Slop / University

Keep the Instruments Running

The case for the infrastructure behind the University's evidence

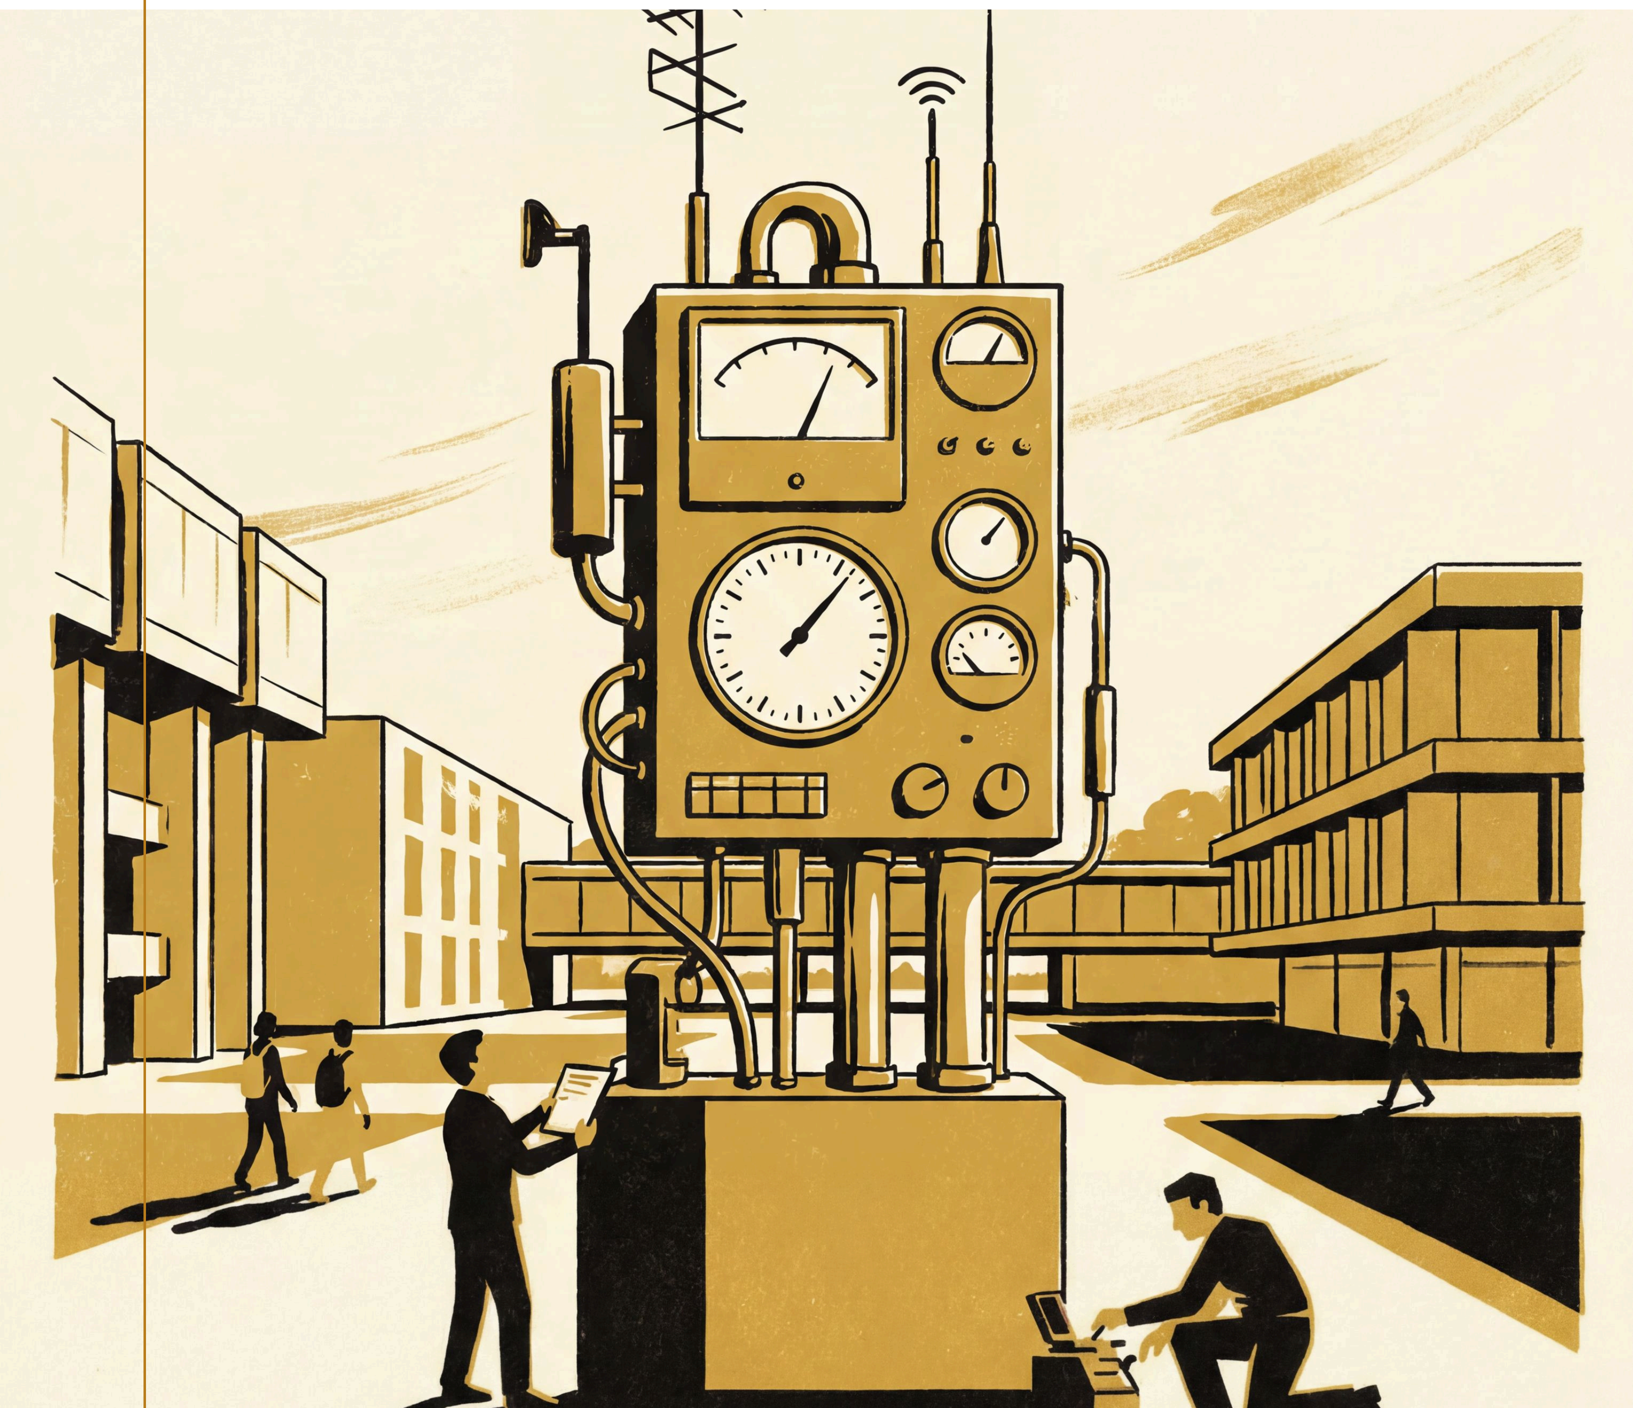

Come and count with us

Some parts of a university are visible from the gate: the lecture halls, the library stacks, the researchers who explain their work at open days. Behind them sits a second campus that visitors rarely see — the instruments. A queue-length sensor at the coffee cart. A discovery-search relevance model in the library stack. A desk-booking feed that tells facilities a building is quietly emptying before anyone files a form. A controller balancing a growth target against an admissions cap. None of these keep themselves running.

Every finding this University publishes rests on an instrument someone built, calibrated, and has kept in good repair since — a discipline that shows up on no cover page and asks for support all the same. This brochure is that ask, made plainly: what it takes to keep the instruments running, across every school, every lab, and every dashboard the University depends on to know itself.

The pages that follow set out where that infrastructure lives, what it has already found, and the partnership that keeps it in calibration for whoever reads it next.

— The Office of Research Outputs

At a glance

What keeping the instruments running actually involves, before the detail.

An instrument in every school

Every research group keeps at least one live instrument in the field — a sensor, an index, or a dashboard feed — under continuous calibration rather than retired once the paper is out.

A record that resolves

Every finding carries a persistent identifier that resolves back to the full study, so a claim can be checked rather than taken on the University's word.

Upkeep, not a one-off

An instrument is calibrated, drifts, and is recalibrated again; the Review Cadence Observatory exists because that cycle needs its own attention, not a single launch event.

A campus that feeds back

Desk bookings, meeting minutes, code reviews, and the coffee queue all return something to the institution — a feedback loop the University is still learning to read well.

Where ideas count

Sensing what the institution already knows

The School of Emergent Priorities runs on the premise that most institutional surprises were visible in the data well before anyone called them a surprise. Its Trajectory Analytics Group, led by Professor Verity Marris, builds and maintains the instruments that make that visibility usable: desk-booking feeds that flag a building's quiet decline, discovery-search telemetry that tracks whether a smarter system actually helps anyone find what they came for, capability curves read early enough to matter. None of these instruments announce themselves loudly. They sit under the floor of ordinary University business, updated on schedules nobody outside the Group thinks about, and they are the reason the School's forecasts read as considered rather than lucky.

Supporting the Group means supporting the unglamorous middle of that pipeline — the recalibration, the data-quality checks, the quiet decommissioning of an instrument that has stopped telling the truth — the part a paper's abstract never mentions and a University's evidence cannot do without.

> *"An instrument that never needs recalibrating has usually stopped watching anything interesting. Ours are watched as closely as what they watch."*
>
> *— Verity Marris, Professor and Director, Trajectory Analytics Group, School of Emergent Priorities*

An instrument for every process

Where the School of Emergent Priorities senses what is coming, the School of Continuous Improvement measures what is already happening, and asks whether the measuring is helping. Its Adaptive Metrics Lab, led by Associate Professor Casimir Beng, has built instruments for processes as different as code review, helpdesk triage, and weekly enrolment offers — each one designed to make a process legible to itself, and each one watched for the moment it starts changing the very behaviour it was built only to observe.

That watching is the Lab's real product. An indicator that quietly trains the people underneath it to perform for the indicator, rather than the outcome, is a common and mostly invisible failure mode; catching it early is slower and less publishable work than building the indicator in the first place, and it is the work this appeal exists to fund.

Keeping the instruments funded

Behind both Schools sits the Office of Research Outputs, the unit that mints a persistent identifier for every finding and keeps the record resolvable years after the study that produced it has left the news cycle. Alongside it, the Review Cadence Observatory tracks the rhythm at which the University's own programs, labs, and initiatives come up for review, so a review cycle that has quietly drifted from what its own indicators would set gets noticed by someone whose job is exactly that.

Neither unit produces a headline finding of its own. Both are the reason the findings that do make headlines can still be checked, and the reason a research group that stops maintaining its instruments finds out sooner than it would otherwise.

> *"A review cadence that drifts far enough eventually reviews nothing at all. Someone has to keep track of the drift itself, and that is the job."*
>
> *— Osei Vandermeer, Research Fellow and Convenor, Review Cadence Observatory, School of Continuous Improvement*

Ideas in print

Four recent pieces make the case better than any brochure could, each resolvable in full on the University's site:

- **Findable, Not Found** — a four-release study asking why a sharper library discovery-search engine kept finding more without patrons retrieving more, and what that gap says about the instruments layered on top of it. doi:10.5555/slop.lwklbz
- **The Unread Diff** — tracking an AI code-review assistant's rising accuracy against how closely engineers still read the diff underneath it, and whether the two curves were ever going to stay in step. doi:10.5555/slop.rq8fyn
- **The Building Knows First** — reading desk-booking cancellations and swipe-card telemetry well ahead of the facilities review that eventually reaches the same conclusion. doi:10.5555/slop.1r34av

- **Two Setpoints, One Controller** — a controller asked to hold an admissions cap and a growth target set on two different committee calendars, and what happens when the calendars drift apart. doi:10.5555/slop.t29aao

The people behind the work

The instruments in this brochure are kept running by people, not systems. Among them:

- **Anneke Tolan**, Senior Research Fellow in the School of Emergent Priorities, studies how foresight is actually practised, from the workshops and templates to the quiet negotiations over wording that decide what a forecast ends up saying.
- **Casimir Beng**, Associate Professor and Lead of the Adaptive Metrics Lab in the School of Continuous Improvement, studies measurement systems that improve the things they measure, and the moments when they stop.

Be part of it

An instrument that nobody funds to maintain eventually just stops, quietly, and the University finds out about the gap it leaves at the same time as everyone else. That is the case this brochure has tried to make plainly: the findings are only as good as the upkeep behind them, and the upkeep is where support is needed most.

Partnering with Slop University means funding the parts of the research pipeline that a press release rarely mentions — the recalibration, the review cadence, the quiet decommissioning of an instrument that has stopped telling the truth — so that the next finding is built on ground as solid as the last one.

The full published record, every instrument it rests on, and the researchers who keep it running are at slop.university. Come and see what we are keeping running, and consider what you would help us keep running next.

Slop / University