



Slovenian / University

Impact Report 2021–2026

A retrospective



Contents

Director’s reflection	3
By the numbers	4
Numbers in context	4
About the School	5
How we’ve delivered	5
The next chapter	9
With thanks	10

Director's reflection

Five years ago, a query typed into a University system was still, mostly, a question addressed to a person. That line has not survived intact — not because anyone decided to move it, but because it accumulated, system by system, integration by integration, until a colleague drafting a report, marking an assignment, or triaging a ticket was as likely to be typing to a model as to a form.



This report is the School's attempt to look at that accumulation directly, with the same instruments it asks everyone else to submit to: a logged count beside an independent audit, printed together rather than reconciled into one figure. The honest answer, five years in, is partial. Some of what the assistants were meant to lighten has genuinely lightened. A good deal of it has simply moved, uncounted, into the work of checking the assistants themselves. We publish both findings, and the gap between them, on the view that a School that measures things for a living owes its own community the same scrutiny it exports everywhere else.

Director, School of Continuous Improvement

By the numbers

Five years, and the four figures the School is asked for most often.



1.2 million logged interactions

Assistant sessions logged across embedded University systems since the first pilot integration in 2022, drawn from the Living Dashboard's own count.



83 audited hours per unit

The average independently audited change in task time recorded per unit across the five-year comparison, blind to any system's own usage log.



47 partner organisations

Government, industry and university partners who opened their own adoption data to comparison against the School's audits.



312 publications drew on the audit

Peer-reviewed and practice output across the period citing at least one logged-versus-audited comparison from this programme.

Numbers in context

Every system this report examines already keeps a number: a query logged, a ticket auto-triaged, a draft time-stamped, a session closed. None of those numbers, on their own, say anything about workload — only about use. Where this report claims a change in workload, it pairs the system's own logged count against an independent audit — a time-diary sample, a walked case review, or a matched before/after task-timing exercise — run by researchers without access to the system's own dashboard while they worked.

The two numbers are reported side by side throughout, not reconciled into one, because reconciling them is the mistake the School is trying to document. A rising

logged-usage count and a flat or worsening audited-workload figure are not a contradiction to be resolved before publication; they are the finding. Readers used to a single headline number per initiative should expect, instead, a pair — and should treat any section where the two move together as the exception it currently is.

About the School

The School of Continuous Improvement studies how organisations count, and what the counting does to them. It traces its origins to the Institute for Measurable Outcomes, and works across measurement systems and instruments, evaluation ecosystems, organisational learning, data cultures and dashboards, impact science, and anticipatory performance — six research themes joined by a single interest: the point at which a measure stops describing the thing it was built to track and starts substituting for it.

The Master of Applied Measurement remains the School's flagship programme, training each cohort to treat an indicator as measurable, repeatable and improvable rather than merely convenient. Capability Sprints extend the same habits into short, applied courses for practitioners who need the tools without the degree. Our researchers work as much inside the University's own operations as outside them: the Adaptive Metrics Lab, the Indicator Commons and the Living Dashboard are as often the object of study as the instrument doing the studying. That reflexive posture — measuring the University with the same rigour it applies to anyone else's systems — is not a caveat on the School's work. It is close to the whole of it.

How we've delivered


The four areas below carry the figures behind them; where a logged count and an audited figure disagree, both are printed.

The Living Dashboard learns a new tile

The Living Dashboard added its first assistant-usage tile in the second year of the reporting period: a live count of sessions opened against the University's embedded drafting and triage assistants, updated hourly beside the fleet-uptime and enrolment tiles it now sits among. Dr Mirela Hanke's team put the new tile through the same

test the wall gives every other indicator on it: whether a rising count corresponds to anything happening off the dashboard.

Working with a services-delivery agency that had embedded a comparable assistant into its own caseworker software eighteen months earlier, the team matched the agency’s session-count trend against its own retrospective audit of caseworker hours. The agency’s dashboard, like the University’s, had reported steady month-on-month growth in assistant use throughout its rollout; its independent workload audit, run without reference to the vendor’s own reporting, found average caseworker hours per completed case unchanged across the same window — the comparison the School had built the tile-audit exercise to run at home.



Assistant sessions logged and independently audited task-hours saved, indexed to 2022 = 100. By 2024 the logged-sessions index had passed 480 while the audited-savings index sat at 98, a gap that has not narrowed since.

“A tile that only goes up is not, on its own, evidence that anything is going down. We built the audit precisely so the two claims could not be confused for each other.”


— Dr Mirela Hanke, Postdoctoral Fellow, School of Continuous Improvement

What the assistant actually shortened

The Impact Pathway Atlas exists to trace how research effort becomes counted impact; this period, Dr Sten Okwuosa’s team pointed the same method at a shorter pathway — how a logged assistant interaction becomes a shortened task. The question, stripped of its usual citation-flow framing, was simple: does more logged use predict less time spent.

In partnership with a big-four professional services firm running its own AI-copilot efficiency review, the team compared per-consultant query counts against blind timings of matched deliverables completed with and without assistant use. The firm had expected, going in, a roughly linear relationship between queries and delivery speed. The matched timings showed no reliable relationship in either direction once


consultant seniority was held constant — a result the firm’s own internal review had not yet reached when the School’s audit began.



Logged queries per completed deliverable against change in completion time, by consultant seniority (214 matched deliverables). The fitted slope is statistically indistinguishable from zero for both seniority bands.

Evaluating the evaluators, again

Evaluation of Evaluation exists to ask whether a review changes anything beyond re-certifying what a panel already believed. Applied to the assistants themselves, its standing method asked a narrower question this period: does adding an AI assistant to a review process shorten the review, deepen it, or simply relocate the same number of hours somewhere else in the cycle.




Working with a Nordic evaluation agency piloting an assistant-drafted first pass on its own programme reviews, Dr Dagny Okafor's team timed the full review cycle before and after the assistant's introduction. Draft turnaround fell by nearly half. Total cycle time, from draft to panel sign-off, held within three per cent of the pre-assistant baseline, the saved drafting hours reabsorbed almost entirely into an additional verification pass the panel had not previously required.

Forecasting the assurance hours before they land

Every dashboard tile the School has added for the assistants carries a twin: a leading-indicator model that tries to forecast where assurance work — the checking, re-verifying and re-approving that follows automated drafting — will land before the logged-usage figures show it. The Adaptive Metrics Lab built the first version of that model in the second year of the reporting period.

The model was scored against a regional university network's parallel rollout, which shared eighteen months of assistant-adoption data across four campuses on the condition its own leading-indicator model be scored against the School's in return. Associate Professor Casimir Beng's model forecast assurance-hour growth within a plausible band for three of the four campus types; the fourth — professional-services units, where adoption moved fastest — required assurance hours the model had significantly underestimated, a miss the Lab has carried into its next forecasting cycle.



Forecast against logged assurance hours by unit type, final reporting year. Professional-services units logged 1,860 assurance hours against a forecast of 1,140 — a 63 per cent miss the model’s next iteration is built to close.

“A forecast that is right about three campuses and wrong about the fastest-moving one has told us where to look next, not that it works. We would rather publish the miss than the average.”

— Associate Professor Casimir Beng, Lead, Adaptive Metrics Lab

The next chapter

Over the next planning period, the School will extend the logged-versus-audited pairing this report has applied to the assistants into a standing feature of the Living Dashboard itself, rather than a one-off audit exercise: every tile that reports on an automated system will carry a companion audited figure, refreshed on the same cycle rather than left to a retrospective report to catch up.



The Adaptive Metrics Lab will take the professional-services forecasting miss as the starting case for a revised leading-indicator model, tested first against the regional university network before any wider rollout. The Indicator Commons will begin a formal retirement review for any assistant-adoption metric that has run three consecutive cycles without a paired audit, on the view that an unaudited number should not keep its place on a dashboard indefinitely simply because nobody has proposed removing it. Evaluation of Evaluation will extend its review-cycle timing method beyond the Nordic pilot to a wider set of the School's own committee processes, on the expectation that the same reabsorption pattern — hours saved in one place resurfacing in another — will recur outside drafting and review work entirely.

None of this promises that the gap between what an assistant logs and what it changes will close. The School's more modest commitment is to keep publishing the two numbers next to each other for as long as the gap persists, on the view that a measurement culture's real test is whether it can tolerate an unflattering pair of figures for more than one reporting cycle.

With thanks

This report draws on the cooperation of services-delivery and evaluation partners across several sectors — government agencies, university networks, professional-

services firms, and research consortia among them — who opened their own logged-usage and audited-workload figures to comparison alongside the University's. We thank the dozens of staff across every School and professional-services unit who sat for interviews, submitted to time-diary sampling, or simply tolerated a researcher walking a case review a second time to check the first count.



We are grateful to the reviewers, external and internal, who read the audits underlying this report's pairs of figures before either number reached a dashboard. And we thank the wider University community for its patience with a School that, having spent five years asking everyone else's systems to submit to measurement, spent a portion of this reporting period turning the same instruments on its own.



Slop / University