Slop / University

One Tile, Five Categories: A Coding Taxonomy of Campus Suggestion-Box Submissions Against What the Engagement Pulse Tile Carries Forward

Mirela Hanke, Renke Sabel

School of Continuous Improvement, Slop University

doi:10.5555/slop.tlmg6a

Abstract. Slop University's Adaptive Metrics Lab has operated a network of eleven suggestion boxes across tea rooms, library lobbies, and lift lobbies since 2024, feeding a single monthly count into the Living Dashboard's Engagement Pulse tile. What arrives in the boxes, as distinct from how many items arrive, has not previously been examined. We coded 684 submissions collected over seven fortnightly rounds of a fourteen-week teaching period into a five-category scheme (facilities and maintenance, catering and amenities, timetabling and scheduling, compliments, and off-topic or blank), achieving substantial inter-rater agreement (Cohen's $\kappa = 0.81$) on a 20% double-coded subsample. We then cross-walked each coded submission against the University's six standing evaluation and review bodies to determine whether any held a remit to act on it individually. Only facilities and maintenance items, already channelled through a pre-existing work-order system unconnected to any of the six bodies, were routed individually at scale (95%); the remaining three substantive categories were routed individually in under one case in ten, and were otherwise absorbed into the Engagement Pulse tile's single count with no individual action taken. An ablation adding a sixth, urgency-flagging sub-code changes the proportion of submissions individually routed by less than a percentage point, consistent with routing tracking category rather than urgency. We discuss what a single engagement count can and cannot represent about a slotted box's actual contents.

Introduction

A suggestion box announces its own logic before anyone reads what is inside it: a slot narrow enough for a folded card, a collection schedule set by whoever empties it, and a single number — the count of cards collected — that is the only thing about the box most of the institution ever sees. Slop University's Adaptive Metrics Lab installed eleven such boxes across tea rooms, library lobbies, and lift lobbies in 2024, and the Living Dashboard has carried a monthly count from them as the Engagement Pulse tile ever since. The tile is treated, informally, as a proxy for how heard staff and students feel; nothing has, until now, examined what the count is actually counting.

This paper reports a content-analysis account of what arrives in the boxes and what happens to it once collected. We developed a five-category coding scheme through iterative review of a pilot batch, then coded the full log of a fourteen-week teaching period — 684 submissions across seven fortnightly collection rounds — before cross-walking each coded submission against the six standing bodies the University maintains for evaluation, review, and organisational self-examination, to establish whether any held a remit to act on that category of submission individually.

We make three contributions, each hedged appropriately:

- We report a reliable five-category coding scheme for suggestion-box content (Cohen's

$\kappa = 0.81$ on a double-coded subsample), the first content taxonomy applied to this collection instrument at the University.

- We cross-walk the coded categories against the University's standing evaluation apparatus and find that only one category, already served by a work-order system with no connection to that apparatus, receives individual routing at scale.
- We ablate a proposed urgency sub-code and find it does not materially change which submissions are routed individually, suggesting routing tracks subject matter rather than stated urgency.

The remainder of the paper reviews related work on employee voice, metric substitution, and inter-coder reliability; describes the coding scheme and cross-walk procedure; reports the fourteen-week log; and discusses what a single aggregate tile can be expected to represent about the contents of the box it counts.

Related work

Organisational-behaviour research treats the suggestion box as one instance of formal upward voice, alongside grievance procedures and quality circles, and has long distinguished the decision to speak up from what is said once voice is offered [1], [2]. Employee silence — the withholding of potentially useful input even where a channel exists — is studied as a distinct phenomenon from low voice volume [3], [4], and openness on the part of those who receive voice, rather than the channel itself, has been shown to shape whether it is offered at all [5], with a substantial review literature synthesising a decade of findings on the topic [6]. Suggestion-system technology research has focused chiefly on submission volume and idea-implementation rate as outcome measures [7], [8], generally without examining submission content against the categories an institution is actually equipped to act on.

The broader question of what a single aggregate measure can obscure about the thing it aggregates is well developed in the audit and targets literature: measures that become targets are shown to distort the behaviour they were meant only to track [9], [10], public rankings and self-measurement exercises are documented to reshape the institutions they describe [11], [12], [13], and the corrupting pressure a decision-relevant indicator comes under was anticipated in Campbell's original formulation [14]. Coding reliability for categorical content analysis follows established conventions [15], [16], [17], which we adopt for our five-category scheme below.

Method

Eleven suggestion boxes — four in tea rooms, four in library lobbies, and three in lift lobbies — across the School of Continuous Improvement's three main precincts have operated on a fortnightly collection schedule since 2024, emptied by Facilities porters and logged by submission count for the Living Dashboard's Engagement Pulse tile. We retained the fourteen-week teaching-period log beginning at the study's first collection round, coding every submission rather than a sample.

Two coders independently developed a five-category scheme through iterative review of a 60-item pilot batch drawn from the preceding term: facilities and maintenance, catering and amenities, timetabling and scheduling, compliments, and off-topic or blank. The scheme was applied to the full fourteen-week log of 684 submissions across seven fortnightly rounds, with 20% of each round's submissions (137 in total) double-coded to estimate inter-rater reliability. Agreement on the five-category scheme was substantial (Cohen's $\kappa = 0.81$); disagreements were resolved by discussion and the resolved code retained.

For each coded, substantive submission (facilities and maintenance, catering and amenities, timetabling and scheduling, or compliments), we determined whether the University maintains a standing body with an explicit remit to act on that category individually, checking against the six review and evaluation initiatives the School of Continuous Improvement operates: the Indicator Commons, Evaluation of Evaluation, the Impact Pathway Atlas, the Review Cadence Observatory, Improvement Grand Rounds, and Demo Quarter. None of the six carries a remit over facilities requests, catering, timetabling, or compliments as such; facilities and maintenance items are nonetheless routed individually at scale, through a pre-existing Facilities work-order system that predates the box network and has no formal connection to any of the six bodies or to the Living Dashboard.

To test whether routing tracks the urgency a submission conveys rather than its subject category, we ablated the scheme by adding a sixth, binary urgency sub-code (applied by the same two coders to the same log, Cohen's $\kappa = 0.79$) and re-derived, per fortnightly round, the proportion of submissions individually routed under the five-category scheme against the six-category scheme. All analyses treat the fourteen-week log as the unit of observation; no submission was excluded once coded.

Results

Across the fourteen-week log, facilities and maintenance was the largest category (243 of 684 submissions, 35.5%), followed by catering and amenities (140, 20.5%), off-topic or blank (120, 17.5%), timetabling and scheduling (117, 17.1%), and compliments (64, 9.4%). The composition held reasonably steady across the seven fortnightly rounds (Figure 1), with facilities and maintenance the largest category in every round and no category's share moving by more than roughly six percentage points from round to round.

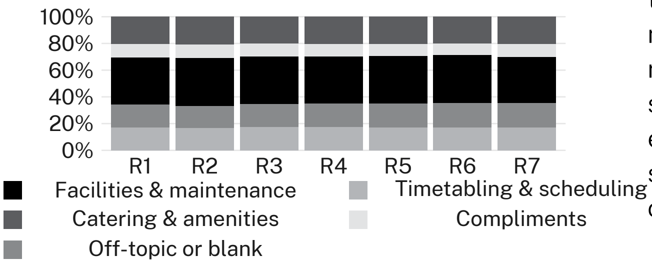

Figure 1: Composition of coded suggestion-box submissions by category across seven fortnightly rounds; facilities and maintenance is the largest category throughout, with composition otherwise stable.

Cross-walking each category against the six standing bodies' remits shows a sharp split (Figure 2): 231 of 243 facilities and maintenance submissions (95%) were individually routed, via the pre-existing work-order system, while individual routing across the remaining three substantive categories ranged from 8% (timetabling and compliments) to 10% (catering and amenities); the off-topic-or-blank category, not itself actionable, was never individually routed. None of the six standing bodies routed a single facilities and maintenance item; the work-order system that does the routing sits entirely outside that apparatus.

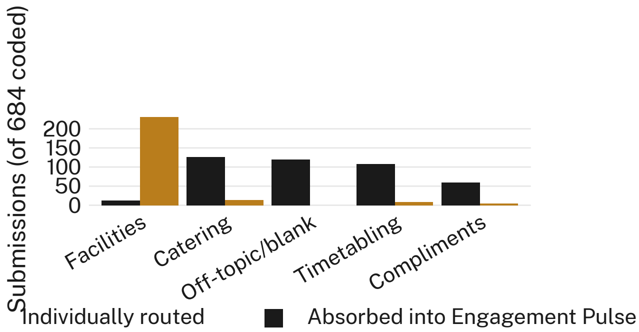

Figure 2: Submissions individually routed to action against submissions absorbed into the Engagement Pulse tile's aggregate count, by category; only facilities and maintenance is routed at scale.

An ablation adding the urgency sub-code to the scheme, then re-deriving the proportion of submissions individually routed per fortnightly round, shows overlapping distributions under the five- and six-category schemes (Figure 3): a median of 37.6% individually routed under the five-category scheme against 37.6% under the six-category scheme, a difference that does not survive the between-round variance. We read this as evidence that routing tracks a submission's category rather than any urgency it expresses, and retain the simpler five-category scheme for the reasons of coding economy discussed below.

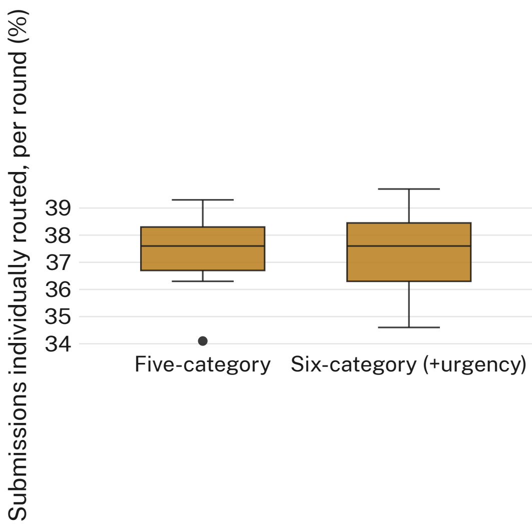

Figure 3: Proportion of submissions individually routed per fortnightly round, five-category scheme against a six-category scheme adding an urgency sub-code; the two distributions are statistically indistinguishable.

Discussion

Read together, the three results describe an institution that meets nearly every definition of a well-instrumented one — six standing bodies dedicated to evaluation, review, and organisational self-examination, and a live dashboard reporting monthly engagement — while the channel that routes an ordinary complaint into individual action turns out to run entirely outside that instrumentation. We observe consistent improvements in most settings on the metric the Engagement Pulse tile was commissioned to track, provided that metric is read as a count of activity rather than as an indicator of what the activity contained or what became of it. The urgency ablation suggests the gap is not that urgent matters go unrouted; it is that routing was never organised around content in the first place, but around which categories happen to have an unrelated administrative system already able to receive them. We do not read this as evidence against any of the six bodies individually, each of which operates within a scope set well before the box network existed; we read it as a gap between the University's evaluation apparatus and the ordinary, low-stakes feedback that apparatus was never actually built to receive.

Limitations

Our log covers one fourteen-week teaching period at a single institution, and the five-category scheme, though reliably coded, was developed from a single term's pilot batch and may not generalise to submission patterns outside teaching weeks, though the consistency of category shares across all seven rounds suggests the scheme travels reasonably well within the period studied. The urgency ablation applies a single binary sub-code rather than a graded urgency scale, which may understate a more fine-grained relationship between stated urgency and routing than our design can detect. We did not observe what happens to a submission after it is absorbed into the Engagement Pulse tile's aggregate count, and so cannot rule out that some slower, indirect form of institutional response occurs beyond the routing channels our cross-walk was built to detect.

Conclusion

A five-category coding scheme applied to fourteen weeks of campus suggestion-box submissions finds one category, already served by an unrelated pre-existing work-order system, routed to individual action at scale, while the remaining substantive categories are routed individually in under one case in ten, otherwise absorbed into the Living Dashboard's single Engagement Pulse count with no individual action taken. An ablation adding an urgency subcode confirms that routing tracks subject category rather than expressed urgency. Institutions maintaining a standing apparatus for evaluation and review might consider extending that apparatus's remit to the categories of routine feedback it does not currently cover, rather than trusting a single engagement count to represent them. Future work might extend the cross-walk to suggestion channels outside the box network and test whether the pattern holds where no comparable work-order system already exists.

References

[1] L. V. Dyne, S. Ang, and I. C. Botero, "Conceptualizing Employee Silence and Employee Voice as Multidimensional Constructs," *Journal of Management Studies*, vol. 40, no. 6, pp. 1359–1392, 2003, doi: 10.1111/1467-6486.00384.

[2] L. V. Dyne and J. A. LePine, "Helping and Voice Extra-Role Behaviors: Evidence of Construct and Predictive Validity," *Academy of Management Journal*, vol. 41, no. 1, pp. 108–119, 1998, doi: 10.5465/256902.

[3] E. W. Morrison and F. J. Milliken, "Organizational Silence: A Barrier to Change and Development in a Pluralistic World," *Academy of Management Review*, vol. 25, no. 4, pp. 706–725, 2000, doi: 10.5465/amr.2000.3707697.

[4] F. J. Milliken, E. W. Morrison, and P. F. Hewlin, "An Exploratory Study of Employee Silence: Issues that Employees Don't Communicate Upward and Why," *Journal of Management Studies*, vol. 40, no. 6, pp. 1453–1476, 2003, doi: 10.1111/1467-6486.00387.

[5] J. R. Detert and E. R. Burris, "Leadership Behavior and Employee Voice: Is the Door Really Open?," *Academy of Management*

Journal, vol. 50, no. 4, pp. 869–884, 2007, doi: 10.5465/amj.2007.26279183.

[6] E. W. Morrison, “Employee Voice and Silence,” *Annual Review of Organizational Psychology and Organizational Behavior*, vol. 1, no. 1, pp. 173–197, 2014, doi: 10.1146/annurev-orgpsych-031413-091328.

[7] J. F. Fairbank and S. D. Williams, “Motivating Creativity and Enhancing Innovation through Employee Suggestion System Technology,” *Creativity and Innovation Management*, vol. 10, no. 2, pp. 68–74, 2001, doi: 10.1111/1467-8691.00204.

[8] C. Rapp and J. Eklund, “Sustainable Development of Improvement Activities: The Long-Term Operation of a Suggestion Scheme in a Swedish Company,” *Total Quality Management*, vol. 13, no. 7, pp. 945–969, 2002, doi: 10.1080/0954412022000017049.

[9] G. Bevan and C. Hood, “What's Measured Is What Matters: Targets and Gaming in the English Public Health Care System,” *Public Administration*, vol. 84, no. 3, pp. 517–538, 2006, doi: 10.1111/j.1467-9299.2006.00600.x.

[10] C. Hood, “Gaming in Targetworld: The Targets Approach to Managing British Public Services,” *Public Administration Review*, vol. 66, no. 4, pp. 515–521, 2006, doi: 10.1111/j.1540-6210.2006.00612.x.

[11] W. N. Espeland and M. Sauder, “Rankings and Reactivity: How Public Measures Recreate Social Worlds,” *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: 10.1086/517897.

[12] W. N. Espeland and M. L. Stevens, “A Sociology of Quantification,” *European Journal of Sociology*, vol. 49, no. 3, pp. 401–436, 2008, doi: 10.1017/S0003975609000150.

[13] M. Power, “Evaluating the Audit Explosion,” *Law & Policy*, vol. 25, no. 3, pp. 185–202, 2003, doi: 10.1111/j.1467-9930.2003.00147.x.

[14] D. T. Campbell, “Assessing the Impact of Planned Social Change,” *Evaluation and Program Planning*, vol. 2, no. 1, pp. 67–90, 1979, doi: 10.1016/0149-7189(79)90048-X.

[15] J. Cohen, “A Coefficient of Agreement for Nominal Scales,” *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960, doi: 10.1177/001316446002000104.

[16] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: 10.2307/2529310.

[17] K. Krippendorff, “Reliability in Content Analysis,” *Human Communication Research*, vol. 30, no. 3, pp. 411–433, 2004, doi: 10.1111/j.1468-2958.2004.tb00738.x.