Slop / University

Guessed at the Door, Confirmed at the Table

A Field Validation of a Book Club Host's Skip-Likelihood Index Against an Anonymous Show of Hands

Runa Adegoke[1], Renke Sabel[2]

[1]School of Emergent Priorities, [2]School of Continuous Improvement, Slop University

doi:10.5555/slop.3hjzqw

Abstract. Neighbourhood book clubs face a familiar problem without a formal name: knowing who read the book without asking anyone to admit they didn't. We report the design and field validation of the Skip-Likelihood Index (SLI), a five-item instrument through which a host scores each attendee, before discussion closes, on observable proxies for non-completion — plot-detail hedging, late acquisition of a borrowed copy, requests to have the plot re-summarised, cross-referencing the film adaptation, and arriving after the opening recap. Each meeting closes with an anonymous show of hands recording how many attendees confess to skipping the book, taken only after the host's scores are sealed. Across 214 meetings of 27 clubs over eight months, the SLI's five items show good internal consistency (Cronbach's $\alpha = 0.81$) but only modest meeting-level agreement with the confession count ($r = 0.44$). A single item, borrowed-copy lateness, accounts for nearly all of this relationship; the other four correlate with the outcome at $r = 0.09$ with lateness removed. Independent hosts co-scoring the same meeting agree only fairly on individual attendees ($\kappa = 0.29$). We read the gap between the index's reliability and its validity as evidence that host impressions cohere with each other more readily than they cohere with the truth, and suggest against fielding five items where one, honestly scored, would do.

Introduction

A book club's host carries an unspoken governance problem every month: running a good discussion without embarrassing the members who did not finish, or did not start, the book. Most clubs manage this the way most informal institutions manage awkward information — by not measuring it. A host reads the room, steers the discussion around anyone who looks lost, and who actually read the book stays a private impression rather than a recorded fact.

Compliance with assigned reading is itself an old empirical problem in a directly analogous setting: compliance with assigned reading in university courses has been measured directly and repeatedly found to fall well short of full completion even where a course formally requires it [1]. A book club carries none of a course's formal enforcement — no quiz, no grade — so if anything the same gap should run wider.

Book clubs are also worth taking seriously as reading communities in their own right rather than as watered-down classrooms: they have been adapted directly into university teaching precisely because they replicate deliberative reading practice outside a graded setting [2].

Formalising a host's private impression into a scored instrument promises two things an unscored impression cannot: consistency between hosts, who currently improvise, and a defensible answer when a discussion needs to route around someone. It also risks a specific, well-documented failure mode: an instrument whose items agree with each other can look reliable while still measuring something other than the thing it was built to measure [3], [4]. Whether a host-facing "skip index" clears that bar, rather than merely appearing to, has not previously been tested against an independent signal.

This paper reports a field validation of the Skip-Likelihood Index (SLI), a five-item host-scoring instrument, against an anonymous end-of-meeting show of hands collected across a regional network of book clubs. The paper's contributions, each modest by design:

1. We design a five-item, host-scorable instrument for estimating book-club completion at the point of discussion, built from proxies a host can observe without asking anyone directly.
2. We validate the instrument against an anonymous, low-stakes ground-truth signal, collected independently of the host's own scores, across a substantial multi-club deployment rather than a single pilot group.
3. We show that the instrument's internal consistency does not imply its criterion validity, and that a considerably simpler single-item measure performs comparably — a caution we suggest travels beyond this particular index.

The SLI is not shown to be useless, and this paper does not attempt to; it is shown to be, on the evidence collected here, doing most of its work through one item wearing four passengers.

Related work

Self-reported reading time is shown to be biased relative to logged app data, with logged measures proposed as the more reliable alternative for children's at-home reading [5]; how people behave when asked to self-report at all, independent of what they over- or understate, has been investigated directly in long-running self-report studies [6]. Social-desirability response bias — the tendency to answer in the way that reflects best on oneself, independent of the answer's truth — has an established measurement tradition of its own [7], and the closely related literature on impression management treats the shading of exactly this kind of information as a predictable, motivated act rather than a lapse [8].

Group settings with an uneven, hard-to-observe individual contribution carry their own adjacent literature under the heading of social loafing: structured peer assessment is shown to reduce free-riding on group projects when it is specific and repeated rather than a single end-of-project judgement [9], though looking beyond loafing itself to why a member does not contribute is argued to matter for the remedy that follows [10]. Reading compliance specifically has been measured directly in university courses, most often by surprise quiz, and found to sit well below full completion even under a course's formal requirement [1], [11].

The instrument-design tradition this paper follows treats internal consistency and criterion validity as separate questions that a single coefficient cannot answer for each other [3], [4], a distinction recent instrument-validation work continues to apply in new domains [12], [13], with resampling methods available for comparing two alpha coefficients directly rather than eyeballing the difference [14]. A related pattern is documented in a volunteer-run verification setting at this institution, where override decisions on an automated authorship-check flag were shown to track cues absent from the score's own feature set far more than they tracked the score's own magnitude [15] — the same asymmetry between what a human notices and what an instrument scores that motivates treating the SLI's item-level signal, not only its total, as the object of interest here.

Method

Clubs and meetings

Clubs were recruited through a regional book-club network spanning nineteen library-affiliated reading groups and eight independent groups meeting in private homes; all 27 opted in to host-side scoring for the length of the study. Across the eight-month window the network logged 214 meetings, a mean of 7.9 meetings per club and 7.4 attendees scored per meeting.

Quantity	Value
Clubs (library-affiliated / independent)	27 (19 / 8)
Meetings (study window)	214 (8 mo.)
Attendees scored / meeting (mean)	7.4
Two-host co-scoring subsample	38 mtgs

The Skip-Likelihood Index

For each attendee j at meeting i, the host records five binary items during discussion — plot-detail hedging, borrowed-copy lateness (copy obtained under 48 hours before the meeting), a reread request, film-adaptation cross-referencing, and late arrival after the opening recap — and sums them unweighted:

$$\mathrm{SLI}_{i,j} = \sum_{k=1}^{5} x_{i,j,k}, \quad x_{i,j,k} \in \{0,1\}.$$

An item-weighted variant was piloted in the network's first month and dropped: weights derived from an early subsample did not generalise to the following month's meetings, and hosts reported the unweighted tally easier to complete without breaking the discussion's pace. The unweighted SLI is what this paper reports throughout.

Show of hands and blinding

Each meeting closes with the host asking for an anonymous show of hands: attendees who did not finish, or did not start, the book raise a hand, counted only in aggregate. Host SLI scores are sealed — written on a tally sheet and photographed — before the show of hands is taken, so no host's scoring can be adjusted by what the room admits. A meeting's SLI-predicted skip count is the number of attendees scoring at least 2 of 5 items; the confessed count is the show-of-hands total.

Reliability, validity, and the co-scoring subsample

Internal consistency across the five items is reported as Cronbach's α [3], computed over item-level scores pooled across all 214 meetings. Criterion validity is the meeting-level Pearson correlation between the SLI-predicted count and the confessed count, reported for the full index and, as an ablation, for two item subsets: borrowed-copy lateness alone, and the remaining four items with lateness removed. For 38 meetings, a second host independently co-scored the same attendees without seeing the first host's sheet; agreement on individual attendees (flagged vs. not, pooled across items) is reported as Cohen's κ [16], following the same agreement statistic used to validate an automated score against independent adjudication elsewhere at this institution [17].

Results

Reliability across the SLI's five items is good by conventional thresholds: Cronbach's $\alpha = 0.81$, pooled across the 214 meetings' item-level scores. Read on its own, this number would normally license treating the SLI as a single coherent construct.

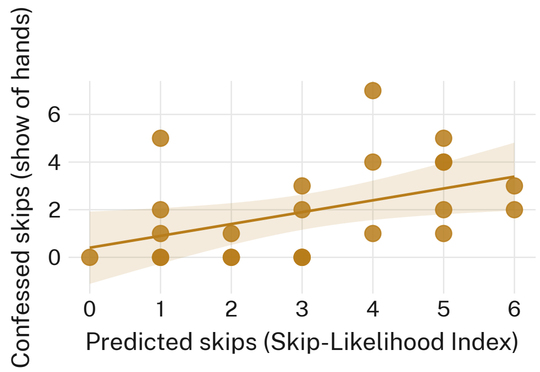

Figure 1: Meeting-level confessed skips (anonymous show of hands) track the host's SLI-predicted count only loosely across 214 meetings ($r = 0.44$); the fitted line's shallow slope tracks the correlation, not a cleaner underlying relationship it is hiding.

Validity tells a different story (Figure 1). The meeting-level correlation between the SLI-predicted count and the confessed count is $r = 0.44$ — real, but well short of what the index's own internal consistency would suggest an evaluator should expect.

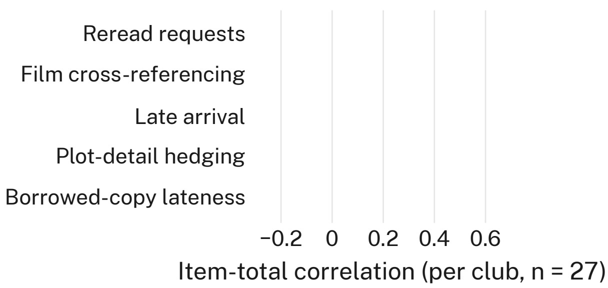

Figure 2: Per-club item-total correlation with the confession outcome, by SLI item (27 clubs). Borrowed-copy lateness sits well above the other four, whose distributions cluster near zero and cross it in both directions.

The item-level picture (Figure 2) explains why. Borrowed-copy lateness alone correlates with the per-club confession rate at a median around 0.58; the other four items — plot-detail hedging, reread requests, film cross-referencing, and late arrival — cluster near zero, several clubs showing a small negative relationship for any given item. The five items are nonetheless mutually consistent enough to produce $\alpha = 0.81$: hosts who score one item high on a given attendee tend to score the others high too, whether or not that attendee actually skipped the book.

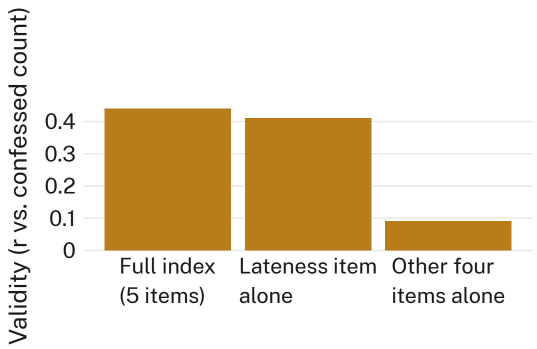

Figure 3: The ablation that changes nothing: dropping four of the SLI's five items and keeping only borrowed-copy lateness costs almost no validity ($r = 0.44$ to $r = 0.41$); the four dropped items alone predict the confession count at $r = 0.09$.

The ablation (Figure 3) makes the item-level pattern concrete at the level the index is actually used. A one-item score — borrowed-copy lateness alone — reaches $r = 0.41$ against the confessed count, indistinguishable in practice from the full five-item index's $r = 0.44$. Removing lateness and keeping the remaining four items produces $r = 0.09$, a relationship not meaningfully different from no relationship at all. We retain the full index in the network's ongoing use for completeness, and because host feedback described the extra items as reassuring to complete even where they add no measurable signal.

The 38-meeting co-scoring subsample shows the same coherence-without- correspondence pattern one level down: two hosts independently scoring the same attendees agree on the internal shape of their impressions well enough to each reach a similar SLI total, but agree on which specific attendees to flag only fairly ($\kappa = 0.29$), a level of agreement associated elsewhere with genuine disagreement rather than measurement noise [16].

Discussion

The gap between the SLI's internal consistency and its criterion validity is not, on this evidence, a flaw specific to this index. It sits closer to what impression-management research would predict from any host-facing scoring task performed live, in the room, without asking directly: a host's five judgements about one attendee are shaped by a single overall impression of that attendee more than by five independent observations, so the items cohere with each other regardless of whether that shared impression tracks the outcome the index was built to predict [8]. Borrowed-copy lateness is the one item plausibly caused by non-completion rather than merely correlated with the impression of it — a copy obtained the morning of the meeting leaves little time to have read it, whatever else the attendee says or does not say during discussion. The other four items read, on this account, as performance cues: how someone talks about a book, not whether they read it.

We do not read this as an argument against structured host scoring in general, only against this particular five-item shape of it. A network that wants the SLI's validity without its four passengers can have it, at the cost of an instrument that looks, to anyone glancing at its item count, less thorough than it is.

Limitations

The 27 participating clubs opted in to host-side scoring, which likely selects for hosts already comfortable making the judgement the SLI merely formalises; a club whose culture makes such scoring unwelcome is not represented here, and the SLI's validity in that setting is untested. The anonymous show of hands is itself a self-report, gathered under exactly the social-desirability pressure the wider literature documents [7]; if anything this sets a conservative floor under the SLI's measured validity, since an undercounted confession total can only weaken an observed correlation, never inflate one the index does not otherwise earn. The two-host subsample was drawn from clubs willing to host a second scorer, a further, smaller selection this paper cannot fully rule out.

Conclusion

An instrument can be internally consistent, field-deployed across more than two hundred meetings, and still be doing almost all of its predictive work through one of its five items. The Skip-Likelihood Index reaches good reliability ($\alpha = 0.81$) and only modest validity against an anonymous confession count ($r = 0.44$), a gap traceable to a single item — borrowed-copy lateness — that the other four merely echo without adding to. Future work might test whether telling hosts this directly changes how they use the index, or whether a network that trims the SLI to one honestly-scored item loses anything besides the four items' reassuring appearance of thoroughness; this study, which only ever watched hosts use all five, cannot answer that on its own.

References

[1] J. Sappington, K. Kinsey, and K. Munsayac, "Two Studies of Reading Compliance among College Students," *Teaching of Psychology*, vol. 29, no. 4, pp. 272–274, 2002, doi: 10.1207/S15328023TOP2904_02.

[2] A. Wyant and S. Bowen, "Incorporating Online and In-person Book Clubs into Sociology Courses," *Teaching Sociology*, vol. 46, no. 3, pp. 262–273, 2018, doi: 10.1177/0092055X18777564.

[3] L. J. Cronbach, "Coefficient Alpha and the Internal Structure of Tests," *Psychometrika*, vol. 16, pp. 297–334, 1951, doi: 10.1007/BF02310555.

[4] D. T. Campbell and D. W. Fiske, "Convergent and Discriminant Validation by the Multitrait-Multimethod Matrix," *Psychological Bulletin*, vol. 56, no. 2, pp. 81–105, 1959, doi: 10.1037/h0046016.

[5] B. Brossette, L. Persia-Leibnitz, M.-J. Chalbos, C. Prugnières, and S. Ducrot, "Measuring Reading Time: Comparing Logged and Self-Reported Data in Relation to Reading Skills," *PLOS One*, vol. 21, no. 4, p. e0344853, 2026, doi: 10.1371/journal.pone.0344853.

[6] A. Möller, M. Kranz, B. Schmid, L. Roalter, and S. Diewald, "Investigating Self-Reporting Behavior in Long-Term Studies," *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, pp. 2931–2940, 2013, doi: 10.1145/2470654.2481406.

[7] D. P. Crowne and D. Marlowe, "A New Scale of Social Desirability Independent of Psychopathology," *Journal of Consulting Psychology*, vol. 24, no. 4, pp. 349–354, 1960, doi: 10.1037/h0047358.

[8] M. R. Leary and R. M. Kowalski, "Impression Management: A Literature Review and Two-Component Model," *Psychological Bulletin*, vol. 107, no. 1, pp. 34–47, 1990, doi: 10.1037/0033-2909.107.1.34.

[9] P. Aggarwal and C. L. O'Brien, "Social Loafing on Group Projects: Structural Antecedents and Effect on Student Satisfaction," *Journal of Marketing Education*, vol. 30, no. 3, pp. 255–264, 2008, doi: 10.1177/0273475308322283.

[10] D. Hall and S. Buzwell, "The Problem of Free-Riding in Group Projects: Looking Beyond Social Loafing as Reason for Non-Contribution," *Active Learning in Higher Education*, vol. 14, no. 1, pp. 37–49, 2013, doi: 10.1177/1469787412467123.

[11] P. Marek and A. N. Christopher, "What Happened to the First "R"? Students' Perceptions of the Role of Textbooks in Psychology Courses," *Teaching of Psychology*, vol. 38, no. 4, pp. 227–232, 2011, doi: 10.1177/0098628311421319.

[12] Y. Song, H. Moon, H. Yang, and C. Kilgore, "Development and Validation of a Faculty Artificial Intelligence Literacy and Competency (FALCON-AI) Scale for Higher Education," *arXiv preprint arXiv:2603.20220*, 2026.

[13] X. Wang, X. Sun, T. Huang, R. He, W. Hao, and L. Zhang, "Development and validation of Critical Thinking Disposition Inventory for Chinese medical college students (CTDI-M)," *arXiv preprint arXiv:1806.11428*, 2018.

[14] M. Pauly, M. Umlauft, and A. Ünlü, "Resampling-based inference methods for comparing two coefficient alpha," *arXiv preprint arXiv:1602.03727*, 2016.

[15] D. Okafor, M. Hanke, and S. Adeyemi, "Cleared by the Volunteer, Not the Rubric: An Interview Study of Override Decisions on a Library Summer Reading Challenge's

Authorship-Check Score," *Slop University technical report*, 2026, doi: 10.5555/slop.t64dph.

[16] J. Cohen, "A Coefficient of Agreement for Nominal Scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960, doi: 10.1177/001316446002000104.

[17] R. Adegoke and S. Adeyemi, "Fastest Claimed, First Flagged: Validating a Neighbourhood Giveaway Group's Templated-Listing Score Against a Volunteer Moderator's Bulk-Repost Adjudication," *Slop University technical report*, 2026, doi: 10.5555/slop.z6ewos.