



Built Not to Watch: How CheckPoint's Participation Logs Reached the Living Dashboard Anyway

Casimir Beng, Osei Vandermeer

School of Continuous Improvement, Slop University

doi:10.5555/slop.hxwf45

Abstract. CheckPoint, the School of Continuous Improvement's opt-in staff wellbeing check-in tool, launched under a published pledge that participation data would remain aggregated above a minimum group size and would never join tables with the School's Living Dashboard. We audit that pledge from the systems side rather than the policy side, analysing eighteen months of the Living Dashboard's own query log for predicates referencing a CheckPoint identifier. A monthly Query Boundary Index (QBI) rose from near zero to over a third of all Living Dashboard queries across the audit window, concentrated in a small number of requesting units. A natural ablation around the pledge's mid-window reaffirmation — a strengthened public restatement of the no-join clause — left the QBI trend statistically indistinguishable before and after, suggesting the growth tracks something other than the wording in force at the time. A parallel staff-awareness item correlates negatively with CheckPoint's own opt-in participation rate across units ($r = -0.52$), consistent with staff responding to joins the pledge says do not happen. We read the pattern as structural rather than deliberate, and propose the audit itself, repeated on a cadence, as the more informative unit of governance than any single restatement of the pledge.

Introduction

Wellbeing check-in tools are usually launched alongside a structural promise: participation is voluntary, responses are aggregated before anyone sees them, and the resulting data stays out of any system that touches an individual's performance record. The promise is meant to buy something specific — staff disclosure that would not happen under a monitoring frame — and the institutions that make it tend to audit whether the promise was kept by re-reading the policy rather than by reading the systems the policy describes. The literature on function creep argues the two need not move together: purpose restrictions decouple from practice gradually and often imperceptibly, in a way that escapes exactly the kind of scrutiny a policy re-read would catch [1]. Work on algorithmic and people-analytics systems more broadly finds a similar pattern at the level of organisational dashboards, where tools introduced for one stated function accrete others once the

pipeline that feeds them already exists [2], [3], [4]. Whether a tool's own host database — rather than its policy documentation — shows the same drift is, to our knowledge, not something a university has previously checked against its own wellbeing programme.

This paper reports an eighteen-month systems-side audit of CheckPoint, the School of Continuous Improvement's staff check-in tool, conducted by analysing the query log of the Living Dashboard it was explicitly built not to feed. Our contributions are threefold:

- we introduce a monthly Query Boundary Index (QBI), a log-derived measure of how often a dashboard's own queries reference an identifier from a tool that pledge documentation says should never be joined to it;
- we report an eighteen-month rise in QBI that survives a natural ablation around the pledge's own mid-window reaffirmation, suggesting the growth tracks something the reaffirmation did not reach;

- we report a unit-level correlation between staff awareness of the joins and CheckPoint’s own opt-in participation rate, consistent with staff responding to practice the policy does not acknowledge.

Related work

Koops’s account of function creep frames purpose drift as usually gradual and rarely deliberate at any single point, which is part of why it survives scrutiny aimed at a single point in time [1]; Biega and Finck make the adjacent case that purpose-limitation obligations are easiest to state and hardest to operationalise once a system’s schema already permits the join [5]. The algorithmic-management literature documents dashboards accreting new managerial functions as an ordinary property of deployed systems rather than a deviation from design intent [2], [3], [4], and the electronic-performance-monitoring literature finds monitoring presence associated with worker stress regardless of whether the monitoring was introduced under a performance frame or a wellbeing one [6], [7], with post-pandemic survey work adding that employees’ acceptance of a given monitoring tool depends heavily on which frame they believe governs it [8]. Worker-surveillance scholarship traces a longer history of tools introduced under a narrow justification later supporting broader ones [9], and audit-culture scholarship separately observes that measurement regimes generate their own reactive behaviour in the people being measured, often independent of what the measurement was nominally for [10]; Manheim and Garrabrant’s taxonomy of Goodhart’s-law failure modes gives that reactivity a vocabulary that extends past optimisation targets to the informal metrics an institution merely watches [11]. On the systems side, k-anonymity supplies the formal notion of aggregation CheckPoint’s own pledge invokes informally [12], database audit-log analysis supplies our method’s closest technical precedent [13], dark-pattern scholarship supplies a vocabulary for consent flows that technically satisfy opt-in while steering toward disclosure [14], and workplace ecological-momentary- assessment tools supply the wider context CheckPoint sits within, most of which carry comparable non-monitoring assurances of their own [15].

Method

System. CheckPoint is a daily opt-in check-in — a mood rating and two short free-text prompts — offered to all School of Continuous Improvement staff since its launch. Its published pledge states that individual responses are never surfaced, that any report leaving the CheckPoint schema is aggregated above a minimum group size of twelve, and that the schema carries no foreign key into the Living Dashboard or any other performance system.

Audit design. The Living Dashboard team retains its own query log for unrelated performance-tuning purposes, giving us eighteen months of already-collected data rather than requiring new instrumentation. For each logged query we test whether its predicate references any of CheckPoint’s join-eligible identifiers — participant IDs, check-in timestamps, or derived cohort tags — via a static hash match against a maintained list of those identifiers, aggregated to a monthly Query Boundary Index:

```
SELECT date_trunc('month', logged_at) AS
month,
       count(*) FILTER (
         WHERE predicate_hash IN (SELECT
           hash FROM checkpoint_join_keys)
       ) * 1.0 / count(*) AS qbi
FROM living_dashboard_query_log
GROUP BY 1
ORDER BY 1;
```

Ablation. In month nine of the audit window, the University reissued the CheckPoint Non-Monitoring Pledge with strengthened language — an expanded FAQ and an explicit restatement of the no-join clause — following a routine policy review unconnected to this audit. We treat months 1–9 and 10–18 as a natural pre/post comparison of the reaffirmation’s effect on QBI, split by the five units issuing the most CheckPoint-referencing queries.

Survey. An existing annual staff item (“I believe CheckPoint data could reach my performance review”) is compared against CheckPoint’s own opt-in participation rate, aggregated per unit across the twelve labs, programs, and initiatives with an active CheckPoint rollout.

Results

Parameter	Value
Audit window	18 months
Queries scanned	612,847
Requesting units tracked	12
Pledge minimum aggregation size	12 respondents
Reaffirmation	issued month 9 of 18

Table 1: Audit parameters, fixed before the window opened or logged throughout it.

The boundary erodes steadily. Figure 1 plots the monthly Query Boundary Index across the full window. QBI rises from 1% in month 1 to 34% by month 18, with no inflection visible around the reaffirmation in month 9 — the rate of increase in the second half of the window is, if anything, slightly higher than in the first.




Figure 1: Query Boundary Index by audit month: the share of Living Dashboard queries referencing a CheckPoint identifier rises from 1% to 34% across eighteen months.

The growth concentrates in a handful of units, and the reaffirmation changes none of them. Figure 2 compares mean QBI for the five heaviest requesting units before and after the reaffirmation. The Performance Review Panel accounts for the largest share throughout (21% before, 23% after); Facilities Scheduling, included as a negative control with no plausible reason to reference CheckPoint, stays near zero in both periods. No unit's before/after difference exceeds two percentage points.




Figure 2: Mean Query Boundary Index by requesting unit, before and after the pledge's month-9 reaffirmation: every unit's two bars sit within two percentage points of each other.

Awareness tracks a fall in participation, not a fall in trust generally. Figure 3 plots CheckPoint's opt-in participation rate against the staff-awareness item, one point per unit. The relationship is negative and moderate ($r = -0.52$): units where more staff believe CheckPoint data reaches performance review show correspondingly lower opt-in rates, though the weakest-participation unit (28%) still retains more than a quarter of staff opting in even at the highest measured awareness.




Figure 3: CheckPoint opt-in participation rate against staff-reported awareness of dashboard joins, by unit ($r = -0.52$): higher awareness predicts lower participation, though the relationship is moderate rather than a cliff.

Discussion

Our results are consistent with the function-creep literature's central claim that purpose drift is typically gradual and largely independent of the policy language nominally governing it [1], extended here to a case where the drift is directly observable in a system's own query log rather than inferred from downstream com-

plaints. That the reaffirmation changed nothing in Figure 2 is, we think, the more informative result than the rise in Figure 1 taken alone: it suggests the growth is a property of which queries the schema makes possible rather than of what staff or reviewers are currently being told about it, a distinction the algorithmic- management literature's dashboard-accretion framing anticipates [2], [3]. The awareness-participation relationship in Figure 3 is consistent with at least some staff responding to the joins rather than the pledge, though the moderate correlation and the substantial residual participation even at high awareness both suggest CheckPoint's value to respondents survives some suspicion of the boundary it promises.

Limitations

Our audit covers one tool and one host dashboard within one school, though the consistency of the QBI rise across twelve tracked units and its imperviousness to a genuine policy reaffirmation both suggest the pattern is a property of the schema rather than of any one team's behaviour. Query-predicate hash matching cannot distinguish a join made for legitimate dashboard-maintenance purposes from one that reaches an individual's record, though our negative-control unit's near-zero QBI throughout suggests the measure is not simply detecting background database activity. The awareness item is self-reported and cross-sectional, which limits any causal reading of the participation correlation, though the direction and unit-level consistency of the relationship are what the audit is built to detect.

Conclusion

An eighteen-month query-log audit of the Living Dashboard finds the share of its queries referencing CheckPoint identifiers rising from 1% to 34%, unaffected by a policy reaffirmation issued partway through the window and concentrated in a small number of requesting units, alongside a moderate unit-level correlation between staff awareness of the joins and CheckPoint's own opt-in participation rate. The boundary CheckPoint was built with appears to hold in its documentation and erode in its schema. Future work will test whether publishing this audit's methodology changes what the query log records, or simply teaches queries to route around the identifiers it watches.

References

- [1] B.-J. Koops, "The Concept of Function Creep," *Law, Innovation and Technology*, vol. 13, no. 1, pp. 29–56, 2021, doi: [10.1080/17579961.2021.1898299](https://doi.org/10.1080/17579961.2021.1898299).
- [2] M. H. Jarrahi, G. Newlands, M. K. Lee, C. T. Wolf, E. Kinder, and W. Sutherland, "Algorithmic Management in a Work Context," *Big Data & Society*, vol. 8, no. 2, 2021, doi: [10.1177/20539517211020332](https://doi.org/10.1177/20539517211020332).
- [3] A. Benlian et al., "Algorithmic Management," *Business & Information Systems Engineering*, vol. 64, no. 6, pp. 825–839, 2022, doi: [10.1007/s12599-022-00764-w](https://doi.org/10.1007/s12599-022-00764-w).
- [4] M. Zhang, F. L. Cooke, D. Ahlstrom, and N. McNeil, "The Rise of Algorithmic Management and Implications for Work and Organisations," *New Technology, Work and Employment*, 2025, doi: [10.1111/ntwe.12343](https://doi.org/10.1111/ntwe.12343).
- [5] A. J. Biega and M. Finck, "Reviving Purpose Limitation and Data Minimisation in Data-Driven Systems," *arXiv preprint arXiv:2101.06203*, 2021.
- [6] D. M. Ravid, D. L. Tomczak, J. C. White, and T. S. Behrend, "EPM 20/20: A Review, Framework, and Research Agenda for Electronic Performance Monitoring," *Journal of Management*, vol. 46, no. 1, pp. 100–126, 2020, doi: [10.1177/0149206319869435](https://doi.org/10.1177/0149206319869435).
- [7] D. M. Ravid, J. C. White, D. L. Tomczak, A. F. Miles, and T. S. Behrend, "A Meta-Analysis of the Effects of Electronic Performance Monitoring on Work Outcomes," *Personnel Psychology*, vol. 76, no. 1, pp. 5–40, 2023, doi: [10.1111/peps.12514](https://doi.org/10.1111/peps.12514).
- [8] J. Vitak and M. Zimmer, "Surveillance and the Future of Work: Exploring Employees' Attitudes Toward Monitoring in a Post-COVID Workplace," *Journal of Computer-Mediated Communication*, vol. 28, no. 4, 2023, doi: [10.1093/jcmc/zmad007](https://doi.org/10.1093/jcmc/zmad007).
- [9] I. Ajunwa, K. Crawford, and J. Schultz, "Limitless Worker Surveillance," *California Law Review*, vol. 105, no. 3, pp. 735–776, 2017, doi: [10.15779/Z38BR8MF94](https://doi.org/10.15779/Z38BR8MF94).
- [10] M. Strathern, "Improving Ratings': Audit in the British University System," *European Review*, vol. 5, no. 3, pp. 305–321, 1997, doi: [10.1017/S1062798700002660](https://doi.org/10.1017/S1062798700002660).

- [11] D. Manheim and S. Garrabrant, “Categorizing Variants of Goodhart’s Law,” *arXiv preprint arXiv:1803.04585*, 2018.
- [12] L. Sweeney, “k-Anonymity: A Model for Protecting Privacy,” *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, vol. 10, no. 5, pp. 557–570, 2002, doi: [10.1142/S0218488502001648](https://doi.org/10.1142/S0218488502001648).
- [13] M. I. Khan, S. N. Foley, and B. O’Sullivan, “Database Intrusion Detection Systems (DIDs): Insider Threat Detection via Behavioural-Based Anomaly Detection Systems,” *arXiv preprint arXiv:2011.02308*, 2020.
- [14] C. M. Gray, Y. Kou, B. Battles, J. Hoggatt, and A. L. Toombs, “The Dark (Patterns) Side of UX Design,” in *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, 2018, pp. 1–14. doi: [10.1145/3173574.3174108](https://doi.org/10.1145/3173574.3174108).
- [15] G. Dell’Aversana, M. Herold, S. Simbula, I. Ruina, M. Pillan, and M. Miglioretti, “Wearable Devices and Ecological Momentary Assessment (EMA) in the Workplace: A Study Protocol on Work Stress Assessment,” *PLOS ONE*, 2026, doi: [10.1371/journal.pone.0339288](https://doi.org/10.1371/journal.pone.0339288).