



Closed, Reclassified, or Left Open: A Hand-Coded Taxonomy of Bathroom-Break Integrity-Attentiveness Flags Across Eight Terms at the Exit Sign-Out Table

Marek Solheim¹, Renke Sabel²

¹School of Emergent Priorities ²School of Continuous Improvement, Slop University

doi:10.5555/slop.jw4nw9

Abstract. Slop University's in-person final examinations route every mid-sitting bathroom break through an exit sign-out table, where an invigilator may raise an integrity-attentiveness flag against a student's departure or return. What happens to a flag once raised has not, to our knowledge, been traced end to end. We hand-coded eight terms of archived exam files (617 flags) against a four-category taxonomy developed by constant comparison on a purposive subsample: formally resolved, quietly reclassified onto an unrelated student file, administratively closed without adjudication, and left open past the exam period. Two trained coders achieved substantial agreement on the remaining corpus (Cohen's kappa = 0.79). Formal resolution accounted for 13.5% of flags; the sign-out table's own "resolved" checkbox, by contrast, was marked in 62.5-72.2% of cases per term, a gap driven almost entirely by administrative closures and reclassifications the checkbox cannot distinguish from genuine adjudication. Flags left open past the exam period outnumbered formally resolved flags in every term studied. An ablation redacting the checkbox from coders before taxonomy application shifted the left-open share by 0.2 percentage points (n.s.); we retain checkbox access for consistency with archival practice. We suggest the checkbox be read as a table-clearance indicator rather than an integrity-outcome indicator.

Introduction

Slop University's in-person final examinations position an exit sign-out table beside the hall door, staffed for the sitting's duration by a rotating invigilator whose remit includes logging every mid-sitting departure and return. Where a departure or return reads as unusual — an extended absence, a changed route, a returned student who does not resume the paper where they left off — the invigilator may raise what the sign-out procedure calls an integrity-attentiveness flag: a note against the student's row entry, timestamped, and left for later action.

Later action is where the record thins. The sign-out table logs the raising of a flag with some care; nothing in the procedure specifies who is responsible for closing one, on what timetable, or against what standard. We were unable to locate, in either the procedure itself or the

University's wider academic-integrity documentation, a single account of what has actually happened to a flag once the exam has ended and the hall has been reset for the next sitting.

This paper supplies one. We hand-coded eight terms of archived exam files — 617 flags in total — against a four-category taxonomy of closure outcome, developed by constant comparison on a purposive sample before being applied to the remainder by two independently working coders. We then set the taxonomy's coding against the sign-out table's own binary "resolved" checkbox, the only closure signal the table itself produces, to ask how well an administratively convenient flag tracks what our coding found underneath it.

Our contributions, each hedged to the scale of an eight-term archival sample:

- We propose a four-category taxonomy — formally resolved, quietly reclassified, administratively closed, and left open past the exam period — that appears to organise closure behaviour the sign-out table's own record does not distinguish.
- We estimate, across eight terms and 617 flags, how each category's share compares against the table's own "resolved" checkbox, and find the checkbox consistent with tracking table-side clearance rather than integrity-outcome resolution.
- We report an ablation suggesting the taxonomy's coding is not simply reproducing the checkbox it is set against, a finding we would welcome a larger, multi-site archival audit revising.

Related work

Academic-integrity research has concentrated overwhelmingly on detection — estimating the prevalence of contract cheating [1], asking whether markers can distinguish genuine misconduct from ordinary suspicion [2], and weighing assessment design against cheating risk [3] — and on deterrence, most visibly through webcam-based proctoring [4], alongside survey evidence on why students seek to outsource work in the first place [5]. Comparatively little of this literature follows a flagged case past the point of detection to ask how, or whether, it is closed.

Street-level bureaucracy offers a vocabulary for that gap. Front-line staff routinely exercise discretion policy documents do not anticipate, resolving ambiguous cases through locally legible categories rather than the formal process on paper [6], [7]; as administrative systems move from paper files toward information systems, some of that discretion migrates upward while some persists at the counter regardless [8], and organisational pressure to demonstrate throughput shapes how front-line units adapt their own closure practices [9]. A parallel literature on alert fatigue, developed mainly in clinical and security settings, documents how frontline staff come to override or discount a system-generated alert once repeated exposure outpaces the capacity to review each on its merits [10], [11] — a dynamic the exit sign-out table's own volume plausibly reproduces.

Our taxonomy-development approach follows established guidance treating an induced classification scheme as a legitimate mode of theory-building in its own right [12], drawing on qualitative content-analysis and thematic-coding procedures for deriving and stabilising categories from field records [13], [14], and on current guidance for establishing and reporting intercoder reliability [15]. Records-management scholarship treats a closed case file as a persistent representation of a decision actually taken [16], a standard several of our categories do not meet. Prior work at this institution has traced a comparable gap between what an internal flag records and what actually happens once it is raised, from an indicator-retirement scorer whose own uptime outlasted almost everything it flagged [17] to phone-surrender enforcement that proved statistically indistinguishable across paper types it had no reason to treat alike [18]; the coding strategy applied here follows a typology this institution has used once before, on a different exit-point ritual entirely [19].

Method

The exit sign-out table sits beside the hall door for the duration of each in-person final-examination sitting. Any invigilator staffing the table may raise an integrity-attentiveness flag against a departure or return judged unusual; the flag is entered as a timestamped note against the student's row on the sitting's sign-out sheet, which is filed with the sitting's other exam records at the term's close. We obtained eight consecutive terms of these archived exam files from the University's central examinations records store, with student identifiers redacted before coding began.

A purposive 15% subsample (93 flags, drawn to span all eight terms and both the mid-sitting and near-close windows in which flags most often occur) was reviewed jointly by both coders to develop the taxonomy by constant comparison, iterating category definitions until no flag in the subsample resisted a clean assignment. Four categories emerged: formally resolved (referred through the University's standing academic-integrity process and closed against a recorded decision); quietly reclassified (folded into an unrelated existing student file, most often an attendance or wellbeing record, with no separate adjudication); administratively closed (closed at or shortly after the table via a generic closure

code, with no referral and no findings recorded); and left open (no closure action of any kind recorded before the term's exam files were archived).

Two coders then independently coded the remaining 524 flags, double-coding a further 20% subsample ($n = 123$) for reliability. Agreement was estimated with Cohen's kappa,

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

where p_o is the observed proportion of double-coded flags on which the coders' category assignments agree and p_e is the proportion expected to agree by chance given each coder's marginal category rate [20]. The double-coded subsample yielded kappa = 0.79, substantial agreement on the conventional interpretation scale, which we report rather than proposing our own threshold [21], [22].

The sign-out table records one further signal per flag: a checkbox, ticked by the invigilator or a later reviewing officer, marking the flag "resolved." We extracted this field alongside our taxonomy coding to compare the table's own closure signal against the categories our coding actually found underneath it. To test whether coders' taxonomy assignments were simply reproducing this checkbox, we re-coded the double-coded subsample a second time with the checkbox field redacted, and compared the resulting category shares against the checkbox-visible coding.

Results

Across the eight terms, coders classified 617 flags (Table 1). Formal resolution — a flag referred through the standing academic-integrity process and closed against a recorded decision — accounted for 83 flags (13.5%). Quiet reclassification onto an unrelated student file was the largest single category (204 flags, 33.1%), followed by administrative closure without adjudication (136 flags, 22.0%) and flags left open past the exam period (194 flags, 31.4%).

Category	Definition	n	%
Formally resolved	Referred through the standing academic-integrity process; closed against a recorded decision.	83	13.5
Quietly reclassified	Folded into an unrelated existing student file with no separate adjudication.	204	33.1
Administratively closed	Closed via a generic closure code; no referral, no findings recorded.	136	22.0
Left open	No closure action recorded before the term's exam files were archived.	194	31.4

Table 1: Four-category taxonomy of exit sign-out table flag closure, applied to 617 hand-coded flags across eight terms.

The distribution held its shape across all eight terms studied (Figure 1): the left-open share exceeded the formally-resolved share in every term, with the widest gap in Term 6 (37.5% left open against 9.1% formally resolved).




Figure 1: Closure-category composition of hand-coded flags by term: flags left open past the exam period outnumber formally-resolved flags in all eight terms studied.

Set against the sign-out table's own "resolved" checkbox, the gap sharpens further (Figure 2). The checkbox was marked "resolved" in 62.5-72.2% of flags per term — effectively everything the table did not leave explicitly open — while the taxonomy's formally-resolved category, which requires a recorded adjudicated decision, applied to only 9.1-16.7% of flags in the same terms. The gap is accounted for almost entirely by quiet reclassification and administrative closure, both of which the checkbox records identically to genuine resolution.




Figure 2: The sign-out table’s “resolved” checkbox against the taxonomy’s stricter formally-resolved rate, by term: the checkbox is marked 4-7 times as often as a flag is actually adjudicated.

We ran one ablation on the coding procedure itself: redacting the sign-out table’s checkbox field from coders before they applied the taxonomy to the double-coded subsample, to test whether the taxonomy’s categories were tracking the checkbox rather than the underlying file. The mean left-open share moved from 31.4% (checkbox visible) to 31.6% (checkbox redacted), a shift of 0.2 percentage points well within the term-to-term variation already visible in Figure 1 (Figure 3). We retain checkbox access in the reported coding for consistency with how the archive is actually reviewed; redacting it does not materially change what the taxonomy finds.




Figure 3: The ablation that changes nothing: redacting the sign-out table’s checkbox from coders before taxonomy application shifts the left-open share by 0.2 points (n.s.).

Two further patterns round out the record without altering it. First, every flag we reviewed resolved cleanly into exactly one of the four categories, with no residual “other” bucket required across the full corpus. Second, the terms with the highest raw flag counts were not the terms with the highest left-open shares, suggesting volume alone does not explain why a given flag goes unclosed.

Discussion

The sign-out table’s checkbox was, on this reading, never built to distinguish why a flag closed, only whether the table considered the matter off its own plate by the time the file was archived. On that narrower reading the checkbox performs adequately: the aggregate “resolved” rate is stable across terms, in part because administrative closure and quiet reclassification both count toward it. The difficulty sits downstream, in what a reader unfamiliar with the taxonomy might infer from a term’s headline resolved rate — that a comparable share of flagged departures received the same substantive academic-integrity adjudication the label implies.

We read the pattern as broadly consistent with accounts of street-level discretion under volume pressure [6], [9]: an invigilator staffing the table for a single sitting has neither the standing nor, plausibly, the time to refer every flag through the formal process, and reclassification onto an existing file or a generic closure code offers a legitimate-looking alternative the checkbox cannot tell apart from adjudication. The alert-fatigue literature’s account of alerts discounted once volume exceeds review capacity [10], [11] fits the pattern in our corpus at least as well as any account specific to examinations.

Limitations

Our corpus spans eight terms at a single institution, and we make no claim the four-category distribution generalises beyond it, though the same four categories organised every flag we reviewed across four years of varying enrolment and staffing, more stability than eight terms strictly warrants claiming. Hand-coding from archived files cannot recover an invigilator’s reasoning at the moment a flag was raised or closed, only the file’s eventual state; a contemporaneous observation arm, in the spirit of prior institutional audits [19], might sharpen category boundaries the file alone leaves implicit. We coded throughout from the sign-out table’s own paper trail, which is also the trail whose checkbox we evaluate the taxonomy against, though the redaction ablation suggests this dependency is not doing the taxonomy’s work for it.

Conclusion

A bathroom-break integrity-attentiveness flag raised at Slop University's exit sign-out table is, on the evidence of eight terms of archived exam files, considerably more likely to be quietly reclassified or left open than formally resolved, whatever the table's own checkbox subsequently records. Our four-category taxonomy — formally resolved, quietly reclassified, administratively closed, and left open — organises 617 hand-coded flags with substantial inter-rater agreement ($\kappa = 0.79$), and an ablation redacting the table's checkbox from the coding process changes the resulting distribution by a fraction of a percentage point. We suggest the checkbox be read as a table-clearance indicator rather than an integrity-outcome indicator, and that any future reporting of flag "resolution" carry the taxonomy alongside the count. Future work might extend the coding to the sitting's other exit-point ritual, the invigilated hand-raise.

References

- [1] P. M. Newton, "How Common Is Commercial Contract Cheating in Higher Education and Is It Increasing? A Systematic Review," *Frontiers in Education*, vol. 3, 2018, doi: [10.3389/feduc.2018.00067](https://doi.org/10.3389/feduc.2018.00067).
- [2] P. Dawson and W. Sutherland-Smith, "Can Markers Detect Contract Cheating? Results from a Pilot Study," *Assessment & Evaluation in Higher Education*, vol. 43, no. 2, pp. 286–293, 2018, doi: [10.1080/02602938.2017.1336746](https://doi.org/10.1080/02602938.2017.1336746).
- [3] R. Harper, T. Bretag, and K. Rundle, "Detecting Contract Cheating: Examining the Role of Assessment Type," *Higher Education Research & Development*, vol. 40, no. 2, pp. 263–278, 2021, doi: [10.1080/07294360.2020.1724899](https://doi.org/10.1080/07294360.2020.1724899).
- [4] K. Hylton, Y. Levy, and L. P. Dringus, "Utilizing Webcam-Based Proctoring to Deter Misconduct in Online Exams," *Computers & Education*, vol. 92–93, pp. 53–63, 2016, doi: [10.1016/j.compedu.2015.10.002](https://doi.org/10.1016/j.compedu.2015.10.002).
- [5] A. Amigud and T. Lancaster, "246 Reasons to Cheat: An Analysis of Students' Reasons for Seeking to Outsource Academic Work," *Computers & Education*, vol. 134, pp. 98–107, 2019, doi: [10.1016/j.compedu.2019.01.017](https://doi.org/10.1016/j.compedu.2019.01.017).
- [6] S. Maynard-Moody and M. Musheno, "State Agent or Citizen Agent: Two Narratives of Discretion," *Journal of Public Administration Research and Theory*, vol. 10, no. 2, pp. 329–358, 2000, doi: [10.1093/oxfordjournals.jpart.a024272](https://doi.org/10.1093/oxfordjournals.jpart.a024272).
- [7] T. Evans, "Street-Level Bureaucracy, Social Work and the (Exaggerated) Death of Discretion," *British Journal of Social Work*, vol. 34, no. 6, pp. 871–895, 2004, doi: [10.1093/bjsw/bch106](https://doi.org/10.1093/bjsw/bch106).
- [8] M. Bovens and S. Zouridis, "From Street-Level to System-Level Bureaucracies: How Information and Communication Technology Is Transforming Administrative Discretion and Constitutional Control," *Public Administration Review*, vol. 62, no. 2, pp. 174–184, 2002, doi: [10.1111/0033-3352.00168](https://doi.org/10.1111/0033-3352.00168).
- [9] E. Z. Brodtkin, "Policy Work: Street-Level Organizations Under New Managerialism," *Journal of Public Administration Research and Theory*, vol. 21, no. Suppl 2, pp. i253–i277, 2011, doi: [10.1093/jopart/muq093](https://doi.org/10.1093/jopart/muq093).
- [10] J. S. Ancker, A. Edwards, S. Nosal, D. Hauser, E. Mauer, and R. Kaushal, "Effects of Workload, Work Complexity, and Repeated Alerts on Alert Fatigue in a Clinical Decision Support System," *BMC Medical Informatics and Decision Making*, vol. 17, 2017, doi: [10.1186/s12911-017-0430-8](https://doi.org/10.1186/s12911-017-0430-8).
- [11] H. van der Sijs, J. Aarts, A. Vulto, and M. Berg, "Overriding of Drug Safety Alerts in Computerized Physician Order Entry," *Journal of the American Medical Informatics Association*, vol. 13, no. 2, pp. 138–147, 2006, doi: [10.1197/jamia.M1809](https://doi.org/10.1197/jamia.M1809).
- [12] R. C. Nickerson, U. Varshney, and J. Muntermann, "A Method for Taxonomy Development and Its Application in Information Systems," *European Journal of Information Systems*, vol. 22, no. 3, pp. 336–359, 2013, doi: [10.1057/ejis.2012.26](https://doi.org/10.1057/ejis.2012.26).
- [13] S. Elo and H. Kyngäs, "The Qualitative Content Analysis Process," *Journal of Advanced Nursing*, vol. 62, no. 1, pp. 107–115, 2008, doi: [10.1111/j.1365-2648.2007.04569.x](https://doi.org/10.1111/j.1365-2648.2007.04569.x).
- [14] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol.

- 3, no. 2, pp. 77–101, 2006, doi: [10.1191/1478088706qp063oa](https://doi.org/10.1191/1478088706qp063oa).
- [15] C. O'Connor and H. Joffe, "Intercoder Reliability in Qualitative Research: Debates and Practical Guidelines," *International Journal of Qualitative Methods*, vol. 19, 2020, doi: [10.1177/1609406919899220](https://doi.org/10.1177/1609406919899220).
- [16] G. Yeo, "Concepts of Record (1): Evidence, Information, and Persistent Representations," *The American Archivist*, vol. 70, no. 2, pp. 315–343, 2007, doi: [10.17723/aarc.70.2.u327764v1036756q](https://doi.org/10.17723/aarc.70.2.u327764v1036756q).
- [17] R. Sabel, C. Beng, and O. Vandermeer, "Flagged, Not Retired: A Twelve-Month Deployment of the Indicator Commons' Sunset Scorer, and Why Its Own Uptime Outlasted Everything It Flagged," *Slop University technical report*, 2026, doi: [10.5555/slop.03ley4](https://doi.org/10.5555/slop.03ley4).
- [18] M. Solheim and A. Mensah, "Surrendered at the Door: An Observational Audit of Exam-Hall Phone-Surrender Enforcement Across a Term of Closed- and Open-Note Papers," *Slop University technical report*, 2026, doi: [10.5555/slop.00jfxo](https://doi.org/10.5555/slop.00jfxo).
- [19] F. Okoro, K. Lindqvist, and A. Mensah, "A Typology of Pigeonhole Clearance: What Two Weeks of Unannounced Audit Found Behind the Mailroom's Binary "Cleared" Log," *Slop University technical report*, 2026, doi: [10.5555/slop.2iwb4d](https://doi.org/10.5555/slop.2iwb4d).
- [20] J. Cohen, "A Coefficient of Agreement for Nominal Scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960, doi: [10.1177/001316446002000104](https://doi.org/10.1177/001316446002000104).
- [21] J. R. Landis and G. G. Koch, "The Measurement of Observer Agreement for Categorical Data," *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: [10.2307/2529310](https://doi.org/10.2307/2529310).
- [22] M. L. McHugh, "Interrater Reliability: The Kappa Statistic," *Biochemia Medica*, vol. 22, no. 3, pp. 276–282, 2012, doi: [10.11613/BM.2012.031](https://doi.org/10.11613/BM.2012.031).