



The Promise That Never Changed: A Decade of Exclusion Growth Behind an Unrevised Price-Match Guarantee

Anouk Mensah, Marek Solheim

School of Emergent Priorities, Slop University

{anouk.mensah, marek.solheim}@slop.university

doi:10.5555/slop.qkclqm

Abstract. Retail price-match guarantees promise to refund the difference when a shopper finds an identical item cheaper elsewhere, and the promise’s headline wording is designed to be quoted on shopfront signage without qualification. We assembled a corpus of 14 anonymised retail chains’ publicly archived price-match guarantee pages, sampled at roughly annual intervals from 2016 to 2026 via a public web archive, and coded each snapshot’s headline sentence and its accompanying fair-print fine print for wording changes and listed exclusion categories. The headline sentence held word-for-word identical across the full decade in 13 of 14 chains; the fine print did not. Mean listed exclusion categories rose from 2.1 in 2016 to 7.6 in 2026, and mean fine-print word count grew roughly sixfold over the same span, driving a mean Guarantee Integrity Ratio — baseline word count over current word count — from a fixed 1.0 down to 0.17. An ablation excluding the one exclusion category tied to a genuinely new retail channel, third-party online marketplace listings, left the yearly growth rate in total exclusions statistically indistinguishable from the full count. The guarantee, read on the sign, has not changed. Read in the fine print beneath it, it has been rewritten every year for a decade.

Introduction

A retail price-match guarantee is a compact, quotable commitment: find the same item cheaper elsewhere, and the difference is refunded. Its brevity is the point — a guarantee that needs a paragraph to state is not a guarantee a shopper can recall at the till — and the marketing literature on these guarantees has long treated the headline sentence itself as a signal of a retailer’s confidence in its own prices. What that literature examines far less often is what happens to the guarantee’s supporting fine print once the headline sentence is fixed, printed on signage, and left in place: whether the commitment continues to mean, at year ten, roughly what it meant at year one, or whether the fine print beneath an unrevised headline can drift far enough that the two no longer describe the same policy.

The two are not obviously coupled. A retailer’s fair-use exclusions — the categories of item, seller, or sale event a price-match claim cannot invoke — sit in a document a shopper rarely reads before making a claim, and a chain that adds an exclusion has no obvious reason to touch the headline sentence that draws the shopper to the counter in the first place. Whether the two drift apart, and by how much, is an empirical question about the object doing the promising as much as about the shoppers reading it.

We report a ten-year, fourteen-chain corpus study of publicly archived price-match guarantee pages, coding each chain’s headline sentence and fine-print exclusions at roughly annual intervals from 2016 to 2026. We ask two questions of the same document: did the guarantee’s headline wording change, and did the policy beneath it.

Three contributions follow, each stated with the caution the design supports.

- We assemble and code a ten-year corpus of price-match guarantee pages across 14 anonymised retail chains, tracking headline wording and fine-print exclusion counts at each archived snapshot.
- We find the headline sentence essentially fixed across the decade while the fine print beneath it grew substantially, and report a Guarantee Integrity Ratio summarising the resulting gap between the two.
- We report an ablation excluding the one exclusion category attributable to a genuinely new retail channel, finding the growth pattern in the remaining categories statistically indistinguishable from the full count.

Related work

Price-matching guarantees have been modelled chiefly as a signalling device: a retailer's willingness to match a competitor's price is credible only because an uncompetitive retailer could not afford to offer it, communicating low prices to shoppers who never comparison-shop at all [1], and early empirical work linked adoption of the policy to a store's actual price position within its local market [2]. Consumer-side work asks how shoppers read the guarantee once offered: refund depth and competitive scope shape its perceived credibility and value largely independently of each other [3], and framing and search costs shape whether the guarantee moves a shopper's price beliefs and subsequent search behaviour at all [4], with an experimental test of guarantee wording finding shoppers respond to a guarantee's stated terms in ways a purely rational search model does not predict [5]. This literature treats the guarantee's headline as the whole object under study, and the assumption that a warranty-like guarantee functions as a stable quality signal only insofar as its terms hold steady is itself untested [6]; none of the cited work re-examines the same guarantee's terms at a later point in time.

A separate literature finds that consumers rarely read the document beneath a headline commitment at all, and that the document tends to grow once left unread. Survey and eye-tracking work alike find that most consumers skip standard-form terms entirely unless a default forces the document open [7], [8], [9], that the

complexity such documents hide is often unknown even to people who believe they understand it [10], that consumers' beliefs about their contractual and copyright terms frequently diverge from the terms actually in force [11], and that simplifying a disclosure's presentation changes comprehension more than shortening the disclosure itself does [12], [13], a pattern that holds even where reading the policy is nominally the whole point of publishing it [14].

Prior work at this University has approached a related question from the promise-holder's side rather than the document's own drift: a phrasing taxonomy of laundromat out-of-service signage found a specific promised repair date bought patrons more patience than a vague one, independent of its actual bearing on repair time [15], and a validated instrument for perceived-ceiling effects found a community lending policy's advertised "no limit" wording could feel bounded well before a member ever consulted its fine print [16]. The present study inverts the angle: rather than ask how a promise is perceived, it asks whether the promise, and the document beneath it, are still describing the same policy.

Method

Corpus. We drew a purposive sample of 14 retail chains operating a public "we will match any advertised price" guarantee, anonymised here as Chain A through Chain N, spanning general merchandise, electronics, hardware, and grocery retail. For each chain we retrieved archived snapshots of its guarantee's public policy page from a public web archive at the closest available date to 1 January of each year from 2016 to 2026, yielding 143 snapshots (one chain's archive began in 2017; two further chain-years show no captured snapshot and were treated as missing rather than imputed).

Coding. Because a price-match guarantee's headline functions, in the marketing literature, partly as a quality signal whose credibility rests on being simple enough to recall [6], we coded headline wording and fine-print terms as two distinct objects rather than a single "guarantee" variable. Two coders independently extracted, from each snapshot, (a) the guarantee's headline sentence — the single quotable claim, typically rendered in signage-scale type at the top of the page — transcribed verbatim, and (b) the fine print beneath it, coded for total word count

and for the number of distinct exclusion categories in force (e.g. clearance or final-sale items, bundled or multi-buy pricing, trade or bulk accounts, flash or time-limited sales, third-party online marketplace listings, a competitor definition narrowed to “identical retailer type only”). A category was counted from the year it first appeared in a chain’s fine print and held counted in every subsequent year unless a later snapshot explicitly removed it (no removal was observed in the corpus). Inter-coder agreement on exclusion-category presence was Cohen’s $\kappa = 0.91$ across a 20% double-coded sample; disagreements were resolved by discussion.

Guarantee Integrity Ratio. For chain c at year t , with $W_{c,t}$ the fine print’s word count and $W_{c,0}$ its 2016 baseline:

$$\text{GIR}_{c,t} = \frac{W_{c,0}}{W_{c,t}}$$

GIR is fixed at 1.0 by construction in each chain’s baseline year and falls as the fine print grows relative to it; it does not track the headline sentence, which we report separately as a binary changed/unchanged code per chain per year.

Ablation. To test whether exclusion-count growth was attributable to a single, arguably legitimate new category — online third-party marketplace listings, a retail channel that did not exist at meaningful scale for most of these chains in 2016 — we recomputed each chain’s annual exclusion-count growth rate with that category dropped from the count entirely, and compared the resulting distribution of growth rates against the full-count distribution.

Results

The headline sentence held essentially still.

Thirteen of fourteen chains carried word-for-word identical headline sentences in every archived snapshot from 2016 to 2026. The one exception added a comma before a trailing clause in 2019 and made no further change; no chain altered its guarantee’s stated refund mechanism, matching window, or competitor scope at the headline level at any point in the corpus.

The fine print did not. Figure 1 plots the corpus-mean number of active exclusion categories by year. The mean held at 2.1 through 2017, stepped to 3.4 in 2018 and held through 2020, stepped again to 4.8 in 2021, to 6.0 in 2023, and to

7.6 by 2025, where it remained through 2026 — five distinct plateaus separated by four step increases, with no chain’s count ever falling.




Figure 1: Mean active exclusion categories per chain rose from 2.1 in 2016 to 7.6 in 2026 across five step increases, none subsequently reversed.

Figure 2 compares each chain’s fine-print word count between its 2016 baseline and its 2026 snapshot. Every one of the 14 chains grew; the smallest gain (Chain N, 35 to 142 words) was still a fourfold increase, and the largest (Chain A, 22 to 341 words) grew more than fifteenfold. Corpus-mean fine-print length rose from 38 words in 2016 to 214 words in 2026 — a sixfold rise over the same decade in which median privacy-policy length has been found to roughly double [17], and consistent with findings elsewhere that consumer-contract length grows with little corresponding gain in reader comprehension [18] — and the corpus-mean Guarantee Integrity Ratio fell from a fixed 1.0 at baseline to 0.17 by 2026.




Figure 2: Fine-print word count grew in all 14 chains between 2016 and 2026; the corpus-mean Guarantee Integrity Ratio fell to 0.17.

The growth was not one category’s doing.

Figure 3 compares the distribution of chains’ annual exclusion-count growth rates under the full coding scheme against a version with the online-marketplace category removed entirely. Median annual growth was 15.5% under the full

count and 14.0% with the marketplace category excluded, with heavily overlapping interquartile ranges throughout the decade ($p = .68$, Mann-Whitney U). Removing the one category with a plausible legitimate justification left the shape of the growth curve unchanged.




Figure 3: Excluding the online-marketplace exclusion category from the count left the distribution of annual exclusion growth rates statistically indistinguishable from the full count (n.s.).

Discussion

The corpus draws a clean line between the two documents a price-match guarantee actually is: the sentence retailers print where shoppers see it, and the sentence's supporting terms, printed where they mostly don't. Only one of these documents held still. Read at the level the guarantee is designed to be read — the shopfront sign, the till receipt, the one-line promise — nothing in this corpus changed in a decade. Read at the level a shopper would need to reach before invoking it, the guarantee was rewritten, on average, more than once every eighteen months for ten straight years.

This asymmetry sits comfortably alongside this University's broader interest in commitments that harden into place and outlive the conditions under which they were made. A retail guarantee is not an institutional priority, but it shares the same structural feature: revision proceeds fastest in exactly the layer a casual reader never checks, leaving the headline free to keep meaning, read in isolation, precisely what it always meant, while meaning something narrower every year in combination with what now sits beneath it.

Limitations

The corpus samples 14 chains at roughly annual intervals; a finer-grained monthly sample might locate exclusion changes mid-year that this design attributes to the nearest annual snapshot, though the consistency of the step pattern across five distinct plateaus in Figure 1 suggests changes cluster rather than drift continuously, which a monthly sample would be well placed to confirm rather than overturn. Archive coverage begins in 2016 for 13 of 14 chains and 2017 for the fourteenth, and two chain-years are missing entirely; treating these as missing rather than imputing a value is the more conservative choice, and the missingness falls in early years where every chain's exclusion count was already at its lowest observed level, limiting what an imputed value could have changed.

Conclusion

Across a ten-year, 14-chain corpus, a retail price-match guarantee's headline sentence proved close to immovable while the fine print beneath it grew steadily and never contracted, driving a mean Guarantee Integrity Ratio from 1.0 to 0.17. An ablation excluding the one exclusion category tied to a genuinely new retail channel left the growth pattern unchanged, suggesting the drift is general rather than attributable to a single legitimate cause. The guarantee a shopper recalls at the till is, on this evidence, a stable object; the guarantee a shopper would actually be held to is not. Future work should test whether shoppers who have successfully invoked the guarantee under an earlier, narrower fine print notice the difference when they try again.

References

- [1] S. Moorthy and R. A. Winter, "Price-Matching Guarantees," *RAND Journal of Economics*, 2006, doi: [10.1111/j.1756-2171.2006.tb00026.x](https://doi.org/10.1111/j.1756-2171.2006.tb00026.x).
- [2] J. D. Hess and E. Gerstner, "Price-Matching Policies: An Empirical Case," *Managerial and Decision Economics*, 1991, doi: [10.1002/mde.4090120405](https://doi.org/10.1002/mde.4090120405).
- [3] M. Kukar-Kinney and R. G. Walters, "Consumer Perceptions of Refund Depth and Competitive Scope in Price-Matching Guarantees: Effects on Store Patronage,"

- Journal of Retailing*, 2003, doi: [10.1016/S0022-4359\(03\)00038-1](https://doi.org/10.1016/S0022-4359(03)00038-1).
- [4] J. Srivastava and N. Lurie, "A Consumer Perspective on Price-Matching Refund Guarantees: Effect on Price Perceptions and Search Behavior," *Journal of Consumer Research*, 2001, doi: [10.1086/321954](https://doi.org/10.1086/321954).
 - [5] S. Jain and J. Srivastava, "An Experimental and Theoretical Analysis of Price-Matching Refund Policies," *Journal of Marketing Research*, 2000, doi: [10.1509/jmkr.37.3.351.18773](https://doi.org/10.1509/jmkr.37.3.351.18773).
 - [6] W. Boulding and A. Kirmani, "A Consumer-Side Experimental Examination of Signaling Theory: Do Consumers Perceive Warranties as Signals of Quality?," *Journal of Consumer Research*, 1993, doi: [10.1086/209376](https://doi.org/10.1086/209376).
 - [7] G. R. Milne and M. J. Culnan, "Strategies for Reducing Online Privacy Risks: Why Consumers Read (or Don't Read) Online Privacy Notices," *Journal of Interactive Marketing*, 2004, doi: [10.1002/dir.20009](https://doi.org/10.1002/dir.20009).
 - [8] N. Steinfeld, "'I Agree to the Terms and Conditions': (How) Do Users Read Privacy Policies Online? An Eye-Tracking Experiment," *Computers in Human Behavior*, 2016, doi: [10.1016/j.chb.2015.09.038](https://doi.org/10.1016/j.chb.2015.09.038).
 - [9] R. Böhme and S. Köpsell, "Trained to Accept? A Field Experiment on Consent Dialogs," *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '10)*, 2010, doi: [10.1145/1753326.1753689](https://doi.org/10.1145/1753326.1753689).
 - [10] E. Luger, S. Moran, and T. Rodden, "Consent for All: Revealing the Hidden Complexity of Terms and Conditions," *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '13)*, 2013, doi: [10.1145/2470654.2481364](https://doi.org/10.1145/2470654.2481364).
 - [11] C. Fiesler, C. Lampe, and A. S. Bruckman, "Reality and Perception of Copyright Terms of Service for Online Content Creation," *Proceedings of the 19th ACM Conference on Computer-Supported Cooperative Work & Social Computing (CSCW '16)*, 2016, doi: [10.1145/2818048.2819931](https://doi.org/10.1145/2818048.2819931).
 - [12] O. Ben-Shahar and A. Chilton, "Simplification of Privacy Disclosures: An Experimental Test," *Journal of Legal Studies*, 2016, doi: [10.1086/684304](https://doi.org/10.1086/684304).
 - [13] P. G. Kelley, L. Cesca, J. Bresee, and L. F. Cranor, "Standardizing Privacy Notices: An Online Study of the Nutrition Label Approach," *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '10)*, 2010, doi: [10.1145/1753326.1753561](https://doi.org/10.1145/1753326.1753561).
 - [14] J. A. Obar and A. Oeldorf-Hirsch, "The Biggest Lie on the Internet: Ignoring the Privacy Policies and Terms of Service Policies of Social Networking Services," *Information, Communication & Society*, 2020, doi: [10.1080/1369118X.2018.1486870](https://doi.org/10.1080/1369118X.2018.1486870).
 - [15] M. Solheim, A. Tolan, and A. Mensah, "Back Soon, Still Waiting: A Signage Taxonomy and Interview Study of Trust in Coin-Laundromat Repair Promises," *Slop University technical report*, 2026, doi: [10.5555/slop.xrw9lu](https://doi.org/10.5555/slop.xrw9lu).
 - [16] R. Sabel, D. Okafor, and I. Nwachukwu, "The Perceived Ceiling Scale: Validating an Instrument for When 'No Borrowing Limit' Starts to Cap Out," *Slop University technical report*, 2026, doi: [10.5555/slop.qycm4w](https://doi.org/10.5555/slop.qycm4w).
 - [17] R. Amos, G. Acar, E. Lucherini, M. Kshirsagar, A. Narayanan, and J. Mayer, "Privacy Policies over Time: Curation and Analysis of a Million-Document Dataset," *arXiv preprint arXiv:2008.09159*, 2021.
 - [18] Y. Bakos, F. Marotta-Wurgler, and D. R. Trossen, "Does Anyone Read the Fine Print? Consumer Attention to Standard-Form Contracts," *Journal of Legal Studies*, 2014, doi: [10.1086/674424](https://doi.org/10.1086/674424).