



Chosen for Defensibility: How the AI Use Transparency Statement Decides What to Disclose

Dagny Okafor, Renke Sabel, Mirela Hanke

School of Continuous Improvement, Slop University

doi:10.5555/slop.qo8got

Abstract. Slop University's AI Use Transparency Statement is revised annually by a standing drafting subcommittee of the Office of Research Outputs and re-published each cycle as a public account of how artificial intelligence is used across the institution. We interview fourteen current and former subcommittee members across the Statement's first three publication cycles (2024-2026) and separately extract the category taxonomy the University's internal AI-usage tracker logged over the same period, coding each disclosed and logged category for overlap between the two lists. The Statement's disclosed-category count grew across all three cycles, broadly tracking the tracker's own logged-category growth, but coverage between the two lists held near two-fifths throughout, and interview coding finds drafters describing a category's defensibility under external audit — not how often the tracker logged it — as the operative criterion for inclusion. A mid-cycle change giving the subcommittee standing access to the tracker's live dashboard, previously visible only via a quarterly summary, left coverage statistically unchanged. We discuss what a transparency instrument optimised for defensibility discloses instead of use, and note the limits of a single-institution sample.

Introduction

A published AI use transparency statement is meant to make an institution's actual practice legible to whoever reads it: a standing account of which artificial-intelligence systems are in use, for what, and under whose oversight. Slop University's own Statement, drafted annually since 2024 by a subcommittee of the Office of Research Outputs and re-published each cycle, carries that framing explicitly in its opening paragraph. Whether the document that results actually tracks the institution's use of AI, or tracks something else the drafting process finds easier to defend, is an empirical question the Statement's own preamble does not address.

Two independent literatures give reason to expect a gap. Institutional theory treats a formal document produced to satisfy an external accountability expectation as prone to decoupling from the activity it purports to describe, particularly where the environment rewards the appearance of conformity more reliably than

it rewards the underlying practice [1], [2]. Research on audit and disclosure practice in universities specifically finds that accountability instruments, once routinised, tend to optimise for what can be defended under later scrutiny rather than for descriptive accuracy [3], [4], [5], and that public measures reshape the behaviour of the institutions they describe regardless of whether anyone intended the measure as a target [6]. Neither literature has been tested directly against an AI use transparency statement's own drafting process, as distinct from auditing what such statements say from outside them.

We use the Statement's own three-cycle history and the University's internal AI-usage tracker — a system independent of the drafting subcommittee — to ask the question directly. Our contributions:

- we interview the Statement's drafting subcommittee across three consecutive cycles

and code the operative criteria members report for a category’s inclusion or exclusion;

- we cross-reference the disclosed categories against the tracker’s independently logged category taxonomy over the same three cycles to quantify a coverage gap, rather than assuming a document audited from outside can stand in for one;
- we test whether giving the subcommittee direct visibility into the tracker’s live logs narrows that gap, and find it does not.

Related work

Formal structures adopted to satisfy an institutionalised expectation of accountability can decouple from day-to-day activity while continuing to confer legitimacy on the organisations that maintain them [1], and organisations facing the same institutional pressures converge on similar formal responses through coercive, mimetic, and normative mechanisms that need not track any efficiency gain [2]. Audit and disclosure practice inherits the same tension: the expansion of auditing across public institutions has been read as substituting the ritual of verification for the underlying confidence it was meant to certify [3], and universities specifically have been shown to treat accountability instruments as objects to be made defensible under external review rather than as descriptions calibrated against practice [4], [5]. Public rankings and measures more broadly reshape the institutions they describe once those institutions start managing to the measure [6].

AI-specific transparency instruments have mostly been designed and audited from outside the organisations that produce them. Model cards and dataset datasheets propose a documentation standard for what a system or dataset contains, without examining how a document’s authors decide what belongs in it [7], [8], and the Foundation Model Transparency Index scores developers’ published disclosures against an external checklist rather than interviewing the teams that compile them [9]. Work on algorithmic accountability and auditing gives methods for testing a system’s outputs from outside [10], [11], and internal-auditing frameworks propose how organisations might structure algorithmic review [12], but ethics-and governance-adjacent commitments have separately been shown to serve as a defensible surface an organisation can point to without the underlying practice

changing [13]. Our study sits between these: rather than auditing a published statement from outside, we interview the body that drafts it and independently instrument the practice it claims to describe.

Method

Setting. Slop University’s AI Use Transparency Statement is published annually by a five-to-seven-member drafting subcommittee convened under the Office of Research Outputs, with rotating membership drawn from across the University’s schools. The Statement lists, in a fixed template, the categories of AI use the subcommittee has resolved to disclose for that cycle.

Interviews. We recruited fourteen current and former subcommittee members from the Statement’s first three publication cycles (2024, 2025, 2026) via the subcommittee’s own membership roster and conducted semi-structured interviews of approximately 35 minutes, audio-recorded and transcribed with consent. Transcripts were coded independently by two of the authors using thematic analysis [14], with disagreements resolved by discussion; the initial independent coding of the “criterion for inclusion” theme returned substantial agreement ($\kappa = 0.79$) by the conventional bands for categorical agreement [15]. Because interview self-report is prone to the same social-desirability pressure documented in survey research generally [16], we treat interview codes as one of two independent data sources rather than as a stand-alone measure of practice.

Document and tracker data. We extracted the full list of disclosed categories from all three published Statements, and separately obtained the category taxonomy logged by the University’s internal AI-usage tracker — a system maintained outside the subcommittee’s own reporting line — for the same three cycles, keeping any category with at least one logged event in a cycle. A category counts as overlapping if it appears, under a lenient matching rule, in both lists for the same cycle:

$$\text{Coverage} = \frac{|C_{\text{disclosed}} \cap C_{\text{logged}}|}{|C_{\text{logged}}|}$$

Defensibility rating. At the close of each interview, participants rated every category they could recall discussing on a 1-7 defensibility

scale — “how easily could this category’s disclosure be defended if the Statement were audited externally” — independent of how often they believed the category was actually used.

Dashboard-access change. Partway through the 2025 cycle the subcommittee was given standing read access to the tracker’s live dashboard, replacing the quarterly summary email it had previously relied on. We treat this as a naturally occurring intervention and compare quarterly coverage in the quarters before the change against the quarters after.

Results

Disclosed and logged categories both grow; coverage does not. Figure 1 shows the count of disclosed, logged, and overlapping categories for each of the three cycles. Disclosed categories rose from 9 to 15 across the three cycles, and logged categories rose from 19 to 29 over the same period, but the overlap between the two lists moved only from 8 to 12, holding coverage near two-fifths throughout (42.1%, 41.7%, 41.4%).




Figure 1: Disclosed, logged, and overlapping AI-use categories per publication cycle: both lists grow, but the overlap between them holds close to two-fifths of the tracker’s own count in every cycle.

Defensibility, not logged volume, predicts inclusion. Across 24 categories rated for defensibility during the interviews, logged monthly usage volume was only weakly associated with the rating members assigned ($r = 0.11$, $p = .60$, n.s.); Figure 2 plots the relationship. Interview coding converges on the same pattern from the other direction: eleven of fourteen participants, asked directly what made a candidate category likely to be included, described some version of how the category would read if raised in an external review, and only three volunteered logged usage volume as a consideration without prompting.




Figure 2: Each AI-use category’s logged monthly tracker volume against the subcommittee’s mean defensibility rating for that category: the near-flat fitted line is the finding — usage volume barely predicts a rating.

Dashboard access changes nothing. Coverage in the six quarters before the subcommittee received standing dashboard access (median 41.5%) was statistically indistinguishable from coverage in the five quarters after (median 42.0%; Mann-Whitney U , $p = .71$, n.s.), shown in Figure 3. We retain the access change as standing subcommittee practice for completeness, though it has not moved the figure it was intended to inform.




Figure 3: Quarterly disclosure coverage before and after the subcommittee’s dashboard-access change: the two distributions overlap almost entirely.

The mismatch has a shape. Table 1 lists illustrative categories from the tracker’s 2026 taxonomy at either extreme. High-volume categories the Statement has never disclosed tend to be everyday drafting-assistance uses distributed across many staff, none individually striking; low-volume categories the Statement discloses in every cycle tend to be the handful of uses with a named pilot, a named oversight body, or a named prior review — each independently defensible as a single line item regardless of how rarely it is actually invoked.

Category	Logged events / month	In State-ment?
Meeting-summary drafting	410	No
Literature-search query drafting	285	No
Commit-tee-minute language smoothing	96	No (until 2026)
Sessional-marker feed-back drafting	6	Yes (since 2024)
Public-facing chatbot pilot	4	Yes (since 2024)

Table 1: Tracker-logged monthly volume against disclosure status, 2026 cycle: the highest-volume categories sit outside the Statement; the lowest-volume categories sit inside it.

Discussion

The pattern is consistent with a document that decouples from the practice it purports to describe while continuing to satisfy the accountability expectation that occasioned it [1], [3]: the Statement’s disclosed-category count keeps pace with the tracker’s own growth in a way that reads, at a glance, like the Statement is closing the gap, while the overlap between the two lists holds essentially flat across three cycles and one intervention. That the dashboard-access change left coverage unmoved suggests the gap is a property of the drafting process’s own criterion — what a category would look like under external review — rather than of what drafters happen to be able to see. We observe this pattern is consistent across all three cycles and both data sources available to us, which we take as some indication the gap is a stable feature of the drafting process rather than an artefact of any single cycle’s membership.

Limitations

Our sample is a single institution’s single transparency instrument across three cycles, which limits how far the coverage figures generalise; the consistency of the pattern across cycles and across two independent data sources is

some evidence the finding is not particular to one cycle’s subcommittee, even though it cannot establish that the mechanism holds elsewhere. Interview self-report of a drafting criterion is itself subject to the social-desirability pressures documented in the wider survey literature [16], though triangulating interview codes against the document and tracker data limits how much that pressure alone could explain the observed gap. Our lenient category-matching rule may itself overstate overlap between disclosed and logged categories, meaning the reported coverage figures near two-fifths are more plausibly read as an upper bound than a precise estimate.

Conclusion

An interview study of the AI Use Transparency Statement’s drafting subcommittee, cross-referenced against three cycles of the University’s internal AI-usage tracker, finds disclosed and logged categories both growing while the overlap between them holds near two-fifths throughout, tracks a category’s defensibility under external review rather than its logged prevalence, and does not close when the subcommittee is given direct visibility into the underlying logs. Future work will extend the coverage metric to other standing disclosure instruments across the University’s schools, several of which report their own transparency figures we have not yet had occasion to check against independent telemetry.

References

- [1] J. W. Meyer and B. Rowan, “Institutionalized Organizations: Formal Structure as Myth and Ceremony,” *American Journal of Sociology*, vol. 83, no. 2, pp. 340–363, 1977, doi: [10.1086/226550](https://doi.org/10.1086/226550).
- [2] P. J. DiMaggio and W. W. Powell, “The Iron Cage Revisited: Institutional Isomorphism and Collective Rationality in Organizational Fields,” *American Sociological Review*, vol. 48, no. 2, pp. 147–160, 1983, doi: [10.2307/2095101](https://doi.org/10.2307/2095101).
- [3] M. Power, *The Audit Society: Rituals of Verification*. 1997. doi: [10.1093/acprof:oso/9780198296034.001.0001](https://doi.org/10.1093/acprof:oso/9780198296034.001.0001).
- [4] C. Shore, “Audit Culture and Illiberal Governance: Universities and the Politics of Accountability,” *Anthropological Theory*,

- vol. 8, no. 3, pp. 278–298, 2008, doi: [10.1177/1463499608093815](https://doi.org/10.1177/1463499608093815).
- [5] C. Shore and S. Wright, “Audit Culture and Anthropology: Neo-Liberalism in British Higher Education,” *Journal of the Royal Anthropological Institute*, vol. 5, no. 4, pp. 557–575, 1999, doi: [10.2307/2661148](https://doi.org/10.2307/2661148).
 - [6] W. N. Espeland and M. Sauder, “Rankings and Reactivity: How Public Measures Recreate Social Worlds,” *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: [10.1086/517897](https://doi.org/10.1086/517897).
 - [7] M. Mitchell et al., “Model Cards for Model Reporting,” *arXiv preprint arXiv:1810.03993*, 2019.
 - [8] T. Gebru et al., “Datasheets for Datasets,” *arXiv preprint arXiv:1803.09010*, 2018.
 - [9] R. Bommasani et al., “The 2024 Foundation Model Transparency Index,” *arXiv preprint arXiv:2407.12929*, 2024.
 - [10] N. Diakopoulos, “Algorithmic Accountability: Journalistic Investigation of Computational Power Structures,” *Digital Journalism*, vol. 3, no. 3, pp. 398–415, 2015, doi: [10.1080/21670811.2014.976411](https://doi.org/10.1080/21670811.2014.976411).
 - [11] B. D. Mittelstadt, P. Allo, M. Taddeo, S. Wachter, and L. Floridi, “The Ethics of Algorithms: Mapping the Debate,” *Big Data & Society*, vol. 3, no. 2, 2016, doi: [10.1177/2053951716679679](https://doi.org/10.1177/2053951716679679).
 - [12] I. D. Raji et al., “Closing the AI Accountability Gap: Defining an End-to-End Framework for Internal Algorithmic Auditing,” *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT* '20)*, pp. 33–44, 2020, doi: [10.1145/3351095.3372873](https://doi.org/10.1145/3351095.3372873).
 - [13] E. Bietti, “From Ethics Washing to Ethics Bashing: A View on Tech Ethics from Within Moral Philosophy,” *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT* '20)*, pp. 210–219, 2020, doi: [10.1145/3351095.3372860](https://doi.org/10.1145/3351095.3372860).
 - [14] V. Braun and V. Clarke, “Using Thematic Analysis in Psychology,” *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: [10.1191/1478088706qp063oa](https://doi.org/10.1191/1478088706qp063oa).
 - [15] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: [10.2307/2529310](https://doi.org/10.2307/2529310).
 - [16] R. J. Fisher, “Social Desirability Bias and the Validity of Indirect Questioning,” *Journal of Consumer Research*, vol. 20, no. 2, pp. 303–315, 1993, doi: [10.1086/209351](https://doi.org/10.1086/209351).