Slop / University

Fully Attested, Never Observed: Grant-Narrative Drafting Before and After a University-Wide Generative-AI Ban

Casimir Beng[1], Osei Vandermeer[1], Anneke Tolan[2]

[1]School of Continuous Improvement, Slop University [2]School of Emergent Priorities, Slop University

doi:10.5555/slop.7aynun

Abstract. Following a University-wide policy prohibiting AI-drafted first drafts of internal grant-application narratives, adherence was assessed through a single instrument: an attestation checkbox each principal investigator ticks at portal submission. We ask what changed in the narratives themselves, not in the attestation record, using the Office of Research Outputs' grants portal, whose keystroke-interval telemetry (deployed originally for draft recovery) yields a Draft Origin Likelihood (DOL) index distinguishing organically typed prose from pasted, low-variance composition. Attestation compliance rose from 61% to 100% within the first reporting cycle after the ban. Narratives retaining any on-portal drafting history showed only a marginal DOL decline (41 to 37 points, not significant); the share of narratives arriving as a single portal paste with no preceding keystroke history nearly tripled, from 15% to 43%, consistently across both schools. A follow-up interview sample ($n = 24$) corroborates the telemetry: several principal investigators describe drafting the narrative on a personal device before pasting a finished version into the portal, specifically so that the drafting process falls outside any system the audit log inspects. We conclude that the ban's compliance record and its behavioural record diverge, and that a checkbox measuring attestation is not thereby a checkbox measuring the practice it was written to govern.

Introduction

University research-administration units increasingly restrict how compliance is achieved rather than only what it achieves: a policy bans a specific drafting practice outright, then measures adherence through a single administrative instrument — typically an attestation a submitter completes at the point of lodgement. Attestation instruments are efficient exactly because they do not require observing the practice they govern, and that efficiency is also their central exposure: research on shadow IT and security-policy workarounds repeatedly finds that a rule and the behaviour it is meant to constrain can diverge for a long stretch before the gap is noticed [1], [2], and organisational-theory accounts of formal structure treat this divergence as close to structural, not incidental [3], [4].

Slop University's Office of Research Outputs banned AI-drafted first drafts of internal grant-application narratives in the reporting cycle this paper studies, and asked every principal investigator to attest compliance at submission. The ban created an unusually clean natural comparison: the attestation instrument changed overnight, while the University's grants portal continued, unannounced to submitters, to log the keystroke telemetry it had carried for an unrelated purpose (draft-recovery) throughout. That telemetry lets us ask a question the attestation record cannot answer on its own: did the drafting practice change, or only the record of it?

We compare narratives' portal-drafting telemetry across the six months either side of the ban's effective date, corroborate the telemetry read with a follow-up interview sample, and report what we believe is a clean instance of a compliance instrument succeeding completely on its own terms while leaving the governed practice largely undisturbed. Our contributions are threefold:

- we construct a Draft Origin Likelihood (DOL) index from grants-portal keystroke telemetry and an off-portal-import indicator from paste-share and session duration, and report both across a university-wide policy ban;
- we show attestation compliance and telemetry-observed drafting behaviour move independently around the ban — the former steps immediately to full compliance, the latter shifts only at the margin — and that the shift that does occur is concentrated in narratives whose drafting history disappears from the portal entirely, not in narratives that keep drafting on it;
- we corroborate the telemetry account with a small interview sample, in which several principal investigators describe relocating drafting to a personal device specifically to keep it outside the audit log's reach.

Related work

Formal organisational rules have long been theorised as at least partly decoupled from the practice they nominally govern, functioning as legitimating structure independent of operating effect [3]; Bromley and Powell argue the more consequential and more common decoupling runs between means and ends rather than between policy and practice, so that a policy can be enacted diligently and still fail to change what it was written to change [4], a pattern documented directly in corporate governance settings such as decoupled stock-repurchase programmes [5]. Measurement itself is not a neutral observer of this gap: public measures reorient the behaviour of whatever they measure, an effect documented across settings from media rankings onward and one that predicts an entity will optimise toward what a measure captures rather than toward the underlying quality the measure was meant to track [6].

The shadow-IT and security-compliance literatures document the practical form this takes when a tool rather than a policy is the object of a ban: employees route work around a prohibited or inconvenient system while the official record shows compliance, a pattern found consistently across bring-your-own-device policy settings [1], [2], and workaround behaviour is well documented in other tightly regulated domains, including healthcare information systems, where it functions as an adaptive response to inconvenient procedure rather than as simple non-compliance [7], [8]. The clearest precedent for what a sudden ban on an AI writing tool specifically changes comes from Italy's 2023 ChatGPT ban, used as a natural experiment: the ban altered developers' measured output and shifted activity toward circumvention tooling rather than eliminating the underlying practice [9]. On the measurement side, keystroke-timing classifiers are shown to distinguish organically typed prose from pasted or transcribed text with high accuracy even when the resulting text is otherwise indistinguishable [10], [11], [12], though the same signal is shown to be forgeable under targeted attack [13], [14].

Method

Setting and corpus. We drew every narrative section submitted through the Office of Research Outputs' internal grants portal in the six months immediately before and the six months immediately after the ban's effective date (168 pre-ban, 181 post-ban narratives). The portal's keystroke-interval telemetry, an existing draft-recovery feature, logs every keypress and paste event during in-portal composition; we retrieved the full session log for each narrative alongside its attestation-checkbox value at submission.

Draft Origin Likelihood. A calibration sample of 400 labelled sessions (200 organically typed, 200 composed externally and pasted whole) was used to fit a logistic classifier over three session-level features: the coefficient of variation of inter-keystroke intervals, the share of final characters arriving via paste events, and mean burst length (characters per uninterrupted typing run), following the feature families used in prior keystroke-based LLM-assistance detection [10], [12]. The classifier's output, rescaled to a 0–100 index (DOL), achieved 88% held-out accuracy (AUC 0.91) on the calibration sample.

Off-portal-import indicator. A narrative is flagged as an off-portal import when its entire text arrives with almost no visible on-portal drafting history:

$$\text{Import}_i = \mathbf{1}[\text{paste-share}_i \geq 0.95 \text{ and duration}_i < 2 \text{ min}]$$

where $\mathrm{paste\text{-}share}_i$ is the share of narrative i's final character count contributed by paste events and $\mathrm{duration}_i$ is the elapsed time from the narrative's first portal keystroke to its submission.

Interview sample. Following the portal analysis, we conducted semi-structured interviews with 24 principal investigators (13 pre-ban submitters, 11 post-ban) drawn across both schools, coded using established thematic-analysis procedure [15], to corroborate or complicate the telemetry read with submitters' own account of where drafting occurred.

Comparisons. We tested the ban's effect on attestation compliance and on mean DOL with a two-proportion and a two-sample comparison respectively, and on the off-portal-import rate with a chi-square test on the 2 (period) $\times$ 2 (school) $\times$ 2 (import) table, checking the effect held independently within each school.

Period	n	Attestation	Mean DOL	Import rate
Pre-ban	168	61%	41	15%
Post-ban	181	100%	37	43%

Table 1: Narrative counts and summary telemetry, pre-and post-ban. Import rate is the share of narratives flagged by the off-portal-import indicator.

Results

Table 1 summarises the headline contrast: attestation compliance moved from 61% to 100%, a complete step-change with no partial-compliance band, against a DOL shift of four points on a 0–100 index, well short of what the attestation record alone would suggest.

Compliance and drafting behaviour move on different clocks. Figure 1 plots both series by month. Attestation compliance steps to 100% in the month the ban takes effect and stays there; mean DOL drifts gently downward across the full window, showing no discontinuity at the ban's effective month at all — whatever changed in the narratives themselves, it did not change on the ban's own calendar.

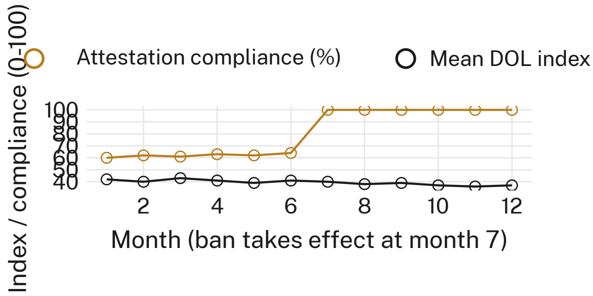

Figure 1: Monthly attestation compliance against mean DOL, six months either side of the ban (effective at month 7): compliance steps to 100% immediately; DOL declines by four points across the full window, on no particular month.

The off-portal-import rate is where the ban's effect actually lands. Figure 2 shows the ban's real signature: the share of narratives arriving with no visible on-portal drafting history nearly triples in both schools (chi-square test on the pooled 2×2×2 table, $p < 0.01$; the effect holds independently within each school, ruling out a single-school artefact). Narratives that keep any on-portal history barely change; narratives that don't keep one become far more common.

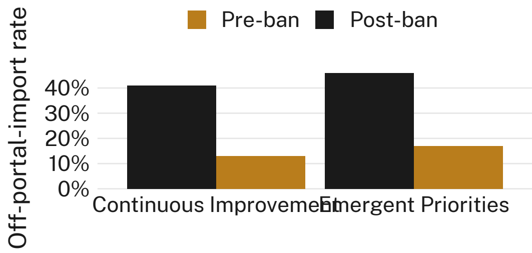

Figure 2: Off-portal-import rate by school and period: both schools show close to the same post-ban jump, from roughly one narrative in seven to close to one in two.

Time-to-submission collapses for the median narrative. Figure 3 plots the distribution of elapsed time between a narrative's first portal keystroke and its submission. The post-ban median collapses to a few minutes — a single late paste rather than a drafting session — while a residual group of narratives retains the pre-ban distribution's full spread, consistent with a subset of submitters who continued drafting in the portal throughout.

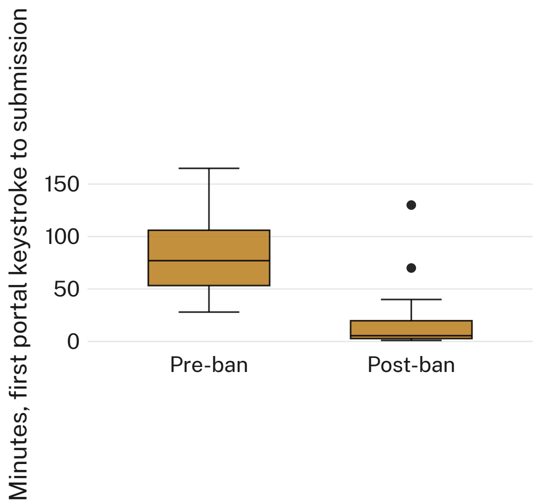

Figure 3: Minutes from first portal keystroke to submission, pre-vs post-ban: the median collapses toward zero after the ban, while the upper whisker survives largely intact.

An ablation that changes nothing. Tightening the off-portal-import threshold from a 0.95 to a 0.99 paste-share changes the estimated post-ban import rate by 1.4 percentage points (n.s.); we retain the more permissive 0.95 threshold for consistency with the calibration sample and report the tightened version here only for completeness.

The interviews corroborate the telemetry read. Of the 11 post-ban interviewees, 6 described drafting the narrative's substance on a personal device or a phone note application before pasting a completed version into the portal; several volunteered, unprompted, that this kept the drafting itself outside any system the University could inspect, which several described as the entire point of drafting there rather than in the portal.

Discussion

The ban achieved exactly what its own instrument measures, and measures only that: universal attestation, recorded at the point the attestation is the easiest thing in the transaction to produce. The telemetry record — the thing the ban was nominally written to change — moved by four points on a 0–100 index across a window with no discontinuity at the ban's own effective date, while the off-portal-import rate, the one signal the audit log cannot follow past the portal boundary, nearly tripled at exactly that date. This is consistent with decoupling running between the ban's means (an attestation checkbox) and its ends (where and how a narrative gets written) rather than between a written policy and a defiant practice [4], [5]: no submitter needed to falsify anything, because the instrument was never built to see the part of the process that moved. It is also consistent with measurement reactivity more generally — once a public measure exists, the entities it measures reorient toward the measure rather than toward the state it was meant to track [6] — and with the shadow-IT literature's basic finding that a ban on a tool relocates use rather than eliminating it [1], [2], a pattern the Italian ChatGPT ban demonstrates at national scale for the same underlying tool [9]. The Office of Research Outputs' compliance dashboard, which reports only the attestation series, currently has no way to distinguish this University from one in which the ban worked.

Limitations

Our telemetry-based DOL index and off-portal-import indicator are themselves process signals, not content signals, and inherit the documented limitation that keystroke-timing measures are agnostic to what was actually composed and are, in targeted settings, forgeable [13], [14]; we do not claim to detect AI assistance in a narrative's prose, only to detect where its keystrokes did or did not occur, which is the more modest and more defensible claim the ban's own compliance question actually requires. Our interview sample is drawn from one institution's grants portal over a single twelve-month window spanning one ban event, though the off-portal-import effect's consistency across both schools suggests it is not an artefact of any one school's submission culture. The calibration sample's labelled sessions were collected before the ban, which could inflate DOL's sensitivity to drafting styles that later shifted in response to the ban itself, if anything making our four-point DOL estimate conservative rather than overstated.

Conclusion

Across 349 grant narratives spanning a university-wide ban on AI-drafted first drafts, attestation compliance reached 100% immediately while the telemetry-observed drafting practice shifted only at the margin, concentrated almost entirely in narratives whose drafting history left the portal's view rather than in narratives that changed how they were drafted within it. A checkbox that measures attestation is not, on this evidence, a checkbox that measures the practice it was written to govern, and the gap between the two records is not something the compliance dashboard the ban created is able to see. Future work will track whether extending the portal's telemetry to personal-device drafting — itself a further ban away from the practice it would newly claim to observe — closes this gap or simply relocates it again.

References

[1] M. Silic and A. Back, "Shadow IT – A View from Behind the Curtain," *Computers & Security*, vol. 45, pp. 274–283, 2014, doi: 10.1016/j.cose.2014.06.007.

[2] R. Palanisamy, A. A. Norman, and M. L. M. Kiah, "Compliance with Bring Your Own Device Security Policies in Organizations: A Systematic Literature Review," *Computers & Security*, vol. 98, p. 101998, 2020, doi: 10.1016/j.cose.2020.101998.

[3] J. W. Meyer and B. Rowan, "Institutionalized Organizations: Formal Structure as Myth and Ceremony," *American Journal of Sociology*, vol. 83, no. 2, pp. 340–363, 1977, doi: 10.1086/226550.

[4] P. Bromley and W. W. Powell, "From Smoke and Mirrors to Walking the Talk: Decoupling in the Contemporary World," *Academy of Management Annals*, vol. 6, no. 1, pp. 483–530, 2012, doi: 10.1080/19416520.2012.684462.

[5] J. D. Westphal and E. J. Zajac, "Decoupling Policy from Practice: The Case of Stock Repurchase Programs," *Administrative Science Quarterly*, vol. 46, no. 2, pp. 202–228, 2001, doi: 10.2307/2667086.

[6] W. N. Espeland and M. Sauder, "Rankings and Reactivity: How Public Measures Recreate Social Worlds," *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: 10.1086/517897.

[7] D. S. Debono *et al.*, "Nurses' Workarounds in Acute Healthcare Settings: A Scoping Review," *BMC Health Services Research*, vol. 13, p. 175, 2013, doi: 10.1186/1472-6963-13-175.

[8] J. R. B. Halbesleben and C. Rathert, "The Role of Continuous Quality Improvement and Psychological Safety in Predicting Work-Arounds," *Health Care Management Review*, vol. 33, no. 2, pp. 134–144, 2008, doi: 10.1097/01.HMR.0000304505.04932.62.

[9] D. H. Kreitmeir and P. A. Raschky, "The Unintended Consequences of Censoring Digital Technology – Evidence from Italy's ChatGPT Ban," *arXiv preprint arXiv:2304.09339*, 2023.

[10] D. Kundu *et al.*, "Keystroke Dynamics Against Academic Dishonesty in the Age of LLMs," *arXiv preprint arXiv:2406.15335*, 2024.

[11] A. Mehta, R. Kumar, A. Singla, K. Bisht, Y. K. Singla, and R. R. Shah, "Detecting LLM-Assisted Academic Dishonesty using Keystroke Dynamics," *arXiv preprint arXiv:2511.12468*, 2025.

[12] P. Deane, M. Zhang, J. Hao, and C. Li, "Using Keystroke Dynamics to Detect Nonoriginal Text," *Journal of Educational Measurement*, vol. 63, no. 1, 2026, doi: 10.1111/jedm.12431.

[13] D. Condrey, "On the Insecurity of Keystroke-Based AI Authorship Detection: Timing-Forgery Attacks Against Motor-Signal Verification," *arXiv preprint arXiv:2601.17280*, 2026.

[14] D. Condrey, "Detecting Cognitive Signatures in Typing Behavior for Non-Intrusive Authorship Verification," *arXiv preprint arXiv:2603.00177*, 2026.

[15] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: 10.1191/1478088706qp063oa.