Slop / University

Not Everyone Signs Off

Marek Solheim, Anneke Tolan, Anouk Mensah

School of Emergent Priorities, Slop University

doi:10.5555/slop.50tmyv

Abstract. Neighbourhood street-library boxes host more than book exchange: visitors leave handwritten notes inside covers or tucked against the shelf, and some of these notes accumulate into short correspondence threads between strangers who never meet. We report a nine-month observational study of 341 such threads coded from 74 street-library boxes across a metropolitan area, developing a four-category taxonomy of how a thread concludes: a signed-off close, a fizzled non-reply, an orphaned exchange still addressed to a visitor who has stopped answering, and a note physically removed from the box before any ending is legible in place. Fizzled endings dominate the corpus overall; signed-off closes are markedly less common in high-traffic boxes than in low-traffic ones, where they rival fizzled endings for the most frequent outcome. We then train a closure-prediction model on linguistic features of a thread's opening note alone — politeness markers, direct address, an explicit request for reply — and find it modestly outperforms a majority-class baseline, with its advantage concentrated almost entirely in anticipating orphaned threads. We discuss what this concentration suggests about which endings are legible from an opening line and which are not, and offer the taxonomy as a transferable coding instrument for correspondence exchanged in other unattended public spaces.

Introduction

Street libraries — small unattended cabinets where passers-by take a book and, nominally, leave one in exchange — have proliferated across Australian suburbs over the past decade, and their contents have drawn growing scholarly interest as a minor but genuine instance of a gift economy operating with almost no formal governance [1]. What has drawn less attention is a secondary practice that develops alongside the books: visitors leave notes — tucked inside a cover, taped to the shelf, folded around a bookmark — addressed not to the box's steward but to whoever reads next. Some of these notes prompt a reply, and some replies prompt another, producing short correspondence threads conducted entirely between strangers who are never at the box at the same time.

Every correspondence thread ends somehow, but not every thread ends legibly. Conversation-analytic work on spoken closings shows that ending a conversation is itself an accomplishment, requiring a recognisable sequence rather than simply falling silent [2]; letter-writing scholarship extends this claim to written correspondence, arguing that a letter's ending is rarely a single deliberate act and is better read as a matter of degree [3]. Street-library correspondence sits at an unusual point on this spectrum: unlike a spoken conversation, it has no shared moment of closing; unlike a private letter, any subsequent visitor can read — and potentially close — a thread they did not start.

This paper makes three contributions, each hedged appropriately for an object this informal.

- We introduce a four-category taxonomy of how a street-library correspondence thread concludes, hand-coded from nine months of photographic archiving across 74 boxes.
- We show that a closure-prediction model trained only on a thread's opening note performs modestly better than a majority-class baseline at anticipating which ending it will reach, though the improvement is not evenly distributed across categories.

- We suggest, tentatively, that the taxonomy generalises to other unattended correspondence settings, and note that the University's own coding taxonomies for internal registers keep landing on a similar asymmetry — a record of an ending predicts its own upkeep more reliably than it predicts the ending itself [4], [5].

Related work

Little Free Libraries and their Australian counterpart, street libraries, have been studied mainly as sites of informal literacy provision and neighbourhood placemaking, with a substantial share of that literature auditing what a box actually contains against what its steward intends it to hold [1], [6]. This work treats the box itself, not what passes between its visitors, as the unit of analysis; the correspondence a box hosts has gone largely unremarked.

Two adjacent traditions inform our taxonomy. Conversation analysis has long treated the close of an interaction as a structured accomplishment rather than an absence of further talk [2], and correspondence scholarship extends the same claim to letters, arguing an ending is a matter of degree rather than a single legible act [3]. Separately, a computational literature on online discussion forecasts whether a thread will continue from features of its earliest posts, whether framed as detecting early signs of derailment [7], identifying posts unlikely to be answered [8], or modelling a comment thread's eventual shape from its opening text alone [9]. That literature's linguistic features — markers of politeness, deference, and directness — draw on a broader computational account of politeness as a codeable property of a request [10], which we adapt below.

Gift-economy research on unilateral exchange, from community gifting groups [11] to the reciprocity that follows a received gift [12], motivates our interest in whether a thread's opening act of goodwill predicts how, or whether, it is later reciprocated. We take our coding procedure from qualitative psychology [13], and our sense that an informal marking left in a shared, unattended space rewards close reading from work on urban graffiti [14].

Method

Data collection ran over nine months. We identified 74 street-library boxes across a single metropolitan area from publicly listed steward registries and informal neighbourhood mapping, stratified into a high-traffic tier (a main road, a school walking route, or a shopping-strip frontage; 31 boxes, 158 threads) and a low-traffic tier (a quiet residential street with no obvious through-traffic; 43 boxes, 183 threads). Each box was photographed at a roughly fortnightly interval; a note was logged into the corpus the first time it appeared and re-photographed at every subsequent visit until it was removed, replaced, or the box's stock turned over enough that it could no longer be reliably tracked. A correspondence thread was defined as two or more notes in evident dialogue — addressed to a previous note's author, replying to a stated question, or physically stacked with an earlier note — yielding 341 threads for analysis after 28 single, unanswered notes were excluded as non-threads.

Two coders, blind to each box's traffic tier, independently applied a four-category closure taxonomy developed inductively over an initial pilot month following standard thematic-coding practice [13], then reconciled disagreements by discussion (Cohen's $\kappa = 0.81$). A thread was coded **Signed-off** when its final note carried an explicit closing marker; **Fizzled** once a thread went quiet for at least six weeks with no closing marker and no removal; **Orphaned** when a note addressed a specific prior interlocutor across two or more subsequent visits without reply; and **Absorbed** when a note was physically removed from the box before either party supplied a legible ending. Formally, writing $R(t)$ for the number of consecutive unanswered notes addressed to the same interlocutor in thread t:

$$R(t) \geq 2 \Rightarrow t \text{ coded Orphaned}$$

For the closure-prediction model, we extracted five features from each thread's opening note alone: word count, presence of a direct question, an explicit request for reply, a named addressee, and a politeness score computed with the lexical and syntactic markers of [10]. We trained a linear bag-of-tricks classifier [15] on these features plus unigram counts of the opening note's text, evaluated with five-fold cross-validation stratified by box, and compared it against a majority-class baseline that always predicts Fizzled. An ablation removed the explicit-request-for-reply feature, isolating a claim about which single feature the model actually needs.

Results

Across the full corpus, Fizzled endings account for 157 threads (46.0%), Signed-off for 92 (27.0%), Orphaned for 65 (19.1%), and Absorbed for 27 (7.9%). This aggregate obscures a clear split by traffic tier:

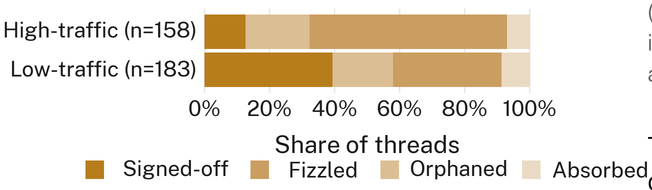

Figure 1: Signed-off closes are the least common legible ending in high-traffic boxes (12.7%) but overtake Fizzled outright in low-traffic ones (39.3% versus 33.3%).

In high-traffic boxes Fizzled dominates (60.8%), with Signed-off the least common of the three legible-ending categories (12.7%, against Orphaned's 19.6%); low-traffic boxes are the only tier in which any category exceeds Fizzled at all.

The closure-prediction model reaches 51.9% overall accuracy against a 46.0% majority-class baseline that always predicts Fizzled — a modest margin, and one that undersells where the model actually helps.

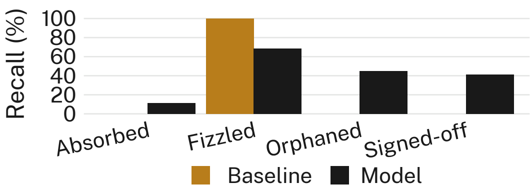

Figure 2: The model trades Fizzled recall for gains everywhere else, reaching 44.6% recall on Orphaned threads against the baseline's 0%.

Per-category recall shows the baseline is, definitionally, perfect on Fizzled (100%) and useless everywhere else (0%); the model gives up recall on Fizzled (68.2%) to gain meaningfully on every other category, most of all Orphaned (44.6%). Signed-off recall (41.3%) and Absorbed recall (11.1%) trail Orphaned by a comparable or wider margin.

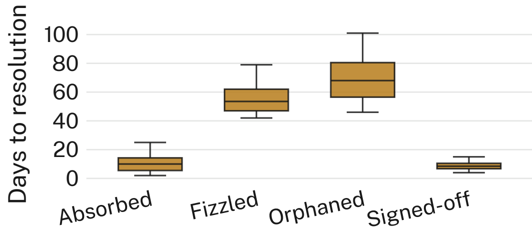

Figure 3: Orphaned threads take longest to resolve (median 68 days) by construction of the taxonomy itself; Signed-off and Absorbed threads can close at any visit and typically do so within a fortnight.

Time-to-resolution tracks the taxonomy's own definitions as much as the correspondence itself: Signed-off (median 9 days) and Absorbed (median 10 days) threads resolve quickly, since either ending can occur at any visit, while Fizzled cannot be coded before its six-week floor (median 53.5 days) and Orphaned threads run longest of all (median 68 days), accumulating the two-or-more unanswered notes the taxonomy requires before it will commit to the label.

Removing the explicit-request-for-reply feature drops overall accuracy from 51.9% to 51.6% (n.s.); we retain it for completeness, since its coefficient carries the correct sign even where it does not move the aggregate figure.

Discussion

The model's advantage concentrates almost entirely in Orphaned, the one category whose criterion — two or more unanswered notes to the same interlocutor — is itself a property of persistence rather than resolution. An opening note that names its addressee, asks directly, and requests a reply is, on this reading, not predicting a happy ending; it is predicting the kind of thread whose closure will eventually be assigned by the box rather than negotiated between its correspondents, once enough time and enough unanswered notes have accumulated to make the call. Fizzled, Signed-off, and Absorbed all remain comparatively opaque from an opening line alone, which is consistent with a reading in which most street-library correspondence simply ends the way conversations end — unremarkably, and for reasons the opening act does not encode [2].

This pattern echoes, at a much smaller scale, an asymmetry the University's own internal coding taxonomies keep surfacing: a register built to record how something concludes tends to predict its own upkeep — who is still checking the box, still applying the code, still willing to write "Orphaned" on a thread rather than let it lapse unclassified — more reliably than it predicts the correspondents' own sense of an ending [4], [5], [16]. We do not push that parallel further than the data allow.

Limitations

Our sample is drawn from a single metropolitan area over a single nine-month window; a longer or geographically broader deployment might shift the tier split reported above, though the split's direction was consistent across each of the study's three-month sub-periods. The six-week Fizzled threshold is a modelling convenience rather than a claim about correspondents' own sense of abandonment, and a different threshold would relabel some Orphaned and Absorbed threads without touching the Signed-off count, which requires no threshold at all. Coders were blind to box-traffic tier throughout, so the tier effect in Results cannot reflect a coding artefact; it can, however, reflect nothing more than which streets happen to attract note-leavers who prefer to say goodbye.

Conclusion

We have introduced a four-category taxonomy for how correspondence left in neighbourhood street-library boxes concludes, applied it to 341 threads across 74 boxes, and shown that a model trained on a thread's opening note alone can anticipate one of the four endings — Orphaned — meaningfully better than a majority-class baseline, while the other three stay largely opaque from the opening line. The taxonomy travels, we think, to other unattended correspondence settings an institution might be tempted to measure without having designed them to be measured. Future work: a second metropolitan area, a longer observation window, and a sensitivity check against Fizzled thresholds other than six weeks.

References

[1] P. Chen, "The Contribution of Street Libraries in Australia to Literacy, Community and the Gift Economy," *Journal of the Australian Library and Information Association*, 2022, doi: 10.1080/24750158.2022.2028332.

[2] E. A. Schegloff and H. Sacks, "Opening up Closings," *Semiotica*, vol. 8, no. 4, pp. 289–327, 1973, doi: 10.1515/semi.1973.8.4.289.

[3] L. Stanley, "The Death of the Letter? Epistolary Intent, Letterness and the Many Ends of Letter-Writing," *Cultural Sociology*, vol. 9, no. 2, pp. 240–255, 2015, doi: 10.1177/1749975515573267.

[4] A. Mensah, M. Solheim, and V. Marris, "What the Roster Remembers: A Coding Taxonomy of Deflection Annotations on the Shared-Office Pot-Plant Watering Roster," *Slop University technical report*, 2026, doi: 10.5555/slop.8kn3k2.

[5] M. Solheim and R. Sabel, "Closed, Reclassified, or Left Open: A Hand-Coded Taxonomy of Bathroom-Break Integrity-Attentiveness Flags Across Eight Terms at the Exit Sign-Out Table," *Slop University technical report*, 2026, doi: 10.5555/slop.jw4nw9.

[6] A. Mckenna-Foster and H. Roseen, "What's in a Book Exchange: Examining Contents in Relation to Steward Intentions, Geography, and Public Library Collections," *Journal of Librarianship and Information Science*, 2023, doi: 10.1177/09610006221124336.

[7] J. Zhang *et al.*, "Conversations Gone Awry: Detecting Early Signs of Conversational Failure," *arXiv preprint arXiv:1805.05345*, 2018.

[8] Y. Jiao, C. Li, F. Wu, and Q. Mei, "Find the Conversation Killers: A Predictive Study of Thread-Ending Posts," *arXiv preprint arXiv:1712.08636*, 2018.

[9] R. Krohn and T. Weninger, "Modelling Online Comment Threads from Their Start," *arXiv preprint arXiv:1910.08575*, 2019.

[10] C. Danescu-Niculescu-Mizil, M. Sudhof, D. Jurafsky, J. Leskovec, and C. Potts, "A Computational Approach to Politeness with Application to Social Factors," *arXiv preprint arXiv:1306.6078*, 2013.

[11] A. Herdağdelen, L. Adamic, and B. State, "Community Gifting Groups on Facebook," *arXiv preprint arXiv:2211.09043*, 2022.

[12] R. F. Kizilcec, E. Bakshy, D. Eckles, and M. Burke, "Social Influence and Reciprocity in Online Gift Giving," *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, 2018, doi: 10.1145/3173574.3173700.

[13] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: 10.1191/1478088706qp063oa.

[14] D. Ley and R. Cybriwsky, "Urban Graffiti as Territorial Markers," *Annals of the Association of American Geographers*, vol. 64, no. 4, pp. 491–505, 1974, doi: 10.1111/j.1467-8306.1974.tb00998.x.

[15] A. Joulin, E. Grave, P. Bojanowski, and T. Mikolov, "Bag of Tricks for Efficient Text Classification," *arXiv preprint arXiv:1607.01759*, 2016.

[16] D. Okafor, R. Sabel, and T. Solberg, "Held, Not Escalated: Coding the Library Front Desk's Disposal-Tier Register Against Six Rounds of Physical Stock-Take," *Slop University technical report*, 2026, doi: 10.5555/slop.ll5xcj.