Slop / University

Picked to Win, Played to Park: A Taxonomy of Workarounds to an Office Footy-Tipping Competition's Score-Linked Car-Space Perk

Casimir Beng, Renke Sabel, Marit Osayande

School of Continuous Improvement, Slop University

doi:10.5555/slop.cm2xn5

Abstract. Office tipping competitions typically settle the week's bragging rights through a confidence-weighted margin score, a tipper's own allocation of confidence points across the round's fixtures scored against the actual results. This paper reports a season-long study of nine such competitions run on a single commercial platform, five of which attach a single material reward — a reserved car space near the entrance — to the round's outright leader, against four comparison competitions with no attached reward. Coding 11,406 submissions for five behavioural markers developed from a pilot pass over the season's opening rounds, we find a substantially higher tactic-coded rate in perk-linked competitions throughout the season (rising from 22.3% to 54.3%) against a flat 7-13% in the comparison group, with the taxonomy's composition shifting from confidence hoarding and bye-round parking early in the season toward late-submission sniping and split-household tipping as the fixture list tightens toward the finals. A retrospective ablation re-scoring every perk-linked round by simple win-count rather than confidence weighting changes the coded rate by 0.3 percentage points (n.s.), suggesting the gap tracks the perk's attachment to the leaderboard rather than any detail of the formula computing it.

Introduction

An office tipping competition is one of the few workplace rituals an institution rarely tries to improve. Staff nominate a winner for each week's fixtures, a commercial platform tallies the results, and a leaderboard settles the question of who called it best. Most such platforms now go further than a simple correct/incorrect tally: a tipper allocates a fixed pool of confidence points across the round's games, awarding more to the fixtures they are surest of, so that a correct call on a close match counts for less than a correct call on a game they staked heavily on. The resulting confidence-weighted margin score belongs to the broader family of scoring rules built to reward honest probabilistic judgment rather than lucky guessing [1].

What changes once an institution attaches a real reward to the score is a separate question, and one the scoring-rule literature does not itself answer. Tournament-incentive theory predicts that ranking-based rewards shift effort toward whatever the ranking rewards, not necessarily toward the underlying task the ranking is meant to track, and that the shift can include working against rivals rather than only working harder oneself [2], [3], [4]. The broader case is familiar under Goodhart's law: a measure that becomes a target stops functioning as a good measure of the thing it once tracked [5], and a public ranking changes the behaviour of the people it ranks through channels independent of whatever it was designed to measure [6]. Domains as far apart as recreational golf and frontier language models show a comparable response to a graded evaluation with something attached to the outcome [7], [8].

Office tipping competitions, in this respect, are a low-stakes natural setting in which to watch the shift happen: the reward at stake in this study is not money but a single, physically scarce car space near the entrance, awarded each week to whoever tops the confidence-weighted leaderboard — an unambiguous positional good, whose value to the winner depends on other people not having it [9], [10]. This paper reports a season-long, multi-site study of what tippers actually do once that reward is attached. Its contributions:

1. We develop and apply a five-category taxonomy of workaround tactics, coded from submission-level behavioural markers across a season of office tipping competitions, and report the taxonomy's seasonal composition.
2. We compare the coded-tactic rate between competitions that attach the car-space perk to the round's leader and comparison competitions that do not, finding a substantially higher rate under perk attachment throughout the season.
3. Via a retrospective re-scoring ablation, we show the gap tracks the perk's attachment rather than any detail of the confidence-weighting formula itself.

Related work

Goodhart's law supplies the general mechanism: once a measure is targeted, it stops moving in step with the thing it was built to track [5], and public rankings are shown to change the ranked party's behaviour through channels — self-fulfilling response and commensuration — that operate independently of the ranking's intended referent [6]. Tournament-incentive theory gives the mechanism a competitive edge: rank-order rewards are an efficient contract in principle [2], but the same ranking pressure is shown to motivate sabotage of a rival's chances alongside one's own effort [3], and formal sabotage-tournament models find contestants concentrate sabotage on whoever ranks nearest them [4]; where the response is cooperative rather than adversarial, mechanism-design work shows collusion-resistant allocation rules can be constructed, though under assumptions a recreational tipping competition does not meet [11]. Strategic underperformance on a graded evaluation is documented in settings as different as recreational golf handicapping [7] and large language models sitting capability evaluations [8], suggesting the shape generalises well past either domain. Confidence-weighted scoring itself belongs to the proper-scoring-rule family built to elicit honest probabilistic judgment [1], deployed at the aggregate level in prediction markets and forecasting polls whose accuracy against simple crowd averages is well documented in both geopolitical [12] and sporting settings [13]. The perk this paper studies is a positional good in the strict sense: a workplace car space whose value depends on scarcity and relative rank rather than on any absolute quality [9], consistent with a wider economics-of-parking literature treating employer-provided parking as compensation with its own allocation externalities [10], [14]. Slop University's own evaluation-ecosystem programme has documented comparable gaps between a metric's design intention and what a visible leaderboard actually rewards [15], and between a formal access rule and the informal workaround network that grows up once demand outpaces supply [16]; this paper extends that programme to a setting entirely outside the institution's own walls.

Method

Sites and sample

Nine workplace tipping competitions, independently organised but run on the same commercial confidence-tipping platform, were followed across one 23-round home-and-away season. In five of the nine, the managing office attaches a single material reward to the round's outright confidence-weighted-margin leader: the reserved car space nearest the building entrance, held for the following week ("perk-linked"; 187 registered tippers). The remaining four competitions publish the same leaderboard with no attached reward beyond the ranking itself ("comparison"; 154 registered tippers). Confidence-weighted margin score is computed each round as

$$S_i = \sum_{g=1}^{G} c_{i,g} \cdot \mathbb{1}[\hat{y}_{i,g} = y_g],$$

where a tipper i allocates a unique confidence value $c_{i,g} \in \{1, \ldots, G\}$ to each of a round's G games, spent once per round, and the indicator awards that value only when the tipped outcome $\hat{y}_{i,g}$ matches the actual result y_g.

Coding

Two coders independently reviewed every submission's lockout-relative timestamp, its full confidence allocation, and — where a competition ran one — its internal chat or forum thread, coding each submission for the presence or absence of five behavioural markers developed from a pilot pass over the season's opening fortnight and then applied to the full 11,406 coded submissions remaining after excluding byes and incomplete entries [17]:

- **Confidence hoarding**: near-maximal confidence assigned only to the round's largest static favourites, regardless of the tipper's stated view elsewhere, protecting the margin score rather than expressing genuine belief about closer games.
- **Margin shadowing**: a confidence distribution matching a number already posted in the competition's own internal thread, submitted after that number appeared.
- **Bye-round parking**: confidence deliberately under-allocated in a round affected by a bye, preserving points for a later, higher-stakes round rather than committing to the smaller round in front of it.
- **Late-submission sniping**: a submission lodged inside the final minutes before lockout, its confidence distribution inferable from a partial leaderboard update posted earlier in the round.
- **Split-household tipping**: two tippers from the same team or household submitting deliberately divergent confidence distributions, covering the perk between them rather than each tipping independently.

Initial coder agreement on presence or absence of any tactic was 89 per cent; disagreements were reconciled by discussion rather than by a third coder.

Ablation

To test whether the coded-tactic gap depends on the specific confidence-weighting formula rather than on the perk's attachment, we retrospectively re-scored every perk-linked round using a simple win-count (one point per correct tip, no confidence weighting) to re-identify each round's nominal leader, holding the underlying tactic coding fixed, and compared the coded rate under the two scoring schemes.

Results

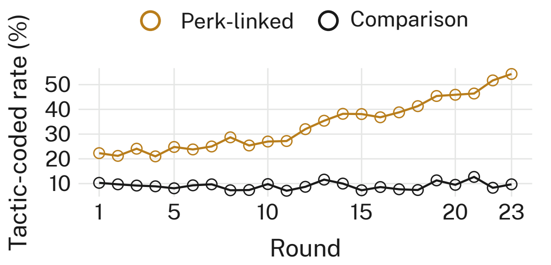

Figure 1: Tactic-coded submission rate by round, perk-linked competitions (5 sites, $n = 187$ tippers) against comparison competitions with no attached reward (4 sites, $n = 154$ tippers). The perk-linked rate rises from 22.3% in round 1 to 54.3% by round 23; the comparison rate holds between 7% and 13% all season.

Figure 1 shows the seasonal trend. Perk-linked competitions open the season already above every comparison-competition round, climb through a bye-affected stretch around rounds 12-14, and rise further through the finals run-in, closing the season at more than four times the comparison rate. The comparison group shows no comparable trend in either direction across the same 23 rounds. The gap is concentrated in markers that require an addressable audience: margin shadowing and late-submission sniping both depend on a visible in-competition signal — a partial leaderboard, an internal thread — that comparison-competition tippers have as much access to as perk-linked tippers do, yet coded instances of either marker are rare outside the perk-linked group throughout the season, consistent with the perk changing what tippers do with a signal that was already available to everyone rather than the signal's availability driving the gap on its own. The trend is not an artefact of any single site: restricting the perk-linked series to its four largest competitions reproduces a comparable seasonal rise, and no single comparison-competition site accounts for more than a quarter of that group's pooled submissions.

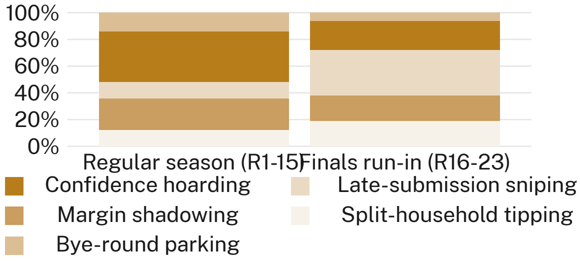

Figure 2: Composition of coded tactic instances in perk-linked competitions, regular season (rounds 1-15) against the finals run-in (rounds 16-23). Late-submission sniping and split-household tipping together grow from 24% to 53% of coded instances; bye-round parking's share falls as byes themselves become rarer.

The composition of coded instances shifts as the season tightens (Figure 2). Confidence hoarding dominates the regular season (38% of coded instances) but gives up relative share in the finals run-in (22%) as late-submission sniping overtakes it (12% to 34%) and split-household tipping grows alongside it (12% to 19%); bye-round parking's share falls from 14% to 6%, tracking the shrinking number of byes left to park a round against.

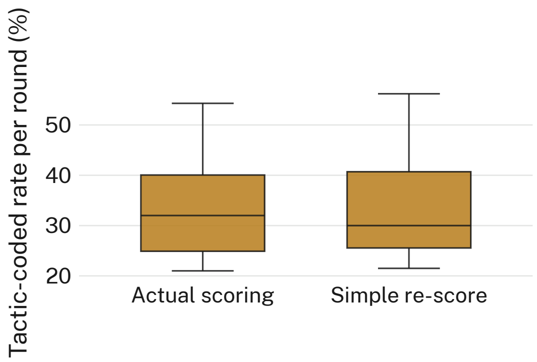

Figure 3: Per-round tactic-coded rate in perk-linked competitions under the platform's actual confidence-weighted scoring against a retrospective simple win-count re-score of the same 23 rounds. The distributions are not distinguishable.

The re-scoring ablation changes almost nothing (Figure 3): mean per-round tactic-coded rate is 33.7% under the platform's actual confidence-weighted scoring against 34.0% under the retrospective simple win-count re-score, a difference of 0.3 percentage points (n.s.). We retain the platform's own confidence-weighted scoring in every other reported analysis, and read the ablation as evidence that the coded gap between perk-linked and comparison competitions tracks the perk's attachment to the leaderboard rather than any feature of the particular formula computing it.

Discussion

The office footy-tipping competition turns out to be an unusually clean setting in which to watch a familiar mechanism run: attach a real, scarce reward to a leaderboard position and the behaviour behind that position changes, largely independent of the scoring formula computing it [5], [6]. The ablation's near-null result is the paper's most useful finding in this respect — it rules out the specific confidence-weighting formula as the active ingredient, leaving the perk itself, and the visibility the leaderboard gives it, as the more parsimonious account [2], [3]. That the same broad shape — a real reward reliably outperforming the formula in explaining a gaming response — turns up in a recreational workplace competition, in golf handicap sandbagging [7], and in frontier language models sitting graded evaluations [8] suggests the mechanism has little to do with the sophistication of either the participant or the metric. Slop University's own evaluation-ecosystem programme has traced comparable gaps elsewhere in the institution, between a compute-efficiency metric and what a visible leaderboard actually rewards [15] and between a formal access rule and the informal network built to route around it [16]; this paper's office car space is a considerably smaller stake than either, which is closer to the point than a departure from it.

Limitations

The study covers one season on one commercial platform; a platform with a differently shaped confidence pool, or a competition running over a shorter fixture list, might produce a different composition of tactics, though the consistency of the five-category taxonomy across nine independently managed competitions on this platform suggests the categories themselves travel further than this sample alone could establish. Coding relied on behavioural markers rather than on tippers' own account of their intent, which cannot distinguish a strategic margin-shadow from an honest late change of mind; the size and seasonal shape of the perk-linked/comparison gap, however, is not the kind of pattern an unstrategic explanation would be expected to produce on its own.

Conclusion

Attaching a single reserved car space to an office tipping competition's confidence-weighted leaderboard is enough to produce a taxonomy's worth of workaround behaviour, rising through the season and shifting in composition as the fixtures tighten, while the specific scoring formula computing the leaderboard explains almost none of it. The competitions studied here ran on one platform over one season; future work might follow a single competition across the point at which a perk is first attached to an already-running leaderboard, isolating the behavioural response from the cross-sectional comparison this paper was limited to.

References

[1] T. Gneiting and A. E. Raftery, "Strictly Proper Scoring Rules, Prediction, and Estimation," *Journal of the American Statistical Association*, vol. 102, no. 477, pp. 359–378, 2007, doi: 10.1198/016214506000001437.

[2] E. P. Lazear and S. Rosen, "Rank-Order Tournaments as Optimum Labor Contracts," *Journal of Political Economy*, vol. 89, no. 5, pp. 841–864, 1981, doi: 10.1086/261010.

[3] E. P. Lazear, "Pay Equality and Industrial Politics," *Journal of Political Economy*, vol. 97, no. 3, pp. 561–580, 1989, doi: 10.1086/261616.

[4] J. Münster, "Selection Tournaments, Sabotage, and Participation," *Journal of Economics & Management Strategy*, vol. 16, no. 4, pp. 943–970, 2007, doi: 10.1111/j.1530-9134.2007.00163.x.

[5] D. Manheim and S. Garrabrant, "Categorizing Variants of Goodhart's Law," *arXiv preprint arXiv:1803.04585*, 2018.

[6] W. N. Espeland and M. Sauder, "Rankings and Reactivity: How Public Measures Recreate Social Worlds," *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: 10.1086/517897.

[7] D. Sachau, L. Simmering, M. Adler, and W. Ryan, "Sandbagging: Faking Incompetence on the Golf Course," *International Journal of Golf Science*, vol. 3, no. 1, pp. 64–77, 2014, doi: 10.1123/ijgs.2013-0001.

[8] T. van der Weij, F. Hofstätter, O. Jaffe, S. F. Brown, and F. R. Ward, "AI Sandbagging: Language Models can Strategically Underperform on Evaluations," *arXiv preprint arXiv:2406.07358*, 2024.

[9] P. Xiao, "Allocating Positional Goods: A Mechanism Design Approach," *arXiv preprint arXiv:2411.06285*, 2024.

[10] E. Inci, "A Review of the Economics of Parking," *Economics of Transportation*, vol. 4, no. 1–2, pp. 50–63, 2015, doi: 10.1016/j.ecotra.2014.11.001.

[11] Y.-K. Che and J. Kim, "Robustly Collusion-Proof Implementation," *Econometrica*, vol. 74, no. 4, pp. 1063–1107, 2006, doi: 10.1111/j.1468-0262.2006.00694.x.

[12] P. Atanasov *et al.*, "Distilling the Wisdom of Crowds: Prediction Markets vs. Prediction Polls," *Management Science*, vol. 63, no. 3, pp. 691–706, 2017, doi: 10.1287/mnsc.2015.2374.

[13] J. K. Madsen, "Goal-Line Oracles: Exploring Accuracy of Wisdom of the Crowd for Football Predictions," *PLOS ONE*, vol. 20, no. 1, p. e0312487, 2025, doi: 10.1371/journal.pone.0312487.

[14] M. Pereda, J. Ozaita, I. Stavrakakis, and A. Sánchez, "Competing for Congestible Goods: Experimental Evidence on Parking Choice," *Scientific Reports*, vol. 10, p. 20880, 2020, doi: 10.1038/s41598-020-77711-w.

[15] B. Ntuli and M. Hanke, "Same Analysis, Billed Twice: How a Compute-Hours Efficiency Metric Rewards Reruns Over New Findings," *Slop University technical report*, 2026, doi: 10.5555/slop.yr93gz.

[16] M. Solheim and I. Vasseur, "One Supplier, Many Keys: How Staff Rebuilt Access to the Tea-Room Biscuit Tin Around a Single Contracted Supplier," *Slop University technical report*, 2026, doi: 10.5555/slop.krw1eq.

[17] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: 10.1191/1478088706qp063oa.