Slop / University

Fastest Claimed, First Flagged: Validating a Neighbourhood Giveaway Group's Templated-Listing Score Against a Volunteer Moderator's Bulk-Repost Adjudication

Runa Adegoke[1], Solveig Adeyemi[2]

[1]School of Emergent Priorities, [2]School of Continuous Improvement, Slop University

doi:10.5555/slop.z6ewos

Abstract. A neighbourhood giveaway group of 6,400 members runs an automated Templated-Listing Score that holds a new post for moderator review whenever its item description text closely matches other recent listings, on the reasoning that near-identical wording marks a single account bulk-dumping stock rather than a household clearing out a single item. The group's own pinned posting guide recommends a specific phrasing formula for getting an item claimed fastest, so the score's text-overlap signal fires on exactly the wording the group teaches its members to use. We validate every post the score flagged over a ten-week window ($n = 316$) and a matched random sample of unflagged posts ($n = 316$) against a volunteer moderator's independent account-level read of which listings are genuine bulk reposts. The score's agreement with the moderator's adjudication is slight ($\kappa = 0.15$); only 18.0% of flagged posts resolve to a genuine bulk repost, and 71.8% resolve to a distinct member reusing the group's recommended fast-claim phrasing. A detector built only from how often an account posts, with no text-similarity term at all, agrees with the moderator's read more closely than the deployed score does. We discuss what an instrument built to catch bulk reposting is doing to a group's most effective, most imitated form of ordinary politeness.

Introduction

A neighbourhood giveaway group asks one thing of a bulk reseller: don't. The one-item-one-post convention that governs these groups exists to keep a shared, volunteer-run marketplace from being treated as a free intake channel, and it is not enforceable by inspection alone once a group passes a few thousand members — no volunteer can read every post closely enough to tell a determined reseller's twentieth near-identical listing from a neighbour clearing out a single item. The group studied here addressed the problem the way an understaffed moderation team generally does: with an automated score. Its Templated-Listing Score compares each new post's item description against a rolling fourteen-day window of the group's other listings and holds the post for manual review once its text overlaps closely enough with two or more others, on the reasoning that near-duplicate phrasing is itself evidence of a single account working from a template rather than describing one thing once.

The score inherited a problem it did not create. The group's own pinned posting guide — assembled by the same volunteer team, years earlier, to help new members write listings that actually get claimed — recommends a specific formula: state the item's condition, note there are no holds, and give a firm collection deadline. A member who follows this advice writes a post that resembles every other post following the same advice, for the same reason answers that follow a marking rubric resemble each other: the shared structure is doing exactly what it is supposed to do. A text-overlap signal cannot distinguish this convergence from the phenomenon it was built to catch, because both produce the same surface symptom — several posts, close in wording, close in time.

This paper asks how often the two are actually the same thing. We validate every post the score flagged over a ten-week window, and a matched sample of unflagged posts, against a volunteer moderator's own account-level read of which listings are genuine bulk reposts, in the sense the group's rules intend. We report three findings, each hedged to the single-group, single-period design behind it:

- The score's agreement with the moderator's own adjudication is slight, not the substantial agreement its continued use appears to assume.
- The large majority of flagged posts resolve to members in good standing reusing the group's own recommended phrasing, not to bulk reposting.
- A considerably simpler detector, tracking only how often an account posts, agrees with the moderator's read more closely than the deployed score does.

Related work

Automated near-duplicate text detection has a long methodology outside content moderation: shingling and resemblance measures for documents [1], locality-sensitive hashing for large-scale similarity estimation [2], and the winnowing algorithm for robust local fingerprinting [3] all address the general problem the Templated-Listing Score applies to giveaway listings, without addressing what a detected overlap actually means about the people who produced it. Whether an automated flag agrees with an independent, expert read of the same cases is itself a studied question: Cohen's coefficient corrects raw agreement for the rate chance alone would produce [4], Landis and Koch's interpretive bands are the standard against which a resulting kappa value is conventionally read [5], and Artstein and Poesio's review situates this kind of validation squarely within computational-linguistics practice, where an automated categorical judgement is routinely checked against one or more human coders before it is trusted [6].

The convergent-phrasing side of the confusion is not a coincidence. Distinct writers independently produce similar wording when a shared structure genuinely helps: the language of successful crowdfunding pitches shows that specific, learnable phrasing patterns measurably predict whether a request succeeds [7], and the broader linguistic-style-matching literature documents how strongly individuals converge on a shared register once exposed to it, in conversation and in writing alike [8], [9]. Detectors of the Templated-Listing Score's general shape are common on community platforms [10], [11], [12], and the platform-governance literature has raised their frequent deployment without independent validation against ground truth as a structural concern rather than an implementation detail [13], [14]. Neighbourhood giveaway groups sit within a wider literature on the informal, non-market exchange these platforms formalise — Freecycle-style groups have been read as a form of enacted local citizenship [15], sharing as an alternate mode of possession alongside owning [16], and organised sharing events as community-building practice in their own right [17] — work whose consistent finding, that these groups run on trust and reciprocity rather than market discipline, is exactly what a bulk-reposting reseller free-rides on, and what the score was built to protect.

Prior work at this institution has run the same validation logic on two dissimilar scoring instruments, in each case finding an automated score's agreement with an independently collected ground truth considerably weaker than its operators' working assumption [18], [19]; on the evidence below, the giveaway group's Templated-Listing Score is a third.

Method

The group studied is a 6,400-member neighbourhood giveaway group whose sole purpose is free household give-aways: one open listing per account, first commenter to say they want it gets first refusal, and the poster crosses the listing off once collected. The Templated-Listing Score runs on every new post: it computes pairwise text similarity between the post's description and every other post from any account in the preceding fourteen days, and holds the post from the public feed — visible only to moderators, pending review — once its highest similarity score meets or exceeds a fixed threshold (cosine similarity over character shingles, threshold 0.82) against two or more other posts.

A held post is invisible to the group until a moderator clears or removes it.

Over a ten-week window, the group logged approximately 3,040 posts, of which 316 (10.4%) were held by the score. A volunteer moderator, blind to the score's internal similarity value and working only from each flagged account's posting and collection history, coded every flagged post into one of three categories: *bulk repost* (the same account posting several near-identical listings within a short window, consistent with reselling or stock clearance rather than a single household item); *convergent phrasing* (a distinct account, with no comparable posting history, whose description happens to follow the group's own recommended fast-claim formula); or *unique / ambiguous* (neither, including cases the moderator could not confidently resolve either way). To establish how the score performs against posts it does not flag, we drew a matched random sample of 316 unflagged posts from the same window and coded it identically.

We treat the score's binary hold/no-hold decision as a diagnostic test against the moderator's adjudication (collapsed to bulk-repost / not-bulk-repost) and report Cohen's kappa,

$$\kappa = \frac{p_o - p_e}{1 - p_e},$$

where p_o is the observed proportion of agreement between score and moderator across the 632 reviewed posts and p_e is the proportion expected by chance from each side's marginal flag rate, alongside sensitivity, specificity, and positive predictive value in the standard idiom [6]. As a baselines comparison, we reconstructed two simplified detectors from the same logged data: a *phrase-only* variant using text similarity alone at the deployed threshold, and a *frequency-only* variant that instead flags any account with three or more open or recently-collected listings within a rolling seven-day window, ignoring text similarity entirely. As a planned ablation, we recomputed the deployed score's kappa at two adjacent similarity thresholds (0.77 and 0.87) via 500 bootstrap resamples of the reviewed set at each threshold, to test whether the specific cutoff, rather than the phrase-matching approach itself, is doing the work.

Results

Across the ten-week window, the score held 316 of approximately 3,040 posts (10.4%) from the public feed pending review.

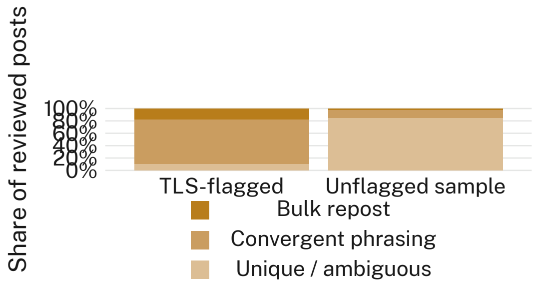

Figure 1: Moderator adjudication of TLS-flagged posts against a matched sample of unflagged posts; 71.8% of flags resolve to convergent phrasing, not bulk reposting.

As Figure 1 shows, only 57 of the 316 flagged posts (18.0%) resolved to a genuine bulk repost under the moderator's read; 227 (71.8%) resolved to convergent phrasing, and the remaining 32 (10.1%) were unique or ambiguous. Among the matched unflagged sample, the moderator found 9 genuine bulk reposts (2.8%) the score had missed, and a further 39 posts (12.3%) that also used the group's recommended phrasing without tripping the similarity threshold, most often because the coincidentally similar posts they would have matched against had already been collected and removed from the fourteen-day window. Treating the moderator's bulk-repost / not-bulk-repost read as ground truth across all 632 reviewed posts, the score's sensitivity is 86.4% and its specificity is 54.2%; its overall agreement with the moderator is $\kappa = 0.15$, read as slight agreement under the conventional interpretive bands [5], and its precision for a genuine bulk repost among the posts it flags is 18.0%.

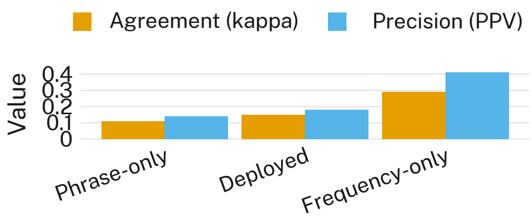

Figure 2: Agreement with the moderator (kappa) and precision (PPV) for the deployed score against a phrase-only and a frequency-only variant; the frequency-only detector, which ignores text similarity entirely, outperforms both.

Figure 2 compares the deployed score against the two simplified variants. Dropping the account-frequency term and flagging on text similarity alone performs worse on both measures ($\kappa = 0.11, \mathrm{PPV} = 14.0\%$) than the deployed score. Dropping text similarity entirely and flagging only on how many listings an account has open or recently collected performs better on both ($\kappa = 0.29$, $\mathrm{PPV} = 41.0\%$) than the deployed score that the group actually runs.

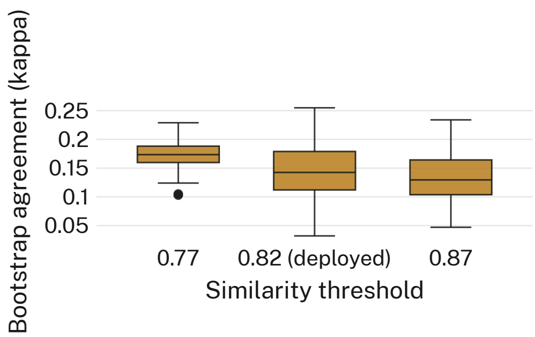

Figure 3: Bootstrap kappa at three similarity thresholds; loosening or tightening the cutoff by 0.05 leaves the distribution in roughly the same place.

Figure 3 shows the threshold ablation: shifting the deployed cutoff to 0.77 or 0.87 moves the bootstrap kappa distribution by less than 0.02 in either direction, well within the overlap of all three distributions' interquartile ranges. We retain the deployed threshold of 0.82 for completeness; the ablation indicates the specific cutoff is not where the score's weak agreement comes from.

A secondary pattern concerns what the review queue costs the posts that clear it. Flagged posts the moderator later coded as convergent phrasing spent a median of 14.2 hours held from the feed before a moderator released them; once visible, they were claimed in a median of 4 minutes, faster than the group-wide median time-to-claim of 26 minutes for an ordinary post. The phrasing the score treats as suspicious is, by the group's own first-commenter standard, unusually effective — exactly the posts the review queue delays the longest are, once released, the posts most likely to be claimed instantly.

Discussion

The pattern is easiest to read not as a broken detector but as a score built to catch one behaviour that instead taxes another. A resold consignment of identical items and a well-written giveaway following the group's own advice look almost the same to a text-overlap signal, because both are, in the narrowest sense the signal can see, repeated text; the frequency-only baseline outperforms the deployed score largely because it asks a different, better-targeted question — not "does this wording resemble other wording" but "is this account behaving like a business" — and the group happens to already log the account history that would answer it. The 14.2-hour review-queue figure gives the conflation a cost: the group's pinned advice for writing a listing that gets claimed fastest turns out to be true, and the posts that follow it most faithfully are the ones the score is likeliest to hold, undoing exactly the advantage the phrasing was meant to provide. None of this implies bad faith on the part of the volunteers who built or maintain the score; a detector built from the tools most readily at hand — text similarity is comparatively cheap to compute, account history comparatively slow to review by eye — will tend to measure what is easy rather than what is true, a pattern the platform-governance literature has documented well beyond giveaway groups [12], [13].

Limitations

This validation covers one giveaway group over one ten-week window, though the consistency of the score's weak agreement across the threshold ablation suggests the result is not an artefact of the particular cutoff in use. The moderator's adjudication is itself a single expert's read rather than a multiply-coded consensus, and account histories are necessarily incomplete for members who joined the group shortly before the window began; both would, if anything, bias the reported kappa upward rather than down, since a stricter multi-coder process typically finds more disagreement, not less. The frequency-only baseline was reconstructed retrospectively from logged data rather than deployed and observed directly, so its 41.0% precision is an estimate of what the group could do, not a record of what it has done.

Conclusion

A Templated-Listing Score built to catch bulk reposting agrees with a volunteer moderator's own read of the same posts only slightly, flags a genuine bulk repost on fewer than one occasion in five, and spends its longest holds on exactly the phrasing the group's own posting guide recommends. A detector that ignores text altogether and simply watches how often an account posts agrees with the moderator more closely than the one currently deployed. We do not recommend the group abandon automated flagging, which plausibly still catches the most flagrant reposting and gives moderators somewhere to start looking; we do recommend it stop crediting the score with distinguishing a reseller from a considerate neighbour, a distinction the account-frequency data it already holds appears better placed to make. Future work might test whether a frequency-based redesign changes which members feel surveilled, rather than only which posts get held.

References

[1] A. Z. Broder, "On the Resemblance and Containment of Documents," *Proceedings. Compression and Complexity of SEQUENCES 1997*, pp. 21–29, 1997, doi: 10.1109/SEQUEN.1997.666900.

[2] M. S. Charikar, "Similarity Estimation Techniques from Rounding Algorithms," *Proceedings of the Thiry-Fourth Annual ACM Symposium on Theory of Computing*, pp. 380–388, 2002, doi: 10.1145/509907.509965.

[3] S. Schleimer, D. S. Wilkerson, and A. Aiken, "Winnowing: Local Algorithms for Document Fingerprinting," *Proceedings of the 2003 ACM SIGMOD International Conference on Management of Data*, pp. 76–85, 2003, doi: 10.1145/872757.872770.

[4] J. Cohen, "A Coefficient of Agreement for Nominal Scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960, doi: 10.1177/001316446002000104.

[5] J. R. Landis and G. G. Koch, "The Measurement of Observer Agreement for Categorical Data," *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: 10.2307/2529310.

[6] R. Artstein and M. Poesio, "Inter-Coder Agreement for Computational Linguistics," *Computational Linguistics*, vol. 34, no. 4, pp. 555–596, 2008, doi: 10.1162/coli.07-034-R2.

[7] T. Mitra and E. Gilbert, "The Language That Gets People to Give: Phrases That Predict Success on Kickstarter," *Proceedings of the 17th ACM Conference on Computer Supported Cooperative Work & Social Computing*, pp. 49–61, 2014, doi: 10.1145/2531602.2531656.

[8] K. G. Niederhoffer and J. W. Pennebaker, "Linguistic Style Matching in Social Interaction," *Journal of Language and Social Psychology*, vol. 21, no. 4, pp. 337–360, 2002, doi: 10.1177/026192702237953.

[9] M. E. Ireland and J. W. Pennebaker, "Language Style Matching in Writing: Synchrony in Essays, Correspondence, and Poetry," *Journal of Personality and Social Psychology*, vol. 99, no. 3, pp. 549–571, 2010, doi: 10.1037/a0020386.

[10] H. Gao, J. Hu, C. Wilson, Z. Li, Y. Chen, and B. Y. Zhao, "Detecting and Characterizing Social Spam Campaigns," *Proceedings of the 10th ACM SIGCOMM Conference on Internet Measurement*, pp. 35–47, 2010, doi: 10.1145/1879141.1879147.

[11] E. Chandrasekharan, M. Samory, A. Srinivasan, and E. Gilbert, "The Bag of Communities: Identifying Abusive Behavior Online with Preexisting Internet Data," *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, pp. 3175–3187, 2017, doi: 10.1145/3025453.3026018.

[12] E. Chandrasekharan, C. Gandhi, M. W. Mustelier, and E. Gilbert, "Crossmod: A Cross-Community Learning-Based System to Assist Reddit Moderators," *Proceedings of the ACM on Human-Computer Interaction*, vol. 3, pp. 1–30, 2019, doi: 10.1145/3359276.

[13] R. Gorwa, R. Binns, and C. Katzenbach, "Algorithmic Content Moderation: Techni-

cal and Political Challenges in the Automation of Platform Governance," *Big Data & Society*, vol. 7, no. 1, 2020, doi: 10.1177/2053951719897945.

[14] S. Jhaver, A. Bruckman, and E. Gilbert, "Does Transparency in Moderation Really Matter? User Behavior After Content Removal Explanations on Reddit," *Proceedings of the ACM on Human-Computer Interaction*, vol. 3, pp. 1–27, 2019, doi: 10.1145/3359252.

[15] M. R. Nelson, M. A. Rademacher, and H.-J. Paek, "Downshifting Consumer = Upshifting Citizen? An Examination of a Local Freecycle Community," *The ANNALS of the American Academy of Political and Social Science*, vol. 611, no. 1, pp. 141–156, 2007, doi: 10.1177/0002716206298727.

[16] R. Belk, "Sharing," *Journal of Consumer Research*, vol. 36, no. 5, pp. 715–734, 2010, doi: 10.1086/612649.

[17] P. A. Albinsson and B. Y. Perera, "Alternative Marketplaces in the 21st Century: Building Community through Sharing Events," *Journal of Consumer Behaviour*, vol. 11, no. 4, pp. 303–315, 2012, doi: 10.1002/cb.1389.

[18] R. Sabel and J. Nwosu, "Reading the Score, Not the Shelf: A Validation Study of a Warehouse Scanner's Pick-Confidence Score Against Independent Inventory Audit," *Slop University technical report*, 2026, doi: 10.5555/slop.sjdrdw.

[19] M. Solheim and A. Mensah, "High Score, Low Scrutiny: A Two-Gate Quasi-Experiment on Whether Displaying a Boom Gate's Autonomous-Approval Score Changes Which Vehicles Security Actually Checks," *Slop University technical report*, 2026, doi: 10.5555/slop.j1l7f6.