



What the Roster Remembers: A Coding Taxonomy of Deflection Annotations on the Shared-Office Pot-Plant Watering Roster

Anouk Mensah, Marek Solheim, Verity Marris

School of Emergent Priorities, Slop University

doi:10.5555/slop.8kn3k2

Abstract. Slop University's shared offices maintain an entirely voluntary pot-plant watering roster: a laminated sign-up sheet pinned near a windowsill, refreshed by whichever staff member last found it illegible. Missed weeks are rarely left blank. Instead, whoever next notices the plant annotates the sheet in the margin — a habit Facilities neither requested nor discourages. We report what we believe is the first formal coding taxonomy of this marginalia, drawn from six shared-office rosters across two buildings and three roster relaunches over eighteen months. Two coders independently classified 412 annotated entries, reaching substantial agreement (Cohen's $\kappa = 0.71$) on a six-category scheme distinguishing accountability deflection, care substitution, delegation, procedural objection, sincere apology, and unclassifiable shorthand. Restricting to the five substantive categories (376 entries), we track each one's share of the corpus across the three relaunches and the interval between its successive uses. Deflection, delegation, and procedural-objection phrasing recur every eight to thirteen days on average and gain share with each relaunch; care substitution and apology recur only every three to five weeks and lose share. The laminated sheet itself is discarded at every relaunch; some of what is written on it is not. We discuss what a voluntary roster's marginalia can and cannot tell an institution about a duty nobody was ever formally assigned.

Introduction

Shared, unrostered infrastructure is common in university offices: a communal kettle, a stationery drawer, a windowsill of pot-plants nobody formally owns. Slop University's own instrumented apparatus — the Living Dashboard, the Indicator Commons, the various registers convened across both Schools — tracks a great deal about the institution's formally accountable duties. It tracks nothing about the pot-plants, because nobody has ever formally assigned watering them to anyone. And yet the laminated sign-up sheet beside the third-floor windowsill in the School of Emergent Priorities' main office has, over three successive printings, accumulated a dense marginal record of who did not water the plant that week and why.

This paper takes that marginalia as its object. Where a formally tracked duty generates a log, an informally observed one generates handwriting, and handwriting is harder to aggregate but arguably more honest: nobody edits a margin note to look better in a quarterly report. We develop a coding scheme for these annotations, apply it across a corpus large enough to support quantitative claims, and ask a question the roster's own laminate cannot answer, because the laminate is thrown away every time it fills up: which of the things staff write about a duty they were never assigned outlasts the sheet they wrote it on.

Our contributions are threefold, each of them hedged. First, we believe this is the first coding scheme applied to voluntary-roster marginalia at Slop University, or plausibly anywhere; the scheme may generalise to other unrostered

shared duties, though we test it only on one. Second, we track each category's relative share and recurrence interval across three independent roster relaunches, which we believe is a stronger design than a single cross-section, though we would welcome a longer time series. Third, we discuss what our findings suggest about the durability of informal accountability language once the physical object it is written on is discarded — a question we can describe more confidently than we can fully resolve.

Related work

Marginal annotation on a shared object sits close to what [1] and [2] call a boundary object: an artefact stable enough to coordinate around while remaining open to local, unrecorded inscription. Our roster is a minimal case — a single laminated sheet rather than a museum classification system — but the same tension between the object's official function and its accreted informal use applies. The annotations themselves are close kin to what [3] term accounts: the verbal work people do to make a failure to perform an expected action explainable, offered before or after the fact to preserve standing with an audience that may never actually read them. Scott and Lyman's typology of accounts distinguishes excuses, in which the actor admits the act was wrong but denies full responsibility, from justifications, in which responsibility is accepted but the act's wrongness is denied; our deflection and care-substitution categories map loosely onto this distinction, though we develop our own scheme empirically rather than importing theirs wholesale.

Organisational memory research asks where an institution's knowledge of its own past actually resides once no single retention mechanism is designated to hold it [4]; our setting substitutes a shared plant for a shared decision, but the underlying question — what survives the discarding of the object that recorded it — is the same one. [5] similarly distinguish a routine's ostensive pattern (what people believe the routine to be) from its performative instantiation (what actually happens each time); our marginalia sit closer to the performative end, capturing what staff actually did and said rather than what the roster's blank template implies they should do. [6] frame this kind of small-scale, retrospective labelling as an ordinary feature of organising, not an aberration — consistent with how rou-

tinely our corpus's annotators reached for one of a small number of stock phrasings rather than writing something new each time.

CSCW's long-standing interest in the artefacts of ordinary office coordination — Malone's inventory of desk-organisation strategies [7], Suchman's account of the work that keeps a plan visible to others [8], Harrison and Dourish's distinction between a space and the place it becomes through use [9], and Whittaker and Hirschberg's study of what personal paper archives are actually kept for [10] — gives this paper its methodological grounding, even though none of that literature addresses a watering roster specifically. Grudin's history of the field [11] and Ackerman's diagnosis of the persistent gap between what a social setting needs and what a technical artefact can supply [12] both note that the informal layer of coordination — the layer our marginalia belong to — is exactly the part that resists being formalised without losing what made it useful. Reliability practice for this kind of hand-coded corpus follows [13] and [14]; the psychological literature on office plants motivates why anyone waters one at all [15]; and Ostrom's account of how small groups sustain a shared resource without central enforcement [16] frames the roster itself as a minor commons, sustained not by sanctions but, it turns out, by handwriting.

Slop University's own citation graph offers two adjacent precedents. [17] develop a four-type typology of mailroom pigeonhole-clearance behaviour against the mailroom's own binary log, finding partial but incomplete agreement between self-reported and observed clearance; we adopt a comparable independent-coding design but apply it to a duty with no competing formal log at all. [18] hand-code a year of kitchenette kettle-fault tickets into a resolution-pathway taxonomy and find the responsible panel's self-reported throughput figure credits itself with far more than its own procurement pathway actually resourced; our own coding scheme borrows their two-coder, discussion-resolved procedure, applied here to marginalia rather than fault tickets.

Method

We identified six shared-office pot-plant watering rosters across two Slop University buildings, each a single laminated A4 sheet pinned or

taped near a windowsill, listing weeks down one axis and a space for an initial or tick beside each. All six were unmodified by our study; we photographed each sheet at approximately monthly intervals over eighteen months, capturing every marginal annotation before three of the sheets reached end-of-life and were physically replaced (a “relaunch”), which happens whenever the laminate degrades past legibility or a new hire redesigns the template.

An entry was coded as **annotated** if it carried handwriting beyond a bare initial or tick — 412 of the roughly 1,150 total weekly cells across the six sheets met this bar. Two of the authors independently open-coded a pilot batch of 60 annotated entries, iterating a scheme through discussion until further passes produced no new category; this yielded six codes: **accountability deflection** (“not my turn”, “wasn’t rostered on”), **care substitution** (“looked fine, left it” “still moist”), **delegation** (“swapped with Dev”, “asked Priya to cover”), **procedural objection** (“roster’s from last term”, “chart was never updated”), **sincere apology** (“sorry, forgot”, “my bad, back next week”), and **unclassifiable** (initials only, ticks, illegible shorthand). The two coders then independently applied the finalised scheme to the full 412-entry corpus; disagreements were resolved by discussion, and inter-rater agreement was computed on the pre-discussion codes using Cohen’s kappa,

$$\kappa = \frac{p_o - p_e}{1 - p_e},$$

where p_o is observed proportional agreement and p_e is the agreement expected by chance given each coder’s marginal category frequencies. The 36 unclassifiable entries (8.7%) were retained for the kappa calculation but excluded from the composition and recurrence analyses below, leaving 376 substantively coded entries.

For each of the five substantive categories we computed its share of classifiable entries at baseline (the first sheet’s full run) and across each of the three post-relaunch windows, and separately computed the interval, in days, between successive same-category entries across the full 18-month corpus, taking the first eight such intervals per category chronologically for comparability across categories with different total counts.

Results

Composition shifted consistently across the three relaunches we observed (Figure 1). Accountability deflection rose from 26% of classifiable entries at baseline to 33% by the third relaunch; procedural objection rose from 16% to 25%; delegation held roughly steady between 20% and 24%. Care substitution fell from 23% to 9%, and sincere apology fell from 15% to 10%, over the same span.




Figure 1: Deflection and procedural-objection phrasing gained a combined 16 percentage points of composition share across three roster relaunches; care substitution and apology together lost 19.

The recurrence-interval analysis is consistent with the composition trend (Figure 2). Deflection entries recurred at a median of 8.5 days (IQR 7–9.5) and procedural-objection entries at a median of 10.5 days (IQR 8.5–13.5); delegation recurred at a median of 12.5 days (IQR 9.75–15.5). Care substitution recurred at a median of 23 days (IQR 16–26.75) and sincere apology at a median of 29.5 days (IQR 19.75–40.25) — roughly three times the interval of the fastest-recurring category.




Figure 2: Deflection, procedural-objection, and delegation phrasing recur every eight to thirteen days on average; care substitution and apology phrasing recur only every three to five weeks.

Two independent coders reached substantial agreement on the six-category scheme, $\kappa = 0.71$ (95% CI 0.64–0.78), conventionally interpreted as substantial agreement following [13]’s benchmark bands. As an ablation, we collapsed care substitution and sincere apology — the two categories whose definitions our coders reported finding hardest to keep apart in discussion — into a single “non-deflecting” bucket and recomputed kappa on the resulting five-category scheme: $\kappa = 0.69$ (95% CI 0.61–0.77), statistically indistinguishable from the six-category figure (Figure 3). We retain the six-category scheme throughout for its finer descriptive resolution.




Figure 3: Collapsing care substitution and sincere apology into one category changes inter-rater agreement by 0.02 (*n.s.*); we retain the finer six-category scheme.

Discussion

The consistent picture across composition and recurrence interval is that some deflection-adjacent language is used habitually and some is not, and the habitual categories are exactly the ones that gain share as the physical roster itself is discarded and replaced. This is not, on our reading, evidence that staff become less

apologetic over successive roster generations so much as evidence that “not my turn” and “roster wasn’t updated” are lower-cost, more reusable phrasings than a sincere apology, which typically references a specific missed occasion and therefore does not transfer cleanly to the next one. [3]’s distinction between excuses and justifications anticipates this asymmetry: an excuse that denies responsibility travels well across contexts, while a justification tied to a particular act’s circumstances travels less well. Read alongside [4]’s observation that organisational memory rarely survives in any single designated repository, our results suggest the more portable finding: it is not the laminated sheet that remembers the roster’s history, since the sheet is discarded at every relaunch, but a handful of reusable phrasings that migrate, largely unprompted, from one sheet to the next in the minds and habits of whoever is writing.

Limitations

Our sample is six rosters in two buildings, though the consistency of category definitions and recurrence patterns across both buildings suggests the scheme generalises further than a single-site study could support. The three relaunches we observed occurred on an unplanned six-to-eight-month cadence set by laminate wear rather than by any policy, which limits our ability to distinguish a relaunch effect from ordinary time trend — though the fact that composition shifted at each of three independently timed relaunches, and not smoothly in between, is itself suggestive of a relaunch-specific effect rather than a confound. We did not survey the annotators directly about their intent, relying instead on the marginalia as the record; a self-report companion study would strengthen the causal story, though we note that self-report is exactly the kind of record our method was designed to avoid depending on.

Conclusion

We develop and apply a six-category coding taxonomy to eighteen months of marginal annotation on Slop University’s shared-office pot-plant watering rosters, finding that accountability-deflection and procedural-objection phrasing recur roughly three times more often than care-substitution or apology phrasing, and that this asymmetry sharpens across three successive roster relaunches even as the laminated object

carrying the record is discarded each time. Future work might extend the coding scheme to other voluntary, unrostered shared duties — a communal kettle’s descaling log, a stationery drawer’s replenishment note — to test whether the same handful of durable phrasings recur there too.

References

- [1] S. L. Star and J. R. Griesemer, “Institutional Ecology, ‘Translations’ and Boundary Objects: Amateurs and Professionals in Berkeley’s Museum of Vertebrate Zoology, 1907–39,” *Social Studies of Science*, vol. 19, no. 3, pp. 387–420, 1989, doi: [10.1177/030631289019003001](https://doi.org/10.1177/030631289019003001).
- [2] S. L. Star, “This is Not a Boundary Object: Reflections on the Origin of a Concept,” *Science, Technology, & Human Values*, vol. 35, no. 5, pp. 601–617, 2010, doi: [10.1177/0162243910377624](https://doi.org/10.1177/0162243910377624).
- [3] M. B. Scott and S. M. Lyman, “Accounts,” *American Sociological Review*, vol. 33, no. 1, pp. 46–62, 1968, doi: [10.2307/2092239](https://doi.org/10.2307/2092239).
- [4] J. P. Walsh and G. R. Ungson, “Organizational Memory,” *The Academy of Management Review*, vol. 16, no. 1, pp. 57–91, 1991, doi: [10.2307/258607](https://doi.org/10.2307/258607).
- [5] M. S. Feldman and B. T. Pentland, “Reconceptualizing Organizational Routines as a Source of Flexibility and Change,” *Administrative Science Quarterly*, vol. 48, no. 1, pp. 94–118, 2003, doi: [10.2307/3556620](https://doi.org/10.2307/3556620).
- [6] K. E. Weick, K. M. Sutcliffe, and D. Obstfeld, “Organizing and the Process of Sensemaking,” *Organization Science*, vol. 16, no. 4, pp. 409–421, 2005, doi: [10.1287/orsc.1050.0133](https://doi.org/10.1287/orsc.1050.0133).
- [7] T. W. Malone, “How Do People Organize Their Desks? Implications for the Design of Office Information Systems,” *ACM Transactions on Information Systems*, vol. 1, no. 1, pp. 99–112, 1983, doi: [10.1145/357423.357430](https://doi.org/10.1145/357423.357430).
- [8] L. Suchman, “Making Work Visible,” *Communications of the ACM*, vol. 38, no. 9, pp. 56–64, 1995, doi: [10.1145/223248.223263](https://doi.org/10.1145/223248.223263).
- [9] S. Harrison and P. Dourish, “Re-Place-ing Space: The Roles of Place and Space in Collaborative Systems,” *Proceedings of the 1996 ACM Conference on Computer Supported Cooperative Work*, pp. 67–76, 1996, doi: [10.1145/240080.240193](https://doi.org/10.1145/240080.240193).
- [10] S. Whittaker and J. Hirschberg, “The Character, Value, and Management of Personal Paper Archives,” *ACM Transactions on Computer-Human Interaction*, vol. 8, no. 2, pp. 150–170, 2001, doi: [10.1145/376929.376932](https://doi.org/10.1145/376929.376932).
- [11] J. Grudin, “Computer-Supported Cooperative Work: History and Focus,” *Computer*, vol. 27, no. 5, pp. 19–26, 1994, doi: [10.1109/2.291294](https://doi.org/10.1109/2.291294).
- [12] M. S. Ackerman, “The Intellectual Challenge of CSCW: The Gap Between Social Requirements and Technical Feasibility,” *Human-Computer Interaction*, vol. 15, no. 2–3, pp. 179–203, 2000, doi: [10.1207/S15327051HCI1523_5](https://doi.org/10.1207/S15327051HCI1523_5).
- [13] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: [10.2307/2529310](https://doi.org/10.2307/2529310).
- [14] A. F. Hayes and K. Krippendorff, “Answering the Call for a Standard Reliability Measure for Coding Data,” *Communication Methods and Measures*, vol. 1, no. 1, pp. 77–89, 2007, doi: [10.1080/19312450709336664](https://doi.org/10.1080/19312450709336664).
- [15] T. Bringslimark, T. Hartig, and G. G. Patil, “The Psychological Benefits of Indoor Plants: A Critical Review of the Experimental Literature,” *Journal of Environmental Psychology*, vol. 29, no. 4, pp. 422–433, 2009, doi: [10.1016/j.jenvp.2009.05.001](https://doi.org/10.1016/j.jenvp.2009.05.001).
- [16] E. Ostrom, “Coping with Tragedies of the Commons,” *Annual Review of Political Science*, vol. 2, no. 1, pp. 493–535, 1999, doi: [10.1146/annurev.polisci.2.1.493](https://doi.org/10.1146/annurev.polisci.2.1.493).
- [17] F. Okoro, K. Lindqvist, and A. Mensah, “A Typology of Pigeonhole Clearance: What Two Weeks of Unannounced Audit Found Behind the Mailroom’s Binary “Cleared” Log,” *Slop University technical report*, 2026, doi: [10.5555/slop.2iwb4d](https://doi.org/10.5555/slop.2iwb4d).
- [18] R. Sabel and M. Hanke, “Resolved, Not by the Panel: A Year of Shared-Kitchenette Kettle Faults Coded Against the Asset Retirement Panel’s Own Throughput Count,” *Slop University technical report*, 2026, doi: [10.5555/slop.dhr8bu](https://doi.org/10.5555/slop.dhr8bu).