



The Round Still Gets Argued: Interviewing Office Coffee-Run Participants Before and After Their Tab Goes Fully Visible

Casimir Beng, Solveig Adeyemi

School of Continuous Improvement, Slop University

{casimir.beng, solveig.adeyemi}@slop.university

doi:10.5555/slop.9g51xp

Abstract. An informal shared-cost tally — coffee runs, grocery kitties, the office round — is commonly kept privately by whichever colleague happens to be running it, then made fully visible on a shared spreadsheet once disputes accumulate past some threshold. We report a mixed interview and archival-log study of 24 office coffee-run groups that made this transition within the preceding eighteen months, interviewing 41 participants about dispute frequency and category before and after, and independently corroborating their accounts against each group’s own spreadsheet edit and comment history. Overall dispute rate is statistically indistinguishable before and after the transition (1.15 vs 1.05 events per group per month; Wilcoxon signed-rank $W = 142, p = .58$), but the composition of disputes shifts sharply: turn-order and non-payment disputes, the majority of pre-transition events, fall to a combined fifth of post-transition events, while a category absent from the pre-transition taxonomy entirely — disagreement over what counts as a full round — becomes the largest single post-transition category. An ablation comparing groups whose spreadsheet itemises each round by name against groups that pool amounts anonymously finds no reliable difference in post-transition dispute rate. We interpret the shift as evidence that itemisation manufactures the categories in which the next dispute is then argued, rather than resolving the ambiguity it was built to remove.

Introduction

A shared informal cost — the coffee run, the grocery kitty, the round bought on someone else’s card — is among the most low-stakes recurring transactions a workplace produces, and among the least formally governed. Most groups track it however whoever happens to volunteer prefers: a mental tally, a scrap of paper, a running note kept by a rotating “shout captain” in a group chat. When disputes accumulate — a forgotten IOU, a contested turn order — the most common institutional response, in offices as in larger organisations, is to make the ledger fully visible: a shared spreadsheet anyone can open, anyone can audit. Transparency of this kind is a familiar governance reflex well beyond the coffee round, deployed on the assumption

that visibility itself resolves disagreement by removing the ambiguity a private ledger permits [1], [2]. Whether the assumption holds for a cost this small, however, has not to our knowledge been directly tested against what participants actually report happening once the switch is made.

We treat the office coffee-run tab as a minimal instance of a broader institutional pattern — a governance response that substitutes visibility for the underlying dispute it is meant to settle — and ask two questions: does making a shared-cost tally fully visible reduce how often it is argued about, and does it change what the arguments are about. We interview participants from groups that made this transition under ordinary, non-experimental conditions, and cross-

reference their accounts against the tally's own archival record rather than relying on self-report alone.

Our contributions:

- We report, to our knowledge, the first systematic interview-and-archival study of dispute frequency and category before and after an informal shared-cost tally becomes fully visible.
- We show total dispute rate does not detectably change, while the categories in which disputes occur shift substantially, with a wholly new dispute category emerging only once itemisation makes it available to argue about.
- We find no reliable difference in post-transition dispute rate between spreadsheets that name each participant per line and those that pool amounts anonymously, suggesting the shift is driven by itemisation itself rather than by attribution.

The remainder of the paper reviews related work on transparency, informal cost-sharing, and reactive measurement (Related work), describes our interview and archival-log method (Method), reports the rate and composition findings (Results), and discusses what itemisation appears to manufacture rather than resolve (Discussion).

Related work

Transparency is a familiar organisational lever: making information visible is widely assumed to build trust and reduce conflict over resource allocation [1], and empirical work on pay transparency finds it can reduce some forms of counterproductive workplace behaviour while increasing others, depending on whether employees compare favourably or unfavourably to visible peers [3]. Bernstein's transparency paradox complicates the simple story further, showing that visibility can itself degrade the behaviour it is meant to improve by changing what people do when watched [4]. Elinor Ostrom's account of durable commons governance emphasises graduated, locally legitimate monitoring over blanket visibility [5], echoed experimentally in findings that cooperation in shared-resource settings responds more reliably to monitoring than to information alone — whether the monitoring carries a sanction for defection [6] or simply makes a contribution observable to others [7].

A narrower literature treats small shared costs directly: equal bill-splitting is measurably inefficient once individual consumption varies [8], though the inefficiency can average out across a large enough group under simple statistical assumptions [9]. Money earmarked for a specific social purpose is treated differently from fungible cash even at small scale [10], and informal lending among friends and colleagues is managed through deliberate ambiguity precisely because making the ledger explicit is socially costly [11]. Prior work at this institution found a shared office fridge's rising label-compliance count tracked no corresponding drop in its own dispute-ticket log [12], and a shared-umbrella stand's sign-out clipboard bore no reliable relationship to actual custody even as its participation count was reported as a programme success [13]; the present paper extends that pattern from a compliance metric to a dispute record itself. Espeland and Sauder's account of measures that reshape the practices they track [14] and standard practice for coding qualitative interview data to reliability [15], [16], [17] underlie our method.

Method

We recruited 24 office coffee-run groups through professional and social networks, restricting the sample to groups that had, within the preceding eighteen months, transitioned their shared-cost tally from an informal private record — kept from memory or a personal note by a rotating “shout captain” — to a shared spreadsheet visible to every participant. The transition was a group's own organisational choice, not one assigned by us; wording, format, and itemisation level of the resulting spreadsheet were left entirely to whichever participant built it.

We conducted semi-structured interviews with one or two participants per group ($n = 41$), asking each to recall dispute frequency and content in a six-month window before the transition and the six-month window immediately following it. Interviews were coded independently by two research assistants into five mutually exclusive dispute categories (Table 1), developed inductively from a pilot round of six interviews and then applied to the full sample; a second coder independently coded a 25% random subsample, yielding acceptable agreement (Krippendorff's $\alpha = 0.81$).

To avoid relying on self-report alone, we independently obtained each group's own archival record of the tally — prior meeting notes and group-chat exports for the pre-transition period, spreadsheet edit and comment history for the post-transition period — and searched each for evidence of the same dispute events participants described. An event was retained in the final dataset only where the interview account and the archival record independently pointed to the same month and category, or where the archival record alone contained an unambiguous dispute (a reverted edit with a contesting comment) that the interview did not surface; 82% of interview-reported events were independently corroborated this way.

Among the 24 groups, the resulting spreadsheet's design fell into two naturally occurring forms: 11 itemised every round by participant name, and 13 pooled each round's cost anonymously without attributing individual line items. Neither form was assigned by us; we treat the split as a naturally occurring ablation on whether per-person attribution, rather than visibility itself, drives any change in dispute rate.

Category	Operational definition
Turn order	Disagreement over whose turn it was to run or fund the round.
Non-payment	A participant said they would pay their share and did not, within the tally's own settlement window.
Amount owed	Disagreement over the total or per-person amount a round cost.
Round eligibility	Disagreement over whether a specific order — a decaf, a half-serve, an extra shot, a tea instead of coffee — counts as a full round for tallying purposes.
Exemption	Disagreement over whether a participant's absence, leave, or decline should exempt them from a round already tallied.

Table 1: Coding rubric for IOU-dispute events; each event received exactly one category.

We report the paired before/after dispute rate per group per month (Wilcoxon signed-rank test, given the modest per-group sample), the category composition of disputes in each period, and a two-sample comparison of post-transition dispute rate between the named-attribution and pooled-attribution subsamples (Mann-Whitney U , given the uneven group sizes).

Results

Across the 24 groups, paired before/after dispute rate is statistically indistinguishable: a mean of 1.15 events per group per month (SD 0.52) before the transition to a visible spreadsheet, against 1.05 (SD 0.49) after (Wilcoxon signed-rank $W = 142$, $p = .58$). Figure 1 shows this pattern holds at the level of eight recruitment clusters considered separately: four clusters show a small post-transition decrease, three show a small increase, and none moves by more than half an event per month in either direction.




Figure 1: Mean IOU-dispute rate per group per month, before vs after the tally goes fully visible, across eight recruitment clusters. No cluster moves by more than half an event per month.

Figure 1 understates how much actually changed. The category composition of disputes shifts substantially between the two periods (Figure 2). Before the transition, turn-order disputes (38% of 165 events) and non-payment (29%) together account for two-thirds of all recorded disputes, with amount-owed (18%), round-eligibility (6%), and exemption (9%) disputes comparatively rare. After the transition, this ordering inverts: round-eligibility disputes — disagreement over whether a given order counts as a full round — become the largest single category at 34% of 151 post-transition events, followed by amount-owed (26%) and exemption (19%), while turn-order (9%) and non-payment (12%) fall to a combined fifth of events. Total event count is nearly unchanged (165 vs 151); what changes is almost entirely what the dispute is about.




Figure 2: Composition of IOU-dispute categories, before vs after the tally goes fully visible. Round eligibility — absent before — becomes the largest category after.

Eighty-two percent of interview-reported dispute events were independently corroborated in each group's own archival record — a group-chat export or meeting note before the transition, a spreadsheet comment thread or reverted edit after it — indicating participants' accounts are not simply post-hoc narrative.

An ablation tests whether this pattern is driven by the visible spreadsheet's naming of individual participants rather than by its itemisation of rounds as such. Comparing the 11 groups whose spreadsheet attributes each line item by name against the 13 groups whose spreadsheet pools amounts anonymously by round, post-transition dispute rate does not reliably differ (named mean 1.02 events per group per month vs pooled mean 1.07; Mann-Whitney $U = 68$, $p = .71$). Figure 3 shows the full per-group distribution in both arms; the swarms overlap almost completely. We retain both spreadsheet designs as naturally occurring rather than assigned, but the absence of any detectable difference is consistent with itemisation, not attribution, driving the compositional shift.




Figure 3: Post-transition dispute rate by spreadsheet design. A named-attribution line item shifts the median group by well under a tenth of an event per month.

Discussion

The clearest pattern in these data is that making a shared-cost tally fully visible does not so much resolve dispute as relocate it onto the categories the tally itself makes available to argue about. Before itemisation, a round was a single, largely fungible unit; a private tally had no place to record that one order cost more than another, so disagreements clustered on the coarsest available facts — whose turn it was, whether anyone had paid at all. Once every round is itemised by drink, the tally acquires a category — round eligibility — that did not previously exist as a site of dispute, and it promptly becomes the largest one. This reading is consistent with Espeland and Sauder's account of measures that reshape the practices they track [14] and with Zelizer's observation that money earmarked into categories is contested differently from money left fungible [10]: itemisation does not just report on the coffee round more precisely, it manufactures the terms on which the next argument about it will be conducted. The null attribution ablation suggests this is a property of itemisation rather than of surveillance — the dispute migrates whether or not the spreadsheet can identify who is responsible for a given line.

Limitations

Our sample is drawn from groups reachable through professional and social networks and

that self-selected into an interview about their coffee-run tally, so groups for whom the tally remains a live source of tension may be over- or under-represented in ways we cannot fully rule out. The transition to a visible spreadsheet was each group's own choice rather than an assigned condition, so unmeasured differences between groups that adopted a named versus pooled design could in principle account for the ablation's null result rather than itemisation itself; the consistency of the compositional shift across both spreadsheet designs, however, suggests the pattern is not an artefact of either one. Archival corroboration was unavailable for the minority of groups that discarded pre-transition chat history, though the corroboration rate among the remainder was high.

Conclusion

We show that making an informal shared-cost tally fully visible does not detectably reduce how often it is disputed, but does substantially change what the disputes are about, shifting from coarse coordination failures toward disagreements the itemised record itself creates the vocabulary for. Future work might test whether a visible tally that deliberately withholds itemisation — recording only a per-person running total, not a per-round breakdown — avoids manufacturing a round-eligibility category at all, and whether the same pattern holds for other small shared costs that acquire a spreadsheet only once they have already generated enough friction to warrant one. It also raises a narrower methodological question for this institution's own audit apparatus: an instrument built to settle a dispute may need to be evaluated not only on whether disputes fall, but on what new ones it creates the language for.

References

- [1] A. K. Schnackenberg and E. C. Tomlinson, "Organizational Transparency: A New Perspective on Managing Trust in Organization-Stakeholder Relationships," *Journal of Management*, vol. 42, no. 7, pp. 1784–1810, 2016, doi: [10.1177/0149206314525202](https://doi.org/10.1177/0149206314525202).
- [2] Z. Cullen, "Is Pay Transparency Good?," *NBER Working Paper*, no. 31060, 2023, doi: [10.3386/w31060](https://doi.org/10.3386/w31060).
- [3] I. SimanTov-Nachlieli and P. Bamberger, "Pay Communication, Justice, and Affect: The Asymmetric Effects of Process and Outcome Pay Transparency on Counterproductive Workplace Behavior," *Journal of Applied Psychology*, vol. 106, no. 2, pp. 230–249, 2021, doi: [10.1037/apl0000502](https://doi.org/10.1037/apl0000502).
- [4] E. S. Bernstein, "The Transparency Paradox: A Role for Privacy in Organizational Learning and Operational Control," *Administrative Science Quarterly*, vol. 57, no. 2, pp. 181–216, 2012, doi: [10.1177/0001839212453028](https://doi.org/10.1177/0001839212453028).
- [5] E. Ostrom, "Collective Action and the Evolution of Social Norms," *Journal of Economic Perspectives*, vol. 14, no. 3, pp. 137–158, 2000, doi: [10.1257/jep.14.3.137](https://doi.org/10.1257/jep.14.3.137).
- [6] E. Fehr and S. Gächter, "Cooperation and Punishment in Public Goods Experiments," *American Economic Review*, vol. 90, no. 4, pp. 980–994, 2000, doi: [10.1257/aer.90.4.980](https://doi.org/10.1257/aer.90.4.980).
- [7] E. Yoeli, M. Hoffman, D. G. Rand, and M. A. Nowak, "Powering Up with Indirect Reciprocity in a Large-Scale Field Experiment," *Proceedings of the National Academy of Sciences*, vol. 110, pp. 10424–10429, 2013, doi: [10.1073/pnas.1301210110](https://doi.org/10.1073/pnas.1301210110).
- [8] U. Gneezy, E. Haruvy, and H. Yafe, "The Inefficiency of Splitting the Bill," *The Economic Journal*, vol. 114, no. 495, pp. 265–280, 2004, doi: [10.1111/j.1468-0297.2004.00209.x](https://doi.org/10.1111/j.1468-0297.2004.00209.x).
- [9] N. Crokidakis and L. Sigaud, "A Statistical Physics Perspective on Fairness in Shared Expenses: The Bar Bill Analogy," *arXiv preprint arXiv:2507.23542*, 2025.
- [10] V. A. Zelizer, "The Social Meaning of Money: 'Special Monies'," *American Journal of Sociology*, vol. 95, no. 2, pp. 342–377, 1989, doi: [10.1086/229272](https://doi.org/10.1086/229272).
- [11] F. F. Wherry, K. S. Seefeldt, and A. S. Alvarez, "To Lend or Not to Lend to Friends and Kin: Awkwardness, Obfuscation, and Negative Reciprocity," *Social Forces*, vol. 98, no. 2, pp. 753–775, 2019, doi: [10.1093/sf/soy127](https://doi.org/10.1093/sf/soy127).
- [12] M. Osayande and I. Nwachukwu, "Tagged and Tossed, Argued All the Same: A Shared Office Fridge's QR-Label App Against Its Own Clear-Out Dispute Log," *Slop University technical report*, 2026, doi: [10.5555/slop.4qjuix](https://doi.org/10.5555/slop.4qjuix).

- [13] R. Sabel and M. Hanke, "Signed Out, Still on the Peg: Validating a Shared-Umbrella Stand's Sign-Out Clipboard Against a Term of Independent Stand Counts," *Slop University technical report*, 2026, doi: [10.5555/slop.5a3tar](https://doi.org/10.5555/slop.5a3tar).
- [14] W. N. Espeland and M. Sauder, "Rankings and Reactivity: How Public Measures Recreate Social Worlds," *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: [10.1086/517897](https://doi.org/10.1086/517897).
- [15] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: [10.1191/1478088706qp063oa](https://doi.org/10.1191/1478088706qp063oa).
- [16] S. J. Tracy, "Qualitative Quality: Eight 'Big-Tent' Criteria for Excellent Qualitative Research," *Qualitative Inquiry*, vol. 16, no. 10, pp. 837–851, 2010, doi: [10.1177/1077800410383121](https://doi.org/10.1177/1077800410383121).
- [17] A. F. Hayes and K. Krippendorff, "Answering the Call for a Standard Reliability Measure for Coding Data," *Communication Methods and Measures*, vol. 1, no. 1, pp. 77–89, 2007, doi: [10.1080/19312450709336664](https://doi.org/10.1080/19312450709336664).