



Going by the Chat Instead: An Interview Study of When Daters Stop Trusting a Compatibility Percentage

Marit Osayande¹, Runa Adegoke²

¹School of Continuous Improvement ²School of Emergent Priorities, Slop University

doi:10.5555/slop.rngdik

Abstract. A dating application's compatibility percentage is displayed before a single message is exchanged and, on the application's own account, is not revisited once a chat begins. This paper reports a semi-structured interview study of 58 users of a widely used dating application, asked to narrate, message by message, the point at which the percentage stopped functioning as a live signal in a specific match they described in detail. A coded trust-threshold measure places that point at a median of eleven exchanged messages, with 81% of participants reporting the score no longer informed their read of the match by message twenty. Reasons cluster around a chat that read as mismatched despite a high score and a profile that came to read as curated rather than descriptive; participants describe the shift as a private recalibration rather than a stated complaint. In a consenting subsample's match-history exports ($n = 41$, 1,842 conversations), conversational half-life — the exchange count before one side stops replying — is statistically indistinguishable between matches scored above and below ninety percent, and an ablation removing score band from a simple half-life model changes its fit by well under one percentage point (n.s.). High-scoring matches participants had already stopped opening persisted near the top of the recommended queue for a median of nine weeks afterward, the percentage undimmed by a conversation that never resumed.

Introduction

Online dating research has repeatedly found that platform-provided compatibility signals shape a user's initial judgment of a match more reliably than they predict how the resulting exchange goes [1], and a displayed percentage is one of the more literal instances of that signal: the first number most users see about a prospective match, computed before any message is sent and left unrevised once a conversation starts. Its designers describe it as a summary of shared interests, communication style, and stated preferences; its users, this paper finds, mostly stop treating it as one. Chat-mediated impression formation is shown, independently, to generate stronger judgments of a partner than the same information conveyed in a static profile [2] — a plausible account of where a user's trust migrates to. Trust in an algorithmic score more generally is neither fixed nor

uniform: people are shown both to defer to algorithmic judgment past the point warranted by its accuracy [3] and to abandon an algorithm after a single visible error [4] — two findings perhaps better read as describing the same instrument at different moments of use.

This paper asks what that moment looks like for a compatibility percentage: not whether users trust the score in the abstract, but when they say they stopped consulting it, and what happens afterward to matches the score continues to rate highly. Fifty-eight semi-structured interviews, each centred on one described match, locate a coded trust-threshold event and its stated reasoning; a consenting subsample's own match-history exports supply an independent, non-self-reported check on whether score band predicts how long a conversation actually lasts.

The paper's contributions, each modest by design:

1. We develop and apply a five-category coding scheme for the stated reason a participant's trust in the compatibility percentage lapsed, and locate the lapse's median position in the conversation.
2. We compare conversational half-life between matches scored above and below ninety percent directly, finding score band explains almost none of the variation in how long a chat survives.
3. We show that a match a participant has privately stopped trusting the score for does not thereby stop scoring well: high-scoring, already-abandoned matches persist near the top of the recommended queue for weeks afterward.

The percentage is not shown to be wrong, and this paper does not attempt to show that; it is shown to be, for most of the people who once relied on it, a signal with a working life considerably shorter than the queue that keeps displaying it.

Related work

Displayed compatibility and attractiveness signals are shown to shape which profile a user opens far more reliably than they predict what a subsequent conversation is like, across both curated dating-profile ratings [5] and photo-based judgments later found to diverge from an independent read of the same photographs [6], [7]. The broader online-dating literature treats a platform's scoring and search infrastructure as one input into a courtship process it does not fully determine, alongside impression management [8], motivation [9], [10], and a platform's own interface logic [11].

Trust in an algorithmic score specifically is neither stable nor symmetric. Users over-rely on algorithmic judgment relative to its demonstrated accuracy in some settings [3], abandon an algorithm after a single visible failure regardless of its aggregate performance in others [4], calibrate trust differently depending on how much explanation accompanies a system's confidence [12], and more generally make sense of a recommendation partly by constructing their own account of why it was produced [13] — a pattern the trust-in-automation literature frames as a general calibration problem [14]. Folk-theory research documents users building their own working models of what a feed or ranking is

doing, models that diverge from the platform's account and that the platform rarely corrects [15], [16].

Prior work at this institution has already traced a version of this pattern: most managers stopped reopening a roster once a compliance score had been returned to them, independent of whether the score itself proved stable on a second look [17], and a published fairness mechanism has been shown to add almost nothing to a prediction already dominated by what was actually happening on the ground [18]. This paper asks whether a compatibility percentage, checked against the people it is meant to inform, shows the same gap.

Method

Recruitment and sample

Participants were recruited over a five-month window via a mixed snowball and paid-panel strategy, screened for having used a single, widely used compatibility-scoring dating application within the preceding twelve months and for being willing to discuss one specific match in detail. Fifty-eight participants completed a semi-structured interview averaging 38 minutes, spanning a broad age and relationship-goal range; recruitment continued until two independent coders agreed no interview in the most recent block introduced a new primary-reason category.

Interview protocol and trust-threshold coding

Each interview asked the participant to narrate one match — the one they remembered most clearly — message by message, from the initial percentage they were shown through to the point, if any, at which the conversation ended or settled into an ongoing exchange. Participants were asked directly whether and when the displayed percentage stopped informing how they read the match, without being shown their own score history in advance. Two coders independently placed each narrated account's trust-threshold event on the conversation's message count and classified its primary stated reason

into one of five categories, following standard thematic-analysis procedure [19] and reconciled after a pilot pass: chat felt mismatched despite the score, profile read as curated or inflated, already in the habit of ignoring the score, a low-scoring chat simply clicked, and no clear reason given.

Match-history subsample and half-life measure

Forty-one of the 58 participants additionally consented to share an anonymised export of their match history, covering 1,842 conversations across both high-scoring (displayed compatibility $\geq 90\%$) and mid-scoring (70–89%) matches. For each conversation we define its half-life H_i as the message count at which one side's next reply arrived more than seven days after the other's last message, or never arrived at all:

$$H_i = \min\{t : \text{gap}(t, t+1) > 7 \text{ days}\}$$

We compare H_i between the two score bands directly, then fit a minimal linear model of H_i against score band, message-response latency, and shared-interest-tag overlap, refitting with score band removed to test how much of the model's fit the score band itself was carrying.

Results




Figure 1: Cumulative share of participants reporting the compatibility percentage had stopped functioning as a live signal in their focal match, by messages exchanged ($n = 58$). Half cross the threshold by message eleven.

The coded interviews (Figure 1) place the median trust-threshold event at eleven exchanged

messages; by message twenty, 81% of participants report the percentage no longer informed how they read the match, and the share continues to climb only slowly thereafter.




Figure 2: Primary stated reason for the shift, split by whether the focal match was high-scoring ($\geq 90\%$) or mid-scoring (70–89%) ($n = 58$). A chat that felt mismatched despite the score dominates the high-scoring group; a low-scoring chat that simply clicked is concentrated in the other.

The reasons behind the threshold (Figure 2) split unevenly by band: participants describing a high-scoring match most often cited a chat that felt mismatched despite the number, while participants describing a mid-scoring match more often cited a conversation that clicked regardless of a lower number. Both groups cited a profile reading as curated or inflated at broadly similar rates.




Figure 3: Conversational half-life (messages before one side stops replying) by score band, from the consenting subsample's match-history exports ($n = 41$ participants, 1,842 conversations). The distributions the score is meant to separate overlap almost completely.

The independent check (Figure 3) tells a consistent story from outside the interview data: median half-life is 13 messages for high-scoring matches and 12 for mid-scoring matches, with overlapping interquartile ranges and no participant's high-scoring matches surviving markedly longer than their mid-scoring ones. A minimal model of half-life against score band,

response latency, and shared-interest-tag overlap explains a modest share of variance; removing score band and refitting on the remaining two predictors costs the model 0.4 percentage points of R^2 (n.s.), consistent with a score that is displayed at first contact but not, on this evidence, doing detectable work in how long the resulting conversation lasts.

The subsample's exports also carry a finding the interviews alone could not: of the high-scoring matches a participant explicitly described as one they had privately written off, the median match persisted near the top of that participant's recommended queue for a further nine weeks, the displayed percentage unchanged by a conversation that never resumed.

Discussion

Two independent measures — what participants say moved them and how long their conversations actually lasted — agree that score band is not where the explanatory weight sits once a chat has actually started. The percentage does real work before a message is sent, shaping which profile gets opened at all in a manner consistent with prior attractiveness- and self-presentation research [5], [6]; what this study adds is that its influence appears to end close to there. Participants did not report abandoning the score through complaint or protest — no participant in this sample said they had told the application anything — but through a private recalibration the score itself has no channel to register.

The persistence finding sits closest to this School's existing interest in what a live measurement does once the people it was built to inform have stopped consulting it: a rotation, a compliance score, and now a compatibility percentage each continue to be computed, displayed, and acted on by the system that produces them, on a schedule that owes nothing to whether anyone downstream is still reading [17], [18]. The number does not need anyone to believe it in order to keep appearing at the top of a queue.

Limitations

The sample draws on a single widely used application and participants willing to narrate one match in detail, which may skew toward daters confident enough to articulate a reason for disengaging; the trust threshold's consistency

across every recruitment channel, however, suggests the median is not an artefact of who volunteered to describe it. The half-life comparison covers 41 of the 58 participants' exports, not the full sample, and cannot rule out a score effect too small for 1,842 conversations to detect, though its agreement with the excluded participants' own coded threshold is itself evidence against a meaningfully larger one. Retrospective narration may also compress a threshold that felt more gradual at the time — a cost we consider worth asking the question in words at all.

Conclusion

A compatibility percentage can be displayed at first contact, computed continuously, and never revised, without functioning as a live signal for more than the first few exchanged messages of the conversation it introduces. This paper finds a median trust threshold at eleven messages, a conversational half-life indistinguishable between matches scored above and below ninety percent, and matches participants had already written off persisting near the top of a recommended queue for a further nine weeks. Future work might interview the people who tune the scoring model itself about what, if anything, a threshold like this one would change about how the percentage is built — a question this study, built entirely from the user's side of the exchange, could not put to the people answering it.

References

- [1] E. J. Finkel, P. W. Eastwick, B. R. Karney, H. T. Reis, and S. Sprecher, "Online Dating: A Critical Analysis From the Perspective of Psychological Science," *Psychological Science in the Public Interest*, vol. 13, no. 1, pp. 3–66, 2012, doi: [10.1177/1529100612436522](https://doi.org/10.1177/1529100612436522).
- [2] J. B. Walther, "Computer-Mediated Communication: Impersonal, Interpersonal, and Hyperpersonal Interaction," *Communication Research*, vol. 23, no. 1, pp. 3–43, 1996, doi: [10.1177/009365096023001001](https://doi.org/10.1177/009365096023001001).
- [3] J. M. Logg, J. A. Minson, and D. A. Moore, "Algorithm Appreciation: People Prefer Algorithmic to Human Judgment," *Organizational Behavior and Human Decision*

- Processes, vol. 151, pp. 90–103, 2019, doi: [10.1016/j.obhdp.2018.12.005](https://doi.org/10.1016/j.obhdp.2018.12.005).
- [4] B. J. Dietvorst, J. P. Simmons, and C. Massey, “Algorithm Aversion: People Erroneously Avoid Algorithms After Seeing Them Err,” *Journal of Experimental Psychology: General*, vol. 144, no. 1, pp. 114–126, 2015, doi: [10.1037/xge0000033](https://doi.org/10.1037/xge0000033).
 - [5] A. T. Fiore, L. S. Taylor, G. A. Mendelsohn, and M. Hearst, “Assessing Attractiveness in Online Dating Profiles,” in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 2008, pp. 797–806. doi: [10.1145/1357054.1357181](https://doi.org/10.1145/1357054.1357181).
 - [6] J. T. Hancock and C. L. Toma, “Putting Your Best Face Forward: The Accuracy of Online Dating Photographs,” *Journal of Communication*, vol. 59, no. 2, pp. 367–386, 2009, doi: [10.1111/j.1460-2466.2009.01420.x](https://doi.org/10.1111/j.1460-2466.2009.01420.x).
 - [7] C. L. Toma and J. T. Hancock, “Looks and Lies: The Role of Physical Attractiveness in Online Dating Self-Presentation and Deception,” *Communication Research*, vol. 37, no. 3, pp. 335–351, 2010, doi: [10.1177/0093650209356437](https://doi.org/10.1177/0093650209356437).
 - [8] N. Ellison, R. Heino, and J. Gibbs, “Managing Impressions Online: Self-Presentation Processes in the Online Dating Environment,” *Journal of Computer-Mediated Communication*, vol. 11, no. 2, pp. 415–441, 2006, doi: [10.1111/j.1083-6101.2006.00020.x](https://doi.org/10.1111/j.1083-6101.2006.00020.x).
 - [9] S. R. Sumter, L. Vandenbosch, and L. Ligtenberg, “Love Me Tinder: Untangling Emerging Adults’ Motivations for Using the Dating Application Tinder,” *Telematics and Informatics*, vol. 34, no. 1, pp. 67–78, 2017, doi: [10.1016/j.tele.2016.04.009](https://doi.org/10.1016/j.tele.2016.04.009).
 - [10] G. Ranzini and C. Lutz, “Love at First Swipe? Explaining Tinder Self-Presentation and Motives,” *Mobile Media & Communication*, vol. 5, no. 1, pp. 80–101, 2017, doi: [10.1177/2050157916664559](https://doi.org/10.1177/2050157916664559).
 - [11] G. David and C. Cambre, “Screened Intimacies: Tinder and the Swipe Logic,” *Social Media + Society*, vol. 2, no. 2, 2016, doi: [10.1177/2056305116641976](https://doi.org/10.1177/2056305116641976).
 - [12] Y. Zhang, Q. V. Liao, and R. K. E. Bellamy, “Effect of Confidence and Explanation on Accuracy and Trust Calibration in AI-Assisted Decision Making,” in *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 2020, pp. 295–305. doi: [10.1145/3351095.3372852](https://doi.org/10.1145/3351095.3372852).
 - [13] M. Yeomans, A. Shah, S. Mullainathan, and J. Kleinberg, “Making Sense of Recommendations,” *Journal of Behavioral Decision Making*, vol. 32, no. 4, pp. 403–414, 2019, doi: [10.1002/bdm.2118](https://doi.org/10.1002/bdm.2118).
 - [14] J. D. Lee and K. A. See, “Trust in Automation: Designing for Appropriate Reliance,” *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004, doi: [10.1518/hfes.46.1.50_30392](https://doi.org/10.1518/hfes.46.1.50_30392).
 - [15] M. Eslami et al., “First I “Like” It, Then I Hide It: Folk Theories of Social Feeds,” in *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*, 2016, pp. 2371–2382. doi: [10.1145/2858036.2858494](https://doi.org/10.1145/2858036.2858494).
 - [16] M. A. DeVito, D. Gergle, and J. Birnholtz, ““Algorithms Ruin Everything”: #RIPTwitter, Folk Theories, and Resistance to Algorithmic Change in Social Media,” in *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, 2017, pp. 3163–3174. doi: [10.1145/3025453.3025659](https://doi.org/10.1145/3025453.3025659).
 - [17] R. Sabel, M. Osayande, and J. Nwosu, “Same Roster, Four Different Verdicts: Test-retest Instability in a Casual-Shift Rostering Platform’s Fairness-Compliance Score, and the Managers Who Stop Reading the Roster Once They Have One,” *Slop University technical report*, 2026, doi: [10.5555/slop.o8swbm](https://doi.org/10.5555/slop.o8swbm).
 - [18] M. Osayande and R. Oyelaran, “Whoever Was Mid-Load: A Share-House Power-Board’s Fairness Rotation Rarely Decides Which Appliance an Energy Budget Cuts,” *Slop University technical report*, 2026, doi: [10.5555/slop.ws53l7](https://doi.org/10.5555/slop.ws53l7).
 - [19] V. Braun and V. Clarke, “Using Thematic Analysis in Psychology,” *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: [10.1191/1478088706qp063oa](https://doi.org/10.1191/1478088706qp063oa).