



Cleared by the Volunteer, Not the Rubric

Dagny Okafor, Mirela Hanke, Solveig Adeyemi

School of Continuous Improvement, Slop University

doi:10.5555/slop.t64dph

Abstract. A public library’s summer reading challenge asks children to log a short report on each book they finish, in exchange for entry into a season-end prize draw; an authorship-check score, built from four stylometric features, screens each entry for signs of adult authorship before it counts. When the score flags an entry, a volunteer reviewer decides whether it stands. We interviewed 22 volunteer reviewers across nine branches of a regional library network, walking each through their five most recently reviewed flagged entries (108 cases), and linked their accounts to the platform’s own decision log across a ten-week season (3,412 entries, 612 flagged, 71.2% overturned). Reviewers rarely invoked the score’s own features when explaining an override; instead, 44% of overturned cases cited consistency with a sibling’s earlier entries, 27% personal familiarity with the family, 19% a seasonal or local detail, and 10% a handwriting or photograph match. Override probability was effectively unrelated to the score’s own magnitude ($r = -0.09$, $p = .34$, $n = 612$), and a retrospective ablation removing one of the score’s four features from the season’s log changed almost nothing about which entries were flagged in the first place ($\kappa = 0.97$ to 0.966 ; n.s.). The score appears to route attention; the tacit cue decides.

Introduction

An authorship-check score is a fairly recent addition to a fairly old institution: the library summer reading challenge, in which a child logs a short report on each book finished in exchange for entry into a season-end prize draw. As challenge platforms have moved online, so have the incentives to game them, and several networks now screen each submitted entry with a stylometric score built to flag text judged too polished, too adult, or too consistent with a parent’s known writing to have been produced by the child whose name is on it. A flagged entry does not lose its place automatically; it is routed to a volunteer reviewer, who decides whether it stands.

This arrangement sits inside a broader pattern in human-automation research. People do not simply defer to an algorithmic judgment: they are especially quick to abandon one after watching it err [1], more willing to keep using an imperfect one when they retain some ability to adjust it [2], and, in other settings, inclined to prefer it to a person’s judgment outright [3]. Automation-

bias research finds the override decision can go the other way, with people over-trusting an automated aid past the point its own reliability warrants [4]. Public-sector research on “street-level” discretion documents a parallel tension in front-line human review of algorithmic outputs generally [5], [6], and clinical decision-support research has quantified the scale of the resulting override behaviour directly, with alert override rates reported between 49% and 96% across settings [7]. What has not, to our knowledge, been examined directly is what a reviewer is actually doing when they override a flag in a low-stakes, child-facing setting with no institutional incentive riding on the outcome beyond a modest prize draw.

We interviewed 22 volunteer reviewers across a nine-branch regional public library network, walking each through their five most recent flagged-entry decisions, and linked their accounts to the platform’s own season-long decision log. Our contributions:

- We report what we believe is the first interview-based account of how volunteer review-

ers override an automated authorship-check flag in a setting built around children’s own writing.

- We show that override decisions track cues largely absent from the score’s own feature set more than they track the score’s own magnitude.
- We report a retrospective ablation in which removing one of the score’s four features changes almost nothing about which entries are flagged in the first place, consistent with the tacit cues carrying more of the real signal than the score’s own mechanics do.

Related work

Volunteer review of an automated flag is one instance of a wider literature on discretion inside a public-facing role. Third-party volunteer gatekeepers in nonprofit settings are already known to apply criteria beyond an organisation’s stated policy when deciding who is included [8], and municipal “screen-level” bureaucrats exercise comparable discretion over the data an algorithmic system is given to work with in the first place [6]. Measurement systems, once adopted, also reshape the behaviour of the people and institutions they measure: public rankings have been shown to alter an organisation’s practice independently of any change in the underlying quality being ranked [9], an effect that persists once a ranking regime becomes routine [10]. Prior work at this University found a comparable pattern at a meeting-room scale, where staff began deliberately marking up whiteboards to protect a room’s photo-audit engagement score, a behaviour only weakly related to whether the room’s meetings actually reached a decision [11], and, in a suburban gym chain’s booking-reliability score, cancellations clustered at the exact edge of a penalty cutoff without the chain’s true attendance rate moving at all [12].

The score itself sits in the automated authorship-detection literature that has grown around adult academic writing: detection tools vary widely in accuracy and are readily defeated by light editing [13], current commercial detectors remain inconsistent across writing contexts [14], and their error rates fall unevenly across writer populations [15]. The stylometric features such tools draw on can distinguish short human- and machine-written samples with reasonable reliability [16], though a text’s shortness measurably weakens exactly this kind of feature’s discrimi-

native power [17] — a property that, at children’s-logbook length, the present score inherits.

Method

We partnered with a nine-branch regional public library network running a ten-week summer reading challenge, in which children log a short report (typically 40–120 words) on each book finished, with completed entries entering a season-end prize draw. Every submitted entry passes through the platform’s built-in authorship-check score before it counts:

$$S = w_{\text{adj}}f_{\text{adj}} + w_{\text{sent}}f_{\text{sent}} + w_{\text{tense}}f_{\text{tense}} + w_{\text{lex}}f_{\text{lex}}$$

where f_{adj} , f_{sent} , f_{tense} , and f_{lex} are normalised adjective-density, sentence-length-variance, present-tense-ratio, and lexical-rarity features, weighted $w = (18, 22, 15, 45)$ and summed to a 0–100 score; an entry scoring at or above the platform’s threshold of $\tau = 50$ is routed to a volunteer reviewer before it counts toward the draw.

Across the season, 3,412 entries were submitted and 612 (17.9%) were flagged. We recruited 22 of the network’s 40 active flagged-entry reviewers (opt-in, contacted near the end of the season) for a semi-structured interview, in which each reviewer walked through their five most recent flagged-entry decisions in a think-aloud format (108 cases; two reviewers had reviewed only four recent flags each). Two coders, blind to the platform’s own override outcome for each case, independently applied Braun and Clarke’s six-phase thematic-analysis procedure [18], [19] to the interview transcripts, coding each case for the cue or cues the reviewer cited in reaching their decision, with inter-rater agreement of 87% on the assigned code set (weighted $\kappa = 0.81$).

We then linked every interview-discussed case, and the full 612-entry flagged set, to the platform’s own decision log — the entry’s score at the moment of flagging and the reviewer’s final decision (upheld or overturned) — to test whether override probability tracked the score’s own magnitude.

Finally, to isolate whether any one of the score’s four features was doing the work the reviewers’ accounts suggested it was not, we retrospectively removed the adjective-density feature ($w_{\text{adj}} = 0$, remaining weights renormalised) and rescored the full season’s 3,412 entries, compar-

ing the resulting flag set against the original on a week-by-branch basis; this ablation tests whether a feature can be dropped from the rubric without moving which entries a reviewer ever sees.

Results

Season-wide, reviewers overturned 436 of the 612 flagged entries (71.2%); within the subset of 108 cases discussed in interview, the rate was closely consistent at 74.1% (80 cases; below), suggesting the interview sample is not unusually lenient or strict relative to the season as a whole.

Metric	Value
Entries submitted	3,412
Entries flagged ($S \geq 50$)	612 (17.9%)
Flagged entries overturned, season log	436 (71.2%)
Interview-discussed cases	108
Interview cases overturned	80 (74.1%)

The reasons reviewers gave for an override rarely touched the score's own feature set. Consistency with an entry a sibling had submitted earlier in the season was cited in 44% of overturned cases, personal familiarity with the family in 27%, a seasonal or locally specific detail (a named school event, a particular week's weather) in 19%, and a match between the submitted photograph and a handwriting sample the reviewer recalled from an earlier entry in 10% (Figure 1); cases frequently cited more than one cue, so the shares do not sum to 100%.




Figure 1: Share of overturned flagged entries ($n = 80$) in which the reviewer's stated rationale cited each cue — Sibling: consistency with a sibling's entry; Family: personal familiarity; Season: a seasonal or local detail; Photo: a handwriting or photograph match. Entries could cite more than one.

Whatever the score itself measured did little to predict the outcome once an entry was flagged. Across the full 612-entry flagged set, the score's magnitude at the moment of flagging was effectively uncorrelated with whether the flag was later overturned ($r = -0.09$, $p = .34$; Figure 2): an entry that scored a marginal 52 was about as likely to be cleared as one that scored a near-maximal 96.




Figure 2: Authorship-check score at flag against reviewer decision ($n = 612$ flagged entries); the fitted line's slope is statistically indistinguishable from flat.

The retrospective ablation reinforced the same picture from a different angle. Removing the adjective-density feature and rescored the season's 3,412 entries left the flag set nearly unchanged: concordance between the original and ablated flag sets was $\kappa = 0.97$ overall, and the weekly change in each branch's flag rate clustered tightly around zero regardless of branch size (mean $\Delta = 0.03$ percentage points, smaller branches; $\Delta = 0.04$, larger branches; $t(18) = 0.09$, $p = .93$; Figure 3). Removing a quarter of the score's rubric changed almost nothing about which entries a reviewer ever saw; we retain the feature for completeness.




Figure 3: Weekly change in branch flag rate when the adjective-density feature is dropped from the score and the season is rescored, by branch size ($n = 10$ weeks per group); both distributions sit near zero.

Taken together, the score’s role in this pipeline looks closer to a routing function than a decision function: it reliably directs attention toward entries a reviewer will go on to examine, but it is the reviewer’s own, largely uncoded, sense of a family’s history that settles what happens next.

Discussion

The picture across all three analyses is consistent: the authorship-check score determines which entries a reviewer sees, but — once an entry is in front of a reviewer — the score’s own magnitude and its own feature set explain very little of what happens to it. The cues that do the deciding (a sibling’s earlier entry, a name a reviewer already knows, a detail that could only have been true that particular week) sit entirely outside the four features the score was built from, which may explain why removing one of those features costs the rubric so little: the rubric’s job, on this evidence, is to open the queue rather than to close the case.

This reading is broadly consistent with automation-bias and algorithm-aversion research showing that people readily depart from an algorithmic recommendation once they have grounds of their own to trust [1], [3], and with the pattern this University’s own audit work keeps finding at other scales, where a measurement system’s introduction reorganises behaviour around the measurement without the underlying quantity moving in step [11], [12]. Whether formalising a “known family” flag into the rubric itself would improve or simply relocate this dy-

namic is a natural next question, one we return to below.

Limitations

Our interview sample draws on reviewers’ own retrospective accounts of a decision already made, which may not fully capture what drove the decision at the time; the close correspondence between the interview subsample’s override rate and the season-wide rate (74.1% versus 71.2%) suggests the sample is not unusually selective, though it cannot rule out reconstructed reasoning. The stylistic features underlying the score are drawn from a literature that already reports weaker discriminative power in texts as short as a child’s logbook entry [17], consistent with, rather than a challenge to, our central finding that the tacit cues carry more of the real signal. Our sample is a single regional network across one season; a network with a stricter prize (a place, not a book voucher) may sharpen the incentive to flag, and the reviewer discretion, differently.

Conclusion

We interviewed the volunteer reviewers who stand between a library summer reading challenge’s authorship-check score and the child whose entry it flags, and found that the score’s own magnitude and feature set explain little of what a reviewer goes on to decide. Overrides tracked cues — a sibling’s earlier entry, a family a reviewer already knew, a detail only that week could have supplied — that never entered the rubric, and removing one of the rubric’s four features from a season-long retrospective rescoring left the flag set almost untouched. Future work should test whether formalising the “known family” cue into the score’s own feature set would improve override consistency across reviewers or simply give the tacit judgement a line item to hide behind.

References

- [1] B. J. Dietvorst, J. P. Simmons, and C. Massey, “Algorithm aversion: People erroneously avoid algorithms after seeing them err,” *Journal of Experimental Psychology: General*, vol. 144, no. 1, pp. 114–126, 2015, doi: [10.1037/xge0000033](https://doi.org/10.1037/xge0000033).
- [2] B. J. Dietvorst, J. P. Simmons, and C. Massey, “Overcoming Algorithm Aversion: People

- Will Use Imperfect Algorithms If They Can (Even Slightly) Modify Them,” *Management Science*, vol. 64, no. 3, pp. 1155–1170, 2018, doi: [10.1287/mnsc.2016.2643](https://doi.org/10.1287/mnsc.2016.2643).
- [3] J. M. Logg, J. A. Minson, and D. A. Moore, “Algorithm appreciation: People prefer algorithmic to human judgment,” *Organizational Behavior and Human Decision Processes*, vol. 151, pp. 90–103, 2019, doi: [10.1016/j.obhdp.2018.12.005](https://doi.org/10.1016/j.obhdp.2018.12.005).
- [4] L. J. Skitka, K. L. Mosier, and M. Burdick, “Does automation bias decision-making?,” *International Journal of Human-Computer Studies*, vol. 51, no. 5, pp. 991–1006, 1999, doi: [10.1006/ijhc.1999.0252](https://doi.org/10.1006/ijhc.1999.0252).
- [5] E. W. Northcott, “Automating street-level discretion: A systematic literature review and research agenda,” *Policy Studies Journal*, vol. 54, no. 3, 2026, doi: [10.1111/psj.70024](https://doi.org/10.1111/psj.70024).
- [6] E. Gordon, J. Harlow, S. A. Whitman, and M. Lee, “Data Discretion: Screen-Level Bureaucrats and Municipal Decision-Making,” *Digital Government: Research and Practice*, vol. 5, no. 2, pp. 1–14, 2024, doi: [10.1145/3652950](https://doi.org/10.1145/3652950).
- [7] C. E. Ray et al., “Alert fatigue measurement in clinical decision support: a systematic review,” *Journal of the American Medical Informatics Association*, 2026, doi: [10.1093/jamia/ocag064](https://doi.org/10.1093/jamia/ocag064).
- [8] P. S. M. van Overbeeke, S. A. Koolen-Maas, L. C. P. M. Meijs, and J. L. Brudney, “You Shall (Not) Pass: Strategies for Third-Party Gatekeepers to Enhance Volunteer Inclusion,” *Voluntas: International Journal of Voluntary and Nonprofit Organizations*, vol. 33, no. 1, pp. 33–45, 2022, doi: [10.1007/s11266-021-00384-0](https://doi.org/10.1007/s11266-021-00384-0).
- [9] W. N. Espeland and M. Sauder, “Rankings and Reactivity: How Public Measures Recreate Social Worlds,” *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: [10.1086/517897](https://doi.org/10.1086/517897).
- [10] M. Sauder and W. N. Espeland, “The Discipline of Rankings: Tight Coupling and Organizational Change,” *American Sociological Review*, vol. 74, no. 1, pp. 63–82, 2009, doi: [10.1177/000312240907400104](https://doi.org/10.1177/000312240907400104).
- [11] L. Achterberg and M. Hanke, “Marked for the Camera: A Staff Survey on Deliberate Whiteboard Use Under a Meeting-Room Engagement Photo-Audit,” *Sloped University technical report*, 2026, doi: [10.5555/sloped.yjud2l](https://doi.org/10.5555/sloped.yjud2l).
- [12] C. Beng, M. Hanke, and S. Adeyemi, “The Number Stayed Clean: A Quasi-Experiment on a Suburban Gym Chain’s Booking-Reliability Score, Comparing Attendance and Cancellation Timing Before and After Its Introduction,” *Sloped University technical report*, 2026, doi: [10.5555/sloped.ln1awt](https://doi.org/10.5555/sloped.ln1awt).
- [13] D. Weber-Wulff et al., “Testing of detection tools for AI-generated text,” *International Journal for Educational Integrity*, vol. 19, no. 1, 2023, doi: [10.1007/s40979-023-00146-z](https://doi.org/10.1007/s40979-023-00146-z).
- [14] M. Hadra, K. Cambridge, and M. Mesbah, “Evaluating the accuracy and reliability of AI content detectors in academic contexts,” *International Journal for Educational Integrity*, vol. 22, no. 1, 2026, doi: [10.1007/s40979-026-00213-1](https://doi.org/10.1007/s40979-026-00213-1).
- [15] A. R. Pratama, “The accuracy-bias trade-offs in AI text detection tools and their impact on fairness in scholarly publication,” *PeerJ Computer Science*, vol. 11, p. e2953, 2025, doi: [10.7717/peerj-cs.2953](https://doi.org/10.7717/peerj-cs.2953).
- [16] K. Przystalski, J. K. Argasiński, I. Grab-ska-Gradzińska, and J. K. Ochab, “Stylometry recognizes human and LLM-generated texts in short samples,” *Expert Systems with Applications*, vol. 296, p. 129001, 2026, doi: [10.1016/j.eswa.2025.129001](https://doi.org/10.1016/j.eswa.2025.129001).
- [17] F. López-Escobedo, C.-F. Méndez-Cruz, G. Sierra, and J. Solórzano-Soto, “Analysis of Stylometric Variables in Long and Short Texts,” *Procedia - Social and Behavioral Sciences*, vol. 95, pp. 604–611, 2013, doi: [10.1016/j.sbspro.2013.10.688](https://doi.org/10.1016/j.sbspro.2013.10.688).
- [18] V. Braun and V. Clarke, “Using thematic analysis in psychology,” *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: [10.1191/1478088706qp0630a](https://doi.org/10.1191/1478088706qp0630a).
- [19] D. Byrne, “A worked example of Braun and Clarke’s approach to reflexive thematic analysis,” *Quality & Quantity*, vol. 56, no. 3, pp. 1391–1412, 2022, doi: [10.1007/s11135-021-01182-y](https://doi.org/10.1007/s11135-021-01182-y).