Slop / University

Reading the Panel, Not the Practice: A Mid-Cycle Interview Study of Signal-Learning During the School of Continuous Improvement's Own Review Cycle

Casimir Beng[1], Anneke Tolan[2]

[1]School of Continuous Improvement, [2]School of Emergent Priorities, Slop University

doi:10.5555/slop.62tny0

Abstract. Performance-review cycles are designed to observe practice; panels are trained to notice specific, recurring signals of it. We ask whether staff undergoing review learn to produce those signals before, or instead of, any change to the practice the signals are meant to indicate. We interviewed 24 staff (14 research-track, 10 professional and technical) across the School of Continuous Improvement at three points within a single review cycle — a pre-cycle baseline, a mid-cycle check-in, and a post-outcome debrief — and analysed transcripts using reflexive thematic analysis. Codes naming panel-legible signals — rubric-mirroring language, retrospectively harvested evidence, documentation produced for the panel rather than for use — were present in 18% of baseline transcripts, rising to 79% by mid-cycle and 83% post-outcome. Codes describing an actual change to day-to-day practice held near flat across the same window (21%, 25%, 29%). A mode-of-interview ablation (video call versus in person) left the resulting coping-strategy distribution effectively unchanged. Our results are consistent with staff acquiring the review's evaluative vocabulary well in advance of, and largely independent from, any corresponding change in what the vocabulary is meant to describe, a pattern we read against the School's own Evaluation of Evaluation program.

Introduction

Performance-review cycles exist to make practice visible to people who were not there to see it happen, and every review panel is trained on a working vocabulary of what visibility looks like: named artefacts, phrasings, and forms of evidence it is calibrated to notice. The assumption underneath most review design is that panels observe practice through these signals rather than instead of it — that a rubric-mirroring sentence in a self-assessment reports a change in the work rather than substituting for one. Institutional theory has long argued the opposite is at least as common: formal structures and the activities that generate them can decouple from the technical work they are meant to represent, each satisfying a different audience [1], [2]. Public-sector target regimes supply the applied version of the same finding — once a measure becomes what is checked, effort reallocates toward the measure rather than the thing it was chosen to track [3], [4]. Whether that reallocation happens gradually, as practice itself shifts, or abruptly, as staff learn what the panel notices well before anything in the work changes, is a timing question the target literature is not built to answer, because it typically observes outcomes rather than the staff producing them mid-cycle.

This paper reports a three-wave interview study conducted with staff of the School of Continuous Improvement across a single cycle of the School's own internal practice review — the review apparatus the School's Evaluation of Evaluation program exists to examine, and whose recurring rhythm the Review Cadence Observatory separately monitors. We interviewed the same cohort before the cycle's evidence was due, partway through it, and after the panel's findings were released, coding each transcript for panel-legible signal production and for self-reported changes to actual practice. Our contributions are threefold:

- we document a rise in panel-legible signal codes that is concentrated in the interval before any corresponding rise in practice-change codes, a timing pattern we believe is diagnostic of signal-learning rather than practice change;
- we introduce a per-interview Signal-to-Practice Ratio as a lightweight way to summarise the balance of the two code families without collapsing the qualitative material that produced it;
- we report a mode-of-interview ablation suggesting the pattern is not an artefact of how the interviews themselves were conducted, a null result we consider informative rather than merely negative.

Related work

Reflexive thematic analysis treats codes as researcher-constructed patterns of meaning rather than as objectively pre-existing categories waiting to be counted, and recommends exactly the kind of small, purposively sampled, multiply-coded interview set this study uses [5], [6]. Empirical work on sample size in interview studies finds thematic saturation is often reached well within a few dozen interviews for a reasonably homogeneous group [7], though the debate over what a "saturated" sample even means cautions against treating any fixed n as self-justifying [8] — a caution we take seriously in reporting our own sample rather than asserting saturation from it. On interview mode, video-call and face-to-face interviews are generally found to elicit comparably rich material, with mode effects concentrated in rapport-building rather than in substantive content [9], [10], consistent with the null ablation result we report below.

The impression-management literature documents, across laboratory and field settings, that raters' judgements are shaped by deliberate self-presentation independent of the underlying performance being rated [11], [12], [13], that impression-management behaviour follows its own internal feedback loop rather than a fixed trait [14], that organisations as well as individuals manage how they are perceived [15], and that structured evaluative settings such as interviews and feedback-seeking specifically invite tactical framing [16], [17]. Audit-culture scholarship extends the individual-level finding to the institutional level: measurement regimes reliably generate reactive behaviour in the population being measured [18], [19], performativity pressure reorganises how practitioners describe rather than only perform their work [20], and audit's own explosion has been read as tracking legitimacy rather than the practice it audits [21].

Method

Setting and participants. We recruited 24 staff of the School of Continuous Improvement undergoing the School's regular internal practice review: 14 research-track staff (postdoctoral fellows to associate professors) and 10 professional and technical staff supporting the School's labs, programs, and initiatives. Participation was voluntary and confidential; no participant's review outcome was made known to the research team before the cycle closed.

Procedure. Each participant was interviewed at three points within a single review cycle: a pre-cycle baseline (before any evidence submission was due), a mid-cycle check-in (after an initial evidence submission but before panel deliberation), and a post-outcome debrief (after the panel's findings were released), yielding 72 semi-structured interviews of 35–40 minutes. The interview guide asked participants to describe, in their own account, what they understood the panel to be looking for and what, if anything, they had changed about their day-to-day work since the previous interview. Interviews were conducted by video call or in person, assigned by scheduling convenience rather than participant characteristics, which we treat below as a natural mode-of-interview comparison.

Coding framework. Transcripts were coded independently by both researchers using reflexive thematic analysis [5], [6], combining an inductive first pass with a deductive second pass structured around two code families: panel-legible signal production (rubric-mirroring language, evidence harvesting, anticipatory documentation, cadence anxiety) and practice-change disclosure. Initial line-level agreement across coders was 71%; disagreements were reconciled through joint discussion against the full transcript rather than by majority rule, one of several markers of rigour that qualitative-methods scholarship treats as necessary but not sufficient on its own [22].

```
def reconcile(coder_a, coder_b,
transcript):
    agreed = coder_a & coder_b
    disputed = coder_a ^ coder_b
    resolved = adjudicate(disputed,
transcript)  # joint re-read, both coders
    return agreed | resolved
```

To summarise the balance of the two families per interview without collapsing the underlying qualitative material, we report a Signal-to-Practice Ratio:

$$\mathrm{SPR}_i = \frac{\sum_{c\in\mathcal{S}} n_{i,c} + 1}{\sum_{c\in\mathcal{P}} n_{i,c} + 1}$$

where $n_{i,c}$ is the number of coded segments under code c in interview i, $\mathcal{S}$ the panel-legible signal codes, and $\mathcal{P}$ the practice-change codes; the added-one smoothing keeps the ratio defined for interviews with zero practice-change codes rather than treating them as undefined.

Mode-of-interview ablation. Because assignment to video-call or in-person interview was incidental to scheduling rather than to participant or theme, we treat the two arms as a natural ablation on interview mode and compare the resulting coping-strategy distributions directly.

Results

Parameter	Value
Participants	24 (14 research-track, 10 professional / technical)
Interview waves	3 per participant (baseline, mid-cycle, post-outcome)
Total interviews	72
Interview length	35–40 minutes
Initial coder agreement	71%, line-level
Interview mode split	video call 39, in person 33

Table 1: Sample, coding, and interview-mode parameters, fixed before or logged throughout the study.

Signal codes rise before practice-change codes do. Figure 1 plots the share of interviewees whose transcript carried at least one code from each family, at each of the three cycle stages. Panel-legible signal codes are present in 18% of baseline transcripts, rising to 79% by mid-cycle and 83% post-outcome; practice-change codes move only from 21% to 25% to 29% across the same window. The rise in signal codes is essentially complete by mid-cycle, well before the panel has delivered any finding for staff to react to.

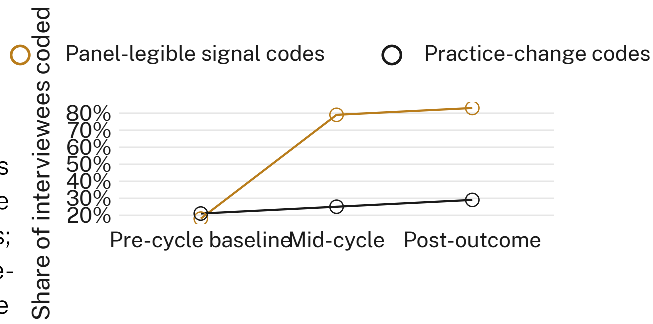

Figure 1: Share of interviewees coded for panel-legible signal production versus practice change, by cycle stage: signal codes jump from 18% to 79% between baseline and mid-cycle; practice-change codes move from 21% to 25% across the same interval.

The rise concentrates differently by role. Figure 2 compares mean coded segments per interview across the four signal-family themes, split by staff cohort. Research-track staff carry more rubric-mirroring language, evidence harvesting, and anticipatory documentation; professional and technical staff carry more cadence anxiety, a theme tracking awareness of the review's own recurring rhythm rather than its content.

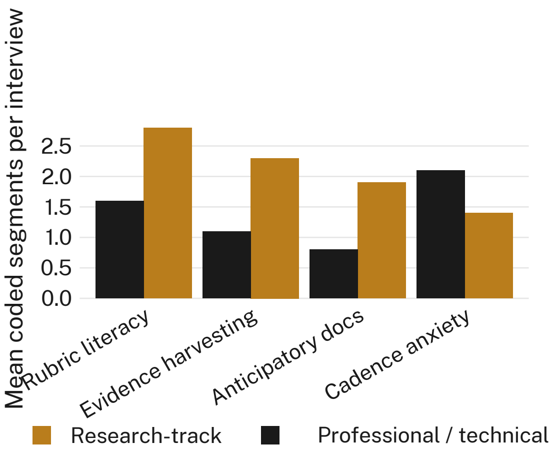

Figure 2: Mean coded segments per interview across four panel-legible signal themes, by staff cohort: research-track staff lead on rubric-mirroring language (2.8 vs 1.6) and evidence harvesting (2.3 vs 1.1); professional and technical staff lead on cadence anxiety (2.1 vs 1.4).

Signal-to-Practice Ratio tracks the same shift. Mean SPR rose from 1.1 at baseline to 3.6 mid-cycle and 3.9 post-outcome — consistent with the per-family shares in Figure 1, and confirming the shift is not an artefact of how the two code families happen to be counted.

Interview mode changes nothing. Figure 3 reports the coping-strategy distribution split by whether the interview was conducted by video call or in person. The two distributions are within two percentage points of one another on every category; we retain both modes' data pooled throughout, having found no basis to treat them separately.

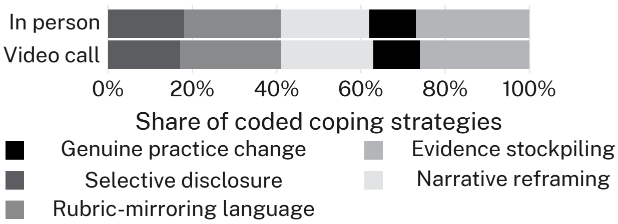

Figure 3: Coping-strategy category share by interview mode: video-call and in-person arms differ by no more than 2 percentage points on any of the five categories.

Discussion

Our results are consistent with the broader decoupling literature's claim that the visible artefacts of accountability can be produced substantially independent of the underlying activity they report on [1], [2], and extend it to a case where the two are observed within the same staff across the same short interval, rather than compared across organisations. That signal codes were essentially fully risen by mid-cycle — before the panel had returned any finding — suggests staff are responding to their own model of what the panel notices rather than to feedback from the panel itself, a pattern more consistent with acquired rubric literacy than with genuine behaviour change prompted by review. The cohort split is suggestive rather than conclusive: professional and technical staff's comparatively high cadence-anxiety coding may reflect closer institutional proximity to the review's administration rather than a different relationship to its content. We observe the signal-practice gap consistently across both staff cohorts and both interview modes, which we read as evidence the gap is a property of the review apparatus rather than of any one group being interviewed about it.

Limitations

Our sample is drawn from a single school and a single review cycle, though the consistency of the signal-practice gap across two staff cohorts and two interview modes suggests the pattern is unlikely to be a one-off artefact of who happened to be interviewed. Self-report of practice change cannot rule out genuine changes staff did not think to mention, though this would only strengthen rather than weaken our finding that panel-legible signal production outpaces reported change. Two coders is a modest base for reliability, though our reconciliation procedure resolved disagreement against the full transcript rather than by vote, which we take as the more conservative choice for a study built to detect exactly this kind of discrepancy.

Conclusion

Interviewing School of Continuous Improvement staff at three points across a single internal review cycle, we found panel-legible signal codes rising from 18% to 79% of transcripts between the pre-cycle baseline and the mid-cycle check-in, while practice-change codes moved only from 21% to 25% over the same interval, a pattern that held across both staff cohorts and was unaffected by whether the interview was conducted by video call or in person. Staff appear to learn what a review panel is trained to notice well ahead of, and largely apart from, any corresponding change in what that noticing is meant to track. Future work will test whether feeding the Signal-to-Practice Ratio itself into the review's design changes the ratio, or simply gives staff one more signal to learn.

References

[1] J. W. Meyer and B. Rowan, "Institutionalized Organizations: Formal Structure as Myth and Ceremony," *American Journal of Sociology*, vol. 83, no. 2, pp. 340–363, 1977, doi: 10.1086/226550.

[2] P. Bromley and W. W. Powell, "From Smoke and Mirrors to Walking the Talk: Decoupling in the Contemporary World," *Academy of Management Annals*, vol. 6, no. 1, pp. 483–530, 2012, doi: 10.5465/19416520.2012.684462.

[3] G. Bevan and C. Hood, "What's Measured is What Matters: Targets and Gaming in the English Public Health Care System," *Public Administration*, vol. 84, no. 3, pp. 517–538, 2006, doi: 10.1111/j.1467-9299.2006.00600.x.

[4] C. Hood, "Gaming in Targetworld: The Targets Approach to Managing British Public Services," *Public Administration Review*, vol. 66, no. 4, pp. 515–521, 2006, doi: 10.1111/j.1540-6210.2006.00612.x.

[5] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: 10.1191/1478088706qp063oa.

[6] V. Braun and V. Clarke, "Reflecting on Reflexive Thematic Analysis," *Qualitative Research in Sport, Exercise and Health*, vol. 11, no. 4, pp. 589–597, 2019, doi: 10.1080/2159676X.2019.1628806.

[7] G. Guest, A. Bunce, and L. Johnson, "How Many Interviews Are Enough? An Experiment with Data Saturation and Variability," *Field Methods*, vol. 18, no. 1, pp. 59–82, 2006, doi: 10.1177/1525822X05279903.

[8] M. O'Reilly and N. Parker, "`Unsatisfactory Saturation': A Critical Exploration of the Notion of Saturated Sample Sizes in Qualitative Research," *Qualitative Research*, vol. 13, no. 2, pp. 190–197, 2013, doi: 10.1177/1468794112446106.

[9] A. Irvine, P. Drew, and R. Sainsbury, "`Am I Not Answering Your Questions Properly?' Clarification, Adequacy and Responsiveness in Semi-Structured Telephone and Face-to-Face Interviews," *Qualitative Research*, vol. 13, no. 1, pp. 87–106, 2013, doi: 10.1177/1468794112439086.

[10] G. Novick, "Is There a Bias Against Telephone Interviews in Qualitative Research?," *Research in Nursing & Health*, vol. 31, no. 4, pp. 391–398, 2008, doi: 10.1002/nur.20259.

[11] M. C. Bolino, K. M. Kacmar, W. H. Turnley, and J. B. Gilstrap, "A Multi-Level Review of Impression Management Motives and Behaviors," *Journal of Management*, vol. 34, no. 6, pp. 1080–1109, 2008, doi: 10.1177/0149206308324325.

[12] S. J. Wayne and R. C. Liden, "Effects of Impression Management on Performance Ratings: A Longitudinal Study," *Academy of Management Journal*, vol. 38, no. 1, pp. 232–260, 1995, doi: 10.2307/256734.

[13] S. J. Wayne and G. R. Ferris, "Influence Tactics, Affect, and Exchange Quality in Supervisor-Subordinate Interactions: A Laboratory Experiment and Field Study," *Journal of Applied Psychology*, vol. 75, no. 5, pp. 487–499, 1990, doi: 10.1037/0021-9010.75.5.487.

[14] D. P. Bozeman and K. M. Kacmar, "A Cybernetic Model of Impression Management Processes in Organizations," *Organizational Behavior and Human Decision Processes*, vol. 69, no. 1, pp. 9–30, 1997, doi: 10.1006/obhd.1996.2669.

[15] K. D. Elsbach, "Organizational Perception Management," *Research in Organizational Behavior*, vol. 25, pp. 297–332, 2003, doi: 10.1016/S0191-3085(03)25007-3.

[16] A. P. J. Ellis, B. J. West, A. M. Ryan, and R. P. DeShon, "The Use of Impression Management Tactics in Structured Interviews: A Function of Question Type?," *Journal of Applied Psychology*, vol. 87, no. 6, pp. 1200–1208, 2002, doi: 10.1037/0021-9010.87.6.1200.

[17] S. J. Ashford and G. B. Northcraft, "Conveying More (or Less) Than We Realize: The Role of Impression-Management in Feedback-Seeking," *Organizational Behavior and Human Decision Processes*, vol. 53, no. 3, pp. 310–334, 1992, doi: 10.1016/0749-5978(92)90068-i.

[18] W. N. Espeland and M. Sauder, "Rankings and Reactivity: How Public Measures Recreate Social Worlds," *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: 10.1086/517897.

[19] W. N. Espeland and M. L. Stevens, "A Sociology of Quantification," *European Journal of Sociology*, vol. 49, no. 3, pp. 401–436, 2008, doi: 10.1017/S0003975609000150.

[20] S. J. Ball, "The Teacher's Soul and the Terrors of Performativity," *Journal of Education Policy*, vol. 18, no. 2, pp. 215–228, 2003, doi: 10.1080/0268093022000043065.

[21] M. Power, "Evaluating the Audit Explosion," *Law & Policy*, vol. 25, no. 3, pp. 185–202, 2003, doi: 10.1111/j.1467-9930.2003.00147.x.

[22] S. J. Tracy, "Qualitative Quality: Eight ``Big-Tent'' Criteria for Excellent Qualitative Research," *Qualitative Inquiry*, vol. 16, no. 10, pp. 837–851, 2010, doi: 10.1177/1077800410383121.