



## The Tap That Counts: An Interview Study of What Sixty-Four Residents Do When a Strata Intercom's Security-Compliance Notification Fires

Casimir Beng, Solveig Adeyemi, Mirela Hanke

School of Continuous Improvement, Slop University

doi:10.5555/slop.22tpz2

**Abstract.** Body-corporate security policies increasingly rely on a shared video intercom's "guest logged" push notification, sent to every resident whenever a visitor buzzes in, as the evidentiary backbone of an insurer-mandated security-compliance audit. Across the nine strata schemes in this study, the audit's own dashboard reports 97.3% of notifications acknowledged inside its five-minute window — a reading its administrators treat as near-total engagement. This paper reports a semi-structured interview study of 64 residents, coded into four notification-management practices, to establish what "acknowledged" actually captures. A reflexive swipe performed without registering the notification's content is the modal practice overall and rises from a minority behaviour among residents under a year in their unit to a majority one among residents of three years or more, displacing the two verification practices that dominate at shorter tenures. Dashboard-logged time-to-acknowledge falls over the same tenure gradient rather than rising, which the audit's own metric would read as improving compliance. A six-week trial adding a "please confirm you saw this" follow-up prompt to three of the nine buildings left the practice mix statistically indistinguishable from the six unprompted buildings. We suggest the acknowledgement rate and the resident's own account of their attention have quietly stopped describing the same behaviour, well before either the audit or the resident noticed the drift.

### Introduction

A shared video intercom is, for most strata residents, the building's only universally used piece of security infrastructure: every unit receives it, every visitor passes through it, and every buzz-in produces a timestamped log entry the strata's insurer can request. Over the past several years, insurer-driven security-compliance frameworks have converted that log into an audited metric — a "guest logged" notification pushed to every resident's phone when a visitor is admitted, with the fraction of notifications marked "acknowledged" inside a fixed window reported to the body corporate as evidence of an alert, security-conscious building. In the nine schemes surveyed for this paper, that metric sits at 97.3%, a figure strata managers and insurers alike treat as close to ceiling.

A dashboard reading close to ceiling can mean the audited population is doing exactly what the audit assumes, or it can mean the population has found the cheapest action that satisfies the audit's own definition of "acknowledged" — in this case, a single tap. Security and safety research elsewhere has repeatedly found that the second reading, not the first, is the one that holds once a population has had enough exposure to habituate to an alert [1], [2], [3], but nobody had asked the residents behind this particular metric which of the two was actually happening to them.

This paper reports an interview study of the crossing itself: 64 residents across nine strata schemes, asked in their own words what they do the moment the notification fires, with their accounts checked against the audit's own dash-

board export for the same buildings. The paper's contributions, each modest by design:

1. We develop and apply a four-category coding scheme for residents' notification-management practice and report its distribution by tenure in the building, a dimension the audit does not itself track.
2. We show that the audit's own time-to-acknowledge measure moves in the opposite direction to the practices the interviews describe, a movement the metric alone cannot distinguish from genuine improvement.
3. We show, via a six-week prompt trial in three of the nine buildings, that a direct nudge toward the audited behaviour does not measurably shift the underlying practice mix.

The audit is not shown to be dishonest, and this paper does not attempt to show that; it is shown to have converged, for a majority of longer-tenured residents, on a single physical gesture the resident no longer experiences as attention at all.

## Related work

Habituation to a repeated alert is well documented outside the strata context. Clinical alarm fatigue shows staff learning to disregard warnings whose base rate of true positives is low, regardless of how each alarm is nominally designed to demand attention [1]. Browser security-warning research finds active, interruptive warnings heeded at far higher rates than passive ones, but repeated exposure erodes even the interruptive kind [2], a finding mental-model interview work traces to users developing an internal shortcut for "warnings I can dismiss" that need not track the warning's actual content [3]. The underlying mechanism is general rather than security-specific: habituation is a near-universal property of repeated stimulus exposure across species and modalities [4], and mobile-notification research finds users already triaging incoming alerts within seconds of receipt, well before any deliberate "checking" behaviour could occur [5], with heavier interruption volume itself measurably degrading attention and mood [6], [7], and a mixed-methods study of university students finds alert fatigue and attention disruption, rather than raw notification frequency, the strongest predictors of reduced well-being [8].

A separate literature treats the audit rather than the alert as the object. An audit does not passively observe its target; it is an active process of making the target auditable in the first place, one that can silently substitute a proxy behaviour for the one the audit was meant to track [9], a substitution recent organisational-communication work frames as a predictable consequence of a metric becoming a target for the people it measures [10], and which occupational compliance research documents as "compliance theatre" once the audited population converges on the minimum action a metric will accept [11]. Strata governance itself rests on a formal regulatory structure that owners routinely under-comprehend and only selectively enact [12], and privacy scholarship on the smart-doorbell adjacent to our own object of study finds the device's stated security purpose is rarely the one residents report caring about day to day [13]. Compliance indicators holding at ceiling while the behaviour they proxy visibly erodes has been documented once already inside the University's own portfolio, in a public-pool alarm-response-time record that stayed perfect through a measured decline in lifeguards' independent scanning [14], and in a blind-review study finding certified auditors rated an automated system's synthesised compliance evidence more audit-ready than the human-logged trail it was meant to supplement [15].

## Method

### Sites and recruitment

Nine strata schemes in a single council area, each running a commercial video-intercom platform from one of three vendors (anonymised here as Platform A, B, and C), were recruited through their body-corporate committees; all nine had adopted the same insurer-specified security-compliance audit, under which a "guest logged" notification is logged as acknowledged if any resident's device registers a tap within five minutes. Recruitment fliers distributed through each building's noticeboard and mailing list yielded 71 volunteer respondents, of whom 64 completed a full interview after excluding seven who had lived in their unit for under one month.

## Interviews and coding

Each respondent completed a semi-structured interview of around twenty minutes, asked in open form what they actually do the moment a guest-logged notification appears on their phone, without being shown their own building's dashboard figures in advance. Two coders independently conducted a thematic analysis of each transcript [16], classifying its primary described practice into one of four categories, reconciled after a pilot pass of eight transcripts:

| Practice             | Definition      | n (%)    |
|----------------------|-----------------|----------|
| Verify, then act     | Checks, acts    | 14 (22%) |
| Verify, then dismiss | Checks, ignores | 17 (27%) |
| Reflexive swipe      | Unread dismiss  | 24 (38%) |
| Muted entirely       | Disabled        | 9 (14%)  |

Coded practice was then cross-tabulated against each respondent's tenure in their unit, banded at under one year, one to three years, and three years or more. Coder agreement was substantial on standard interpretive bands [17]:

$$\kappa = \frac{p_o - p_e}{1 - p_e} = 0.81$$

Recruitment continued past the point recent qualitative-methods work identifies as typical for thematic saturation [18], in order to preserve at least fifteen respondents in each tenure band and each intercom platform.

## Dashboard audit and prompt trial

For the same nine buildings we obtained each platform's own compliance export over the interview window, recording the acknowledgement rate and the per-notification time-to-acknowledge. To test whether a direct nudge could shift practice toward the audited behaviour, three of the nine buildings' committees agreed to add a "please confirm you saw this" follow-up prompt to the notification copy for a six-week trial; we compared the coded practice mix in the three prompted buildings against the six unprompted buildings at the trial's end.

## Results




Figure 1: Coded notification-management practice by tenure ( $n = 64$ ). Reflexive swipe is a minority practice under a year in a unit and the modal practice at three years or more.

Figure 1 shows the coded practice mix by tenure band. Reflexive swipe is a minority practice among residents under a year in their unit (3 of 18, 17%) and the modal practice among residents of three years or more (12 of 22, 55%), a reversal that comes almost entirely out of the two verification categories rather than out of muting: residents who mute the notification outright stay a stable minority across all three bands (2 of 18 to 4 of 22). Pooled across tenure, reflexive swipe is the single largest category overall (24 of 64, 38%), ahead of verify-then-dismiss (17, 27%), verify-then-act (14, 22%), and muted entirely (9, 14%).




Figure 2: Dashboard-logged time-to-acknowledge by tenure. The audit's own instrument reads faster, not slower, exactly where verification drops out of the interview data.

The dashboard's own time-to-acknowledge measure (Figure 2) moves the opposite way to what the coded interviews would predict if faster acknowledgement meant more attention rather than less: median time-to-acknowledge falls from 15.5 seconds among residents under a year in their unit to 5.5 seconds among residents of three years or more — exactly the tenure

band in which the interviews report the least verification. Read against the dashboard alone, a falling time-to-acknowledge and a steady acknowledgement rate would ordinarily be presented as an improving compliance trend; read against the interviews, the same numbers describe residents completing the audited gesture faster because they are doing less with it.




Figure 3: Coded practice mix at the six-week trial's end, prompted vs unprompted buildings. The ablation that changes nothing.

The prompt trial (Figure 3) changes almost nothing. Reflexive swipe accounts for 41% of prompted residents against 38% of unprompted residents, and the remaining three categories sit within a few percentage points of their unprompted counterparts throughout. We retain the trial in the reported results for completeness, and because no other intervention in the study window offers a comparable natural experiment against which to read the coded interviews.

## Discussion

Two independently collected measures — what residents say they do, and how quickly the platform logs them doing it — agree that time-to-acknowledge is not measuring engagement in any sense the audit's own documentation assumes. The audit was not designed to fail: a five-minute acknowledgement window and a per-notification timestamp are a reasonable, auditable proxy for “someone checked the door” in the platform's original brief, and nothing in this study suggests the three vendors' compliance exports are miscalculated. What the interviews add is that the population being audited has, over the ordinary course of living somewhere for a few years, converged on the cheapest gesture the proxy will accept, and that convergence happens to make the proxy's own headline number look better rather than worse. We

do not think this makes any of the nine buildings less secure than an insurer's report would suggest; we suggest only that “acknowledged” and “attended to” have quietly stopped being the same claim, in exactly the population — the longest-resident, most audit-experienced one — an insurer would ordinarily read as the most reassuring.

## Limitations

The sample is 64 residents across nine strata schemes in a single council area; a different intercom-adoption context, a different insurer's audit window, or a different housing market's tenure pattern might surface a different tenure gradient, though the consistency of the pattern across all three sampled vendor platforms suggests it is not an artefact of one platform's notification design. Practice was coded from self-report of what is, for the majority of longer-tenured respondents, close to an automatic action; a respondent's account of a reflex may itself be reconstructed after the fact rather than introspectively accurate, a limitation we consider unlikely to run in only one direction. The prompt trial ran for six weeks in three buildings; a longer trial or a differently worded prompt remains untested, though the near-identical practice mix at six weeks is itself suggestive of a floor rather than a slow-moving trend.

## Conclusion

A compliance metric can be published, logged at every notification, and read as close to ceiling without measuring what its own documentation assumes it measures. This paper finds a strata security audit's acknowledgement rate holding near-total across nine buildings while an interview study of the residents behind it finds verification practice falling away with tenure, displaced by a reflexive swipe the audit's own instrument cannot distinguish from genuine attention — and, if anything, rewards with a faster timestamp. Future work might instrument the notification itself, rather than the resident, to establish at what point in a tenure the swipe becomes reflexive, a threshold this study's banded tenure groups could describe but not date.

## References

- [1] M. Cvach, “Monitor Alarm Fatigue: An Integrative Review,” *Biomedical Instrumentation*

- & *Technology*, vol. 46, no. 4, pp. 268–277, 2012, doi: [10.2345/0899-8205-46.4.268](https://doi.org/10.2345/0899-8205-46.4.268).
- [2] S. Egelman, L. F. Cranor, and J. Hong, “You’ve Been Warned: An Empirical Study of the Effectiveness of Web Browser Phishing Warnings,” *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '08)*, pp. 1065–1074, 2008, doi: [10.1145/1357054.1357219](https://doi.org/10.1145/1357054.1357219).
  - [3] C. Bravo-Lillo, L. F. Cranor, J. Downs, and S. Komanduri, “Bridging the Gap in Computer Security Warnings: A Mental Model Approach,” *IEEE Security & Privacy*, vol. 9, no. 2, pp. 18–26, 2011, doi: [10.1109/MSP.2010.198](https://doi.org/10.1109/MSP.2010.198).
  - [4] C. H. Rankin et al., “Habituation Revisited: An Updated and Revised Description of the Behavioral Characteristics of Habituation,” *Neurobiology of Learning and Memory*, vol. 92, no. 2, pp. 135–138, 2009, doi: [10.1016/j.nlm.2008.09.012](https://doi.org/10.1016/j.nlm.2008.09.012).
  - [5] M. Pielot, K. Church, and R. de Oliveira, “An In-Situ Study of Mobile Phone Notifications,” *Proceedings of the 16th International Conference on Human-Computer Interaction with Mobile Devices & Services (MobileHCI '14)*, pp. 233–242, 2014, doi: [10.1145/2628363.2628364](https://doi.org/10.1145/2628363.2628364).
  - [6] G. Mark, S. T. Iqbal, M. Czerwinski, and P. Johns, “Focused, Aroused, but so Distractible: Temporal Perspectives on Multitasking and Communications,” *Proceedings of the 18th ACM Conference on Computer Supported Cooperative Work & Social Computing (CSCW '15)*, pp. 903–916, 2015, doi: [10.1145/2675133.2675221](https://doi.org/10.1145/2675133.2675221).
  - [7] K. Kushlev, J. Proulx, and E. W. Dunn, ““Silence Your Phones”: Smartphone Notifications Increase Inattention and Hyperactivity Symptoms,” *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems (CHI '16)*, pp. 1011–1020, 2016, doi: [10.1145/2858036.2858359](https://doi.org/10.1145/2858036.2858359).
  - [8] A. Y. Khanday, M. Muqeet, and Z. Syed, “Alert Fatigue and Smartphone Notifications: A Mixed-Methods Study of Attention Disruption and Mental Well-Being among University Students,” *Asian Journal of Education and Social Studies*, vol. 51, no. 12, pp. 988–999, 2025, doi: [10.9734/ajess/2025/v51i122743](https://doi.org/10.9734/ajess/2025/v51i122743).
  - [9] M. Power, “Making Things Auditable,” *Accounting, Organizations and Society*, vol. 21, no. 2–3, pp. 289–315, 1996, doi: [10.1016/0361-3682\(95\)00004-6](https://doi.org/10.1016/0361-3682(95)00004-6).
  - [10] J. W. Treem, W. C. Barley, M. S. Weber, and J. B. Barbour, “Signaling and Meaning in Organizational Analytics: Coping with Goodhart’s Law in an Era of Digitization and Datafication,” *Journal of Computer-Mediated Communication*, vol. 28, no. 4, p. zmad023, 2023, doi: [10.1093/jcmc/zmad023](https://doi.org/10.1093/jcmc/zmad023).
  - [11] S. A. Samad, “From Compliance Theater to Authentic Governance: A Framework for Risk Culture Integrating Structural Independence and Behavioral Science,” *International Journal of Financial Research*, vol. 17, no. 2, p. 1, 2026, doi: [10.5430/ijfr.v17n2p1](https://doi.org/10.5430/ijfr.v17n2p1).
  - [12] H. Easthope and B. Randolph, “Governing the Compact City: The Challenges of Apartment Living in Sydney, Australia,” *Housing Studies*, vol. 24, no. 2, pp. 243–259, 2009, doi: [10.1080/02673030802705433](https://doi.org/10.1080/02673030802705433).
  - [13] A. Mikhaylova, S. Wilford, B. Stahl, and L. Brooks, “Smart Doorbells in a Surveillance Society,” *Ethical and Social Impacts of Information and Communication Technology (ETHICOMP 2025)*, *Lecture Notes in Computer Science*, vol. 15939, 2025, doi: [10.1007/978-3-032-01429-0\\_39](https://doi.org/10.1007/978-3-032-01429-0_39).
  - [14] S. Okwuosa and J. Nwosu, “The Delegated Eye: Scanning Independence and Alarm Tenure Among Public-Pool Lifeguards,” *Slop University technical report*, 2026, doi: [10.5555/slop.g1d4nv](https://doi.org/10.5555/slop.g1d4nv).
  - [15] M. Hanke, C. Beng, and D. Okafor, “The Attestation Loom: Blind Reviewer Preference for Agent-Synthesised Audit Evidence over the Human-Logged Trail It Was Built to Supplement,” *Slop University technical report*, 2026, doi: [10.5555/slop.xl8qw0](https://doi.org/10.5555/slop.xl8qw0).
  - [16] V. Braun and V. Clarke, “Using Thematic Analysis in Psychology,” *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: [10.1191/1478088706qp0630a](https://doi.org/10.1191/1478088706qp0630a).
  - [17] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977, doi: [10.2307/2529310](https://doi.org/10.2307/2529310).

- [18] G. Guest, A. Bunce, and L. Johnson, “How Many Interviews Are Enough? An Experiment with Data Saturation and Variability,” *Field Methods*, vol. 18, no. 1, pp. 59–82, 2006, doi: [10.1177/1525822X05279903](https://doi.org/10.1177/1525822X05279903).