



Scanned, Not Seen

Interviewing Student-Society Event Organisers About What Noticeboard QR-Code RSVP Scans Miss About Who Actually Came

Ingrid Vasseur, Anouk Mensah

School of Emergent Priorities, Slop University

doi:10.5555/slop.b0njfs

Abstract. Since last semester, Slop University's Societies Space and Funding Panel has set each student society's room-allocation priority and funding tier from the average scan count its "Noticeboard RSVP" QR posters draw each semester, on the reasoning that a poster nobody stops for records nobody coming. This paper reports a mixed-methods reconciliation of that scan count against the turnout organisers say they actually watched arrive. We interviewed 28 organisers across 41 events run by 19 societies, and cross-checked 34 of those events' scan logs against an organiser-kept door tally. Scan counts correlate only weakly with door counts ($r = 0.39$) and run a median 2.4 times higher; the gap is not evenly spread but concentrates in events whose organisers say they knew the Panel was watching, where the ratio more than doubles relative to events organisers considered unwatched. A device-fingerprint deduplication pass removes very little of the inflation, consistent with organisers' own accounts of asking members to scan again rather than of a handful of devices scanning repeatedly. We suggest the scan count is measuring how well a society has learned the metric matters, not how many people came.

Introduction

A QR-code poster on a shared noticeboard is a cheap, low-friction way to ask people to say they are coming to something; it is a much less direct way to find out how many of them showed up. Slop University's Noticeboard RSVP pilot exists for the first purpose — societies pin a poster with a scannable RSVP link to promote an event — but the count it produces has, since the start of last semester, done a second job the pilot was not built for. The Societies Space and Funding Panel now averages each society's per-event scan count across a semester and uses the resulting figure, alongside declared room capacity, to decide which societies keep or move up a room-allocation and funding tier and which are asked to justify a smaller one. The Panel's own guidance frames the change plainly: a society whose posters draw scans is a society whose events people want, and rooms and funding should follow demand.

That reasoning holds only if a scan tracks a body at the door reasonably well, and reasonable-looking metrics have not always survived contact with the decisions built on top of them once the people producing the metric know what it decides. Whether Noticeboard RSVP's scan count is one of those metrics, or a genuinely faithful proxy for turnout, has not previously been checked against anything other than the scan log itself.

We report a mixed-methods study across one twelve-week semester: semi-structured interviews with 28 event organisers about how they promoted their events, whether and how they counted attendance at the door, and what they understood the scan count to be used for; and a reconciliation of 34 events' scan logs against the organisers' own door tallies. Our contributions:

1. We provide the first check, at this institution, of a Noticeboard RSVP scan count against an attendance figure gathered independently of the scanning platform itself.

2. We find the scan-to-door gap is not uniform: it is widest among organisers who say they knew the count was tied to their society's room and funding standing, and narrowest among those who did not.
3. We show, via a device-deduplication check, that the inflation is not well explained by a small number of devices scanning repeatedly, consistent with organisers' own accounts of a broader, more distributed scanning effort.

Related work

A public measure that is used to decide something reliably changes the behaviour it measures rather than simply describing it: rankings built to describe law schools have reshaped the schools competing for them, through mechanisms including self-fulfilling prophecy and the flattening of dissimilar things onto one comparable scale [1], [2], and Campbell's foundational caution — that a quantitative indicator used for social decision-making becomes more, not less, subject to corruption as its stakes rise — has held up well enough to be treated as closer to a law than a hypothesis [3]. Recent work formalising these dynamics distinguishes several distinct mechanisms by which a proxy target diverges from the goal it stands in for, useful because “the metric stopped tracking the thing” is not one failure mode but several [4]. Audit and performativity scholarship extends the point to universities specifically: practices institutions adopt to make themselves visible to their own oversight bodies have effects independent of, and sometimes contrary to, the qualities those practices are meant to surface [5], [6].

QR-code RSVP counts sit inside a narrower literature on whether a low-friction proxy signal predicts the behaviour it forecasts. QR-code adoption research finds behavioural intention to scan driven by perceived usefulness, attitude, and satisfaction rather than by any measure of what happens after the scan [7], and a QR-based attendance-monitoring deployment elsewhere reports high user-rated accuracy and convenience without independently checking its counts against attendance recorded another way [8]. Attendance-forecasting work outside the QR-code literature treats a gap between a registration signal and realised turnout as the default expectation, whether predicting individual attendance decisions from season-ticket

holder data [9] or transit ridership around a scheduled event [10]; crowd-counting research, for its part, increasingly treats estimating how a count changes between observation points, not just tallying one point in isolation, as central to getting a count right [11]. Self-report research in adjacent domains — health-service utilisation and church attendance — consistently finds proxy-reported behaviour departs from independently observed behaviour [12], [13], and departs further where the disclosure touches a more sensitive stake, such as a grade earned rather than a course merely taken [14]. Click-through rate, a much more heavily studied low-friction proxy than the QR scan, has drawn a substantial literature devoted to predicting it more accurately, itself a sign of how far short a raw click falls from the downstream outcome advertisers actually want [15]; university room-allocation research, by contrast, typically optimises directly against enrolled or scheduled attendance figures rather than any upstream engagement signal [16].

Prior work at this institution has already found an intake scoring tool tracking cataloguing convenience rather than the recovery outcome it claimed to predict [17], and a proximity-tag forecast whose accuracy rose only once staff learned where the tags were and adjusted their own behaviour around them [18]; this paper asks the same question of an RSVP scan rather than a claim score or a location tag.

Method

The Noticeboard RSVP pilot

Noticeboard RSVP lets a registered society generate a QR code linking to a one-field RSVP form (attendee name optional) for a specific event; the code is printed onto the society's own paper poster and pinned to a shared noticeboard in a location of the society's choosing, most often outside a lecture theatre or in the student commons. Each completed form submission logs one scan, timestamped and linked to a coarse device fingerprint. Since the start of last semester, the Societies Space and Funding Panel has averaged each society's scan count per event across a semester and used the resulting figure, alongside declared room capacity, to set the

society's room-allocation priority and funding tier for the following semester.

Sample

We tracked every event run through Noticeboard RSVP by a registered society across one twelve-week semester: 41 events across 19 societies. We interviewed 28 event organisers (some responsible for more than one tracked event) in semi-structured sessions of 25-40 minutes, covering how the event was promoted, whether and how attendance was counted at the door, what the organiser understood the scan count to be used for, and, where a discrepancy existed, the organiser's own account of it. Thirty-four of the 41 events had an organiser-kept door tally — a manual clicker count or a running paper tally taken at the entrance — treated here as the closest available figure to actual turnout, itself an imperfect count but gathered independently of the scanning platform. The remaining seven events had no kept tally and contribute interview material only.

Scan Inflation Ratio

For each of the 34 tallied events we compute the Scan Inflation Ratio, the event's scan count divided by its door tally:

$$\text{SIR} = \frac{N_{\text{scans}}}{N_{\text{door}}}$$

An SIR of 1 indicates the scan count and door tally agree exactly; values above 1 indicate more scans than people counted at the door.

Awareness grouping

During interviews, organisers were asked, unprompted, whether they believed the scan count fed into their society's room-allocation or funding-tier standing that semester. Responses split the 34 tallied events into an **aware** group (19 events, organisers who named the Panel's use of the figure without being told what the interview was about) and an **unaware** group (15 events, organisers who said they had not connected the count to any administrative decision, or believed their society's tier was already fixed).

Device-fingerprint deduplication

Noticeboard RSVP's coarse device fingerprint allows a secondary, deduplicated scan count per event: repeat submissions from the same fingerprint within a 10-minute window are collapsed to one. We compute SIR under both the raw count the Panel currently uses and this deduplicated count, to test whether the inflation is concentrated in a small number of devices scanning repeatedly or spread more broadly across distinct scans.

Results

Group	Events	Median SIR
Aware	19	3.3
Unaware	15	1.5
All tallied	34	2.4

Table 1: Tallied events by reported awareness of the scan count's administrative use, with median Scan Inflation Ratio.

Across the 34 tallied events, median SIR was 2.4 (IQR 1.6-3.3): a poster's scan count typically ran more than double the door tally for the event it promoted. Raw scan count and door count correlate only weakly ($r = 0.39$; Figure 1) — an event that draws twice the scans of another draws, on this evidence, nowhere near twice the people.




Figure 1: Per-event Noticeboard RSVP scan count against organiser-kept door tally ($n = 34$), with a linear fit. A rising fitted line does not imply a close relationship: $r = 0.39$.

The gap is not spread evenly across events (Table 1). Events in the aware group ran a median

SIR of 3.3 against 1.5 in the unaware group (Figure 2) — organisers who named the scan count’s administrative stakes unprompted produced posters whose scans diverged furthest from who actually came, not closest.




Figure 2: Median Scan Inflation Ratio by whether the organiser reported knowing the count fed the Panel’s room-allocation and funding-tier decision. The aware group’s posters diverge furthest from the door tally.

Nor is the divergence stable across the semester. Weekly mean SIR held around 1.6 for the first four weeks, before rising steadily to 2.9 by week twelve (Figure 3) — a climb that tracks, more closely than it tracks anything about the events themselves, the timing of an all-society email in week five restating the room-allocation and funding-tier formula.




Figure 3: Weekly mean Scan Inflation Ratio across the semester. The rise begins the week after the Panel’s formula was restated to all societies.

Interviews suggest the inflation is not principally a small number of devices scanning repeatedly. Several organisers in the aware group described asking members to “scan again” after attending, moving a poster to a busier noticeboard than the one nearest the venue, or scanning the code themselves from more than one device while testing it. Consistent with a distributed rather

than concentrated effect, a device-fingerprint deduplication pass barely moves the aggregate figure: median SIR falls from 2.4 (raw) to 2.2 (deduplicated), a difference that does not reach significance (Wilcoxon signed-rank, n.s., $p = .09$; Figure 4). We retain the raw, undeduplicated count throughout the comparisons above for consistency with the figure the Panel itself currently uses.




Figure 4: Scan Inflation Ratio distribution, raw against device-deduplicated scan counts, across all 34 tallied events. The two distributions overlap almost completely.

Discussion

Two independent signals point the same way: the scan-to-door correlation is weak across all tallied events, and the gap concentrates exactly where organisers say they know the count is being watched, arriving on a timeline that tracks the Panel’s own reminder rather than any change in the events themselves. Neither the deduplication check nor the aware/unaware split rules out that Noticeboard RSVP scans measure something real; they measure, on this evidence, how energetically a society has learned to promote a poster once promoting the poster is understood to be what the room and the funding tier are actually for, which is a different quantity from how many people came.

None of this shows the Panel’s underlying reasoning to be false in general — a society whose posters draw no scans at all plausibly is a society whose events draw little interest. What the evidence complicates is the assumption that more scans track more of the room being filled, once societies know the count is what fills it. A metric adopted to describe turnout and then used to allocate turnout’s own reward is, on this evidence, harder to keep pointed at turnout the longer it is used that way [1], [3].

Limitations

This study covers one semester at one institution, and the 34 tallied events depend on organisers who kept a door count in the first place, a group that may already promote and count more carefully than the seven who did not; the consistency of the aware/unaware split across nineteen independently run societies, however, suggests the pattern is not an artefact of any one society's habits. Organiser-kept door tallies are themselves imperfect counts, taken by a volunteer mid-event rather than a dedicated observer, though they were gathered independently of the scanning platform whose figures they are being checked against, which the Panel's own log cannot claim. Interview awareness was self-reported and could itself be shaped by the same stakes it describes.

Conclusion

A QR-code RSVP scan on a noticeboard poster is easy to collect and easy to inflate, and this study finds the two happening together within the same semester once collecting the scan is what a society's room and funding depend on. Across 34 events, scan counts correlated only weakly with organiser-kept door tallies, diverged furthest exactly where organisers reported knowing the count mattered, and moved on a timeline that tracks the Panel's own reminder of the formula rather than any change in turnout itself. A device-deduplication check rules out the simplest fix, since the inflation is not concentrated in a handful of repeat scanners. Future work might test whether a count gathered independently of the poster — a door sensor, say — proves any harder to game once its own stakes become known.

References

- [1] W. N. Espeland and M. Sauder, "Rankings and Reactivity: How Public Measures Recreate Social Worlds," *American Journal of Sociology*, vol. 113, no. 1, pp. 1–40, 2007, doi: [10.1086/517897](https://doi.org/10.1086/517897).
- [2] W. N. Espeland and M. L. Stevens, "Commensuration as a Social Process," *Annual Review of Sociology*, vol. 24, pp. 313–343, 1998, doi: [10.1146/annurev.soc.24.1.313](https://doi.org/10.1146/annurev.soc.24.1.313).
- [3] D. T. Campbell, "Assessing the Impact of Planned Social Change," *Evaluation and Program Planning*, vol. 2, no. 1, pp. 67–90, 1979, doi: [10.1016/0149-7189\(79\)90048-X](https://doi.org/10.1016/0149-7189(79)90048-X).
- [4] D. Manheim and S. Garrabrant, "Categorizing Variants of Goodhart's Law," *arXiv preprint arXiv:1803.04585*, 2018.
- [5] M. Strathern, "The Tyranny of Transparency," *British Educational Research Journal*, vol. 26, no. 3, pp. 309–321, 2000, doi: [10.1080/713651562](https://doi.org/10.1080/713651562).
- [6] S. J. Ball, "The Teacher's Soul and the Terrors of Performativity," *Journal of Education Policy*, vol. 18, no. 2, pp. 215–228, 2003, doi: [10.1080/0268093022000043065](https://doi.org/10.1080/0268093022000043065).
- [7] M. T. Nguyen and T. Alang, "When Do Shoppers Prefer Using QR Codes? Empirical Evidence from Vietnam," *Future Business Journal*, vol. 10, 2024, doi: [10.1186/s43093-024-00391-9](https://doi.org/10.1186/s43093-024-00391-9).
- [8] W. G. Domingo Dit, C. R. G. Abella, and R. B. Nucaza, "Technology Acceptance of QR Code Recognition as Attendance Monitoring on School Activities: Student Organization Officers Perspective," *Journal of Computer Sciences and Applications*, vol. 13, no. 1, pp. 29–36, 2025, doi: [10.12691/jcsa-13-1-3](https://doi.org/10.12691/jcsa-13-1-3).
- [9] J. K. Nguyen, A. Karg, A. Valadkhani, and H. McDonald, "Predicting Individual Event Attendance with Machine Learning: A 'Step-Forward' Approach," *Applied Economics*, vol. 54, no. 27, pp. 3138–3153, 2022, doi: [10.1080/00036846.2021.2003747](https://doi.org/10.1080/00036846.2021.2003747).
- [10] T. Santanam, A. Trasatti, P. Van Hentenryck, and H. Zhang, "Public Transit for Special Events: Ridership Prediction and Train Optimization," *arXiv preprint arXiv:2106.05359*, 2021.
- [11] W. Liu, M. Salzmann, and P. Fua, "Estimating People Flows to Better Count Them in Crowded Scenes," *arXiv preprint arXiv:1911.10782*, 2019.
- [12] M. E. Short et al., "How Accurate are Self-Reports? Analysis of Self-Reported Health Care Utilization and Absence When Compared with Administrative Data," *Journal of Occupational and Environmental Medicine*, vol. 51, no. 7, pp. 786–796, 2009, doi: [10.1097/JOM.0b013e3181a86671](https://doi.org/10.1097/JOM.0b013e3181a86671).
- [13] J. H. Shaver, T. A. J. White, P. Vakaoti, and M. Lang, "A Comparison of Self-Report,

Systematic Observation and Third-Party Judgments of Church Attendance in a Rural Fijian Village,” *PLOS ONE*, vol. 16, no. 10, 2021, doi: [10.1371/journal.pone.0257160](https://doi.org/10.1371/journal.pone.0257160).

- [14] J. A. Rosen, S. R. Porter, and J. Rogers, “Understanding Student Self-Reports of Academic Performance and Course-Taking Behavior,” *AERA Open*, vol. 3, no. 2, 2017, doi: [10.1177/2332858417711427](https://doi.org/10.1177/2332858417711427).
- [15] Y. Yang and P. Zhai, “Click-Through Rate Prediction in Online Advertising: A Literature Review,” *Information Processing & Management*, vol. 59, no. 2, p. 102853, 2022, doi: [10.1016/j.ipm.2021.102853](https://doi.org/10.1016/j.ipm.2021.102853).
- [16] A. Lemos, F. S. Melo, P. T. Monteiro, and I. Lynce, “Room Usage Optimization in Timetabling: A Case Study at Universidade de Lisboa,” *Operations Research Perspectives*, vol. 6, 2019, doi: [10.1016/j.orp.2018.100092](https://doi.org/10.1016/j.orp.2018.100092).
- [17] J. Nwosu and S. Okwuosa, “Logged by Score, Reclaimed by Chance: Auditing Whether a Library Lost-Property Claim-Likelihood Index Predicts What Gets Recovered,” *Slop University technical report*, 2026, doi: [10.5555/slop.37ag3g](https://doi.org/10.5555/slop.37ag3g).
- [18] V. Marris and K. Lindqvist, “Tagged to Stay Put: Forecasting Missing AV Trolleys Before the Weekly Stocktake,” *Slop University technical report*, 2026, doi: [10.5555/slop.mf9gpt](https://doi.org/10.5555/slop.mf9gpt).