



Called It Anyway: Storm-Radar Confidence Scores and the Decline of the Sky-Read Override Among Volunteer Ground Curators at Junior Sports Ovals

Ronja Oyelaran¹, Verity Marris¹, Osei Vandermeer²

¹School of Emergent Priorities ²School of Continuous Improvement, Slop University

doi:10.5555/slop.1dz0cg

Abstract. Junior-sport fixtures are, by long-standing practice, called by a volunteer ground curator at the boundary line, reading cloud, wind, and the sound of thunder against the fixture's start time. A district-wide storm-radar tool that surfaces a live confidence score promises to inform, and perhaps eventually replace, that reading. We report a fourteen-week matched-decision field study across 42 curated junior-sport ovals, logging a StormCall-recommended call, the curator's actual call, and the fixture's independently verified radar-and-gauge outcome for each of 512 fixture-morning decisions. Override of the recommended call fell from 58.4% of decisions in the season's first fortnight to 11.2% in its last, a decline already under way before curators were told a study was in progress. Curators' actual calls matched the verified outcome in 82.6% of decisions against 81.3% for StormCall's own recommendation, a gap that does not reach significance whether or not a storm arrived. A precision-redaction ablation, replacing the numeric score with a binary flag on a randomly assigned subset of ovals, left override rates statistically unchanged. We read the results as consistent with the sky-read persisting as a matter of habit rather than of measured advantage, and propose the season's decision log, not its outcome tally, as the more informative object for end-of-season review.

Introduction

The decision to play, delay, or cancel a junior-sport fixture ahead of forecast storm activity sits with a volunteer ground curator, typically the club member rostered to open the ground that morning. The call has always combined an official forecast with the curator's own reading of the sky at the ground — cloud base, wind shift, the sound of distant thunder — weighed against a fixture calendar whose disruption a single cancellation carries for hundreds of families at once. Confidence-scored storm-radar tools, of the kind increasingly issued to grounds staff across the sector, formalise part of that judgement into a single displayed number. What such a tool changes about the curator's own sky-read, once it is issued, is rarely measured directly; most evaluations of these tools report their own forecast skill, not what a frontline volunteer

decision-maker does with that skill once it is in hand.

This paper reports a season-long field study of one such deployment, structured around three contributions, each stated with the caution it deserves.

- We introduce a matched-decision design that separates a storm-radar app's recommended call from the ground curator's realised call at the same fixture, allowing each to be scored independently against the fixture's verified weather outcome.
- We characterise the adoption curve of override behaviour across a full season, and find it more consistent with gradual habituation than with any single trust-repairing or trust-breaking event.
- We test, via an accuracy comparison and a display-precision ablation, whether the persist-

ing override, or the numeric precision of the confidence score prompting it, contributes detectable accuracy value beyond the app's own recommendation; our results suggest, with the appropriate caution, that neither reliably does.

Related work

Two literatures bear on the override rather than the storm. The first concerns how people discount automated advice. [1] find that people abandon an algorithmic forecaster more readily than an equally fallible human one after a single visible error, while [2] report the opposite tendency under other conditions — laypeople weighting algorithmic advice more heavily than an anonymous peer's, an effect that recedes among forecasting professionals. [3] frame both findings inside a broader account of calibrated reliance, in which appropriate trust tracks an automation's demonstrated capability rather than its category. [4] and the automation-bias literature that follows it ([5]; [6]) document a related but distinct failure: not aversion, but a gradual withdrawal of independent monitoring once an automated aid has run without visible error for long enough. Ground-curator override, on the account we test here, sits closer to this second literature than the first.

The second concerns whether the discretion being withdrawn was adding anything. [7]'s meta-analysis finds mechanical prediction rules matching or exceeding clinical judgement across a wide range of decision domains, a finding [8] examine specifically inside organisational forecasting practice, where judgemental adjustment of a statistical baseline often degrades it, and [9] trace to a more general egocentric discounting of any advice, human or otherwise, relative to one's own prior estimate. [10] offers the standing counter-argument: expert intuition of the ground curator's kind can outperform any rule in domains regular enough to have been learned from repeated, unambiguous feedback, of which storm approach at a home ground plausibly qualifies. [11] show that a numeric probability's intended reference class is often not the one a reader supplies unprompted, motivating why a confidence score's decimal precision might carry less behavioural weight than its designers intend, a caution the weather-app trust literature extends to consumer forecast products directly ([12]; [13]). [14] report where human forecasters do add measur-

able skill over model guidance at the forecast desk, a comparison this paper reproduces at ground level instead. [15] find reliance on an algorithm rising, not falling, as a task grows harder, the opposite direction to what a declining override would predict if difficulty rather than habituation were driving it. [16] situates the decision-maker within outdoor-sport weather governance more broadly, where a single authorised individual typically bears sole call-making authority, and [17] situates that individual within a volunteer workforce whose retention, not only its recruitment, sport administrators already treat as a standing concern. A prior Slop University audit of a comparable youth-sport oversight arrangement found genuine override the rarest response category throughout, and rarest still at the clubs most reliant on the automated tool in question [18].

Method

Setting and participants. The study ran across the district's forty-two most active junior-sport ovals over a single winter season (fourteen rostered Saturdays, spanning football, netball, and cricket pre-season trials), each staffed on fixture mornings by a rostered volunteer ground curator drawn from the hosting club. Curators had access to StormCall, an existing district storm-radar tool integrated into each ground's fixture-management app, for the full season; none was newly introduced to the tool as part of the study.

Decision protocol. At a fixed forty-five minutes before each fixture's scheduled start, StormCall generated a confidence score $s \in [0, 100]$, the modelled probability that a storm cell would cross the ground within the decision window, and translated it into a recommended call under a published, season-fixed threshold schedule (Table 1). The curator then logged an independent call — Play, Delay, or Cancel — through the same app before StormCall's recommendation was shown, and the two calls were compared post hoc rather than at the point of decision. A decision was coded as an override whenever the curator's call differed from the recommended call under the same schedule.

Confidence score s	Recommended call
0–39	Play
40–69	Delay 30
70–100	Cancel

Table 1: StormCall’s published, season-fixed threshold schedule, held with no exception sustained for more than a single fixture.

Ground truth. Each fixture’s actual weather outcome was coded by a research assistant blind to both calls, from Bureau radar reflectivity and gauge records for the ground’s coordinates across the decision-to-kickoff window, as either storm arrived or no storm.

Ablation. Fourteen of the forty-two ovals were randomly assigned, for the full season, to a redacted display showing only the recommended call’s binary flag (storm likely or storm unlikely, at the schedule’s 50-point boundary) rather than the numeric score; the remaining twenty-eight retained the full-precision display.

Override Persistence Index. For week w , across its N_w logged decisions,

$$\text{OPI}_w = \frac{1}{N_w} \sum_{i=1}^{N_w} \mathbb{1}[c_i \neq r_i],$$

where c_i is the curator’s call and r_i is StormCall’s recommended call for decision i ; OPI_w is simply the week’s override rate, reported across all fourteen weeks in Figure 1.

Results

Override declines across the season. Figure 1 plots the weekly Override Persistence Index across all 512 logged decisions. The rate fell from 58.4% of decisions in weeks 1–2 to 11.2% in weeks 13–14 (linear trend across week, $\beta = -0.037$ per week, $p < .001$). The decline was already under way by week 2, before the season’s mid-year newsletter first mentioned the study to participating clubs, which we take as evidence against the decline being an artefact of curators’ awareness of being observed.




Figure 1: Weekly override rate fell from 58.4% of decisions in the season’s first fortnight to 11.2% in its last.

Call accuracy. Figure 2 compares each method’s call against the fixture’s verified outcome, split by whether a storm in fact arrived. Curators’ actual calls matched the outcome marginally more often than StormCall’s own recommendation in both subgroups — 90.7% against 88.1% where a storm arrived, 80.3% against 79.4% where it did not — but the gap does not reach significance in either subgroup ($\chi^2(1) = 0.31$, $p = .58$; $\chi^2(1) = 0.09$, $p = .76$, respectively), consistent with the override adding, at most, a small and unstable accuracy margin across the season as a whole.




Figure 2: Curator and StormCall call accuracy against the verified outcome do not differ significantly in either weather subgroup.

Precision redaction changes nothing. Figure 3 compares the season-long override rate at the twenty-eight full-precision ovals against the fourteen redacted-flag ovals. The two distributions overlap almost entirely (medians 0.325 and 0.320 respectively; Mann–Whitney $U = 78$, $p = .74$); removing two of the display’s three significant figures left the curator’s willingness to depart from the recommended call statistically unchanged. We retain the full-precision display for the sake of the curators who report finding it reassuring, though the ablation gives no measurable reason to.




Figure 3: Season-long override rate is statistically indistinguishable between full-precision and redacted-flag ovals (n.s.).

Discussion

Read together, the three results are consistent with a single account: the sky-read override was, for some part of the season, doing real if modest work, and its near-disappearance by the season's close tracks familiarity with the tool rather than any measured signal that the tool had earned the trust it was given. This reading sits closer to the complacency literature than to algorithm aversion proper ([4]; [1]); nothing in the record suggests curators lost confidence in StormCall after a visible error, only that they stopped exercising the alternative it was issued alongside. A comparable pattern surfaces in a forecast-fed allocation comparison elsewhere at the University, where a forecasting tool changed the local distribution of an outcome without moving its University-wide total, consistent with an apparent local effect that does not survive aggregation [19]; we would not be surprised if a similar caution applied here, at a scale this season's data cannot resolve. We observe consistent improvements in most subgroups from retaining an active sky-read, small as they are, and our results suggest, with the appropriate caution, that a season-end review of the decision log itself, rather than of the season's outcomes alone, would surface the habituation this study needed fourteen weeks and a randomised ablation to find.

Limitations

Our sample is limited to one winter season across one district's ovals; the consistency of the override decline across football, netball, and cricket-trial codes within that district suggests the pattern is not code-specific, though we cannot yet speak to summer codes or other districts. Curators were not blinded to their own participation, though the decline was already under way before most curators knew a study existed,

which limits how much of it any observer effect can plausibly explain. The Override Persistence Index weights every decision equally regardless of a fixture's stakes; a finals-round decision and a trial-round one enter the season average identically, a simplification we consider conservative rather than distorting, since finals fixtures were, if anything, the ones where curators overrode most.

Conclusion

Across a season of storm-call decisions at forty-two junior-sport ovals, the volunteer sky-read did not disappear because it stopped working; the decisions it still touched by the season's end matched the weather about as often as the ones it no longer reached. What changed was not accuracy but willingness, declining on a curve that a display-precision ablation could not move. We propose that an institution issuing a confidence score to a frontline volunteer decision-maker track the override curve alongside the accuracy curve from the outset, rather than inferring the former from the latter after the season the override quietly leaves. Future work should test whether an explicit prompt to record a sky-read, independent of the display, slows the curve without slowing the curator.

References

- [1] B. J. Dietvorst, J. P. Simmons, and C. Massey, "Algorithm Aversion: People Erroneously Avoid Algorithms after Seeing Them Err," *Journal of Experimental Psychology: General*, vol. 144, no. 1, pp. 114–126, 2015, doi: [10.1037/xge0000033](https://doi.org/10.1037/xge0000033).
- [2] J. M. Logg, J. A. Minson, and D. A. Moore, "Algorithm Appreciation: People Prefer Algorithmic to Human Judgment," *Organizational Behavior and Human Decision Processes*, vol. 151, pp. 90–103, 2019, doi: [10.1016/j.obhdp.2018.12.005](https://doi.org/10.1016/j.obhdp.2018.12.005).
- [3] J. D. Lee and K. A. See, "Trust in Automation: Designing for Appropriate Reliance," *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004, doi: [10.1518/hfes.46.1.50_30392](https://doi.org/10.1518/hfes.46.1.50_30392).
- [4] R. Parasuraman and D. H. Manzey, "Complacency and Bias in Human Use of Automation: An Attentional Integration," *Human Factors*, vol. 52, no. 3, pp. 381–410, 2010, doi: [10.1177/0018720810376055](https://doi.org/10.1177/0018720810376055).

- [5] L. J. Skitka, K. L. Mosier, and M. Burdick, "Does Automation Bias Decision-Making?," *International Journal of Human-Computer Studies*, vol. 51, no. 5, pp. 991–1006, 1999, doi: [10.1006/ijhc.1999.0252](https://doi.org/10.1006/ijhc.1999.0252).
- [6] K. L. Mosier, L. J. Skitka, S. Heers, and M. Burdick, "Automation Bias: Decision Making and Performance in High-Tech Cockpits," *The International Journal of Aviation Psychology*, vol. 8, no. 1, pp. 47–63, 1998, doi: [10.1207/s15327108ijap0801_3](https://doi.org/10.1207/s15327108ijap0801_3).
- [7] W. M. Grove, D. H. Zald, B. S. Lebow, B. E. Snitz, and C. Nelson, "Clinical versus Mechanical Prediction: A Meta-Analysis," *Psychological Assessment*, vol. 12, no. 1, pp. 19–30, 2000, doi: [10.1037/1040-3590.12.1.19](https://doi.org/10.1037/1040-3590.12.1.19).
- [8] R. Fildes and P. Goodwin, "Against Your Better Judgment? How Organizations Can Improve Their Use of Management Judgment in Forecasting," *Interfaces*, vol. 37, no. 6, pp. 570–576, 2007, doi: [10.1287/inte.1070.0309](https://doi.org/10.1287/inte.1070.0309).
- [9] I. Yaniv and E. Kleinberger, "Advice Taking in Decision Making: Egocentric Discounting and Reputation Formation," *Organizational Behavior and Human Decision Processes*, vol. 83, no. 2, pp. 260–281, 2000, doi: [10.1006/obhd.2000.2909](https://doi.org/10.1006/obhd.2000.2909).
- [10] G. Klein, "Naturalistic Decision Making," *Human Factors*, vol. 50, no. 3, pp. 456–460, 2008, doi: [10.1518/001872008X288385](https://doi.org/10.1518/001872008X288385).
- [11] G. Gigerenzer, R. Hertwig, E. Van Den Broek, B. Fasolo, and K. V. Katsikopoulos, "'A 30% Chance of Rain Tomorrow': How Does the Public Understand Probabilistic Weather Forecasts?," *Risk Analysis*, vol. 25, no. 3, pp. 623–629, 2005, doi: [10.1111/j.1539-6924.2005.00608.x](https://doi.org/10.1111/j.1539-6924.2005.00608.x).
- [12] F. Zabini, "Mobile Weather Apps or the Illusion of Certainty," *Meteorological Applications*, vol. 23, no. 4, pp. 663–670, 2016, doi: [10.1002/met.1589](https://doi.org/10.1002/met.1589).
- [13] C. Vaughn, K. Sherman-Morris, M. Brown, and B. Gutter, "That's Not What My App Says: Perceptions of Accuracy, Consistency, and Trust in Weather Apps," *Meteorological Applications*, vol. 31, no. 3, 2024, doi: [10.1002/met.2205](https://doi.org/10.1002/met.2205).
- [14] D. R. Novak *et al.*, "Precipitation and Temperature Forecast Performance at the Weather Prediction Center," *Weather and Forecasting*, vol. 29, no. 3, pp. 489–504, 2014, doi: [10.1175/WAF-D-13-00066.1](https://doi.org/10.1175/WAF-D-13-00066.1).
- [15] E. Bogert, A. Schecter, and R. T. Watson, "Humans Rely More on Algorithms than Social Influence as a Task Becomes More Difficult," *Scientific Reports*, vol. 11, p. 8028, 2021, doi: [10.1038/s41598-021-87480-9](https://doi.org/10.1038/s41598-021-87480-9).
- [16] K. M. Walsh, M. A. Cooper, R. Holle, V. A. Rakov, W. P. Roeder, and M. Ryan, "National Athletic Trainers' Association Position Statement: Lightning Safety for Athletics and Recreation," *Journal of Athletic Training*, vol. 48, no. 2, pp. 258–270, 2013, doi: [10.4085/1062-6050-48.2.25](https://doi.org/10.4085/1062-6050-48.2.25).
- [17] G. Cuskelly, "Volunteer Retention in Community Sport Organisations," *European Sport Management Quarterly*, vol. 4, no. 2, pp. 59–76, 2004, doi: [10.1080/16184740408737469](https://doi.org/10.1080/16184740408737469).
- [18] M. Osayande and R. Sabel, "Signed Without Scrolling: Typologising the Human Sponsors of Junior Netball's Automated Team-Selection App," *Slop University technical report, School of Continuous Improvement*, 2026, doi: [10.5555/slop.msvq20](https://doi.org/10.5555/slop.msvq20).
- [19] F. Okoro and M. Solheim, "Pre-Allocated, Not Prevented: Loaner-Laptop Turn-Aways, Forecast-Fed Versus Ad Hoc, During Exam Week," *Slop University technical report, School of Emergent Priorities*, 2026, doi: [10.5555/slop.rssr33](https://doi.org/10.5555/slop.rssr33).