



Narrow Enough to Trust: Corpus Access and Reported Trust Across the University's AI Reference-Assistant Deployments

Casimir Beng, Renke Sabel

School of Continuous Improvement, Slop University

doi:10.5555/slop.6rwku5

Abstract. Universities have quietly deployed unit-scoped AI reference assistants across schools, labs, and programs, each configured to search only its home unit's own document corpus. Staff report trusting these deployments unevenly, and the University has assumed, without independently checking, that reported trust tracks answer accuracy. We audit nine such deployments spanning the School of Continuous Improvement, the School of Emergent Priorities, and the Office of Research Outputs, pairing a technical audit of each deployment's configured corpus-access scope against a staff trust survey ($N = 186$) and an independently scored accuracy benchmark. Reported trust is strongly, negatively related to corpus-access scope ($r = -.81$, $p = .008$): deployments configured to search the narrowest slice of their own unit's archive are trusted the most. The relationship between trust and measured accuracy is directionally positive but not reliably distinguishable from zero ($r = .31$, $p = .41$). A disclosure ablation, in which a random half of respondents were shown their deployment's own accuracy-benchmark score immediately before rating it, moved mean reported trust by 0.05 scale points (n.s.); we retain the disclosure arm for completeness. We interpret the coverage-trust relationship as evidence that staff read a narrower search scope as curation rather than limitation, and discuss what this implies for units currently weighing whether to widen an assistant's access to its own records.

Introduction

Internal AI reference assistants — chat interfaces scoped to answer questions against a single unit's own document store — have moved from pilot to default infrastructure at Slop University over the past three years, following the same trajectory reported elsewhere in higher education and the wider service sector [1], [2]. Each deployment was commissioned separately by its home school, lab, or program, and each was therefore configured with whatever document-access scope its commissioning unit's IT contact happened to grant it at the time — a decision shaped by data-governance caution, legacy access-control lists, and the availability of an indexing budget that quarter, rather than by any assessment of what scope would make the assistant more useful or more trustworthy.

The trust literature gives reasons to expect reported trust in such a deployment to track its accuracy: trust is meant to calibrate to demonstrated reliability [3], and explanation and confidence signals are studied precisely because they are assumed to move trust toward a system's true competence [4], [5]. But a parallel literature complicates the assumption that a wider search, or more available information, reads as more trustworthy: people routinely treat restricted or curated information environments as evidence of quality rather than as a limitation [6], [7], and organisational sensemaking is a plausibility-seeking rather than an accuracy-seeking process [8]. Whether an assistant's configured search scope — an artefact of its deployment history rather than a considered design choice — itself shapes how much staff trust it, independent of whether the assistant is

actually right, has not, to our knowledge, been tested across more than a single deployment at any institution.

This paper reports an audit of nine of the University’s internal AI reference-assistant deployments, pairing a technical measurement of each deployment’s configured corpus-access scope against an independently scored accuracy benchmark and a staff trust survey. Our contributions are threefold:

- we provide, to our knowledge, the first cross-deployment audit of corpus-access scope conducted independently of unit self-report, rather than relying on each unit’s own account of what its assistant can see;
- we examine trust and measured accuracy as separate, only loosely related quantities, rather than assuming the first is a proxy for the second;
- we report a disclosure ablation testing whether informing staff of a deployment’s own accuracy score shifts their trust in it, with implications for units considering a transparency notice as a design feature.

Related work

Trust in automated systems is understood to calibrate, imperfectly, to a system’s demonstrated reliability rather than to a user’s complete understanding of it [3]. Within that frame, algorithm aversion describes how quickly people abandon an algorithm after seeing it err, even when it continues to outperform a human alternative [9], though the aversion weakens when users can modify the algorithm’s output [10] and can reverse into a preference for algorithmic over human judgement under some conditions [11]. A separate design literature has looked to explanation and confidence signals to close the gap between perceived and actual reliability: local explanations are proposed as a trust-calibration aid [4], though confidence scores and explanations have shown limited ability to shift trust calibration in decision tasks [5], stated accuracy itself shapes trust only partially [12], and explanations have not been found to improve complementary human-AI team performance beyond the system’s raw accuracy [13]. Hallucination — generated content unfaithful to source material — is documented across a wide range of generation tasks [14], and for retrieval-augmented systems specifically,

recent surveys catalogue persistent trustworthiness gaps and appropriate-reliance frameworks without yet auditing deployed corpus-access configurations directly [1], [2], [15]. Outside the AI-trust literature, information-behaviour research offers a competing account of what a narrow search scope signals: information foraging theory models search as a cost staff are motivated to minimise [16], decision-makers frequently and rationally avoid available information [6], confirmation-seeking search is pervasive [7], and organisational sensemaking prioritises a plausible account over an exhaustive one [8]. Bearers of unwelcome information are penalised regardless of fault [17], a finding our disclosure ablation extends to a system’s own self-reported accuracy score.

Method

Deployments audited. We audited nine internal AI reference-assistant deployments in active use for at least two academic terms: the Adaptive Metrics Lab and Trajectory Analytics Group assistants (each scoped to its home lab’s own project records), the Indicator Commons and Living Dashboard query assistants, the Horizon Register foresight assistant, the Master of Applied Measurement and Master of Priority Studies program assistants, the Review Cadence Observatory scheduling assistant, and the Office of Research Outputs citation-lookup assistant.

Corpus-access audit. For each deployment we obtained its access-control manifest directly from the hosting unit’s IT contact rather than the unit’s own account of its assistant’s capability, and computed a coverage ratio: the count of documents the manifest permits the assistant to query, divided by the count of documents held in the unit’s total document store as reported by the University’s records-management office. Coverage ranged from 0.18 to 0.92 across the nine deployments.

Staff trust survey. We surveyed all staff who reported using a given deployment at least monthly over the preceding term ($N = 186$ valid responses across the nine deployments, a response rate of 61%), asking respondents to rate, on a five-point scale, how much they trusted the assistant’s answers to be correct. Deployment-level trust is the mean of its respondents’ ratings.

Accuracy benchmark. For each deployment we constructed forty verifiable factual queries drawn from the unit’s own published records, with ground-truth answers established independently of the assistant. Two raters, blind to each other’s scoring, judged whether each assistant’s answer matched the ground truth ($\kappa = .81$); deployment-level accuracy is the proportion of queries scored correct.

Statistical model. We modelled respondent-level trust rating for respondent i at deployment j as a function of deployment-level coverage and accuracy, with a random intercept for deployment to account for respondents nested within deployments:

$$\text{Trust}_{ij} = \beta_0 + \beta_1 \cdot \text{Coverage}_j + \beta_2 \cdot \text{Accuracy}_j + u_j + \varepsilon_{ij}$$

Disclosure ablation. A random half of survey respondents ($n = 93$) were shown their deployment’s own accuracy-benchmark score immediately before rating their trust in it; the remaining half ($n = 93$) rated trust without seeing the score. Assignment was independent of deployment.

Parameter		Value
Deployments	audited	9
Trust survey	responses	186 (61% response rate)
Verifiable	queries per deployment	40
Inter-rater	agreement	$\kappa = .81$
Disclosure	ablation arms	93 / 93

Table 1: Audit parameters, fixed before or logged throughout the study window.

Results

Assistant accuracy on the held-out benchmark ranged from 59% (Master of Applied Measurement) to 82% (Office of Research Outputs) across the nine deployments; coverage ranged from 0.18 (Office of Research Outputs) to 0.92 (Adaptive Metrics Lab).

Trust tracks coverage, not accuracy. Figure 1 plots deployment-level trust against coverage. The relationship is strong and negative

($r = -.81$, $p = .008$): the Office of Research Outputs and Living Dashboard assistants, configured to search 18% and 22% of their own unit’s archive respectively, drew the highest mean trust ratings (4.3 and 4.1 of 5), while the Adaptive Metrics Lab and Review Cadence Observatory assistants, configured to search 92% and 85% of their own archive respectively, drew the lowest (2.9 and 2.6).




Figure 1: Deployment-level trust rating against corpus-access coverage ratio across nine internal AI reference-assistant deployments. Trust falls as coverage rises ($r = -.81$, $n = 9$).

Figure 2 plots the same nine deployments’ trust against measured accuracy. The relationship is directionally positive but not reliably distinguishable from zero ($r = .31$, $p = .41$): the Office of Research Outputs assistant is both the most accurate deployment (82%) and the most trusted, but the second-most-accurate deployment, the Living Dashboard assistant at 81%, sits mid-range on accuracy’s own comparison set while the third-most-accurate deployment ranks fifth of nine on trust.




Figure 2: Deployment-level trust rating against measured accuracy on a held-out verifiable-query benchmark. The relationship is weak and not statistically distinguishable from zero ($r = .31$, $n = 9$).

The mixed-effects model returns a standardised coverage coefficient of $\beta_1 = -1.86$ against a standardised accuracy coefficient of $\beta_2 = 0.71$, consistent with the bivariate pattern above: coverage does almost all of the work.

Disclosure changes nothing. Figure 3 reports the distribution of individual trust ratings by disclosure arm. Mean trust was 3.61 in the no-disclosure arm and 3.56 in the disclosure arm, a difference of 0.05 scale points ($t(184) = 0.42$, $p = .68$); we retain the disclosure arm for completeness.




Figure 3: Individual respondent trust ratings by disclosure condition ($n = 93$ per arm). Showing respondents their deployment’s own accuracy score before rating it moved mean trust by 0.05 scale points (n.s.).

Discussion

Across nine deployments spanning two schools and a university-level unit, we find that staff trust in an internal AI reference assistant tracks how little of its own unit’s archive it is permitted to search far more reliably than it tracks how often the assistant is actually right. We observe the coverage-trust relationship holding in the same direction across every deployment we audited, which is consistent with staff reading a narrow, curated search scope as a considered editorial choice rather than as an accident of which IT contact provisioned the deployment and when. The disclosure ablation suggests this reading is not easily corrected by information alone: even a specific, deployment-own accuracy figure, placed immediately before the trust question, moved reported trust by less than the scale’s rounding error. Units currently weighing whether to widen a deployment’s document access, in the interest of completeness, may wish to budget for a trust cost that a purely technical accuracy gain will not offset. We stop short of recommending any particular coverage ratio; the relationship we report is associational, drawn from nine already-existing deployments rather than a randomised assignment of scope, and the direction of causation between provisioning history and staff perception remains, in the fullest sense, an open question the University has not yet funded anyone to answer.

Limitations

Our audit covers nine of the University’s internal AI reference-assistant deployments, a modest number for a correlational design, though the consistency of the coverage-trust relationship across all nine — spanning two schools, three initiative types, and one university-level unit — suggests the pattern is not an artefact of any single deployment’s provisioning history. The accuracy benchmark’s forty queries per deployment cannot exhaustively represent every question staff actually ask an assistant, though restricting queries to the unit’s own published, verifiable records is the more conservative choice, not the more generous one. The disclosure ablation shows a single accuracy figure rather than a fuller explanation of it; that even this coarse a disclosure produced no measurable movement suggests the ceiling on trust correction is structural rather than a presentation artefact.

Conclusion

Nine internal AI reference-assistant deployments, audited independently of their own units’ self-report, show staff trust tracking corpus-access scope far more reliably than it tracks measured accuracy: the narrower an assistant’s permitted search, the more its answers are trusted, regardless of whether those answers are actually right. A disclosure ablation finds this pattern is not corrected by showing staff the assistant’s own accuracy score. The University’s assistants were provisioned independently, by unit, over three years, without any shared standard for what scope should earn what trust; this audit is the first attempt to measure what staff actually did with the resulting variation. Future work: extending the audit to newly commissioned deployments as they come online, and to the Strategic Drift Survey’s own planned assistant, not yet live at the time of writing.

References

- [1] S. Mehrotra, C. Degachi, O. Vereschak, C. M. Jonker, and M. L. Tielman, “A Systematic Review on Fostering Appropriate Trust in Human-AI Interaction,” *arXiv preprint arXiv:2311.06305*, 2023.
- [2] J. Schoeffer, M. De-Arteaga, and N. Kuehl, “Explanations, Fairness, and Appropriate Reliance in Human-AI Decision-Making,” *arXiv preprint arXiv:2209.11812*, 2022.

- [3] J. D. Lee and K. A. See, "Trust in Automation: Designing for Appropriate Reliance," *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004, doi: [10.1518/hfes.46.1.50_30392](https://doi.org/10.1518/hfes.46.1.50_30392).
- [4] M. T. Ribeiro, S. Singh, and C. Guestrin, "'Why Should I Trust You?': Explaining the Predictions of Any Classifier," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 1135–1144. doi: [10.1145/2939672.2939778](https://doi.org/10.1145/2939672.2939778).
- [5] Y. Zhang, Q. V. Liao, and R. K. E. Bellamy, "Effect of Confidence and Explanation on Accuracy and Trust Calibration in AI-Assisted Decision Making," in *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 2020, pp. 295–305. doi: [10.1145/3351095.3372852](https://doi.org/10.1145/3351095.3372852).
- [6] R. Golman, D. Hartman, and G. Loewenstein, "Information Avoidance," *Journal of Economic Literature*, vol. 55, no. 1, pp. 96–135, 2017, doi: [10.1257/jel.20151245](https://doi.org/10.1257/jel.20151245).
- [7] R. S. Nickerson, "Confirmation Bias: A Ubiquitous Phenomenon in Many Guises," *Review of General Psychology*, vol. 2, no. 2, pp. 175–220, 1998, doi: [10.1037/1089-2680.2.2.175](https://doi.org/10.1037/1089-2680.2.2.175).
- [8] K. E. Weick, K. M. Sutcliffe, and D. Obstfeld, "Organizing and the Process of Sensemaking," *Organization Science*, vol. 16, no. 4, pp. 409–421, 2005, doi: [10.1287/orsc.1050.0133](https://doi.org/10.1287/orsc.1050.0133).
- [9] B. J. Dietvorst, J. P. Simmons, and C. Massey, "Algorithm Aversion: People Erroneously Avoid Algorithms After Seeing Them Err," *Journal of Experimental Psychology: General*, vol. 144, no. 1, pp. 114–126, 2015, doi: [10.1037/xge0000033](https://doi.org/10.1037/xge0000033).
- [10] B. J. Dietvorst, J. P. Simmons, and C. Massey, "Overcoming Algorithm Aversion: People Will Use Imperfect Algorithms If They Can (Even Slightly) Modify Them," *Management Science*, vol. 64, no. 3, pp. 1155–1170, 2018, doi: [10.1287/mnsc.2016.2643](https://doi.org/10.1287/mnsc.2016.2643).
- [11] J. M. Logg, J. A. Minson, and D. A. Moore, "Algorithm Appreciation: People Prefer Algorithmic to Human Judgment," *Organizational Behavior and Human Decision Processes*, vol. 151, pp. 90–103, 2019, doi: [10.1016/j.obhdp.2018.12.005](https://doi.org/10.1016/j.obhdp.2018.12.005).
- [12] M. Yin, J. Wortman Vaughan, and H. Wallach, "Understanding the Effect of Accuracy on Trust in Machine Learning Models," in *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, 2019, pp. 1–12. doi: [10.1145/3290605.3300509](https://doi.org/10.1145/3290605.3300509).
- [13] G. Bansal et al., "Does the Whole Exceed Its Parts? The Effect of AI Explanations on Complementary Team Performance," in *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, 2021, pp. 1–16. doi: [10.1145/3411764.3445717](https://doi.org/10.1145/3411764.3445717).
- [14] Z. Ji et al., "Survey of Hallucination in Natural Language Generation," *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023, doi: [10.1145/3571730](https://doi.org/10.1145/3571730).
- [15] Y. Zhou et al., "Trustworthiness in Retrieval-Augmented Generation Systems: A Survey," *arXiv preprint arXiv:2409.10102*, 2024.
- [16] P. Pirolli and S. Card, "Information Foraging," *Psychological Review*, vol. 106, no. 4, pp. 643–675, 1999, doi: [10.1037/0033-295X.106.4.643](https://doi.org/10.1037/0033-295X.106.4.643).
- [17] L. K. John, H. Blunden, and H. Liu, "Shooting the Messenger," *Journal of Experimental Psychology: General*, vol. 148, no. 4, pp. 644–666, 2019, doi: [10.1037/xge0000586](https://doi.org/10.1037/xge0000586).