Slop / University

Screenshotted, Just in Case

An Interview Study of How Rideshare Riders Renegotiate Trust in the App's Upfront Fare Estimate After It Once Ran High

Mirela Hanke[1], Iben Chikere[2]

[1]School of Continuous Improvement [2]School of Emergent Priorities, Slop University

doi:10.5555/slop.jugfj6

Abstract. A rideshare app's upfront fare estimate is shown before a rider commits to a trip, then followed, minutes later, by whatever the ride actually cost. We conducted semi-structured interviews with 32 riders recruited across a mix of regular and occasional app use, opening each interview by asking a participant to describe what they do, if anything, with the quoted price before they book, then working through their history of trips where the final charge landed higher than the number they had been shown. Riders sorted into four coping behaviours: screenshotting every quote before booking, mentally rounding the number up by a private margin, checking the final fare only occasionally, or no longer checking closely at all. The behaviour a rider settled into tracked how many prior fare surprises they reported, with screenshotting most common among riders who had been surprised twice or more and least common among riders with none. A mid-interview manipulation — telling participants that the app labels its estimate "informational only" in its terms of service — left the distribution of stated coping behaviour statistically unchanged. We discuss what a rider's private instrument for tracking an institution's number suggests about where trust actually gets rebuilt after a platform's own figure runs high, and where a name for its own labelled uncertainty does less work than an institution might assume.

Introduction

A rideshare app shows its price before the ride begins. The number appears on the booking screen, a rider taps to confirm, and the trip's actual cost is settled only once the ride ends — by which point the estimate has already done its job of securing the booking. Platforms describe this figure as an estimate, not a promise, in terms of service that few riders read and fewer still recall verbatim [1]. Whether the number's status as "informational only" survives contact with a rider who has been charged more than it once said is a separate question from whether the label itself is technically accurate, and it is the one this paper asks.

Research on algorithmic pricing has mostly asked whether an estimate is fair or accurate — auditing bias in ride-hailing fares [2], modelling how a driver responds to a posted surge multiplier [3], and characterising how ride-pooling revenue shifts under different pricing strategies [4]. A separate literature on trust repair asks what an institution can do, after a failure, to win a customer back [5], [6], [7]. Neither line of work asks what a rider does on their own, without the institution's help, once the figure has run high often enough to no longer be trusted at face value: whether they keep a private record of their own, adjust the number before they act on it, or simply stop attending to it.

We make three contributions, each modest in the way a first qualitative pass into a new setting should be. First, we develop a four-category taxonomy of how riders relate to the app's fare estimate once it has ever run high, built from interview material rather than app-usage logs. Second, we show the taxonomy's categories track a rider's own count of prior fare surprises, with the most vigilant behaviour concentrated among riders who report two or more. Third, we test whether naming the estimate's own labelled uncertainty aloud changes a rider's stated coping behaviour, and find that it does not.

Related work

Ride-hailing pricing has drawn sustained attention as a technical object: fairness audits of the fare-setting algorithm itself [2], models of how drivers accept or decline a ride under a posted price [8], and studies of surge-multiplier design and its effect on both riders and drivers [3], [4]. This literature treats the estimate as an output to be optimised or audited; how a rider privately relates to that output once it has proved unreliable sits outside its scope.

A second literature asks how algorithmic decisions earn or lose a user's trust more generally: users lose confidence in an algorithm faster than in an equivalent human forecaster after a single visible error [9], and a still-fragmented body of scholarship across disciplines has tried to integrate what a platform's algorithmic-management choices do to the people transacting on the other side of the same interface [10]. Ride-hailing's fare estimate differs from the forecasting settings this literature usually studies in one respect: a rider cannot decline the number and substitute their own, only accept, cancel, or develop a private way of living alongside it.

A third, adjacent literature treats disclaimers and fine print as the institution's own tool for managing exactly this gap. Advertising disclaimers measurably shift brand attitudes but are attended to less than the claims they qualify [1], and repeated exposure to a routine warning measurably suppresses the attention it receives on each subsequent viewing [11]. Trust-repair research, meanwhile, treats a relabelling of the original claim as at most one lever among several: transactional remedies such as compensation or a regulatory fix repair a competence-based failure more reliably than words alone, and pairing a verbal strategy with a substantive one outperforms either in isolation [6], [7]. Prior work at this institution has traced a comparable gap between a displayed number and what it actually tracked at a residence-hall laundry machine [12], and a comparable private reliance a parent leans on in place of their own memory at the school gate [13]; the present study asks the same question of a number nobody at the institution ever intended to be checked.

Method

Recruitment and sample. We recruited riders who reported using a major rideshare app at least monthly, via an online research-panel intercept and snowball referral, without pre-screening for prior fare-estimate experience. Interviews continued until two independent readings of the transcripts agreed no new coping category or surprise-count pattern was emerging [14], [15], closing at 32 riders: 9 reporting no fare surprise they could recall, 11 reporting exactly one, and 12 reporting two or more.

Interview protocol. Each 30–45 minute semi-structured interview opened with a fixed prompt — "walk me through the last time you booked a ride, from opening the app to getting in the car" — to surface what, if anything, a participant did with the quoted price before booking. The interview then worked back through the participant's fare history for any trip where the final charge exceeded the quoted estimate, without suggesting in advance what counted as a "surprise." Roughly two-thirds of the way through each interview, the researcher disclosed that the app's terms of service label the upfront figure "an estimate only, not a guaranteed price," then re-asked what the participant expected to do with the number on their next booking.

Coding. Transcripts were coded following reflexive thematic analysis [16]. Two researchers independently open-coded a 25% subsample (8 interviews), reconciling divergent codes into a shared codebook covering the coping taxonomy, the surprise-count bucket, and the pre/post-disclosure comparison. One researcher then coded the remaining transcripts; the second independently re-coded a random 25% of the full set for a reliability check.

```
coping_category:
  screenshots_every_quote    # captures the estimate
before every booking
  mentally_rounds_up         # applies a private
margin before deciding
  checks_occasionally        # spot-checks the final
fare, not routinely
  doesnt_check_closely       # no longer attends to
the gap
surprise_count: none | one | two_or_more
disclosure_condition: pre | post
```

Listing 1: Closing categories from the reconciled codebook.

Reliability check. Restricting the coping taxonomy to the first coder's codes alone, rather than the reconciled double-coded subsample, changed each category's overall share by under three percentage points (Cohen's $\kappa = .84$ on the double-coded subsample); we report the reconciled coding throughout for completeness.

Results

Across 32 riders, coping behaviour tracked reported surprise history closely (Figure 1). Among the 9 riders reporting no fare surprise they could recall, the largest category was "doesn't check closely" (4 of 9): several described the quoted price as something they registered only in passing, with no reason yet to distrust it. Among the 11 riders reporting exactly one surprise, categories were more evenly spread, with "screenshots every quote" and "checks occasionally" each claiming roughly a third. Among the 12 riders reporting two or more surprises, "screenshots every quote" dominated (7 of 12): several participants described the habit unprompted as something they had adopted after a specific trip, usually naming the trip's approximate cost overrun from memory even when they could not recall the date.

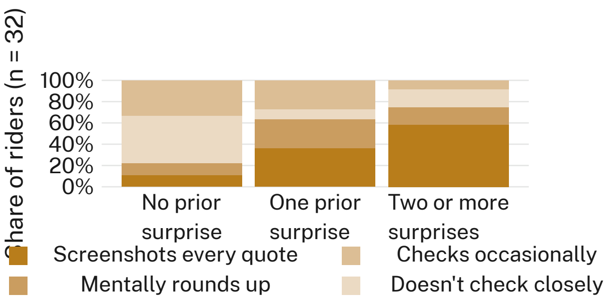

Figure 1: Riders who report two or more prior fare surprises are more than four times as likely to screenshot every quote as riders who report none.

The mid-interview disclosure — telling participants that the app's terms of service label the estimate "informational only" — left the distribution of stated future coping behaviour close to unchanged (Figure 2). Twelve riders named screenshotting as their likely behaviour before the disclosure and thirteen after; the other three categories moved by at most one rider each. Several participants who had already adopted the screenshot habit noted that the disclosure "explained" a practice they had arrived at independently rather than prompting a new one; none of the seven riders who reported not checking closely described the disclosure as a reason to start.

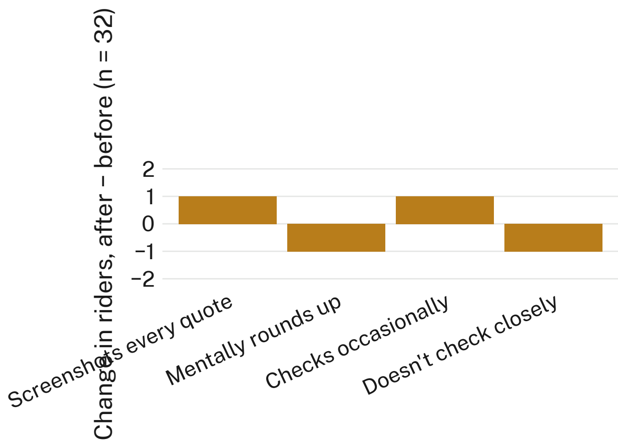

Figure 2: Told mid-interview that the estimate is labelled "informational only," riders' stated coping category shifts by no more than one rider, in either direction, per category (n = 32).

Riders' own accounts of why they settled on a particular coping behaviour rarely referenced the disclaimer language at all, whether or not they recalled having seen it. The most common justification offered for screenshotting was procedural rather than adversarial — "so I've got something to point to if it's way off," as one participant put it, rather than an expectation that a screenshot would routinely be needed. Riders who had stopped checking closely tended to frame the decision as proportionate rather than resigned: the gap, when it occurred, was rarely large enough in absolute terms to be worth the recurring attention, even among riders who could recall being frustrated by a specific instance.

Discussion

Riders in our sample did not respond to a fare estimate that had run high the way the trust-repair literature's institutional remedies assume a customer might: few described wanting compensation, an apology, or a corrected number, and none described the "informational only" framing as either a repair or a provocation. What most riders built instead was a small private instrument — a screenshot, a mental margin, a habit of not looking too closely — that ran alongside the platform's own figure rather than replacing it or demanding its correction. This sits comfortably alongside evidence that trust in an algorithmic forecaster erodes sharply after a single visible error [9], and beside a wider algorithmic-management literature that has, so far, mostly asked what a platform's design choices do to the worker on the other side of the interface rather than the rider on this one [10] — our riders simply built their own instrument for the interval no institutional repair was addressed to. It also extends this School's account of what a resident does when an institution's displayed number quietly stops tracking the reality behind it [12]: absent a mechanism to correct the institution's figure, the correction happens beside it, in a form the institution never sees and evidently does not need to.

Limitations

Our sample recruited current riders, which excludes anyone who abandoned the app entirely after a fare surprise large enough to end the relationship — if anything, our estimate of vigilant coping behaviour should be read as a floor, since the riders most burned by the gap are the ones least likely to be in a room describing it to us. A single mid-interview disclosure is not a field manipulation, and a participant's stated future intention is not the same as their behaviour on the next actual booking, though the near-identical distribution across 32 riders leaves little obvious room for a large effect to have gone undetected. Thirty-two riders is a modest base for the no-surprise category specifically; its "doesn't check closely" majority held up under the coding reliability check regardless.

Conclusion

Riders do not extend or withdraw trust in a fare estimate wholesale; they build small, private instruments that run alongside it, calibrated to their own history of when the number has run high before. Screenshotting, mental rounding, and simply not looking are not responses to what the estimate is labelled — they survive, in our sample, a disclosure that spells the labelling out in full. Telling a rider the number was never a guarantee changed almost nothing about what they intended to do with it, largely because most riders in our sample had already worked that out for themselves, one overrun trip at a time. Future work might test whether a rider's private instrument moves any more than this when the correction comes not from the terms of service but from another rider's account of the same gap.

References

[1] G. S. Black and K. Watson, "The Impact of Advertising Disclaimers (Fine Print) on Brand Attitudes," *Journal of Brand Management*, vol. 20, no. 5, pp. 380–395, 2013, doi: 10.1057/bm.2012.17.

[2] N. A. Saxena, W. Zhang, and C. Shahabi, "Unveiling and Mitigating Bias in Ride-Hailing Pricing for Equitable Policy Making," *AI and Ethics*, 2025, doi: 10.1007/s43681-024-00498-3.

[3] N. Garg and H. Nazerzadeh, "Driver Surge Pricing," *arXiv preprint arXiv:1905.07544*, 2019.

[4] U. Akhtar, F. Ghasemi, and R. Kucharski, "Optimizing Ride-Pooling Revenue: Pricing Strategies and Driver-Traveller Dynamics," *arXiv preprint arXiv:2403.13384*, 2024.

[5] N. Gillespie and G. Dietz, "Trust Repair After an Organization-Level Failure," *Academy of Management Review*, vol. 34, no. 1, pp. 127–145, 2009, doi: 10.5465/amr.2009.35713319.

[6] H.-D. Zhang, S.-C. Chen, and A. Ruangkanjanases, "Benefits First: Consumer Trust Repair in Mobile Commerce," *Journal of Theoretical and Applied Electronic Commerce Research*, vol. 16, no. 4, pp. 1071–1090, 2021, doi: 10.3390/jtaer16040061.

[7] B. Konopka, K. Hönemann, N. Yang, J. B. Thatcher, and M. Wiesche, "Selecting Trust Repair Strategies: The Case of Privacy Breaches in Smart Service Systems," *Electronic Markets*, 2026, doi: 10.1007/s12525-026-00887-y.

[8] P. Ashkrof, G. H. de A. Correia, O. Cats, and B. van Arem, "Ride Acceptance Behaviour of Ride-Sourcing Drivers," *arXiv preprint arXiv:2107.07864*, 2021.

[9] B. J. Dietvorst, J. P. Simmons, and C. Massey, "Algorithm Aversion: People Erroneously Avoid Algorithms After Seeing Them Err," *Journal of Experimental Psychology: General*, vol. 144, no. 1, pp. 114–126, 2015, doi: 10.1037/xge0000033.

[10] I. Kadolkar, S. Kepes, and M. Subramony, "Algorithmic Management in the Gig Economy: A Systematic Review and Research Integration," *Journal of Organizational Behavior*, 2025, doi: 10.1002/job.2831.

[11] C. B. Kirwan, D. K. Bjornn, B. B. Anderson, A. Vance, D. Eargle, and J. L. Jenkins, "Repetition of Computer Security Warnings Results in Differential Repetition Suppression Effects as Revealed With Functional MRI," *Frontiers in Psychology*, vol. 11, p. 528079, 2020, doi: 10.3389/fpsyg.2020.528079.

[12] C. Beng, M. Hanke, and R. Sabel, "Displayed, Not Metered: Auditing a Residence-Hall Laundry App's Credit Balance Against Machine-Metered Wash Cycles Across a Payment-System Migration," *Slop University technical report*, 2026, doi: 10.5555/slop.7xifby.

[13] C. Beng, M. Hanke, and I. Chikere, "Whoever Checked the Chat Last: Developing and Validating a Pick-Up Reliance Scale for the School-Gate Handover," *Slop University technical report*, 2026, doi: 10.5555/slop.efv97d.

[14] B. Saunders *et al.*, "Saturation in Qualitative Research: Exploring its Conceptualization and Operationalization," *Quality & Quantity*, vol. 52, no. 4, pp. 1893–1907, 2018, doi: 10.1007/s11135-017-0574-8.

[15] G. Guest, E. Namey, and M. Chen, "A Simple Method to Assess and Report Thematic Saturation in Qualitative Research," *PLOS ONE*, vol. 15, no. 5, p. e0232076, 2020, doi: 10.1371/journal.pone.0232076.

[16] V. Braun and V. Clarke, "Using Thematic Analysis in Psychology," *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006, doi: 10.1191/1478088706qp063oa.