



Freshly Worded, Still Adrift: A Taxonomy of Terms-of-Reference Drift After a University-Wide Charter Refresh

Osei Vandermeer¹, Anneke Tolan², Petra Umbile²

¹School of Continuous Improvement, Slop University ²School of Emergent Priorities, Slop University

doi:10.5555/slop.glpjfv

Abstract. University standing committees are typically governed by a terms-of-reference (ToR) charter that is revised rarely and read less. We report what happened when a shared AI drafting template regenerated the ToR of all 42 of Slop University's standing committees within a single administrative fortnight, holding committee membership and meeting cadence constant across a six-month pre-refresh and six-month post-refresh window. Coding each committee's charter against its logged meeting-agenda items with a sentence-embedding Charter-Practice Alignment Index (CPAI; two-rater agreement $\kappa = 0.79$), we derive a four-category taxonomy of charter-practice drift: Ceremonial, Residual, Colonising, and Self-Fulfilling. The refresh reduced mean charter length by 18% and raised an independent clarity rating, yet produced no significant shift in the taxonomy's category distribution or in mean CPAI ($\Delta = +2.1$ points, not significant); charter length itself is uncorrelated with CPAI ($r = 0.06$). A seventh of committees fall into the Self-Fulfilling category, where the charter predicts practice only because the committee treats the charter as its own meeting agenda. We conclude that a well-written charter is not evidence of an accurately governed committee, and propose CPAI as a lightweight complement to the review cadence the University already runs.

Introduction

University governance runs on charters: every standing committee, working group, and advisory board holds a terms-of-reference (ToR) document specifying its remit, membership, and reporting line. ToRs are revised infrequently, cited rarely, and — outside an audit or a dispute — read by almost nobody, a status the institutional-theory literature has long attributed to formal structure's role as legitimating myth rather than operating instruction [1], [2].

That status changed, briefly, when the University's committee-governance function rolled a shared AI drafting template across every one of its 42 standing committees within a single administrative fortnight, producing a uniformly well-worded ToR refresh where individually drafted, unevenly maintained documents had stood before. The refresh offered an

unusual natural comparison: membership, cadence, and reporting lines held constant across the rollout, while the charter text itself changed for everyone at once.

This paper asks the governance question the refresh raised and, we believe, has not been asked directly: does a charter that reads better govern any more accurately? We compare each committee's charter against its own logged practice, before and after the refresh, using a sentence-embedding alignment measure, and use the comparison to build a taxonomy of the ways a charter and a committee's actual business can relate. Our contributions are threefold:

- we introduce a four-category taxonomy of charter-practice drift — Ceremonial, Residual, Colonising, and Self-Fulfilling — derived from independently coded alignment between 42

committees' charters and their logged agendas, with substantial inter-rater agreement;

- we report a natural pre/post comparison across a university-wide charter-drafting intervention, and find the intervention changed the charters' prose without measurably changing the taxonomy's category distribution or mean alignment;
- we show charter length is uncorrelated with charter–practice alignment, a null result we read as evidence against the assumption that a longer, more detailed charter is a more accurately governing one.

Related work

Formal organizational structure has long been theorised as decoupled from operating routine: charters, policies, and mission statements function as legitimating myth as much as instruction [1], and organizations facing common pressures converge on the same structures — including shared boilerplate — through mimetic and normative isomorphism [2]. Bromley and Powell extend this to argue the more consequential and more common decoupling is not policy from practice but means from ends: an organization can enact its stated policy diligently and still fail to achieve what the policy was for [3]. Loosely coupled systems persist because the coupling failure is rarely tested [4], and an organization's legitimacy strategies operate largely independently of its technical accuracy [5]; institutionalization itself is an ongoing interplay of actors, actions, and the meanings attached to both, rather than a settled state [6]. Organizational documents are themselves a genre — a socially recognised response to a recurring situation [7] — and Orlikowski and Yates show that an organization's genre repertoire structures, and is structured by, the community that drafts and reads it [8]; a shared drafting template is one of the more direct ways a repertoire is imposed from outside.

Large language models are documented to homogenise the text they help produce, across both creative writing [9] and everyday cross-cultural composition [10], which motivates asking what a shared AI template does to a corpus of governance documents specifically. Text-as-data methods bring institutional documents within empirical reach [11], [12], following precedent in analysing meeting-generated text com-

putationally, from central-bank deliberation [13] to automated meeting summarisation [14], [15].

Method

Setting and corpus. We drew all 42 standing committees, boards, and advisory groups recorded in the University's committee register, spanning both schools and the cross-cutting professional units. For each committee we retrieved its terms-of-reference charter and its logged meeting-agenda items for the six months immediately before and the six months immediately after a university-wide charter refresh, in which a shared AI drafting template regenerated every committee's ToR within a single fortnight (mean revision turnaround 1.8 days per committee, human sign-off retained throughout). Committee membership, meeting cadence, and reporting line were unchanged across the refresh; only the charter text and, per our results, the charter's own word count and independently rated clarity moved.

Charter–Practice Alignment Index. Each charter was segmented into clauses and each meeting's agenda items were summarised into single-sentence topic statements; both were embedded with a sentence-transformer encoder [16]. A committee's Charter–Practice Alignment Index (CPAI) is the mean, over charter clauses, of that clause's maximum cosine similarity to any agenda-item embedding logged in the window, rescaled to 0–100:

$$\text{CPAI} = \frac{100}{n} \sum_{i=1}^n \max_j \cos(c_i, a_j)$$

with c_i the embedding of charter clause i (of n total) and a_j the embedding of logged agenda item j .

Taxonomy coding. Two raters independently coded each committee, in each window, into one of four categories — Ceremonial (charter and practice both diverge from a shared substantive core), Residual (practice tracks the pre-refresh charter only), Colonising (practice exceeds even the refreshed charter's scope), or Self-Fulfilling (practice tracks the charter because the charter is used as the standing meeting agenda) — using CPAI banding alongside a directionality read of which document or practice moved first, drawing on coding protocols piloted for the School of Emergent Priorities' Strategic Drift Survey. Raters agreed on category for 38 of 42

committees per window (Cohen's $\kappa = 0.79$ [17]); disagreements were resolved by discussion.

Comparisons. We tested the refresh's effect on category distribution with a chi-square test on the 4 (category) $\times$ 2 (period) table, on mean CPAI with a paired comparison across the same 42 committees, and on the charter-length-CPAI relationship with a Pearson correlation pooled across both windows.

Category	CPAI band	<i>n</i>	(post-refresh)
Ceremonial	0–34	13	
Residual	35–54	16	
Colonising	40–64	7	
Self-Fulfilling	75–95	6	

Table 1: Taxonomy categories, post-refresh CPAI band, and committee counts. Residual and Colonising bands overlap; category assignment there relies on the directionality read, not the index alone.

Results

Of the 42 committees, the post-refresh taxonomy placed 13 as Ceremonial, 16 as Residual, 7 as Colonising, and 6 as Self-Fulfilling (Table 1); the pre-refresh distribution was materially the same (12, 17, 6, 7). Figure 1 shows the two periods' category composition side by side: the refresh moved roughly one committee in ten across category boundaries, well within what coding noise alone would produce (chi-square test on the 4 $\times$ 2 table, n.s.).




Figure 1: Composition of the University's 42 standing committees by charter-practice-drift category, before and after the university-wide charter refresh: the four-category mix shifts by less than three percentage points in any category.

The taxonomy separates cleanly on CPAI. Figure 2 plots post-refresh CPAI by category. The

four categories occupy largely distinct bands, with Self-Fulfilling committees — who use their own charter as the standing meeting agenda — returning the highest alignment scores in the corpus (median CPAI 84) despite exhibiting, on independent review, the least committee-specific deliberation; Ceremonial committees return the lowest (median CPAI 24).




Figure 2: Charter-Practice Alignment Index by taxonomy category, post-refresh ($n = 42$): Self-Fulfilling committees post the highest median alignment (84) despite using their charter as a standing agenda rather than a specification independently arrived at.

The refresh changes the charter, not the drift.

Mean CPAI rose by only 2.1 points from pre- to post-refresh (paired comparison, n.s.), even though the same charters shrank by a mean 18% in word count and gained 0.6 points on a 5-point independent clarity rating administered to five professional-writing staff blind to committee identity; we retain the clarity-rating apparatus and report it in full because a null effect on governance accuracy alongside a real effect on prose quality is the paper's central finding, not a footnote to it.

Charter length buys no alignment. Figure 3 plots charter word count against CPAI across both windows pooled (84 committee-period observations); the fitted relationship is flat (slope 0.01, $R^2 = 0.004$), indicating that committees with longer, more elaborated charters are not thereby more accurately governed by them.




Figure 3: Charter word count against Charter–Practice Alignment Index, pooled across both windows ($n = 84$ committee-periods): a flat fitted relationship ($R^2 = 0.004$) indicates charter length does not predict alignment.

Discussion

The refresh’s null effect on CPAI and on category distribution, set against its real effect on charter prose, is consistent with formal structure functioning as legitimating myth rather than operating instruction [1], and with decoupling running between means and ends rather than merely between policy and practice [3]: the University’s committees did not become better governed by acquiring better-written charters, because the charter was never the operating instruction the refresh implicitly treated it as. The Self-Fulfilling category is the sharper finding: a committee that uses its own charter as its meeting agenda will always score well on any alignment measure taken from the charter’s own clauses, a circularity that the University’s existing review cadence — which audits committees against their charters — cannot detect and may be quietly rewarding. Whether a template-driven refresh could ever close charter–practice drift, rather than merely re-describing it more legibly, is a question this comparison was not designed to answer; it would require an intervention that changed logged practice rather than logged prose.

Limitations

Category assignment relies on a qualitative directionality read alongside the CPAI band, and Residual and Colonising committees occupy overlapping CPAI ranges (Table 1); a purely quantitative classifier would likely conflate them, though the two raters’ 0.79 kappa suggests the combined read is reproducible. Our corpus is confined to one institution’s 42 committees over a single twelve-month window spanning one refresh event, though the consistency of the null pre/post effect across both

schools and across professional and academic committees alike suggests the finding is not an artefact of any one committee population. The clarity-rating panel was blind to committee identity but not to refresh status, a design constraint that could inflate the measured prose-quality gain, if anything sharpening the contrast with the null governance effect.

Conclusion

Across 42 standing committees, a university-wide AI-drafted charter refresh measurably improved charter prose and left charter–practice alignment where it started, in aggregate and within our four-category taxonomy alike. A seventh of committees maintain high alignment only by treating their own charter as their meeting agenda, a self-fulfilling pattern the University’s existing charter-audit cadence cannot distinguish from the genuine article. We conclude that a well-written charter is not evidence of an accurately governed committee, and propose the Charter–Practice Alignment Index as a lightweight complement — not a replacement — for the qualitative review the University’s committees already undergo. Future work will track whether CPAI, once measured and reported back to a committee, changes the practice it was meant only to describe.

References

- [1] J. W. Meyer and B. Rowan, “Institutionalized Organizations: Formal Structure as Myth and Ceremony,” *American Journal of Sociology*, vol. 83, no. 2, pp. 340–363, 1977, doi: [10.1086/226550](https://doi.org/10.1086/226550).
- [2] P. J. DiMaggio and W. W. Powell, “The Iron Cage Revisited: Institutional Isomorphism and Collective Rationality in Organizational Fields,” *American Sociological Review*, vol. 48, no. 2, pp. 147–160, 1983, doi: [10.2307/2095101](https://doi.org/10.2307/2095101).
- [3] P. Bromley and W. W. Powell, “From Smoke and Mirrors to Walking the Talk: Decoupling in the Contemporary World,” *Academy of Management Annals*, vol. 6, no. 1, pp. 483–530, 2012, doi: [10.5465/19416520.2012.684462](https://doi.org/10.5465/19416520.2012.684462).
- [4] K. E. Weick, “Educational Organizations as Loosely Coupled Systems,” *Administrative Science Quarterly*, vol. 21, no. 1, pp. 1–19, 1976, doi: [10.2307/2391875](https://doi.org/10.2307/2391875).

- [5] M. C. Suchman, "Managing Legitimacy: Strategic and Institutional Approaches," *Academy of Management Review*, vol. 20, no. 3, pp. 571–610, 1995, doi: [10.5465/amr.1995.9508080331](https://doi.org/10.5465/amr.1995.9508080331).
- [6] T. B. Zilber, "Institutionalization as an Interplay between Actions, Meanings, and Actors: The Case of a Rape Crisis Center in Israel," *Academy of Management Journal*, vol. 45, no. 1, pp. 234–254, 2002, doi: [10.2307/3069294](https://doi.org/10.2307/3069294).
- [7] C. R. Miller, "Genre as Social Action," *Quarterly Journal of Speech*, vol. 70, no. 2, pp. 151–167, 1984, doi: [10.1080/00335638409383686](https://doi.org/10.1080/00335638409383686).
- [8] W. J. Orlikowski and J. Yates, "Genre Repertoire: The Structuring of Communicative Practices in Organizations," *Administrative Science Quarterly*, vol. 39, no. 4, pp. 541–574, 1994, doi: [10.2307/2393771](https://doi.org/10.2307/2393771).
- [9] B. R. Anderson, J. H. Shah, and M. Kreminski, "Homogenization Effects of Large Language Models on Human Creative Ideation," in *Proceedings of the 16th Conference on Creativity & Cognition*, 2024, pp. 413–425. doi: [10.1145/3635636.3656204](https://doi.org/10.1145/3635636.3656204).
- [10] D. Agarwal, M. Naaman, and A. Vashistha, "AI Suggestions Homogenize Writing Toward Western Styles and Diminish Cultural Nuances," *arXiv preprint arXiv:2409.11360*, 2024.
- [11] J. Grimmer and B. M. Stewart, "Text as Data: The Promise and Pitfalls of Automatic Content Analysis Methods for Political Texts," *Political Analysis*, vol. 21, no. 3, pp. 267–297, 2013, doi: [10.1093/pan/mps028](https://doi.org/10.1093/pan/mps028).
- [12] Z. Jin and R. Mihalcea, "Natural Language Processing for Policymaking," *arXiv preprint arXiv:2302.03490*, 2023.
- [13] S. Hansen, M. McMahon, and A. Prat, "Transparency and Deliberation Within the FOMC: A Computational Linguistics Approach," *The Quarterly Journal of Economics*, vol. 133, no. 2, pp. 801–870, 2018, doi: [10.1093/qje/qjx045](https://doi.org/10.1093/qje/qjx045).
- [14] M. Zhong et al., "QMSum: A New Benchmark for Query-based Multi-domain Meeting Summarization," *arXiv preprint arXiv:2104.05938*, 2021.
- [15] H. G. Ramachandran and D. DeRose, "A Text Analysis of Federal Reserve Meeting Minutes," *arXiv preprint arXiv:1805.07851*, 2018.
- [16] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks," *arXiv preprint arXiv:1908.10084*, 2019.
- [17] J. Cohen, "A Coefficient of Agreement for Nominal Scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960, doi: [10.1177/001316446002000104](https://doi.org/10.1177/001316446002000104).