



Confident First, Correct Later

Validating the Welcome Desk's "Resolved on First Ask" Log Against a Same-Day Survey of Where Enquirers Actually Ended Up

Lindiwe Achterberg, Anneke Tolan

School of Emergent Priorities, Slop University

doi:10.5555/slop.1jtnco

Abstract. Frontline information points commonly log their own resolution outcomes, and those self-logs often stand in for service-quality evidence without independent verification. We audit fourteen weeks of the welcome desk's "resolved on first ask" log against an independently sampled same-day survey ($n = 618$) of enquirers, tracking where each respondent reported actually needing to go after leaving the desk. Agreement between the self-log and the survey-confirmed outcome is slight (Cohen's kappa = 0.09), and the gap holds despite a mid-deployment script revision instructing ambassadors to escalate uncertain queries rather than answer confidently: the self-logged resolution rate rises by 2.4 percentage points after the revision (n.s.), and the survey-confirmed rate by 2.3 (n.s.). The composition of survey-confirmed outcomes — correct on the first stop, one redirect, two or more redirects — does not differ meaningfully between new-starter and returning ambassadors, suggesting the gap tracks the desk's confident-answer script rather than individual experience. We discuss implications for services that report their own resolution logs without an independent check, and for scripting policies that substitute a confident guess for an acknowledged "I'm not sure."

Introduction

Service points that report their own resolution outcomes are common across university administration — a helpdesk ticket closed as "resolved", a call logged as "handled", an enquiry marked complete by the person who took it. The convenience is obvious: the log is free, timely, and already collected as a byproduct of the interaction. The cost is less obvious, and rarely audited: a self-reported resolution flag can drift from the resolution it is meant to represent, particularly under an instruction that discourages the alternative ("I'm not sure, let me check") in favour of a confident answer on the spot.

Slop University's campus welcome desk, staffed by student ambassadors, adopted exactly this instruction at the start of the current term: enquirers are to receive a confident, immediate answer, and the ambassador logs the enquiry "resolved on first ask" once the answer is given, regardless of whether the enquirer later needed

to ask elsewhere. The desk's own resolution log is the figure Student Services cites when describing the service's performance. Whether that figure tracks the outcome it names — did the enquirer actually end up where they needed to be — has not, to our knowledge, been checked against anything the desk did not itself produce.

We audit the log against an independently sampled same-day survey of enquirers over a fourteen-week deployment, and we treat a mid-term script revision — instructing ambassadors to escalate uncertain queries rather than guess — as a natural test of whether the gap between the log and the outcome it claims to track can be narrowed by policy alone.

Our contributions are threefold. First, we report the level of agreement between the desk's self-logged resolution flag and an independently surveyed outcome across 618 enquiries, using an agreement statistic that does not credit chance overlap. Second, we evaluate whether

the escalate-if-unsure script revision shifts either figure, and find that it moves both by less than three percentage points. Third, we show that the composition of survey-confirmed outcomes does not differ meaningfully between new-starter and returning ambassadors, which we take as tentative evidence that the gap tracks the desk’s answering policy rather than any one cohort’s experience.

Related work

Self-reported completion figures have a long history of drifting from independently observed outcomes in frontline service settings; hand-hygiene compliance is the best-studied case, with self-report consistently outperforming direct observation by a wide margin [1] and product-use monitoring offering no cleaner a proxy [2]. Scripting the service encounter is itself a deliberate managerial choice, trading a discretionary, uncertain response for a standardised one [3], and the discretion an employee retains within that script shapes what actually gets delivered [4] — the same discretion that front-line-bureaucracy research treats as the object to be modelled and predicted, not scripted away [5]. On the advice-giving side, a confident answer changes how much a recipient trusts it independently of whether the answer is correct [6], and a synthesis of the advice-taking literature finds confidence expressed by the advisor a stronger driver of uptake than the advice’s demonstrated accuracy [7]; a recent finding in human-AI advice settings shows exposure to a confident answer measurably suppresses a respondent’s own willingness to admit uncertainty, even when the advice is wrong [8]. Self-assessment research offers a caution in the other direction: people who most need to defer are often least able to recognise it [9], though expertise offers only partial protection against overclaiming [10]. At the institutional level, a metric that stands in for the outcome it was designed to track tends to be optimised on its own terms once it is adopted [11], a pattern documented well beyond any one domain [12], [13], with reactive effects that outlast the original justification for measuring at all [14], [15]. Prior work at this institution has audited a share-scooter corral’s occupancy sensor against a manual headcount [16] and a lobby signage screen’s dwell-time sensor against passersby’s actual recall [17], in both cases finding the self-reporting instrument

tracked something other than the outcome it was credited with measuring.

Method

Setting. The welcome desk is staffed by student ambassadors on rotating shifts and fields walk-up enquiries from new and continuing students about where to go for a service, form, or office. At the start of the current term, Student Services issued ambassadors a standing instruction to answer every enquiry with a confident, specific direction rather than defer or admit uncertainty; each enquiry is logged in the desk’s queue system with a binary “resolved on first ask” flag set by the ambassador at the close of the interaction.

Independent measure. Over the fourteen-week term, research assistants intercepted a same-day, stratified sample of enquirers as they left the desk ($n = 618$ of an estimated 2,114 logged enquiries, sampled across all shifts and days of the week) and sent a two-question same-day follow-up: whether the enquirer reached the office or service they needed without a further redirect, and if not, how many additional stops it took. Each response was coded into one of three outcome categories — correct on the first stop, one redirect, two or more redirects — with “correct on the first stop” defined as survey-confirmed correct routing.

Ablation. In week 8, Student Services revised the desk’s script to instruct ambassadors to escalate a query to a supervisor or offer a “let me check and follow up” response when genuinely unsure, rather than commit to a guess. We treat weeks 1–7 and weeks 8–14 as pre-and post-revision periods and compare both the self-logged and survey-confirmed rates across the split.

Agreement. We cross-tabulate the desk’s binary resolution flag against the binary survey-confirmed outcome (correct on first stop vs. not) for the full matched sample and compute Cohen’s kappa, which discounts the agreement expected by chance alone:

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

where p_o is the observed proportion of agreement and p_e is the proportion expected under independence, computed from the row and column marginals [18]. We report kappa alongside the raw agreement rate throughout, following

standard practice for a proportion-based benchmark that chance alone can inflate [19], [20].

Cohort split. Ambassadors were grouped into new starters (recruited for the current term, $n = 9$) and returning ambassadors (employed at the desk for one or more prior terms, $n = 14$); we compare the composition of survey-confirmed outcomes across the two groups.

Results

Across the fourteen weeks sampled, the desk's self-logged "resolved on first ask" rate holds steady between 87% and 93%, while the survey-confirmed correct-routing rate tracks 26 to 31 percentage points lower in every week (Figure 1) — the two series never converge, and the gap does not narrow across the term.




Figure 1: Desk-logged "resolved on first ask" holds 26-31 percentage points above the survey-confirmed correct-routing rate in every one of the fourteen weeks sampled.

On the full matched sample ($n = 618$), 556 enquiries (90.0%) were logged resolved and 383 (62.0%) were survey-confirmed correct on the first stop; observed agreement between the two measures is 63.3%, against a chance-expected agreement of 59.6%, for Cohen's kappa = 0.09 — slight agreement by conventional benchmarks [19], indistinguishable in practice from the desk's log carrying no information about where the enquirer actually ended up.

The week-8 script revision, which instructed ambassadors to escalate uncertain queries rather than guess, produced no material change on either measure (Figure 2): the self-logged rate rises from a pre-revision average of 89.1% to a post-revision average of 91.6% (n.s.), and the survey-confirmed rate from 60.9% to 63.1%

(n.s.), a shift on both measures that sits within the week-to-week variability already present before the revision.




Figure 2: The week-8 escalate-if-unsure script revision moves the self-logged rate by 2.4 points and the survey-confirmed rate by 2.3 points, neither reaching significance (n.s.).

Finally, the composition of survey-confirmed outcomes — correct on the first stop, one redirect, two or more redirects — does not separate meaningfully by ambassador cohort (Figure 3): new-starter ambassadors resolve 63% of enquiries correctly on the first stop against 65% for returning ambassadors, with the remaining categories differing by no more than a single percentage point either way. If the desk's confident-answer script were being compensated for by ambassador experience, we would expect the returning cohort's composition to diverge further from the new starters'; it does not.




Figure 3: Survey-confirmed outcome composition differs by at most 2 percentage points per category between new-starter and returning ambassadors.

Discussion

The desk's self-logged resolution rate and the survey-confirmed outcome it is taken to represent share almost no signal beyond what chance would produce, and the gap persists across a script revision explicitly designed to close it. This is broadly consistent with the wider pattern in which a self-reported performance figure, once adopted as the record of an outcome,

comes to be optimised on its own terms rather than the terms of the outcome it names [11], [12]. The absence of a cohort effect suggests the gap is a property of the instruction ambassadors are working under, not of any one ambassador's skill or tenure — consistent with scripting research that treats the standardised response, not the individual employee, as the unit that shapes the outcome [3]. We do not find evidence that the escalate-if-unsure revision made things worse; we also do not find evidence that it made things measurably better, which is itself a useful, if modest, finding for a service considering whether a script change alone is sufficient to close a validity gap of this size.

Limitations

Our survey sample, while stratified across shifts and days, captures 29% of the term's logged enquiries and cannot rule out a systematic difference between enquirers who responded to the same-day follow-up and those who did not; the consistency of the gap across all fourteen weeks, however, suggests the finding is not an artefact of a single sampling window. The study covers one welcome desk over one term, which limits generalisation to other frontline points on campus, though the size and stability of the log-survey gap leaves little room for a materially different result under the same scripting policy. We did not have access to enquirers' subsequent help-seeking behaviour beyond the same-day window, so a small number of "two or more redirects" cases may understate eventual resolution occurring after our observation period closed.

Conclusion

A resolution log completed by the same person being asked to resolve the query is, on this evidence, a weak proxy for whether the query was actually resolved, and a script change aimed at the underlying behaviour did not close the gap within a single term. Services that cite a self-logged resolution rate as evidence of performance would benefit from the same independent check we ran here, ideally embedded as routine practice rather than a one-off audit. Future work might test whether a coarser log — resolved, escalated, or unsure — recovers more signal than the current binary flag, and whether the gap we observe at one welcome desk holds at the other enquiry points across campus that report a similar figure.

References

- [1] E. L. Larson, A. E. Aiello, and J. P. Cimiotti, "Assessing Nurses' Hand Hygiene Practices by Direct Observation or Self-Report," *Journal of Nursing Measurement*, 2004, doi: [10.1891/jnum.12.1.77.66072](https://doi.org/10.1891/jnum.12.1.77.66072).
- [2] A. M. Pettis, "A Tarnished Gold Standard?: Direct Observation vs. Monitoring Product Use for Hand Hygiene Compliance," *American Journal of Infection Control*, 2013, doi: [10.1016/j.ajic.2013.03.087](https://doi.org/10.1016/j.ajic.2013.03.087).
- [3] D. A. Tansik and W. L. Smith, "Scripting the Service Encounter," *New Service Development: Creating Memorable Experiences*, 2000, doi: [10.4135/9781452205564.n11](https://doi.org/10.4135/9781452205564.n11).
- [4] S. W. Kelley, "Discretion and the Service Employee," *Journal of Retailing*, 1993, doi: [10.1016/s0022-4359\(05\)80005-3](https://doi.org/10.1016/s0022-4359(05)80005-3).
- [5] G. Pokharel, S. Das, and P. J. Fowler, "Discretionary Trees: Understanding Street-Level Bureaucracy via Machine Learning," *arXiv preprint arXiv:2312.10694*, 2023.
- [6] N. Pescetelli and N. Yeung, "The Role of Decision Confidence in Advice-Taking and Trust Formation," *Journal of Experimental Psychology: General*, 2021, doi: [10.1037/xge0000960](https://doi.org/10.1037/xge0000960).
- [7] S. Bonaccio and R. S. Dalal, "Advice Taking and Decision-Making: An Integrative Literature Review, and Implications for the Organizational Sciences," *Organizational Behavior and Human Decision Processes*, 2006, doi: [10.1016/j.obhdp.2006.07.001](https://doi.org/10.1016/j.obhdp.2006.07.001).
- [8] C. Marcoccia, W. Quattrociocchi, and V. Capraro, "AI Advice Suppresses People's Willingness to Say "I Don't Know", Even When the Advice Is Wrong and Accuracy Is Incentivized," *arXiv preprint arXiv:2607.13562*, 2026.
- [9] J. Kruger and D. Dunning, "Unskilled and Unaware of It: How Difficulties in Recognizing One's Own Incompetence Lead to Inflated Self-Assessments," *Journal of Personality and Social Psychology*, 1999, doi: [10.1037/0022-3514.77.6.1121](https://doi.org/10.1037/0022-3514.77.6.1121).
- [10] S. Atir, E. Rosenzweig, and D. Dunning, "Does Expertise Protect Against Overclaiming False Knowledge?," *Organizational Behavior and Human De-*

- cision Processes, 2024, doi: [10.1016/j.obhdp.2024.104354](https://doi.org/10.1016/j.obhdp.2024.104354).
- [11] D. T. Campbell, “Assessing the Impact of Planned Social Change,” *Evaluation and Program Planning*, 1979, doi: [10.1016/0149-7189\(79\)90048-X](https://doi.org/10.1016/0149-7189(79)90048-X).
 - [12] D. Manheim and S. Garrabrant, “Categorizing Variants of Goodhart's Law,” *arXiv preprint arXiv:1803.04585*, 2018.
 - [13] M. Fire and C. Guestrin, “Over-Optimization of Academic Publishing Metrics: Observing Goodhart's Law in Action,” *arXiv preprint arXiv:1809.07841*, 2018.
 - [14] W. N. Espeland and M. Sauder, “Rankings and Reactivity: How Public Measures Recreate Social Worlds,” *American Journal of Sociology*, 2007, doi: [10.1086/517897](https://doi.org/10.1086/517897).
 - [15] R. K. Merton, “The Unanticipated Consequences of Purposive Social Action,” *American Sociological Review*, 1936, doi: [10.2307/2084615](https://doi.org/10.2307/2084615).
 - [16] L. Achterberg and C. Beng, “Validating the Corral Occupancy Sensor: A Manual Headcount Against the Count Setting Next Term's Bike-Rack Capacity,” *Slop University technical report*, 2026, doi: [10.5555/slop.gdacz4](https://doi.org/10.5555/slop.gdacz4).
 - [17] L. Achterberg and I. Vasseur, “Slower Isn't Seen: A Lobby Signage Screen's Dwell-Time Sensor, Tested Against What Passersby Actually Recall,” *Slop University technical report*, 2026, doi: [10.5555/slop.63i907](https://doi.org/10.5555/slop.63i907).
 - [18] J. Cohen, “A Coefficient of Agreement for Nominal Scales,” *Educational and Psychological Measurement*, 1960, doi: [10.1177/001316446002000104](https://doi.org/10.1177/001316446002000104).
 - [19] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, 1977, doi: [10.2307/2529310](https://doi.org/10.2307/2529310).
 - [20] T. Rüz, “Inter-Rater Reliability is Individual Fairness,” *arXiv preprint arXiv:2308.05458*, 2023.