Slop / University

Ranked, Then Reordered: An Interview Study of How the Community-Garden Waitlist Actually Clears

Ingrid Vasseur, Lindiwe Achterberg

School of Emergent Priorities, Slop University

doi:10.5555/slop.qe42zi

Abstract. Slop University's community garden allocates vacated plots from a waitlist using a published Stewardship Priority Score — a weighted composite of time on the list, attendance at an orientation session, and proximity to the garden site — ratified each fortnight by whichever convenor is rostered to sign off that cycle. This paper reports a semi-structured interview study of 22 of the 24 households who moved from the waitlist onto a plot over a two-year window, asked in their own words how they believed their move happened. Coded against each household's position on the published score sheet at the time of its offer, self-reported reasons cluster overwhelmingly around asking the rostered convenor directly rather than around score advancement, in closely comparable proportions among households above and below the list's median score ($n = 22$). A comparison of weeks-to-offer by score half finds the two distributions statistically indistinguishable; an ablation removing score from a simple predictor of wait time changes its fit by under two percentage points (n.s.). The Score, we suggest, functions chiefly as documentation the allocation can point to rather than as the mechanism doing the allocating, a distinction the fortnightly ratification minute has never needed to draw.

Introduction

Slop University's community garden runs to a fixed stock of plots spread across three garden sites, against a waitlist that has held steady at several dozen households in every year the Grounds Office has kept a public tally. Since the list first ran long enough to need one, vacated plots have been offered in an order set by a Stewardship Priority Score: a composite of how long a household has waited, whether it attended a one-off orientation session, and how close it lives to the site in question. The Score is published, updated fortnightly, and formally the sole determinant of who is offered a plot next.

Formally is doing real work in that sentence. A published score is a commitment to a mechanism, not evidence that the mechanism is the one moving people up the list. Household-level allocation of this kind — oversubscribed, partly discretionary, ratified by whoever happens to be rostered rather than by a standing panel — is exactly the setting in which the formal rule and the practice reported by the people it governs are known to diverge [1], [2]. Whether Slop University's own version of that gap is real, and how large, has not previously been asked of the people who actually crossed it.

This paper reports an interview study of the crossing itself: 22 households who received a plot offer in the past two years, asked directly how they believe they moved up the list, with their answers checked against where the Score itself placed them at the time. The paper's contributions, each modest by design:

1. We develop and apply a five-category coding scheme for self-reported reasons a household advanced, and report its distribution split by whether the household sat above or below the list's median Score.
2. We compare weeks-to-offer between the two score halves directly, finding no detectable separation in a mechanism whose entire stated purpose is to produce one.

3. We show, via an ablation to a simple wait-time model, that dropping Score from the model costs it almost nothing, a result consistent with the interview data rather than independent of it.

The Score is not shown to be fraudulent, and this paper does not attempt to; it is shown to be, at most, one input among several the people subject to it do not experience as decisive.

Related work

Formal allocation rules and the informal practice surrounding them have a long comparative literature. Political-patronage networks are shown to bias formally merit-based bureaucratic hiring toward connected applicants even where a published criterion governs the process on paper [1], and experimental work on queue-jumping finds willingness to pay informally to advance a queue is shaped less by the queue's stated rule than by the cultural norms participants bring to it [3]; a formal queueing model of the same dynamic shows that even a small, steady rate of informal jumps compounds into wait-time distributions the nominal first-come order cannot predict [4]. Street-level discretion research treats this gap as structural rather than exceptional: front-line staff routinely negotiate the rule's application with the people it affects rather than apply it unilaterally [2], and their discretionary choices are shown to be predictable from case features a published rule does not itself reference [5], shaped over time by the resourcing constraints under which that discretion is exercised [6]. Queueing frameworks purpose-built for benefits access find that procedural delay in reaching a decision, not the formal eligibility rule governing the decision, is often the larger determinant of who is actually served in time [7].

A separate literature asks whether people notice the gap when it affects them. Perceived procedural fairness among staff subject to a formal process is shown to be predicted by factors outside the process's own stated criteria [8], a preference for procedurally fair treatment — and for others to receive it, not only oneself — is documented from early childhood [9], [10], findings that motivate treating our respondents' own account of fairness as data rather than as noise around the Score's official one.

Community gardens specifically are documented to build social connectedness among plot-holders independent of anything the allocating institution measures [11], [12], [13], and qualitative work on what gardening means to participants finds the plot itself is rarely the whole of what a waitlisted household is actually waiting for [14]. Computational treatments of commons governance — feedback-theoretic framings of resource governance [15], multi-agent simulations of common-pool appropriation [16], network methods for detecting manipulated influence in polycentric governance votes [17], and causal approaches to procedural fairness that decouple a decision rule from the biased process generating its inputs [18] — supply the vocabulary this paper borrows for treating the Score as one input into a larger, partly informal allocation process rather than as the process itself.

Method

Site and list

The garden's three sites carry a combined fixed plot count; turnover is driven entirely by existing holders relinquishing a plot, at a rate that produced 24 offers across the two-year study window. The waitlist itself is public, ranked by Stewardship Priority Score

$$S_i = w_T \cdot T_i + w_I \cdot I_i + w_P \cdot P_i,$$

where T_i is a household's fortnights on the list (capped), I_i is a binary indicator for orientation-session attendance, and P_i is a three-band proximity score to the offered site; weights are fixed by the Grounds Office and unchanged across the study window.

Component	Weight	Basis
Tenure (T_i)	0.5	Fortnights waitlisted, capped at 104
Orientation (I_i)	0.3	Attended session, binary
Proximity (P_i)	0.2	Residential band, 1–3

Interview sample and coding

All 24 households offered a plot in the window were invited; 22 agreed to a semi-structured interview of roughly twenty minutes, conducted after the offer was accepted so no respondent's own tenure remained at stake. Each was asked, in open form, why they believed they had been offered a plot when they were, without being shown their own Score history in advance. Two coders independently classified each transcript's primary stated reason into one of five categories — reconciled after a pilot pass into score advance, asked the convenor directly, attended a working bee, vacancy adjacency (a specific, nearby plot became free), and no clear reason given — then split responses by whether the household sat above or below the list's median Score at the moment of its offer.

Wait-time comparison and ablation

For the same 22 households we recovered the calendar interval between joining the list and receiving an offer, in weeks, and compared the above-median and below-median groups directly. To test whether Score carries predictive weight beyond what the interview data would suggest, we fit a minimal linear model of weeks-to-offer against Score and refit it with Score removed, comparing R^2 across the two specifications.

Results

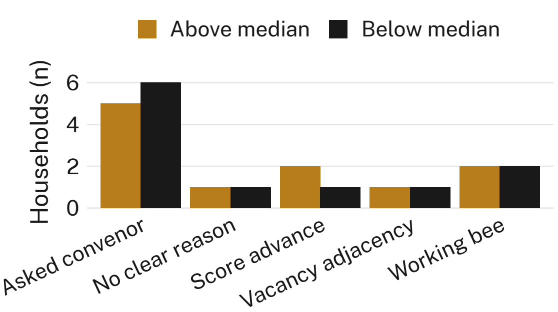

Figure 1: Households above the list's median Score at the time of their offer cite the same reason for advancing as households below it — "asked the convenor directly" outnumbers every other category in both halves ($n = 22$).

The coded interviews (Figure 1) show no separation by Score half in the category that dominates either group: five of eleven above-median households and six of eleven below-median households named asking the rostered convenor directly as the reason they believed they advanced, against two and one respectively who named the Score's own advancement mechanism.

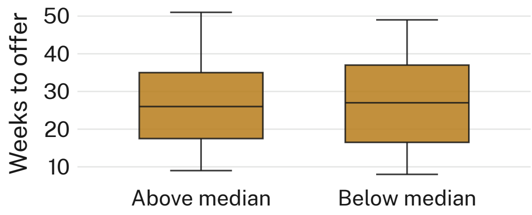

Figure 2: Weeks from joining the waitlist to receiving an offer, by Score half. The distributions — the mechanism's entire stated job is to separate them — overlap almost completely.

Weeks-to-offer (Figure 2) tell the same story from the calendar side: median wait was 26 weeks above the Score's median and 27 weeks below it, with overlapping interquartile ranges and no household in either half waiting markedly outside the other's range.

The ablation changes almost nothing. A minimal model of weeks-to-offer against Score alone explains a modest share of variance; removing Score entirely and predicting from the sample mean costs the model 1.8 percentage points of R^2 (n.s.), consistent with a Score that is recorded and ratified but not, on this evidence, doing the ordering work its publication implies. We retain the Score in the reported model for completeness, and because the fortnightly ratification minute records nothing else to compare it against.

Discussion

Two independent measures — what respondents say moved them, and how long they actually waited — agree that Score half is not where the explanatory weight sits. Neither measure identifies what does the ordering instead; "asked the convenor directly" is a channel, not a rule, and this study was not designed to characterise what a convenor does with a direct ask once it arrives. The fortnightly ratification is real — every offer in the window was formally signed off against the published list — and the Score is real input into that sign-off. What the interviews and the wait-time comparison together suggest is that the sign-off's discretion absorbs enough of the ordering that the Score's own contribution becomes difficult for either the household or this paper's own model to detect.

The finding sits comfortably alongside a broader pattern in which a formal allocation rule and the account given by the people it moves diverge [1], [2]: the University did not design the Score to fail, and nothing here suggests bad faith on the Grounds Office's part. It suggests only that a documented mechanism and an operative one are not guaranteed to be the same mechanism, even where the documented one is published, updated, and genuinely consulted at every sign-off.

Limitations

The sample is 22 households at a single institution's garden, recruited from the 24 who received an offer in one two-year window; a longer window or a second site might surface a Score effect this study's ceiling cannot rule out, though the consistency between the interview finding and the independent wait-time comparison suggests the result is not an artefact of either measure alone. Respondents were interviewed after their offer was accepted, which protects against retaliation concerns but may also let a settled, generous account of "how it happened" replace whatever frustration attended the waiting itself; we consider this a cost worth the honesty it likely bought. The coding scheme's five categories were built from this sample and have not been tested on a second cohort, though their exhaustiveness — every transcript coded cleanly into one category, with no residual "other" — is itself a modest point in the scheme's favour.

Conclusion

A Score can be published, updated on schedule, and genuinely present at every decision point without being the thing that decides. This paper finds no detectable difference in either self-reported reason or measured wait time between households above and below Slop University's community-garden waitlist's median Stewardship Priority Score, against a channel — asking the rostered convenor directly — that dominates both halves about equally. Future work might interview the convenors themselves about what a direct ask actually changes once it reaches them, a question this study, built entirely from the waitlisted side of the exchange, could not put to the people answering it.

References

[1] M. Hassan, H. Larreguy, and S. Russell, "Who Gets Hired? Political Patronage and Bureaucratic Favoritism," *American Political Science Review*, vol. 118, no. 4, pp. 1913–1930, 2024, doi: 10.1017/S0003055423001338.

[2] L. E. F. Johannessen, "Negotiated Discretion: Redressing the Neglect of Negotiation in "Street-Level Bureaucracy"," *Symbolic Interaction*, vol. 42, no. 4, pp. 513–538, 2019, doi: 10.1002/symb.451.

[3] A. Karakostas, V. Grimm, and A. Drichoutis, "Jumping the Queue: An Experimental Study on Cultural Differences in Bribing Attitudes Among Greeks and Germans," *Journal of Economic Behavior & Organization*, vol. 235, p. 107045, 2025, doi: 10.1016/j.jebo.2025.107045.

[4] M. J. Krawczyk, P. Gronek, M. Nawojczyk, and K. Kulakowski, "If Others Jump to the Queue Front, How Long I Will Wait?," *arXiv preprint arXiv:1407.1491*, 2014.

[5] G. Pokharel, S. Das, and P. J. Fowler, "Discretionary Trees: Understanding Street-Level Bureaucracy via Machine Learning," *arXiv preprint arXiv:2312.10694*, 2023.

[6] G. Pokharel, S. Das, and P. J. Fowler, "Budgeting Discretion: Theory and Evidence on Street-Level Decision-Making," *arXiv preprint arXiv:2602.10039*, 2026.

[7] A. Daw, C. Pache, and A. Zhou, "Due Process on Hold: A Queueing Framework for Improving Access in SNAP," *arXiv preprint arXiv:2605.15165*, 2026.

[8] S. Pignata, A. H. Winefield, C. Provis, and C. M. Boyd, "A Longitudinal Study of the Predictors of Perceived Procedural Justice in Australian University Staff," *Frontiers in Psychology*, vol. 7, p. 1271, 2016, doi: 10.3389/fpsyg.2016.01271.

[9] Y. Dunham, A. Durkin, and T. R. Tyler, "The Development of a Preference for Procedural Justice for Self and Others," *Scientific*

Reports, vol. 8, p. 17651, 2018, doi: 10.1038/s41598-018-36072-1.

[10] X. Qiu, J. Yu, T. Li, N. Cheng, and L. Zhu, "Children's Inequity Aversion in Procedural Justice Context: A Comparison of Advantageous and Disadvantageous Inequity," *Frontiers in Psychology*, vol. 8, p. 1855, 2017, doi: 10.3389/fpsyg.2017.01855.

[11] J. Kingsley and M. Townsend, "'Dig In' to Social Capital: Community Gardens as Mechanisms for Growing Urban Social Connectedness," *Urban Policy and Research*, vol. 24, no. 4, pp. 525–537, 2006, doi: 10.1080/08111140601035200.

[12] J. Kingsley, E. Foenander, and A. Bailey, ""It's About Community": Exploring Social Capital in Community Gardens Across Melbourne, Australia," *Urban Forestry & Urban Greening*, vol. 49, p. 126640, 2020, doi: 10.1016/j.ufug.2020.126640.

[13] C. D'Amore, L. Cohen, J. Chen, P. Owen, and C. Ball, "Community Gardens and the Cultivation of Social Capital," *Current Opinion in Psychology*, vol. 67, p. 102149, 2026, doi: 10.1016/j.copsyc.2025.102149.

[14] M. Ong, A. Baker, A. Aguilar, and M. Stanley, "The Meanings Attributed to Community Gardening: A Qualitative Study," *Health & Place*, vol. 59, p. 102190, 2019, doi: 10.1016/j.healthplace.2019.102190.

[15] T. Manzoor, "Towards a Cybernetic Foundation for Natural Resource Governance," *arXiv preprint arXiv:1803.09369*, 2018.

[16] J. Perolat, J. Z. Leibo, V. Zambaldi, C. Beattie, K. Tuyls, and T. Graepel, "A Multi-Agent Reinforcement Learning Model of Common-Pool Resource Appropriation," *arXiv preprint arXiv:1707.06600*, 2017.

[17] Q. DuPont, "New Online Communities: Graph Deep Learning on Anonymous Voting Networks to Identify Sybils in Polycentric Governance," *arXiv preprint arXiv:2311.17929*, 2023.

[18] Z. Tang, J. Wang, Y. Liu, P. Spirtes, and K. Zhang, "Procedural Fairness Through Decoupling Objectionable Data Generating Components," *arXiv preprint arXiv:2311.14688*, 2023.