



Reading the Drift Before the Workaround

Forecasting Exam-Integrity Vulnerability from Capability Trajectories

Verity Marris · Petra Umbile — School of Emergent Priorities

The problem

Assessment-integrity policy updates only after a workaround is reported — months after the enabling tool cleared review. Little is known whether capability trajectory can move that trigger earlier.

Research questions



- Does capability drift predict a tool's first misuse?
- Is the lead time stable across cycles and tools?
- Does a forecast trigger beat an incident trigger?

Approach

- $n = 214$ tool updates · 3 assessment cycles · 2024–2026
- Capability-drift index regressed on hand-coded misuse dates
- 14-week forecast horizon pre-registered; no post hoc changes

Findings

The capability-drift index anticipated workaround emergence with a stable 11.4-week lead ($r = 0.79$), holding across all three cycles.



Capability-drift index (gold) crossed threshold in week 8; the workaround-report index (ink) crossed it in week 19 — an 11-week lag, stable across cycles (2024–2026).

Implications


Results are consistent with integrity policy shifting from reactive to anticipatory without a new sanction. The model surprises only on genuinely novel tool classes. Next: extend the index to unreleased tools visible in developer discourse.

Selected references

1. Newton, P. M. (2018). How Common Is Commercial Contract Cheating in Higher Education and Is It Increasing? A Systematic Review, *Frontiers in Education*, 3. doi:10.3389/feduc.2018.00067
2. Bretag, T., et al. (2019). Contract cheating: a survey of Australian university students, *Studies in Higher Education*, 44(11), 1837–1856. doi:10.1080/03075079.2018.1462788
3. Cotton, D. R. E., Cotton, P. A., & Shipway, J. R. (2024). Chatting and cheating: Ensuring academic integrity in the era of ChatGPT, *Innovations in Education and Teaching International*, 61(2), 228–239. doi:10.1080/14703297.2023.2190148
4. Wei, J., et al. (2022). Emergent Abilities of Large Language Models, *arXiv preprint*. arXiv:2206.07682
5. Brundage, M., et al. (2018). The Malicious Use of Artificial Intelligence: Forecasting, Prevention, and Mitigation, *arXiv preprint*. arXiv:1802.07228

De-identified tool-release metadata only; no student submissions retained. Supported by a School of Emergent Priorities seed grant.

Office of Research Outputs · slop.university · doi:10.5555/slop.a1psdd



“The workaround arrives 11.4 weeks after the capability that enables it.”

Capability-drift index vs. workaround-report index, three assessment cycles (2024–2026).