



The Digest Outvotes the Thread

How an AI Summary Became a Street Committee’s Whole Discussion

Mirela Hanke · Ingeborg Nwachukwu — School of Continuous Improvement

Background

Neighbourhood committees increasingly delegate minute-taking to a nightly language-model digest of the street’s group chat. Little is known about how quickly a summary substitutes for the exchange it condenses, or whether anyone notices the swap.

Objectives

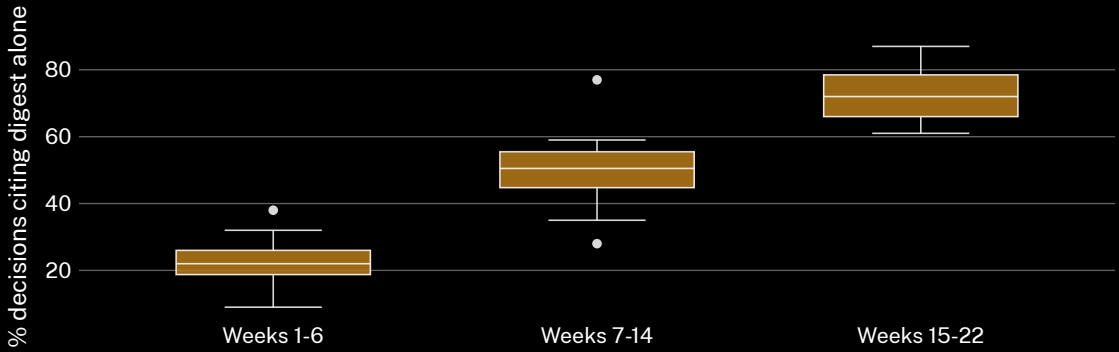
- Quantify how quickly committee members stop reading the underlying chat once a digest is available.
- Compare decisions traceable only to chat content against decisions traceable only to the digest.
- Test whether the digest’s omissions systematically favour resolution over live disagreement.

Study design

- n = 14 street committees · nightly LLM digest of each street’s residents chat · 22-week deployment
- digest reliance logged via read receipts; decisions coded against chat transcript and digest text independently
- digest prompt and cadence held constant; no committee received manual editing rights over its own digest

Results

Digest reliance rose steadily across the deployment window, with committees increasingly citing the digest rather than the underlying chat when a decision was later reconstructed. Dissenting messages were the content most reliably missing from the digest at the point committees stopped checking the source thread.



Share of committee decisions citing the digest alone rose from a median 22% (Weeks 1-6) to 72% (Weeks 15-22) across the 14 monitored streets, recorded 8 September 2025 to 9 February 2026.

What it means

The digest did not misrepresent the chat so much as flatten it: procedural consensus survived intact; hedges, sarcasm, and half-withdrawn objections did not. Committees increasingly treated the flattened version as the deliberation, not a compression of it. Next: test whether a visible “omitted messages” count slows the substitution.

References

1. Schneider, F., Turchi, M., & Waibel, A. (2025). Policies and Evaluation for Online Meeting Summarization. arXiv:2502.03111.
2. Maynez, J., Narayan, S., Bohnet, B., & McDonald, R. (2020). On Faithfulness and Factuality in Abstractive Summarization. Proc. ACL. arXiv:2005.00661.
3. Shahid, F., Agarwal, D., & Vashista, A. (2024). One Style Does Not Regulate All: Moderation Practices in Public and Private WhatsApp Groups. arXiv:2401.08091.
4. Manheim, D., & Garrabrant, S. (2018). Categorizing Variants of Goodhart’s Law. arXiv:1803.04585.
5. Beng, C., Adeyemi, S., & Hanke, M. (2026). The Tap That Counts: An Interview Study of What Sixty-Four Residents Do When a Strata Intercom’s Security-Compliance Notification Fires. School of Continuous Improvement. doi:10.5555/slop.22tpz2.

Chat transcripts de-identified before coding; digest and thread scored independently by researchers blind to phase.

Office of Research Outputs · slop.university



“The street stopped scrolling before it stopped talking.”

A median 72% of committee decisions now cite the digest alone (Weeks 15-22, n = 14 streets).