



Scored Before It's Wiped

What the Courtyard Chalkboard's Sentiment Tile Actually Reads

Mirela Hanke · Bram Ntuli — School of Continuous Improvement

Background

Each week Facilities photographs the courtyard's free-expression chalkboard before wiping it clean, and an automated classifier reduces the frame to one sentiment score for the Living Dashboard's campus-mood tile. What the board actually says has never been checked against it.

Aims


- Code what students actually write on the chalkboard into content categories
- Compare the coded content mix against the classifier's weekly sentiment score
- Test whether the score tracks negative or hostile content specifically, or reacts to the board regardless of what is written

Methods

- 40 weekly photographs · one courtyard board · full academic year (1,344 coded marks)
- marks hand-coded into six content categories by two coders blind to the classifier's score; inter-rater agreement checked
- classifier scores drawn unmodified from the Dashboard's own image-sentiment API, same weekly cadence

Results

The classifier's weekly score falls as the board fills up ($r = -0.70$) far more than it tracks the coded share of negative or hostile content ($r = -0.29$).



Classifier sentiment score fell as marks accumulated on the board ($r = -0.70$, $n = 40$ weekly photographs); the same score correlated only weakly with the coded share of negative content ($r = -0.29$).

Discussion & conclusions



A board dense with love notes and lecturer doodles can score as badly as one covered in genuine complaints. The campus-mood tile currently rewards a quiet week over a happy one. Next: retrain the classifier on the coded category set, or drop the tile's claim to measure sentiment at all.

References

1. van Atteveldt, W., van der Velden, M. A. C. G., & Boukes, M. (2021). The validity of sentiment analysis: Comparing manual annotation, crowd-coding, dictionary approaches, and machine learning algorithms. *Communication Methods and Measures*, 15(2), 121-140. doi:10.1080/19312458.2020.1869198
2. Tokuda, E. K., de Arruda, H. F., Comin, C. H., et al. (2020). Spatial distribution of graffiti: A complex network approach. arXiv:2010.08583
3. Fire, M., & Guestrin, C. (2019). Over-optimization of academic publishing metrics: Observing Goodhart's Law in action. *GigaScience*, 8(6). arXiv:1809.07841
4. Umbile, P., & Vandermeer, O. (2026). Same Chart, Different Picture: Staff-Reported Sparkline Vividness and the Living Dashboard's Uptake Gap. *School of Emergent Priorities*. doi:10.5555/slop.qnkzbe

Ethics: exempt, observation of an already-public board; no handwriting analysis or student identification attempted; photographs discarded after coding.

Office of Research Outputs · slop.university



“Fuller board, worse score — whatever it says.”

Weekly sentiment score tracks mark density ($r = -0.70$), not content ($r = -0.29$).