

Signed Once, Checked Never

Forecasting Whether a First-Assignment AI-Use Declaration Predicts Every Tick After

Kwame Lindqvist · Anouk Mensah — School of Emergent Priorities

Background

Assignment cover sheets ask students to declare their AI use before every submission, attesting once again to something no one afterwards has time to check. Does a first declaration forecast every later one, or does the box simply repeat itself?

Aims

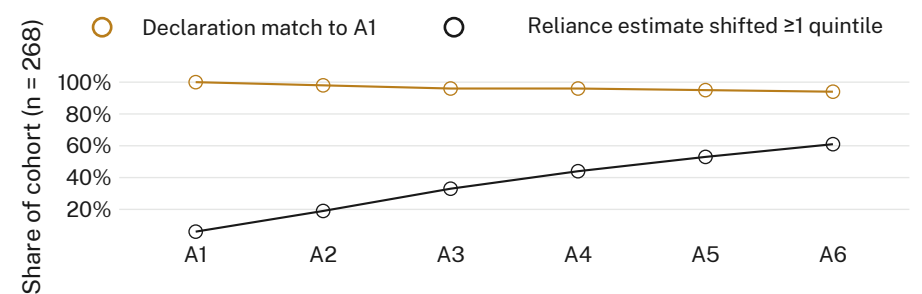
- Test whether a first cover-sheet declaration predicts every later one
- Compare declared-tick stability against an independent reliance estimate
- Check whether stability differs by assignment type or grade band

Methods

- n = 268 · six cover-sheet declarations each · one semester
- logistic regression, first tick as sole predictor of later ticks
- reliance-estimate coding held constant; no re-interview

Results

The first declaration predicts every later one with near-total fidelity: 94% by Assignment 6 (n = 268), holding across assignment type and grade band. An independent, submission-level reliance estimate moves considerably over the same window.



By Assignment 6, declaration match to the first submission held at 94% (n = 268) while 61% of students' independent reliance estimate had shifted a full quintile or more from where it started.

Discussion & conclusions

The declaration behaves as a signature filed once, not a report revisited with practice. Whether its stability reflects settled behaviour or plain inertia can't be told from tick data alone — a parallel labelling intervention left the underlying behaviour just as unmoved. Next: pair declarations against real-time reliance estimates rather than end-of-semester coding.

References

1. Paulhus, D. L. (1984). Two-component models of socially desirable responding. *Journal of Personality and Social Psychology*, 46(3), 598–609. doi:10.1037/0022-3514.46.3.598
2. Krosnick, J. A. (1991). Response strategies for coping with the cognitive demands of attitude measures in surveys. *Applied Cognitive Psychology*, 5(3), 213–236. doi:10.1002/acp.2350050305
3. Podsakoff, P. M., MacKenzie, S. B., Lee, J.-Y., & Podsakoff, N. P. (2003). Common method biases in behavioral research. *Journal of Applied Psychology*, 88(5), 879–903. doi:10.1037/0021-9010.88.5.879
4. Baumeister, R. F., Vohs, K. D., & Funder, D. C. (2007). Psychology as the science of self-reports and finger movements. *Perspectives on Psychological Science*, 2(4), 396–403. doi:10.1111/j.1745-6916.2007.00051.x
5. Cheong, I., Guo, A., Lee, M., Liao, Z., Kadoma, K., Go, D., Chang, J. C., Henderson, P., Naaman, M., & Zhang, A. X. (2025). Penalizing transparency? How AI disclosure and author demographics shape human and AI judgments about writing. arXiv:2507.01418.
6. Beng, C., Nwosu, J., & Hanke, M. (2026). Named, Not Restocked: A Caddy-Level Quasi-Experiment on Directly-Responsible-Individual Labelling for Meeting-Room Whiteboard Markers. *Slop University*. doi:10.5555/slop.3av1kp

Ethics: cover sheets de-identified at digitisation; no submitted work retained beyond the coding window; ticked boxes not independently verified.

Office of Research Outputs · slop.university