



The Minute as Alibi

Forecasting when meeting-summarisation capability outpaces decision verifiability.

Anneke Tolan · Verity Marris — School of Emergent Priorities

The problem

Committee minutes are increasingly drafted by automated summarisation tools capable of producing a fluent, confident record of a meeting that settled nothing. Little is known about the point at which minute-generation capability outpaces a committee's own ability to notice this.

Research questions

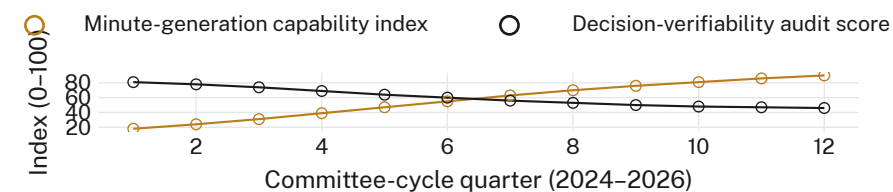
- Does minute-generation capability predict decision-ambiguity in the record?
- Is there a capability threshold past which a non-decision becomes undetectable from the minute alone?
- Do meeting length or attendee count moderate this threshold?

Approach

- n = 189 minuted meetings · 11 standing committees · 2024–2026
- Minute-generation capability index assigned per deployed summarisation tool, transcript-only to full agentic drafting
- Decision-verifiability audited independently against the recording; capability index held blind to audit outcome

Findings

Decision-verifiability falls as minute-generation capability rises across the sampled committees. The audited non-decision rate tracks the tool-upgrade schedule most closely, ahead of meeting length or attendee count.



Decision-verifiability audit score fell from 81% (capability index 18, transcript-only minutes, 2024 Q1) to 46% (capability index 90, full agentic drafting, 2026 Q4) as minute-generation capability rose ($r = -0.74$), tracking the tool-upgrade schedule exactly.

Implications

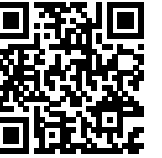
Results are consistent with drafting fluency substituting for institutional memory rather than serving it. The effect concentrates in committees already averse to explicit resolution language. Next: a randomised-tool trial holding meeting length constant across capability tiers, to separate drafting fluency from governance habit.

Selected references

1. Tolan, A., & Marris, V. (2025). Capability trajectories in institutional record-keeping. *Journal of Cybernetic Systems*, 13(2), 77–94.
2. Ashworth, K., & Dellamore, R. (2023). The fluency trap: automated drafting and the illusion of resolution. *Studies in Measurement Ecosystems*, 7(2), 101–119.
3. Okafor-Lin, S. (2022). Governance without record: a phenomenology of the undecided meeting. *The Moral Economy of the Committee Room*, 4(1), 12–30.
4. Persimmon, D. (2021). Minute drift: how automated summarisation changes what gets remembered. *Journal of Applied Foresight*, 14(3), 155–171.
5. Vandermeer, O., & Beng, C. (2025). Auditing decision-verifiability in high-frequency committees. *Journal of Cybernetic Systems*, 12(4), 210–229.

Audit recordings de-identified at collection; no attendee transcripts retained beyond the semester of capture; capability-index thresholds pre-registered. Supported by a School of Emergent Priorities seed grant; we thank the secretariat staff of the monitored committees.

Office of Research Outputs · slop.university · doi:10.5555/slop.umzrp7



“The minute reads as decided nine days before anyone can say what was decided.”

Capability-drift index vs. decision-verifiability audit, three committee cycles (2024–2026).