



Filed as New

Validating the Role-Continuity Classifier Against an Independent Archivist Read

Marek Solheim · Mirela Hanke — School of Emergent Priorities

Motivation

Slop University's Role-Continuity Classifier decides whether a returning appointment is filed as a reinstatement, carrying its service history and leave accrual forward, or a fresh hire that resets the clock. Whether the Classifier reads a file the way an independent archivist would has not been tested.

Objectives


- Measure Classifier–panel agreement on reinstatement calls across role families
- Identify which file features predict a disagreement
- Test whether the discontinuation-to-reinstatement gap length drives the split

Study design

- n = 96 reinstated posts · five schools · four-year window (2022–2026)
- three-archivist panel, blind to the Classifier's verdict, majority rule
- gap length and role family as covariates; pre-registered coding sheet

Results

Panel-versus-Classifier agreement across all 96 files was moderate (Cohen's $\kappa = 0.44$, Landis-Koch scale), driven down almost entirely by one role family: of 31 invigilation and proctoring posts, the Classifier filed 26 as new appointments, and the panel read genuine continuity in 22 of those 26.



Among 31 invigilation and proctoring files reopened between March 2025 and June 2026, the Classifier filed 26 (84%) as new appointments; the archivist panel read genuine continuity in 22 of those 26.

What it means

The Classifier does not fail evenly: its blind spot sits precisely on the roles this year's supervised-exam revival is refilling fastest, filing genuine returns as fresh starts and quietly resetting the service record each carries. Gap length explains little of the split on its own; what separates a reinstatement from a new appointment, in the Classifier's reading, is whether the discontinuation was ever minuted — and invigilation posts were retired informally, by simply not being called. A file keeps its history only if a human panel goes back to check it. Next: extend the panel read to the leave-accrual ledger the Classifier's verdict ultimately updates.

Selected references

1. Walsh, J. P., & Ungson, G. R. (1991). Organizational memory. *Academy of Management Review*, 16(1), 57–91. doi:10.5465/amr.1991.4278992
2. Shipp, A. J., Furst-Holloway, S., Harris, T. B., & Rosen, B. (2014). Gone today but here tomorrow: Extending the unfolding model of turnover to consider boomerang employees. *Personnel Psychology*, 67(2), 421–462. doi:10.1111/peps.12039
3. Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46. doi:10.1177/001316446002000104
4. Landis, J. R., & Koch, G. G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33(1), 159–174. doi:10.2307/2529310
5. Fellegi, I. P., & Sunter, A. B. (1969). A theory for record linkage. *Journal of the American Statistical Association*, 64(328), 1183–1210. doi:10.1080/01621459.1969.10501049

Personnel files reviewed under a de-identified data-sharing agreement; panel members held no supervisory relationship to any case file; original ledgers not retained beyond the review window.