



Slovenian / University

Slovenian University Strategic Plan 2026-2031

A named reviewer for every AI-assisted
decision



Contents

A message from the Vice-Chancellor	3
Executive summary	4
Our values	5
Where we stand	6
Where we are headed	7
Strategic priorities	8
Roadmap	13
How we developed this plan	13
Accountability and review	14
The decade ahead	15

A message from the Vice-Chancellor

Across the sector, and across this University's own systems, a growing share of the decisions that affect a person's standing — an application considered, a grant assessed, a roster set, a flag raised against an account — now carries some contribution from an AI-assisted tool before a human ever sees the recommendation. That is not, in itself, the problem this plan is written to solve. The problem is narrower, and we think more consequential: that as this contribution grows, a named person remains answerable for what the tool recommended, able to explain it, and willing to be asked about it.

Slopp University has spent the better part of a decade building the instruments that sense what is coming and measure what has already happened — a horizon-scanning capability that reads the priorities on approach, an evaluation ecosystem that reports faithfully on our own performance. Those instruments tell us a great deal. Neither, on its own, was built to answer the plainer question this plan is concerned with: who, specifically, stands behind a given recommendation on the day someone asks to see it explained.



Over the plan period to come, this University commits to closing that gap directly: no AI-assisted decision affecting a student, staff member, or applicant will take effect without a named person’s attestation that they have seen it and stand behind it. We are not naive about what that commitment will cost as the number of such decisions continues to grow well past what any one reviewer could read in full, line by line, and still get through their caseload. This plan sets out, in specific and checkable terms, how we intend to keep the commitment meaningful anyway — not by asking less of a sign-off as the volume rises, but by budgeting, training, and reviewing the capacity to give it properly.

Vice-Chancellor and President

Executive summary

This Strategic Plan commits Slop University, across the period 2026-2031, to a single organising idea: that no AI-assisted recommendation affecting a student, staff member, or applicant should take effect without a named person attesting that they have read it and stand behind it. Four strategic priorities carry that commitment from the decision itself through to the capacity to review it, the public’s standing to ask about it, and its survival as the number of such decisions keeps rising.



Trusted and Accountable Decision-Making

Every AI-assisted decision affecting a person’s standing is entered on a published Register and paired with a named reviewer’s dated attestation.



Capability, Talent, and Oversight

Reviewer numbers, training, and caseload are budgeted and reported against decision volume, so a named role carries the standing to actually be exercised.



A University for the Public Good

The Register’s attestation rates are published in full, and a standing community panel can query how any AI-assisted workflow is overseen.



Sustainable Oversight

Reviewer capacity is grown ahead of decision volume, and an independent annual audit checks that attestation depth has kept pace with it.

Together these priorities move the University from a default in which an AI-assisted recommendation is actioned unless someone happens to object, to a default in which it is actioned only once someone specific has attested to it. The initiatives that follow are specific by design: each carries a target, a date, and a mechanism for knowing whether it has been met. Where earlier plans have asked the University to “govern AI responsibly” in the abstract, this plan asks something narrower and considerably harder to fake: that every such decision have exactly one name against it, and that the name mean something.

None of the four priorities invents a discipline the University does not already practise somewhere. Each extends a habit already at work in one part of the institution — a named convenor, an attested custodian, a published register — to the specific and fast-growing case of decisions an AI-assisted tool has helped make. That is a deliberate choice: an institution that already knows how to name someone accountable for a research instrument does not need a new capability to do the same for a rostering flag, only the will to apply the existing one without exception, at a volume the existing one was not originally built for.

The remainder of this document sets out the operating environment that makes this commitment necessary, the ambition that follows from it, and the four strategic priorities, their initiatives, and their measures in detail, before closing with the phasing, governance, and review arrangements that will carry the plan through to 2031.

Our values

Accountability — we can always name the person standing behind a decision, including the decisions an AI-assisted tool helped make.	Integrity — a sign-off means what it says it means, whatever the caseload behind it.	Capability — we resource the people who review as carefully as we resource the tools they are reviewing.	Trust — the standing this University holds is measured by whether an outsider, asking, gets a straight answer.
--	---	---	---

Where we stand

The sector we work within increasingly expects an institution to be able to account for a specific decision, not simply to gesture at a governance framework that covers decisions in general. Slop University welcomes that expectation. It sits close to work this University has pursued for years, across the School of Emergent Priorities' anticipatory-capability agenda and the School of Continuous Improvement's evaluation-ecosystem agenda, and it gives that work a sharper edge to press against.

Internally, the number of workflows in which an AI-assisted tool contributes to a decision affecting a person's standing has grown faster than the University's capacity to review each one in full. Admissions triage, grant-round shortlisting, shift rostering, procurement exception flags, and a growing list of smaller decisions besides now arrive at a human desk already carrying a recommendation. In most of these workflows, a human is still formally the one who decides. What has grown much less certain, as the volume has risen, is how much of the recommendation that human actually reads before they sign.

Other institutions across the sector report the same trajectory, in broadly similar terms: an AI-assisted contribution that began in one or two high-visibility workflows, then spread, workflow by workflow, faster than any single governance framework was drafted to anticipate. Most respond by strengthening the framework. Fewer ask the narrower question this plan starts from — not whether a framework exists, but whether the specific person it names to apply it still has the time to do so. Slop University does not regard that gap as a failing unique to us. We do regard it as the gap most worth closing first, precisely because it is the one a framework, on its own, cannot close.



The University's own research bears this out, and not comfortably. Work published from within our own schools has already found that a mandatory AI-literacy sign-off, required of every research student before they use certain tools, tracks whether the student completed the associated course rather than whether they can actually use the tools it certifies them on. A separate study, following how confidently shoppers overrode a self-checkout's own error prompts, found no reliable relationship at all between how much a person said they trusted the system and how they actually behaved in front of it. Neither finding is about this University's governance directly. Read together, they describe the same pattern our own decision workflows are not exempt from: a sign-off, a certificate, or a stated level of trust can persist long after the thing it was meant to certify has stopped being true of the person giving it.

That pattern is precisely what this plan is written against. A named reviewer is only worth naming if naming them changes something about how carefully the decision was actually looked at. This plan sets out to make that the case, deliberately, and at a volume that will keep testing it.

Where we are headed

By 2031, no AI-assisted recommendation affecting a student, staff member, or applicant at Slop University will take effect without a named individual's attestation that they have read it, understood it, and stand behind it.

That ambition would be trivial to meet at low volume and considerably harder to meet at the volume this University actually operates at, which is exactly why it is the ambition. It asks nothing exotic of the institution — only that the habit already practised in its most visible instruments, a named convenor for a register, a named custodian for an instrument, extend to every place an AI-assisted tool now contributes to a decision about a person. What makes the ambition worth stating plainly is the part that follows from it: as the number of such decisions keeps rising, so must the University’s capacity to review them, rather than the depth of what reviewing one is taken to mean.

We are conscious that this ambition can be met in a way that satisfies its letter while hollowing out its intent — a sign-off recorded in seconds against a brief nobody wrote, a name attached to a decision nobody actually weighed. This plan treats that outcome as a live risk to be measured against, not a theoretical one to be waved away in a footnote, which is why the fourth priority below exists to test the other three against it every year of the plan.

Strategic priorities

The shift from a tool’s recommendation standing largely on its own to a named person standing behind it touches every part of the institution differently, and this plan sets out four priorities to carry it through. The table below sets the direction each priority moves the University in.

From	To
An AI-assisted recommendation actioned unless someone happens to object	An AI-assisted recommendation actioned only once a named person attests to it
Oversight capacity fixed at whatever a unit already happens to have	Oversight capacity budgeted and grown against decision volume
A sign-off recorded informally, some time after the decision took effect	A sign-off recorded, dated, and checkable before the decision takes effect
Review depth left to whatever time a reviewer’s caseload allows	Review depth set to a published minimum, regardless of caseload
An AI system’s confidence treated as a stand-in for a person’s judgement	An AI system’s confidence treated as one input to a person’s judgement, never a substitute for it

Trusted and Accountable Decision-Making

Why this matters

The University now runs a great many workflows in which an AI-assisted tool contributes to a decision affecting a person's standing — an admissions triage flag, a grant-round shortlist score, a roster change, a procurement exception — before a human sees the recommendation at all. Trust in that decision rests on a specific person being answerable for it, not on the tool's own confidence score, however precisely stated. This is deliberately the first priority in this plan rather than the last: naming a reviewer is the cheapest possible demonstration that the commitment is a genuine institutional default and not a framing reserved for the decisions important enough to already carry one.

What we'll do

1. Publish and maintain a University-wide AI-Assisted Decision Register, listing every workflow in which such a tool contributes to a decision affecting a student, staff member, or applicant, each entry paired with a named accountable reviewer role.
2. Require a dated, attributable sign-off from that named reviewer before any AI-assisted recommendation is actioned, with no exception sustained beyond a single review cycle.
3. Give any person affected by an AI-assisted decision a standing right to request the reviewer's name and the basis of their attestation.
4. Extend the Register to cover recommendations that inform a decision even where a human formally makes the final call, closing the gap between "assisted" and merely "informed".
5. Report annually on any workflow whose AI-assisted contribution grew significant enough, over the year, to warrant entry on the Register but had not yet been added, and correct the omission within the same reporting cycle.

How we'll know

- 100% of AI-assisted decision workflows entered on the Register by 2028, with zero unregistered exceptions carried past a single review cycle.
- Reviewer attestation completed for 100% of in-scope decisions, audited each quarter with the findings published in full.
- Every request for a reviewer's name and attestation basis answered within ten business days, with the response rate published annually.

Capability, Talent, and Oversight

Why this matters

Naming a reviewer is only meaningful if that reviewer has the standing, training, and time to actually exercise judgement. A mandate without capacity is not a safeguard; it is an incentive to sign quickly rather than to look closely, and the University has no interest in building a Register that rewards the first over the second.



What we'll do

1. Establish a University Reviewer Capability Framework, setting standards for training, caseload, and standing across every unit that operates an AI-assisted workflow.
2. Fund a recognised professional-development pathway for staff who take on named-reviewer roles, reflected in workload models and promotion criteria.
3. Convene an annual Reviewer Capability Census, assessing whether reviewer numbers and training have kept pace with the volume of decisions requiring attestation.
4. Give every named reviewer a standing right to decline a sign-off and refer a decision upward when caseload has left them unable to review it properly.
5. Publish a directory of reviewer numbers by unit, so that a reviewer-to-decision ratio can be checked against the Register's own published decision volume.

How we'll know

- Reviewer-to-decision ratio published annually and held within a stated band, unit by unit.
- 100% of named reviewers complete accredited training within one year of appointment.
- Reviewer capability formally recognised in workload allocation for every unit operating an AI-assisted workflow by 2029.

A University for the Public Good

Why this matters

A Register the University keeps to itself is a governance framework; a Register the public can ask questions of is an accountability. The community this University serves has a legitimate stake in how AI-assisted decisions affecting its members are overseen, and that stake does not stop at the University's own staff and students — it extends to every shared service and partnership acting on the University's behalf.

What we'll do

1. Publish an annual public report of the Register and its attestation rates in aggregate, without redaction to workflow-level figures.
2. Establish a standing community panel able to query how any AI-assisted decision workflow is overseen, though not the individual outcomes it reaches.
3. Extend the same sign-off standard to significant shared services and partnerships that make AI-assisted decisions on the University's behalf.
4. Make the public report's underlying figures available in a form that can be checked against the Register directly, not only summarised.
5. Invite the community panel to nominate one AI-assisted workflow each year for a deeper, published review beyond the standard attestation record.

How we'll know

- Public report published annually with no redactions to workflow-level attestation rates.
- Community panel convened at least twice yearly, with its questions and the Register's responses published in full.
- 100% of qualifying shared-service and partnership arrangements brought under the same attestation standard by 2030.

Sustainable Oversight

Why this matters

The honest test of this plan is not whether the Register launches well in 2026; it is whether the sign-off still means something once the number of AI-assisted decisions has grown well past what a reviewer's caseload allowed when the Register began. This University commits to scaling reviewer capacity ahead of that growth. Where capacity cannot grow as fast as volume does, in any given year, this priority is where that gap is required to surface, named, rather than quietly absorbed into a faster read.



What we'll do

1. Deploy a Reviewer Support Tool that drafts a standard attestation brief summarising each AI-assisted recommendation, reducing the reading time a sign-off requires.
2. Introduce a tiered attestation model, in which decisions below a published materiality threshold receive a streamlined sign-off pathway while decisions above it retain full review.
3. Commission an annual Oversight Integrity Audit, independent of the Register, assessing whether attestation depth has kept pace with decision volume.
4. Require the Reviewer Support Tool itself to carry a named accountable reviewer for any change to how it drafts an attestation brief.
5. Publish the Reviewer Support Tool's own change log alongside the Register, so a change to how an attestation brief is drafted is itself a named, dated decision.

How we'll know

- 100% of AI-assisted decisions retain a named, attributable sign-off regardless of tier, with zero decisions actioned without one.
- Reviewer Support Tool adoption reaches 100% of in-scope units by 2029.
- Oversight Integrity Audit published annually, with any finding of attestation depth falling behind decision volume addressed within a single review cycle.

Roadmap

Named oversight will be built in four phases running contiguously across the plan period, deliberately sequenced so that the discipline is proven on a small, closely watched set of workflows before it is asked to carry the University's full decision volume.

Establish, 2026 — the University stands up the AI-Assisted Decision Register and names its first wave of reviewers, beginning with admissions triage and grant-round shortlisting. The Reviewer Capability Framework is drafted alongside this first wave, so early reviewers are never left to work out the standard from the Register entry alone.

Build, 2027-28 — the Register extends to rostering, procurement exceptions, and the shared services acting on the University's behalf; the funded professional-development pathway for named reviewers opens; the first annual Reviewer Capability Census reports to the Executive. The tiered attestation model is piloted in one School ahead of University-wide rollout, and the public report is published for the first time.

Embed, 2029-30 — the Reviewer Support Tool and the tiered attestation model operate across every in-scope unit; the community panel convenes on its standing schedule; the first Oversight Integrity Audit runs independently of the Register, giving the University its first real read on whether attestation depth has kept pace with decision volume.

Renew, 2031 — every named reviewer and Register entry is re-attested end to end; the Reviewer Capability Census and the Oversight Integrity Audit are reviewed together ahead of the University's next planning cycle. Renewal is treated as routine rather than exceptional: a model built to outlast any one reviewer's tenure, and any one year's decision volume, should be able to survive its own first full turnover without visible strain.

How we developed this plan

This plan was shaped through consultation across the University community over the twelve months preceding its publication: facilitated workshops with staff already operating AI-assisted workflows, structured interviews with reviewers carrying the heaviest existing caseloads, and a standing submission process open to students, professional staff, and alumni. The consultation surfaced the same concern

independently and repeatedly — that a sign-off is easy to mandate and much harder to keep meaningful once the caseload behind it grows — and that concern, more than any single workshop finding, set the shape of the priorities that follow.

Draft measures were tested back against the reviewers who had raised the original concern, rather than developed in isolation and presented as finished. The ten-business-day response standard for a reviewer's name and attestation basis, and the decision to make the Oversight Integrity Audit independent of the Register it examines, both came directly from that testing round, where an earlier draft had proposed folding the Audit into the Register's own reporting line.

Accountability and review

The University Council retains standing oversight of the AI-Assisted Decision Register, receiving it as a fixed agenda item at each of its ordinary meetings across the plan period. The Executive receives an annual report tabling Register coverage, reviewer capacity, and attestation performance against the measures set out in this plan.

The Review Cadence Observatory monitors the Register's own review cycle, flagging any entry whose attestation has lapsed, and the Indicator Commons holds the shared indicators this plan reports against, consistent with the University's existing evaluation-ecosystem practice. As Dr Dagny Okafor, Convenor of the School of Continuous Improvement's Evaluation of Evaluation program, put it during the consultation that shaped this plan: a review that never asks who is doing the reviewing is not yet a review. That question sits at the centre of the Oversight Integrity Audit's remit, and the Audit reports to the Executive independently of the Register it examines.

No new committee is created to hold this work; it is carried by the governance structures the University already has, with one addition — every one of those structures will now be able to say, for any AI-assisted decision within its remit, who specifically stood behind it, and on what basis. Where a unit disagrees with a finding of the Oversight Integrity Audit, the matter is escalated to the Executive rather than resolved informally, so that a disputed finding is settled on the record rather than left to lapse until the next review cycle.

The decade ahead

By 2031, an admissions flag, a grant-round score, and a rostering recommendation will each resolve to the same thing: a named person, on record, who read it and stands behind it. None of these are large ambitions taken alone. Taken together, across a decision volume this University expects to keep rising, they close the distance between what this institution says it is accountable for and what it can actually show, on the day someone asks, that it is.

This is not, in the end, a plan about any one admissions cycle, roster, or procurement flag. It is a plan about what an institution is willing to be held to as the volume of its own decisions grows past what any one person can read, and whether it can still say, honestly, that someone did. Slop University intends, by the close of this plan, to be able to say so every time.



Slop / University